

SRM Exercise

Bartu Kadir Tamer

2025-06-27

Uploading the Related Packages:

```
required_packages <- c(
  "ggplot2", "Hmisc", "car", "caret", "MASS", "randomForest",
  "xgboost", "tidyverse", "tidyr", "dplyr", "rpart", "rpart.plot",
  "ipred", "e1071"
)

# Install missing packages only
install_if_missing <- function(pkg) {
  if (!requireNamespace(pkg, quietly = TRUE)) {
    install.packages(pkg, dependencies = TRUE)
  }
}

# Loop through and install if needed
invisible(lapply(required_packages, install_if_missing))
```

```
library(ggplot2)
library(Hmisc) #Used for descriptive statistics
```

```
##
## Attaching package: 'Hmisc'

## The following objects are masked from 'package:base':
##
##   format.pval, units
```

```
library(car)
```

```
## Loading required package: carData
```

```
library(caret)
```

```
## Loading required package: lattice
```

```
library(MASS)
library(randomForest)
```

```
## randomForest 4.7-1.2
```

```
## Type rfNews() to see new features/changes/bug fixes.
```

```
##
```

```
## Attaching package: 'randomForest'
```

```
## The following object is masked from 'package:ggplot2':
```

```
##
```

```
##     margin
```

```
library(xgboost)
library(tidyverse)
```

```
## -- Attaching core tidyverse packages ----- tidyverse 2.0.0 --
```

```
## v dplyr      1.1.4      v readr      2.1.5
```

```
## v forcats    1.0.0      v stringr    1.5.1
```

```
## v lubridate  1.9.3      v tibble     3.2.1
```

```
## v purrr      1.0.2      v tidyr      1.3.1
```

```
## -- Conflicts ----- tidyverse_conflicts() --
```

```
## x dplyr::combine()      masks randomForest::combine()
```

```
## x dplyr::filter()       masks stats::filter()
```

```
## x dplyr::lag()          masks stats::lag()
```

```
## x purrr::lift()         masks caret::lift()
```

```
## x randomForest::margin() masks ggplot2::margin()
```

```
## x dplyr::recode()       masks car::recode()
```

```
## x dplyr::select()       masks MASS::select()
```

```
## x dplyr::slice()        masks xgboost::slice()
```

```
## x purrr::some()         masks car::some()
```

```
## x dplyr::src()          masks Hmisc::src()
```

```
## x dplyr::summarize()     masks Hmisc::summarize()
```

```
## i Use the conflicted package (<http://conflicted.r-lib.org/>) to force all conflicts to become errors
```

```
library(tidyr)
library(dplyr)
library(rpart)
library(rpart.plot)
library(ipred)
library(e1071)
```

```
##
```

```
## Attaching package: 'e1071'
```

```
##
```

```
## The following object is masked from 'package:Hmisc':
```

```
##
```

```
##     impute
```

To understand working directory of the Rstudio. It is important to have the file which we want to analyze in same directory.

```
getwd()
```

```
## [1] "/Users/ben/Downloads/SRM Exercise/SRM Exercise"
```

If the Downloads file contains the dataset you can copy the path and paste it into the `read.csv()` function.

Uploading and Viewing the Dataset:

```
df=read.csv("EIB2023_dummies_filtered.csv",header = T)
View(df)
```

Data Summary:

Variables :

`progr_binary`: Support for progressive taxation on carbon consumption (rich pay more, poor pay less on fuel taxes)

`ctax_binary` (Response variable for this exercise): Support for taxing environmentally harmful profits (e.g., from fossil fuels)

`subsidies_binary`: Support for removing fossil fuel subsidies and redirecting funds to renewables

`income_scale`: Country-normalized income score (standardized by mean and SD)

`any_cc_last2year_factor`: Whether the respondent's region had major climate disasters in the last 2 years (yes/no)

`gender`: Respondent's gender

`educationsecondary`: Dummy = 1 if respondent completed secondary education

`educationtertiary`: Dummy = 1 if respondent completed tertiary education

`age`: Respondent's age (can be used directly or scaled as `age_scale`)

`urbanizationtown`: Dummy = 1 if respondent lives in a town/suburban area

`has_children`: Dummy = 1 if respondent has children

`trust`: Respondent's trust in their country's ability to fight climate change and maintain social equity Levels: Very confident, Rather confident, Not really confident, Not confident at all

`lr_scale`: Left-right political ideology scale (0 = far left, 10 = far right)

`new_ccknowledge_index`: Composite score of climate change knowledge

`countryBelgium`, `countryFrance`, ..., `countrySweden`, etc. 0/1 dummies for each country

`str()` function helps to understand dimension and structures of the objects in the dataset.

```
str(df)
```

```
## 'data.frame': 22729 obs. of 47 variables:
## $ progr_binary : int 1 0 1 1 1 1 1 1 1 0 ...
## $ ctax_binary : int 1 1 1 1 1 1 1 1 1 0 ...
## $ subsidies_binary : int 1 1 1 1 1 1 1 1 1 0 ...
## $ income_scale : num -1.403 NA -0.446 0.484 1.057 ...
## $ age_scale : num -9.54 -31.54 -17.54 5.46 20.46 ...
## $ trust : chr "No really confident" "Rather confident" "No really confident"
## $ LR_scale_scale : num 4.348 1.348 0.348 1.348 -1.652 ...
## $ new_ccknowledge_index_scale: num 0.0249 -0.3084 0.0249 0.136 0.0249 ...
## $ any_cc_last2year_factor : int 0 0 0 0 0 0 0 0 0 0 ...
## $ regional_heterogeneity : int 1 1 1 1 1 1 1 1 1 1 ...
## $ LR_scale : int 10 7 6 7 4 3 4 NA 2 NA ...
## $ income : chr "Less than 7 070 €" "Prefer not to say" "14 230 € to under 16 5
## $ new_ccknowledge_index : num 0.667 0.333 0.667 0.778 0.667 ...
## $ gender : chr "male" "male" "female" "female" ...
## $ country_w : num 0.962 1.134 0.934 0.911 1.009 ...
## $ educationprimary : int 0 0 0 0 0 0 0 0 0 0 ...
## $ educationsecondary : int 1 1 1 0 1 0 0 0 0 1 ...
## $ educationtertiary : int 0 0 0 1 0 1 1 1 1 0 ...
## $ urbanizationtown : int 0 1 1 0 0 1 1 0 1 1 ...
## $ urbanizationrural : int 1 0 0 0 1 0 0 0 0 0 ...
## $ has_childrenyes : int 0 0 0 0 0 1 1 0 0 0 ...
## $ countryBelgium : int 0 0 0 0 0 0 0 0 0 0 ...
## $ countryBulgaria : int 0 0 0 0 0 0 0 0 0 0 ...
## $ countryCroatia : int 0 0 0 0 0 0 0 0 0 0 ...
## $ countryCyprus : int 0 0 0 0 0 0 0 0 0 0 ...
## $ countryCzech.Republic : int 0 0 0 0 0 0 0 0 0 0 ...
## $ countryDenmark : int 0 0 0 0 0 0 0 0 0 0 ...
## $ countryEstonia : int 0 0 0 0 0 0 0 0 0 0 ...
## $ countryFinland : int 0 0 0 0 0 0 0 0 0 0 ...
## $ countryFrance : int 0 0 0 0 0 0 0 0 0 0 ...
## $ countryGermany : int 0 0 0 0 0 0 0 0 0 0 ...
## $ countryGreece : int 0 0 0 0 0 0 0 0 0 0 ...
## $ countryHungary : int 0 0 0 0 0 0 0 0 0 0 ...
## $ countryIreland : int 0 0 0 0 0 0 0 0 0 0 ...
## $ countryItaly : int 1 1 1 1 1 1 1 1 1 1 ...
## $ countryLatvia : int 0 0 0 0 0 0 0 0 0 0 ...
## $ countryLithuania : int 0 0 0 0 0 0 0 0 0 0 ...
## $ countryLuxembourg : int 0 0 0 0 0 0 0 0 0 0 ...
## $ countryMalta : int 0 0 0 0 0 0 0 0 0 0 ...
## $ countryPoland : int 0 0 0 0 0 0 0 0 0 0 ...
## $ countryPortugal : int 0 0 0 0 0 0 0 0 0 0 ...
## $ countryRomania : int 0 0 0 0 0 0 0 0 0 0 ...
## $ countrySlovakia : int 0 0 0 0 0 0 0 0 0 0 ...
## $ countrySlovenia : int 0 0 0 0 0 0 0 0 0 0 ...
## $ countrySpain : int 0 0 0 0 0 0 0 0 0 0 ...
## $ countrySweden : int 0 0 0 0 0 0 0 0 0 0 ...
## $ countryThe.Netherlands : int 0 0 0 0 0 0 0 0 0 0 ...
```

Counting the number of NA values.

```
sum(is.na(df))
```

```
## [1] 7064
```

We have 7064 NA values in the dataset. Some of the ML function that we are using the packages causes error when we want to deal with NA generated dataset. Therefore, it is better to getting rid of them.

```
df_clean=na.omit(df)
sum(is.na(df_clean))
```

```
## [1] 0
```

We will use the ctax_binary as our response in this analysis, so we can delete the other binary responses.

```
df_clean=df_clean[, -c(1,3,12)]
```

We can check the descriptive statistics of the dataset for understanding the first insights of the dataset.

```
describe(df_clean)
```

```
## df_clean
##
## 44 Variables      18698 Observations
## -----
## ctax_binary
##      n missing distinct      Info      Sum      Mean
##  18698         0         2    0.524   14479    0.7744
##
## -----
## income_scale
##      n missing distinct      Info      Mean  pMedian      Gmd
##  18698         0      270         1    0.03855  0.03021    1.143
##      .05      .10      .25      .50      .75      .90      .95
## -1.3884036 -1.3140793 -0.9228093  0.0004226  1.0398237  1.3125329  1.4025175
##
## lowest : -1.77838 -1.64296 -1.61668 -1.51722 -1.49367
## highest:  1.57411  1.61239  1.66954  1.67938  1.79275
## -----
## age_scale
##      n missing distinct      Info      Mean  pMedian      Gmd      .05
##  18698         0      1767         1    1.145    1.105    19.2 -25.4740
##      .10      .25      .50      .75      .90      .95
## -22.4740 -12.4740  0.9247  15.5260  23.1522  26.4630
##
## lowest : -31.5352 -31.474 -30.593 -30.5669 -30.561
## highest:  45.4678  46.1922  47.463  51.4331  53.664
## -----
## trust
##      n missing distinct
##  18698         0         4
##
## Value      No confident at all No really confident      Rather confident
## Frequency              3016              8373              5763
## Proportion              0.161              0.448              0.308
##
## Value      Very confident
```

```

## Frequency          1546
## Proportion        0.083
## -----
## LR_scale_scale
##      n missing distinct      Info      Mean  pMedian      Gmd      .05
##    18698      0      270        1  0.01449 -0.05924    2.395  -3.6809
##      .10      .25      .50      .75      .90      .95
##    -2.6809  -1.1092  -0.3348  1.3482  2.9913  4.2339
##
## lowest : -5          -4.878   -4.84746 -4.83556 -4.82637
## highest: 4.70789    4.89079  4.94488  4.98488  4.99129
## -----
## new_ccknowledge_index_scale
##      n missing distinct      Info      Mean  pMedian      Gmd      .05
##    18698      0      412        1  0.01096  0.01894    0.1795  -0.27900
##      .10      .25      .50      .75      .90      .95
##   -0.21584 -0.09509  0.02644  0.12713  0.20194  0.24044
##
## lowest : -0.533278 -0.510613 -0.508222 -0.499557 -0.497397
## highest: 0.380189  0.381778  0.381889  0.39527   0.40071
## -----
## any_cc_last2year_factor
##      n missing distinct      Info      Sum      Mean
##    18698      0        2    0.738    8174    0.4372
##
## -----
## regional_heterogeneity
##      n missing distinct      Info      Sum      Mean
##    18698      0        2    0.587   13701    0.7328
##
## -----
## LR_scale
##      n missing distinct      Info      Mean  pMedian      Gmd      .05
##    18698      0        10    0.966    5.546     5.5    2.387     2
##      .10      .25      .50      .75      .90      .95
##        3        4        5        7        8        10
##
## Value          1      2      3      4      5      6      7      8      9     10
## Frequency      750    702   1584   1847   5612   2599   1985   1767    799   1053
## Proportion  0.040  0.038  0.085  0.099  0.300  0.139  0.106  0.095  0.043  0.056
##
## For the frequency table, variable is rounded to the nearest 0
## -----
## new_ccknowledge_index
##      n missing distinct      Info      Mean  pMedian      Gmd      .05
##    18698      0        17    0.989    0.6643  0.6667    0.1826  0.3889
##      .10      .25      .50      .75      .90      .95
##    0.4444  0.5556  0.6667  0.7778  0.8333  0.8889
##
## 0.1111111111111111 (1, 0.000), 0.166666666666667 (18, 0.001), 0.222222222222222
## (90, 0.005), 0.277777777777778 (275, 0.015), 0.333333333333333 (485, 0.026),
## 0.388888888888889 (713, 0.038), 0.444444444444444 (954, 0.051), 0.5 (1354,
## 0.072), 0.555555555555556 (1576, 0.084), 0.611111111111111 (2021, 0.108),
## 0.666666666666667 (2400, 0.128), 0.722222222222222 (2565, 0.137),

```

```

## 0.7777777777777778 (2399, 0.128), 0.8333333333333333 (2085, 0.112),
## 0.8888888888888889 (1057, 0.057), 0.9444444444444444 (544, 0.029), 1 (161, 0.009)
##
## For the frequency table, variable is rounded to the nearest 0
## -----
## gender
##      n missing distinct
##  18698      0      2
##
## Value      female      male
## Frequency   9244   9454
## Proportion  0.494  0.506
## -----
## country_w
##      n missing distinct      Info      Mean  pMedian      Gmd      .05
##  18698      0      6967      1  0.9945  0.9657  0.2611  0.6620
##      .10      .25      .50      .75      .90      .95
##  0.7496  0.8595  0.9545  1.0721  1.2468  1.4222
##
## lowest : 0.158708 0.160538 0.174153 0.237714 0.243678
## highest: 3.73711 3.75487 3.77212 4.85878 6.25698
## -----
## educationprimary
##      n missing distinct      Info      Sum      Mean
##  18698      0      2  0.359  2596  0.1388
##
## -----
## educationsecondary
##      n missing distinct      Info      Sum      Mean
##  18698      0      2  0.724  7595  0.4062
##
## -----
## educationtertiary
##      n missing distinct      Info      Sum      Mean
##  18698      0      2  0.744  8507  0.455
##
## -----
## urbanizationtown
##      n missing distinct      Info      Sum      Mean
##  18698      0      2  0.711  7217  0.386
##
## -----
## urbanizationrural
##      n missing distinct      Info      Sum      Mean
##  18698      0      2  0.476  3696  0.1977
##
## -----
## has_childrenyes
##      n missing distinct      Info      Sum      Mean
##  18698      0      2  0.672  6327  0.3384
##
## -----
## countryBelgium
##      n missing distinct      Info      Sum      Mean

```

```

##      18698      0      2      0.119      775      0.04145
##
## -----
## countryBulgaria
##      n missing distinct      Info      Sum      Mean
##      18698      0      2      0.13      848      0.04535
##
## -----
## countryCroatia
##      n missing distinct      Info      Sum      Mean
##      18698      0      2      0.124      809      0.04327
##
## -----
## countryCyprus
##      n missing distinct      Info      Sum      Mean
##      18698      0      2      0.055      352      0.01883
##
## -----
## countryCzech.Republic
##      n missing distinct      Info      Sum      Mean
##      18698      0      2      0.129      842      0.04503
##
## -----
## countryDenmark
##      n missing distinct      Info      Sum      Mean
##      18698      0      2      0.126      820      0.04385
##
## -----
## countryEstonia
##      n missing distinct      Info      Sum      Mean
##      18698      0      2      0.065      414      0.02214
##
## -----
## countryFinland
##      n missing distinct      Info      Sum      Mean
##      18698      0      2      0.129      842      0.04503
##
## -----
## countryFrance
##      n missing distinct      Info      Sum      Mean
##      18698      0      2      0.12      781      0.04177
##
## -----
## countryGermany
##      n missing distinct      Info      Sum      Mean
##      18698      0      2      0.136      887      0.04744
##
## -----
## countryGreece
##      n missing distinct      Info      Sum      Mean
##      18698      0      2      0.133      868      0.04642
##
## -----
## countryHungary

```



```

##          n missing distinct      Info      Sum      Mean
##    18698          0          2    0.118      770  0.04118
##
## -----
## countryIreland
##          n missing distinct      Info      Sum      Mean
##    18698          0          2    0.132      860  0.04599
##
## -----
## countryItaly
##          n missing distinct      Info      Sum      Mean
##    18698          0          2    0.121      789  0.0422
##
## -----
## countryLatvia
##          n missing distinct      Info      Sum      Mean
##    18698          0          2    0.064      406  0.02171
##
## -----
## countryLithuania
##          n missing distinct      Info      Sum      Mean
##    18698          0          2    0.065      412  0.02203
##
## -----
## countryLuxembourg
##          n missing distinct      Info      Sum      Mean
##    18698          0          2    0.056      354  0.01893
##
## -----
## countryMalta
##          n missing distinct      Info      Sum      Mean
##    18698          0          2    0.027      167  0.008931
##
## -----
## countryPoland
##          n missing distinct      Info      Sum      Mean
##    18698          0          2    0.128      834  0.0446
##
## -----
## countryPortugal
##          n missing distinct      Info      Sum      Mean
##    18698          0          2    0.131      854  0.04567
##
## -----
## countryRomania
##          n missing distinct      Info      Sum      Mean
##    18698          0          2    0.133      866  0.04632
##
## -----
## countrySlovakia
##          n missing distinct      Info      Sum      Mean
##    18698          0          2    0.065      412  0.02203
##
## -----

```

```
## countrySlovenia
##      n missing distinct    Info    Sum    Mean
##  18698      0         2    0.066    420 0.02246
##
## -----
## countrySpain
##      n missing distinct    Info    Sum    Mean
##  18698      0         2    0.134    876 0.04685
##
## -----
## countrySweden
##      n missing distinct    Info    Sum    Mean
##  18698      0         2    0.129    844 0.04514
##
## -----
## countryThe.Netherlands
##      n missing distinct    Info    Sum    Mean
##  18698      0         2    0.12    779 0.04166
##
## -----
```

EDA Part:

Plotting the Dataset:

What are the distributions of the variables in the dataset ?

```
library(ggplot2)

distribution_plot <- function(data) {
  for (v in names(data)) {

    # Skip constant or all-NA variables
    if (all(is.na(data[[v]])) || length(unique(data[[v]])) <= 1) next

    # Histogram for continuous numeric variables
    if (is.numeric(data[[v]]) && length(unique(data[[v]])) > 2) {
      p <- ggplot(data, aes_string(x = v)) +
        geom_histogram(bins = 30, fill = "purple", color = "white") +
        labs(title = paste("Histogram of", v), x = v, y = "Count") +
        theme_minimal()

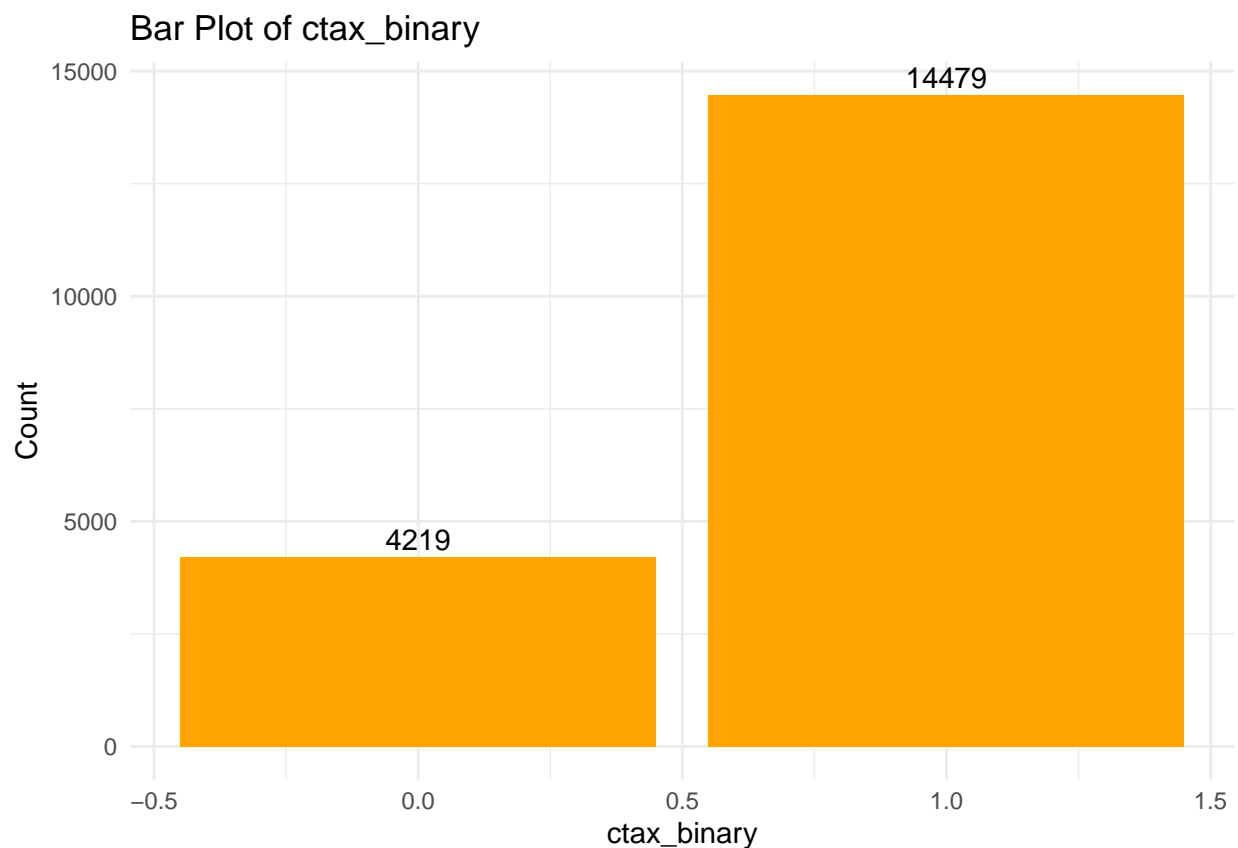
      # Bar plot for categorical or binary variables
    } else {
      p <- ggplot(data, aes_string(x = v)) +
        geom_bar(fill = "orange") +
        geom_text(stat = "count", aes(label = ..count..), vjust = -0.3) +
        labs(title = paste("Bar Plot of", v), x = v, y = "Count") +
        theme_minimal()
    }
  }

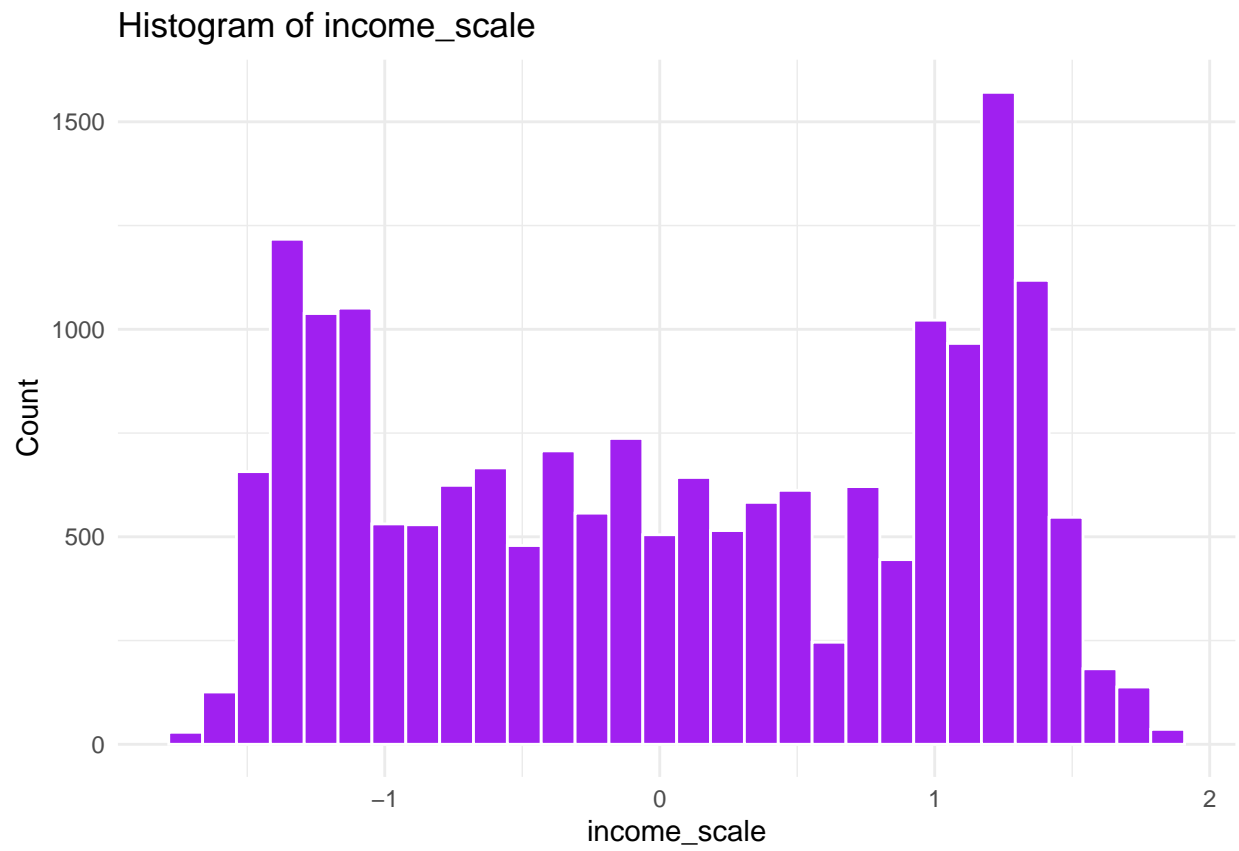
  print(p)
}
```

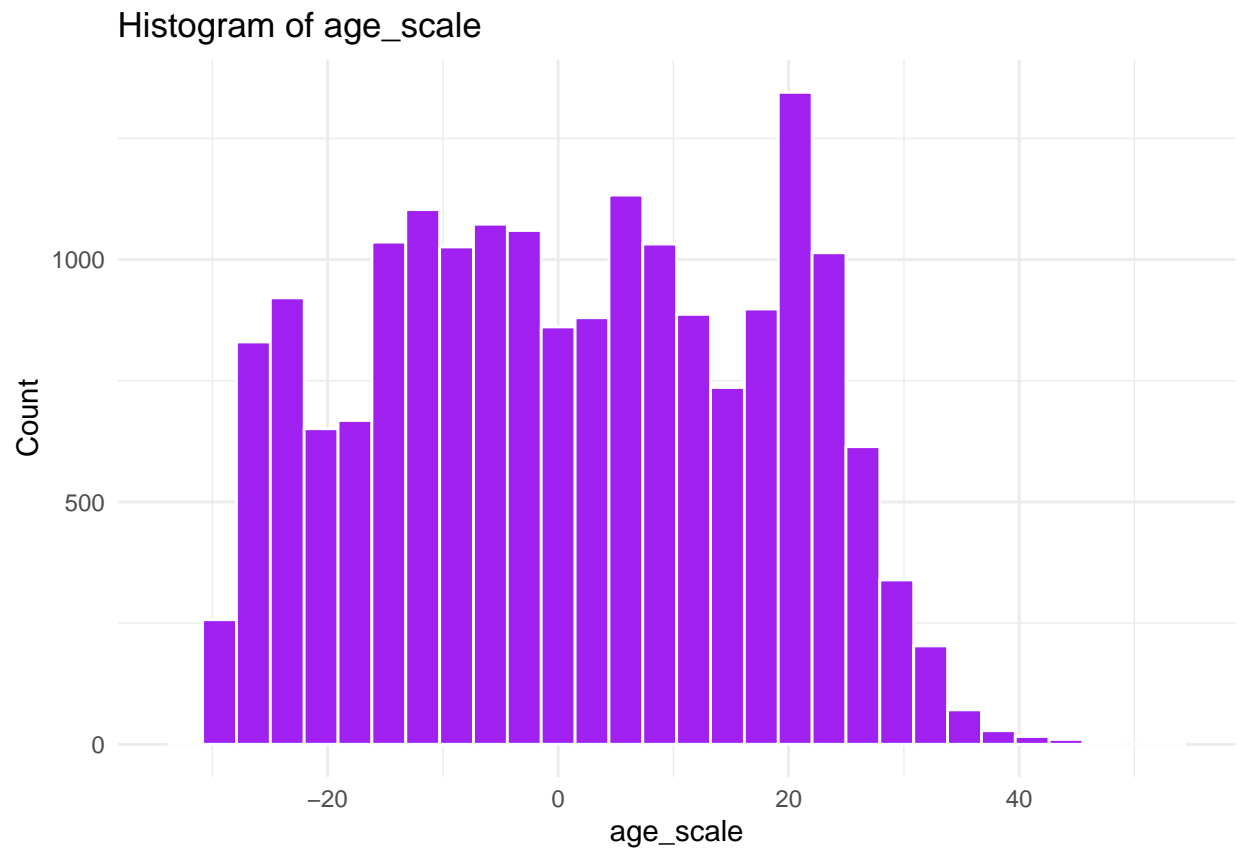
```
}  
}
```

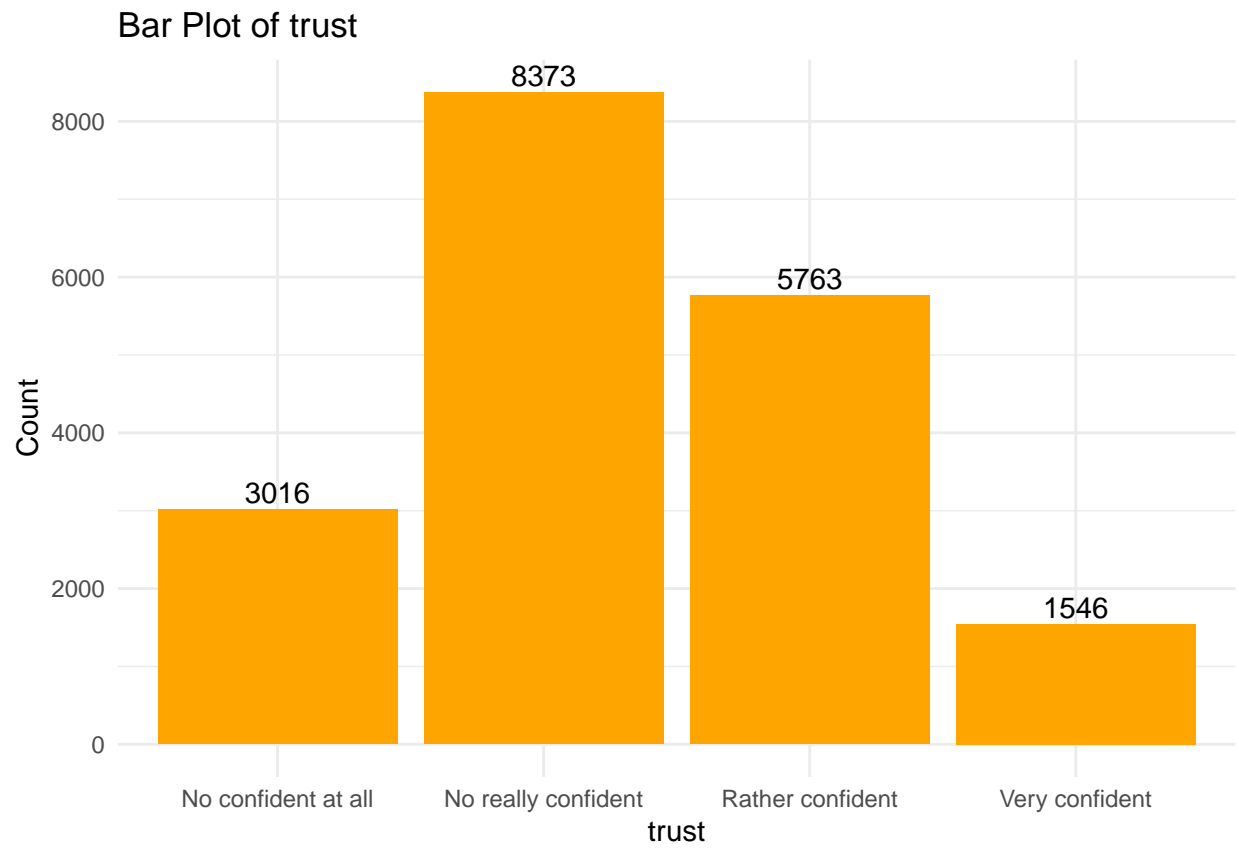
```
distribution_plot(df_clean)
```

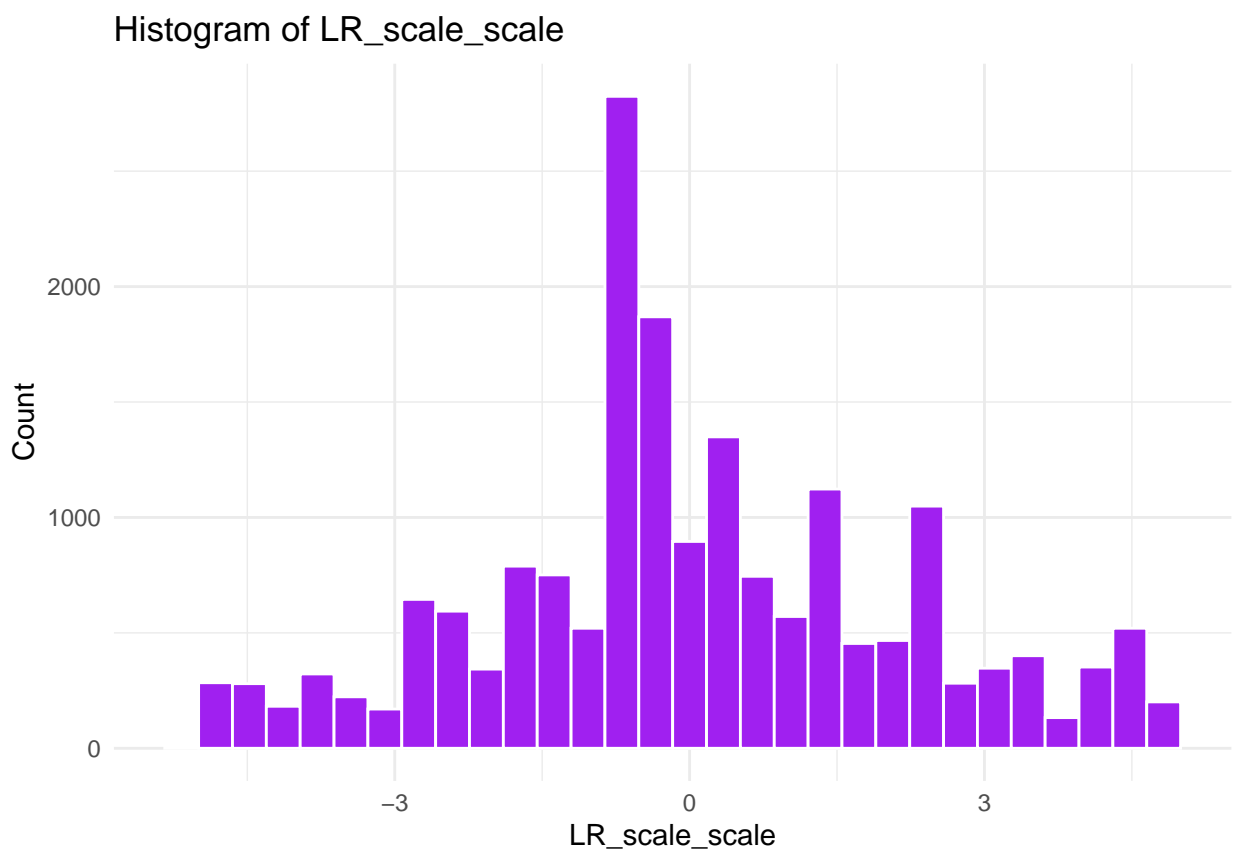
```
## Warning: 'aes_string()' was deprecated in ggplot2 3.0.0.  
## i Please use tidy evaluation idioms with 'aes()'.  
## i See also 'vignette("ggplot2-in-packages")' for more information.  
## This warning is displayed once every 8 hours.  
## Call 'lifecycle::last_lifecycle_warnings()' to see where this warning was  
## generated.  
  
## Warning: The dot-dot notation ('..count..') was deprecated in ggplot2 3.4.0.  
## i Please use 'after_stat(count)' instead.  
## This warning is displayed once every 8 hours.  
## Call 'lifecycle::last_lifecycle_warnings()' to see where this warning was  
## generated.
```



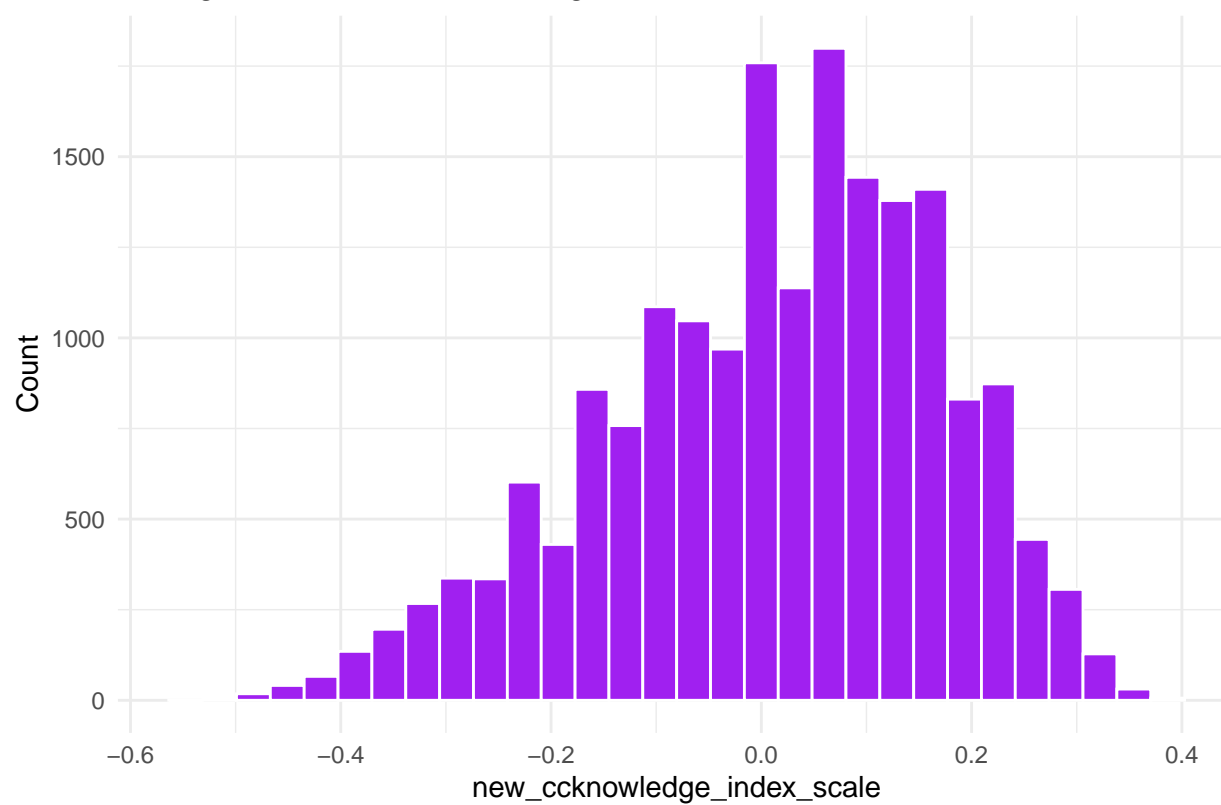


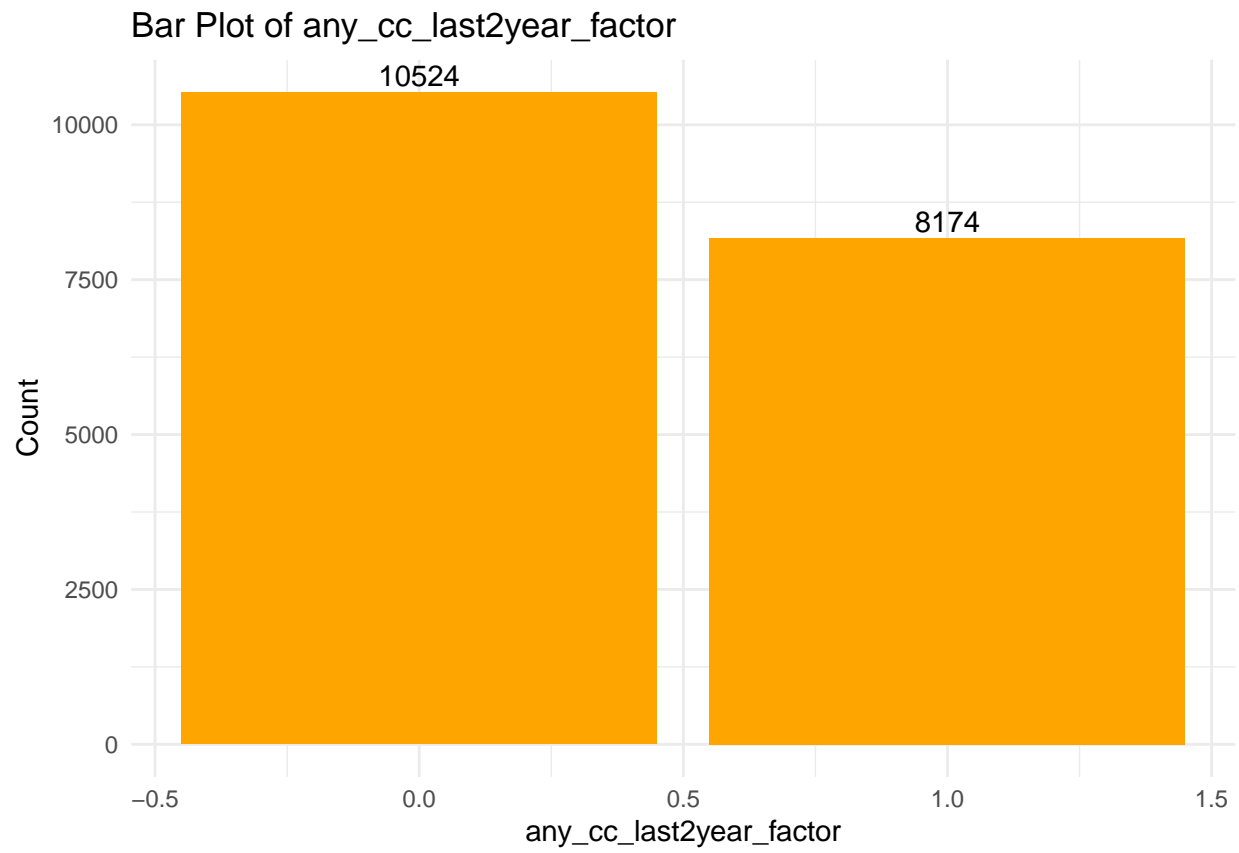


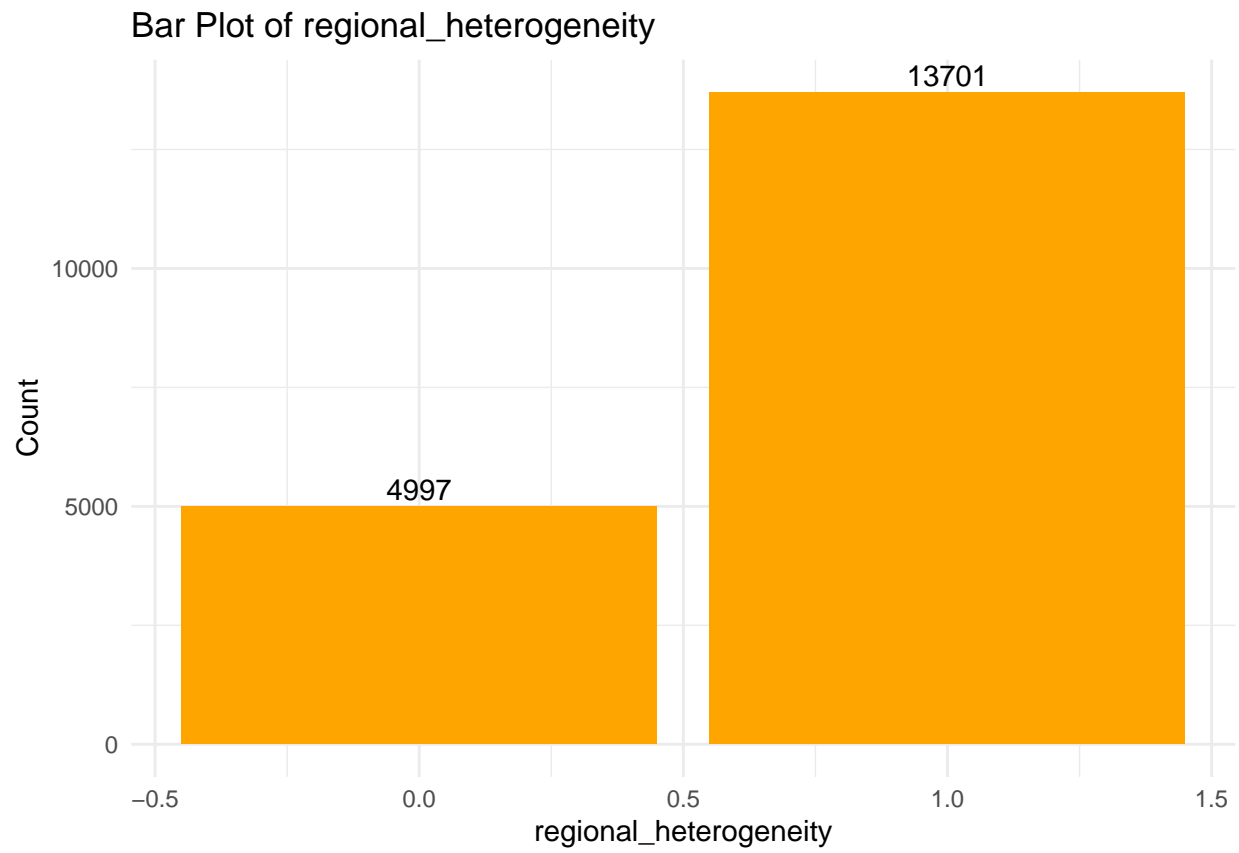


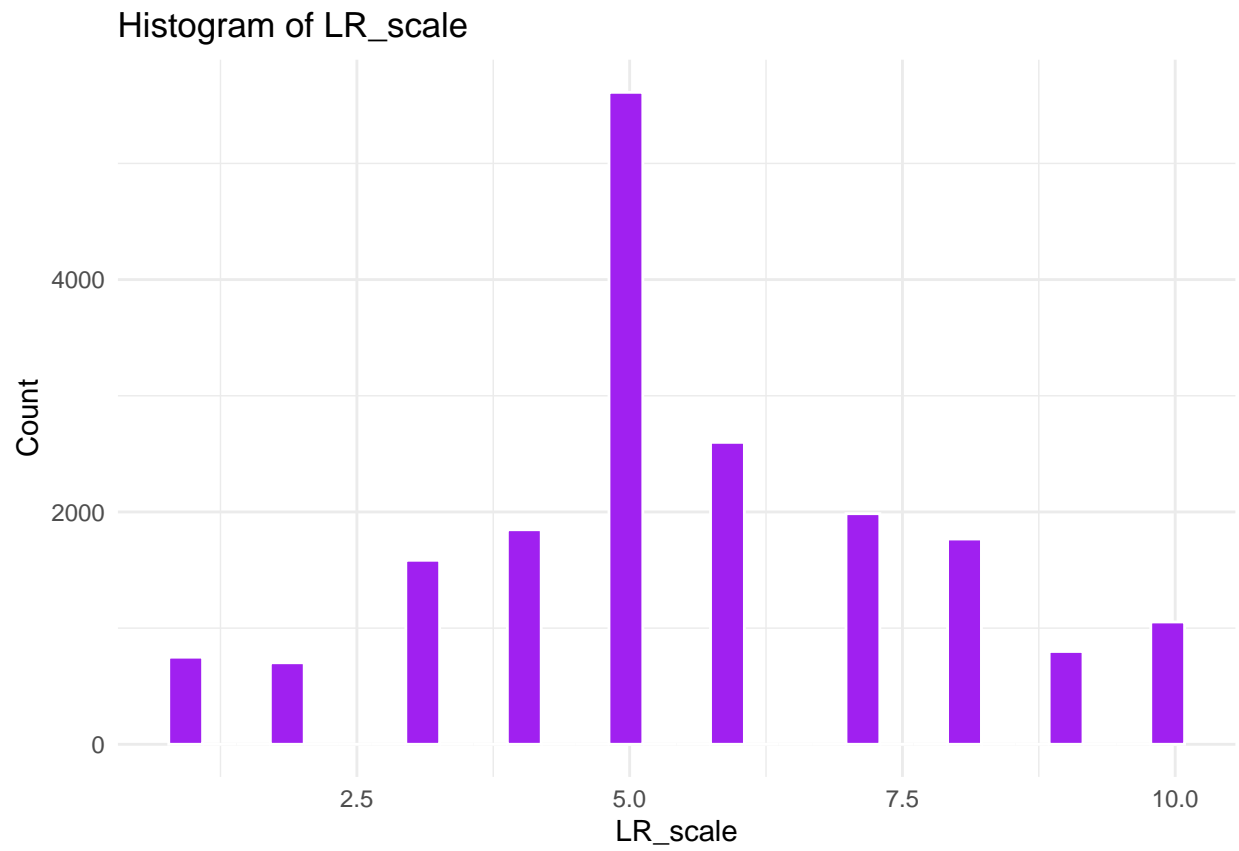


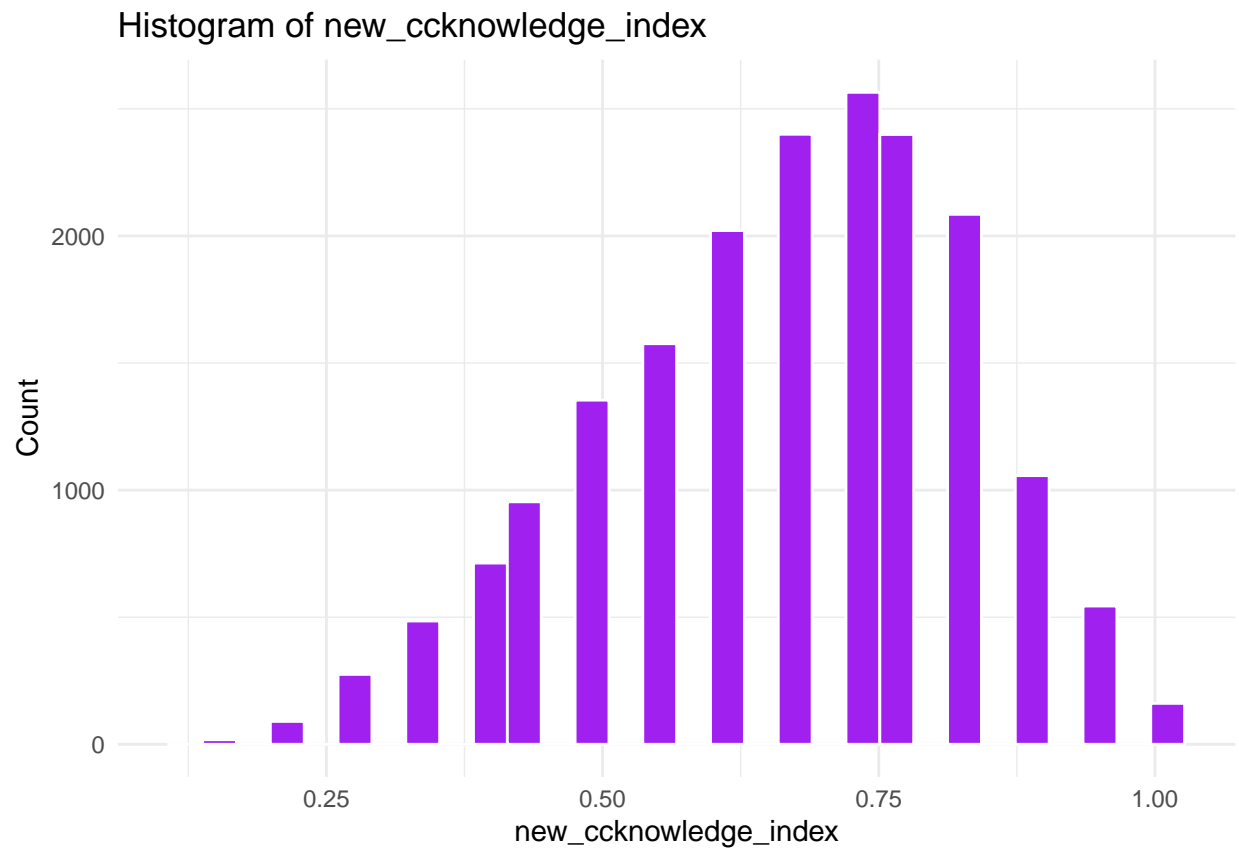
Histogram of new_ccknowledge_index_scale

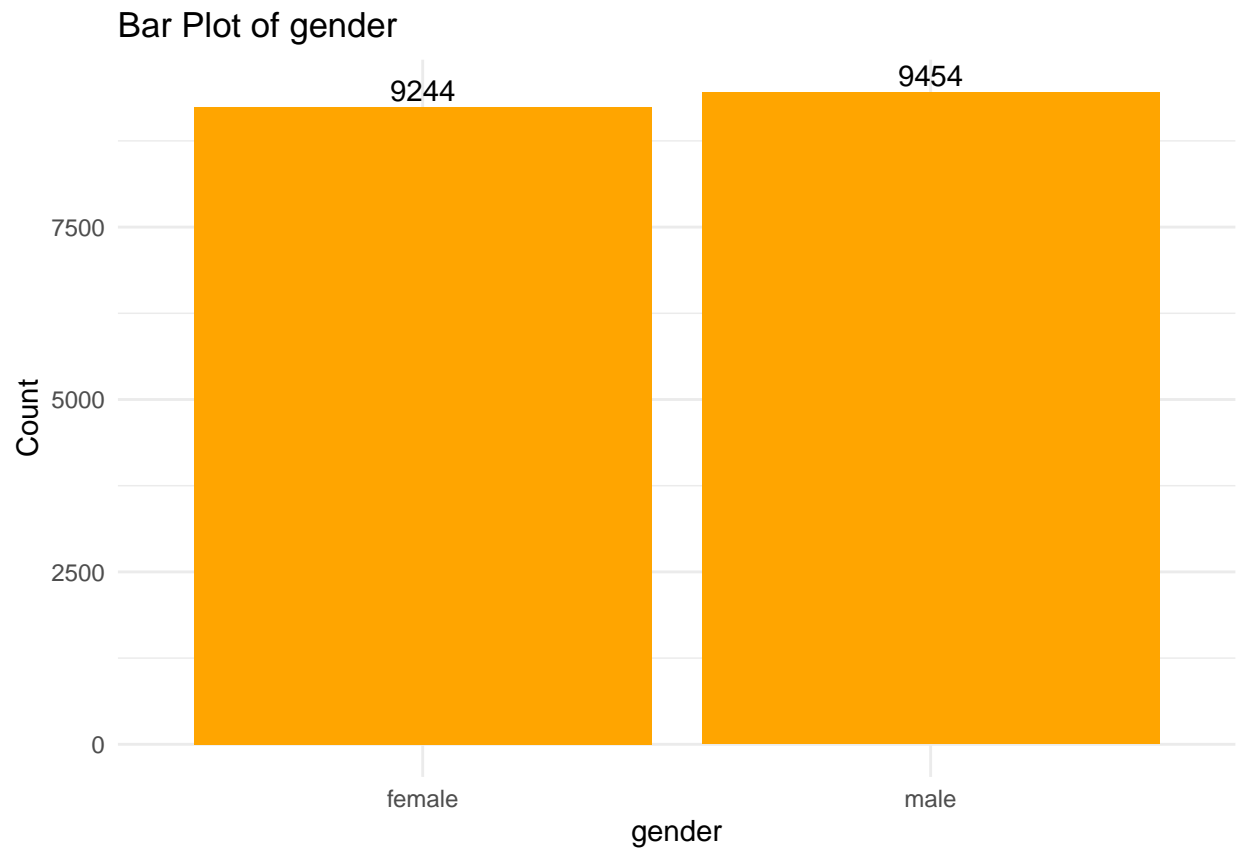


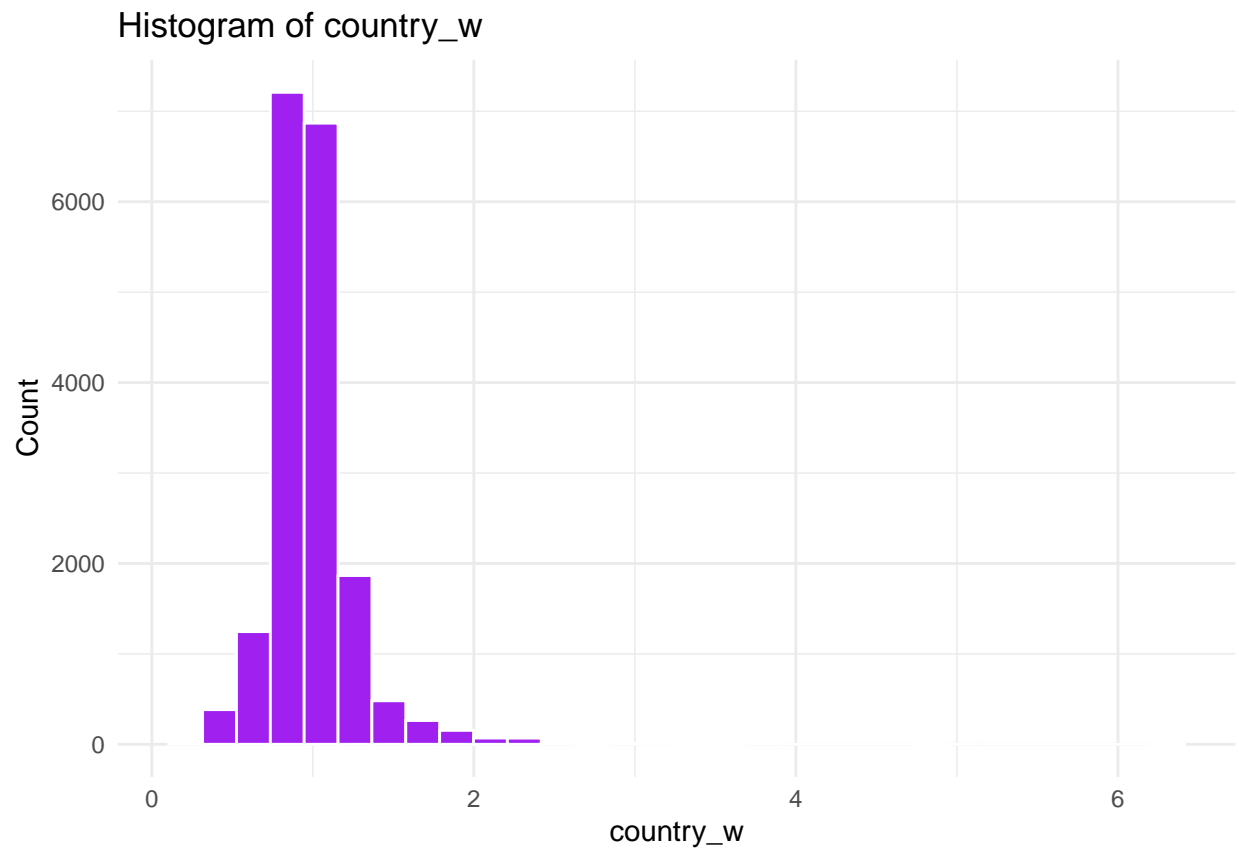


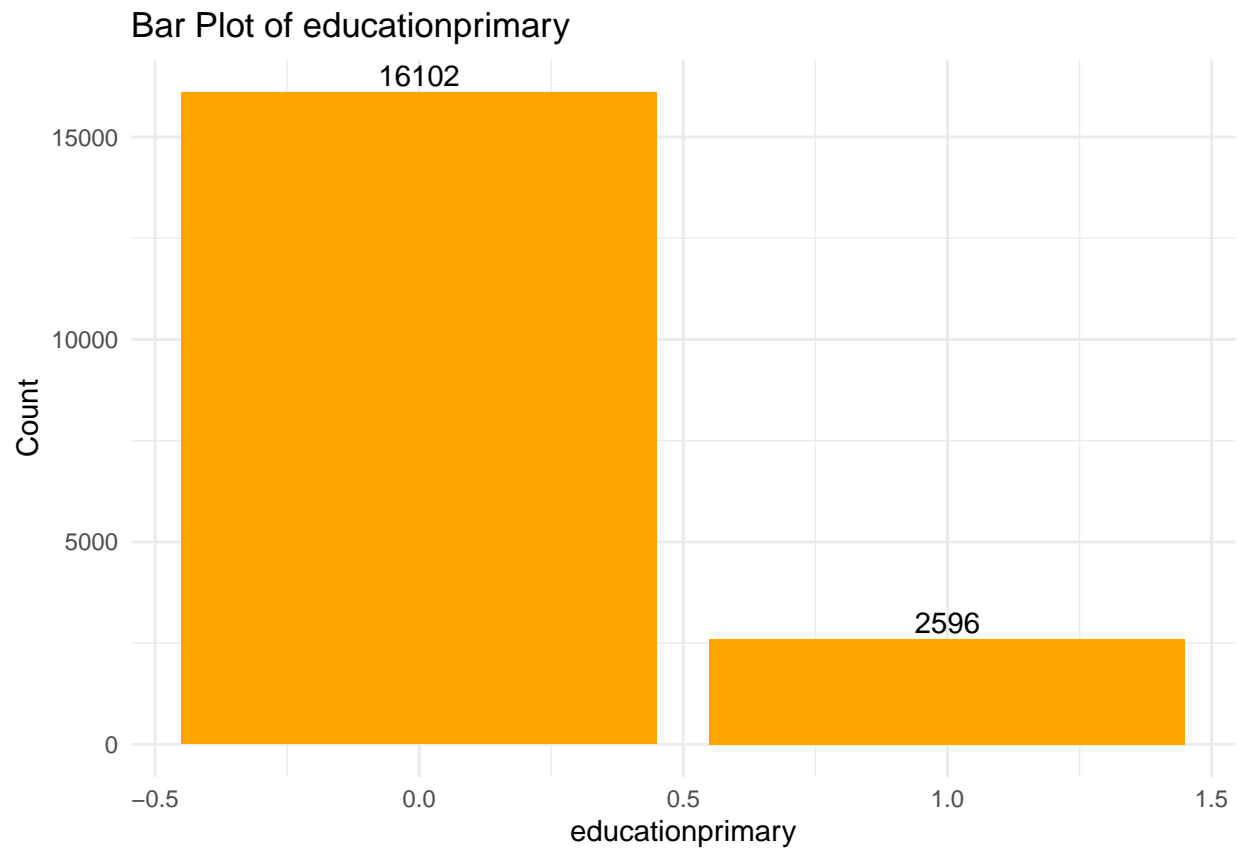


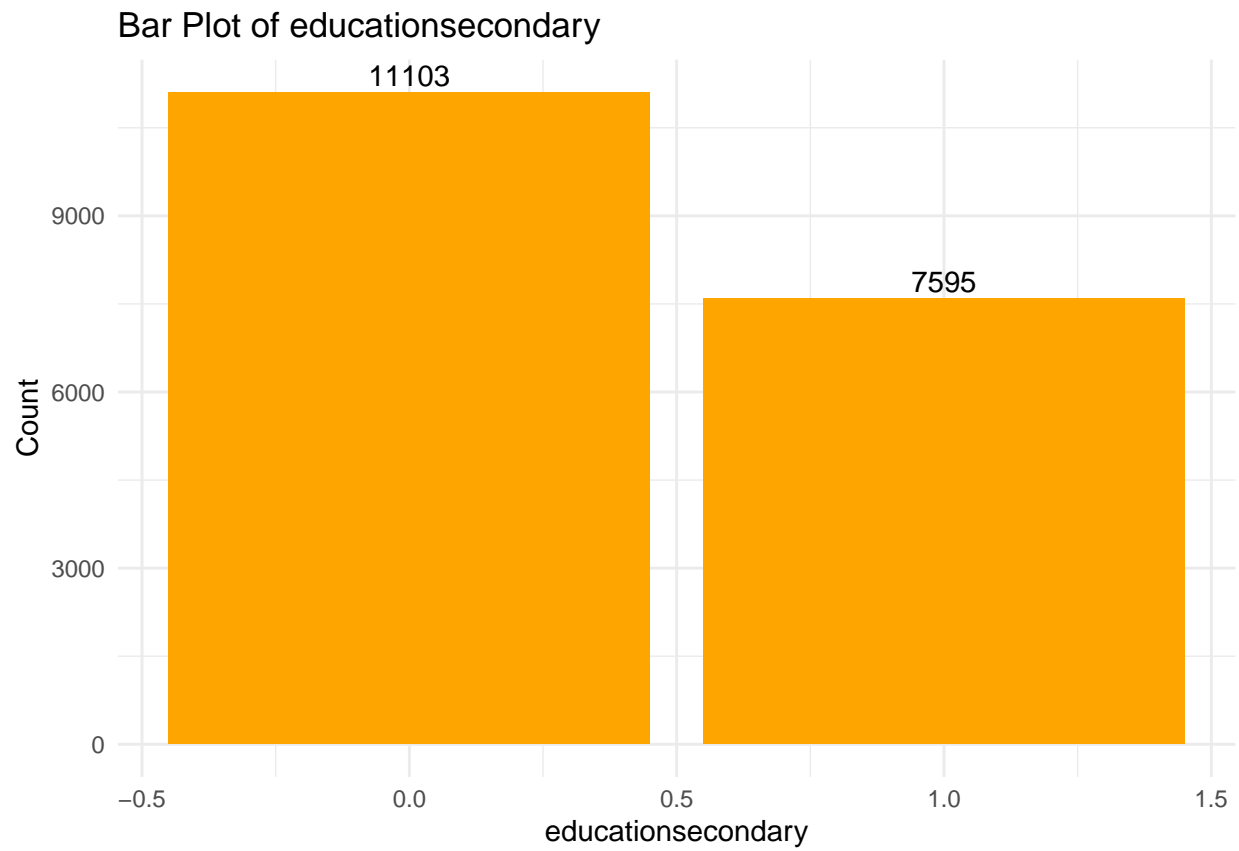


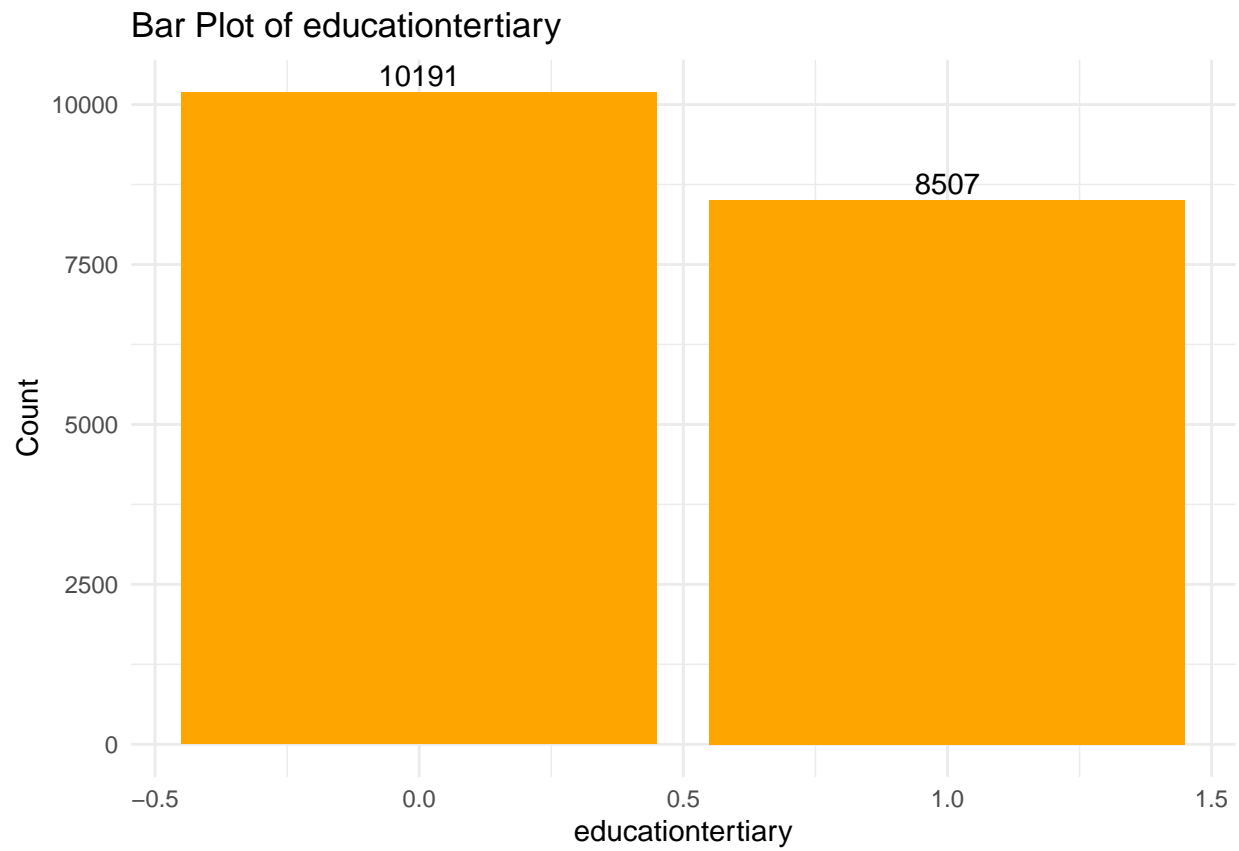


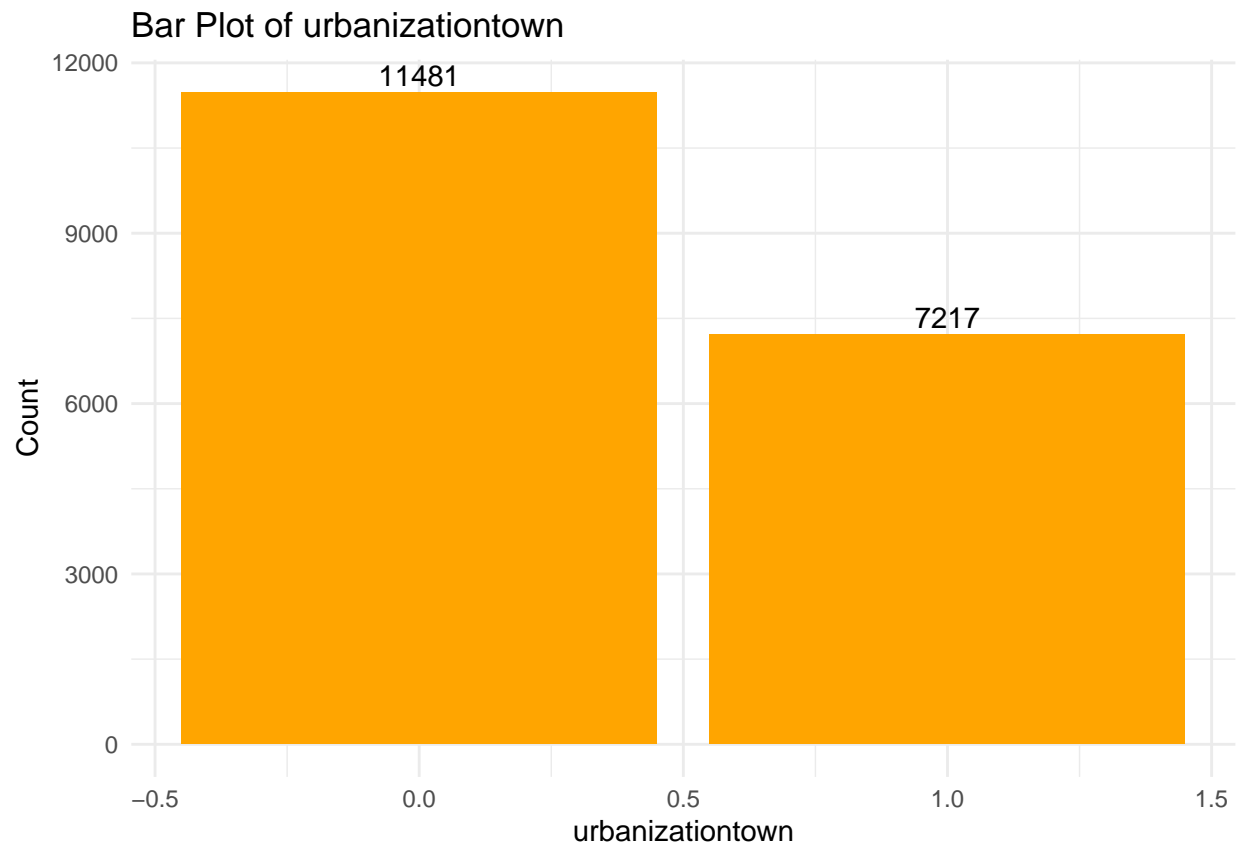


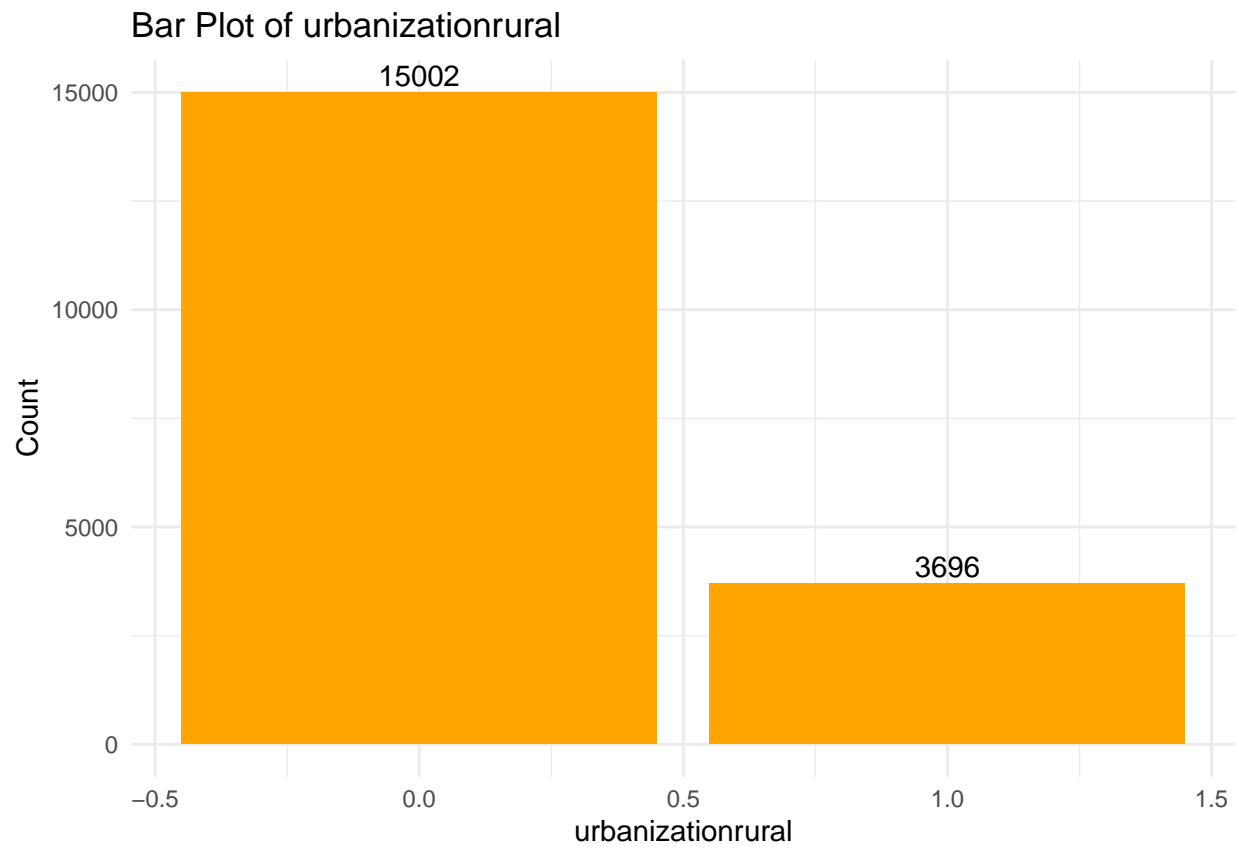


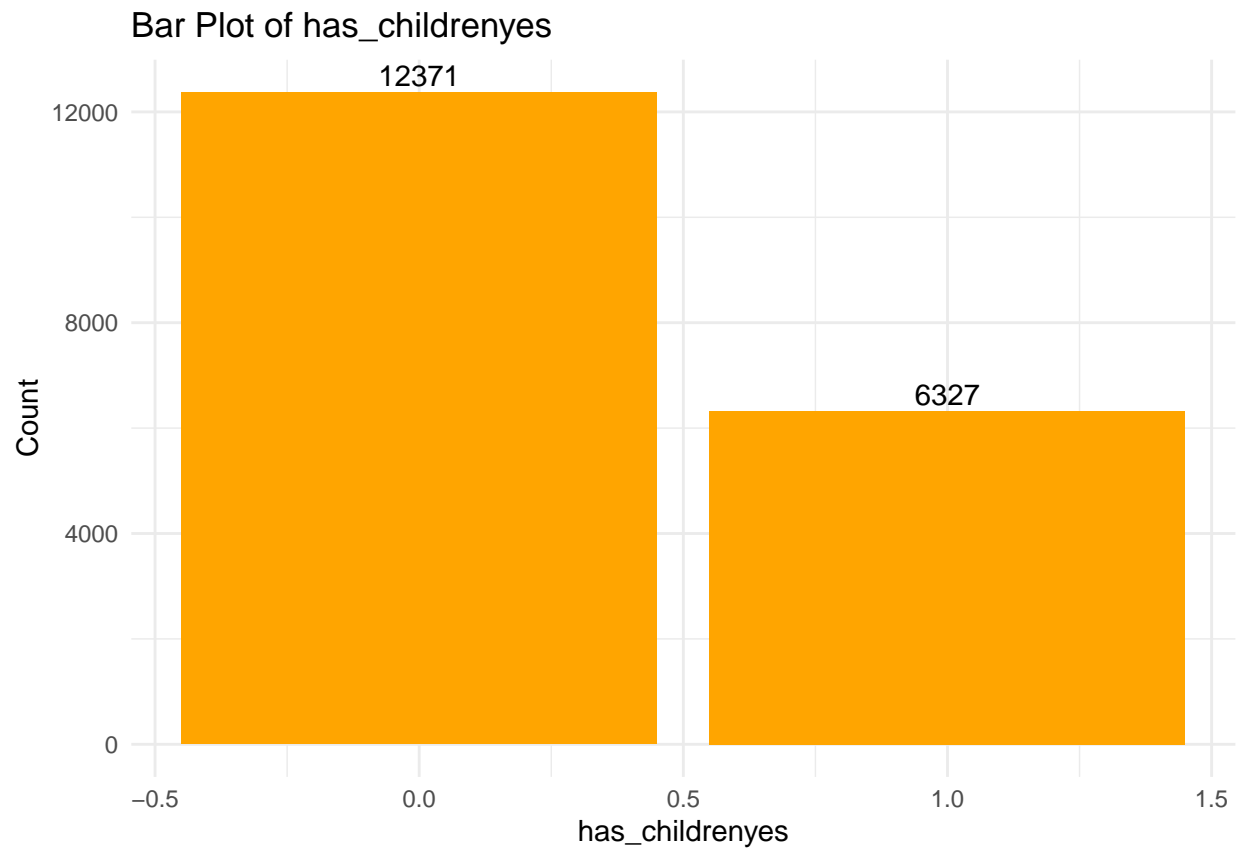


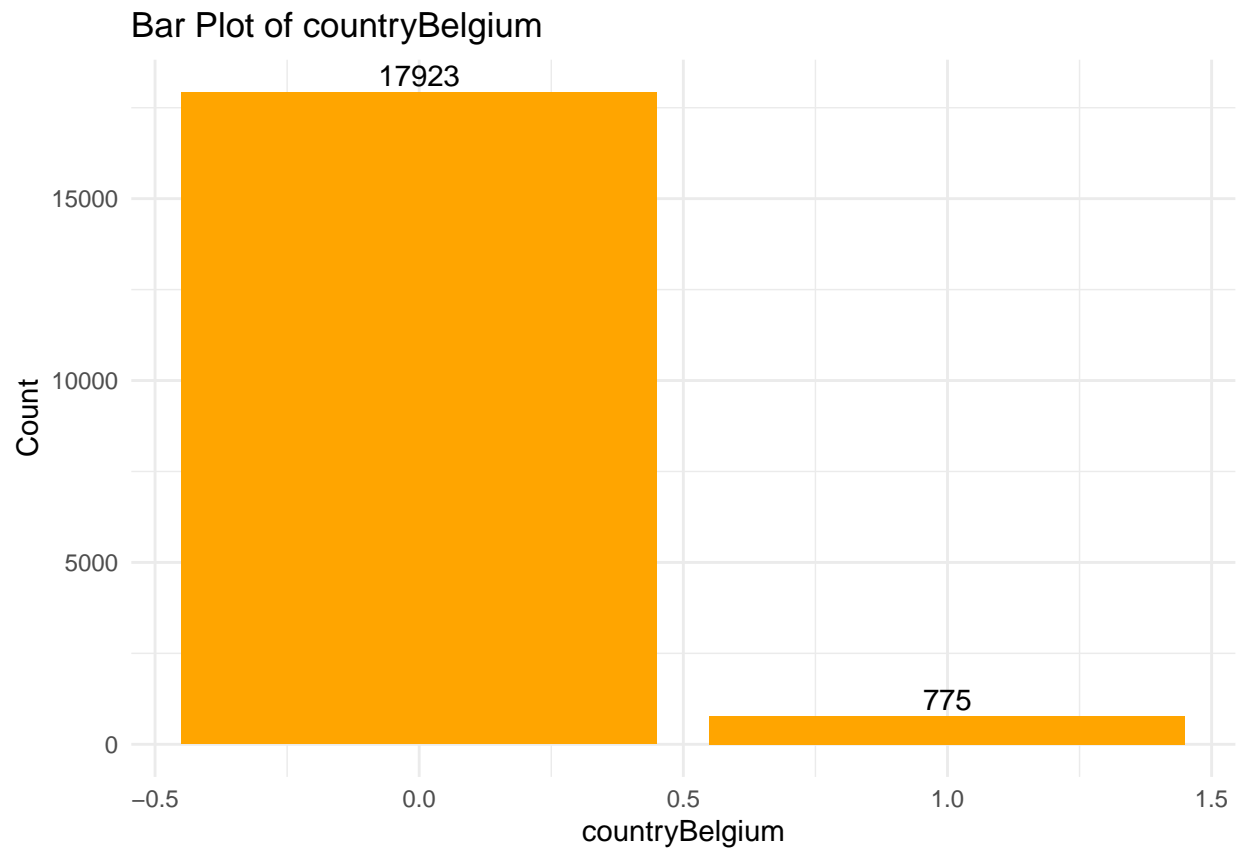


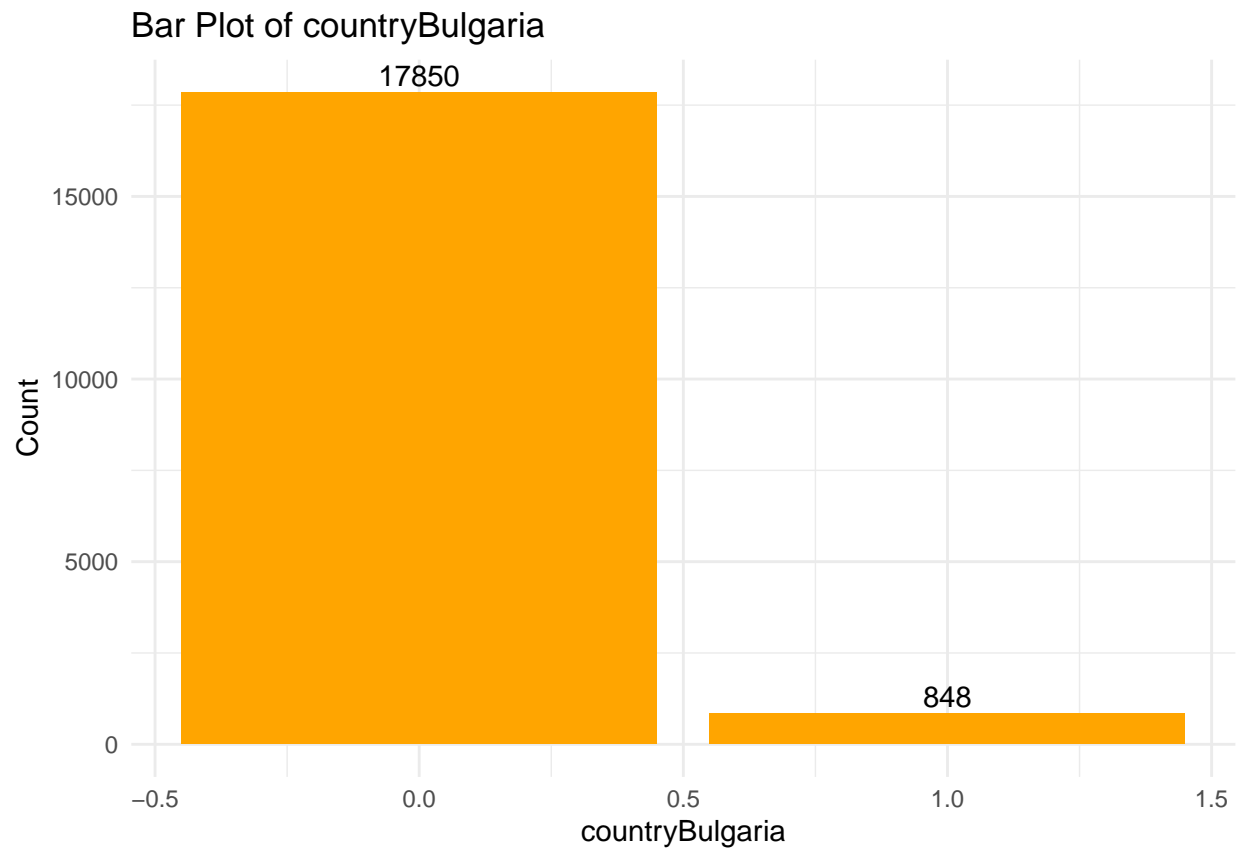


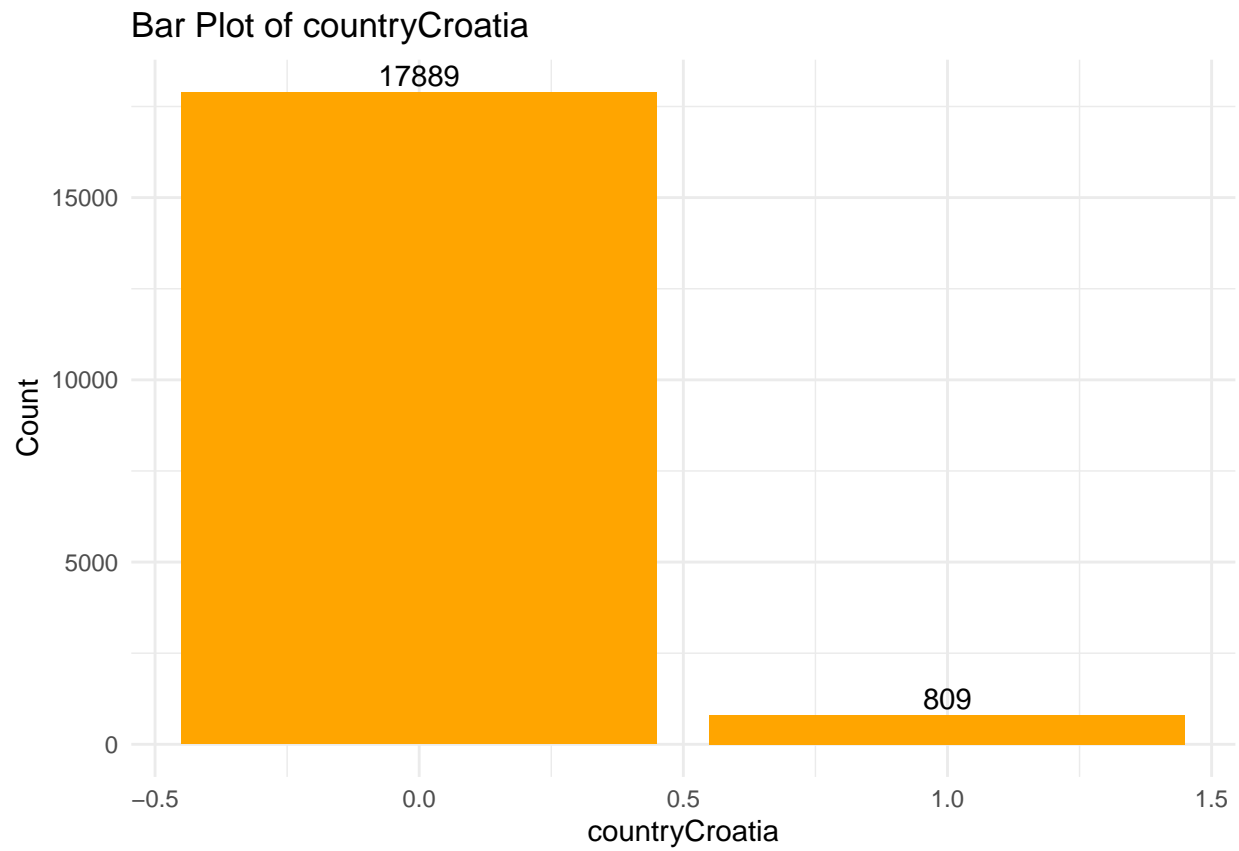


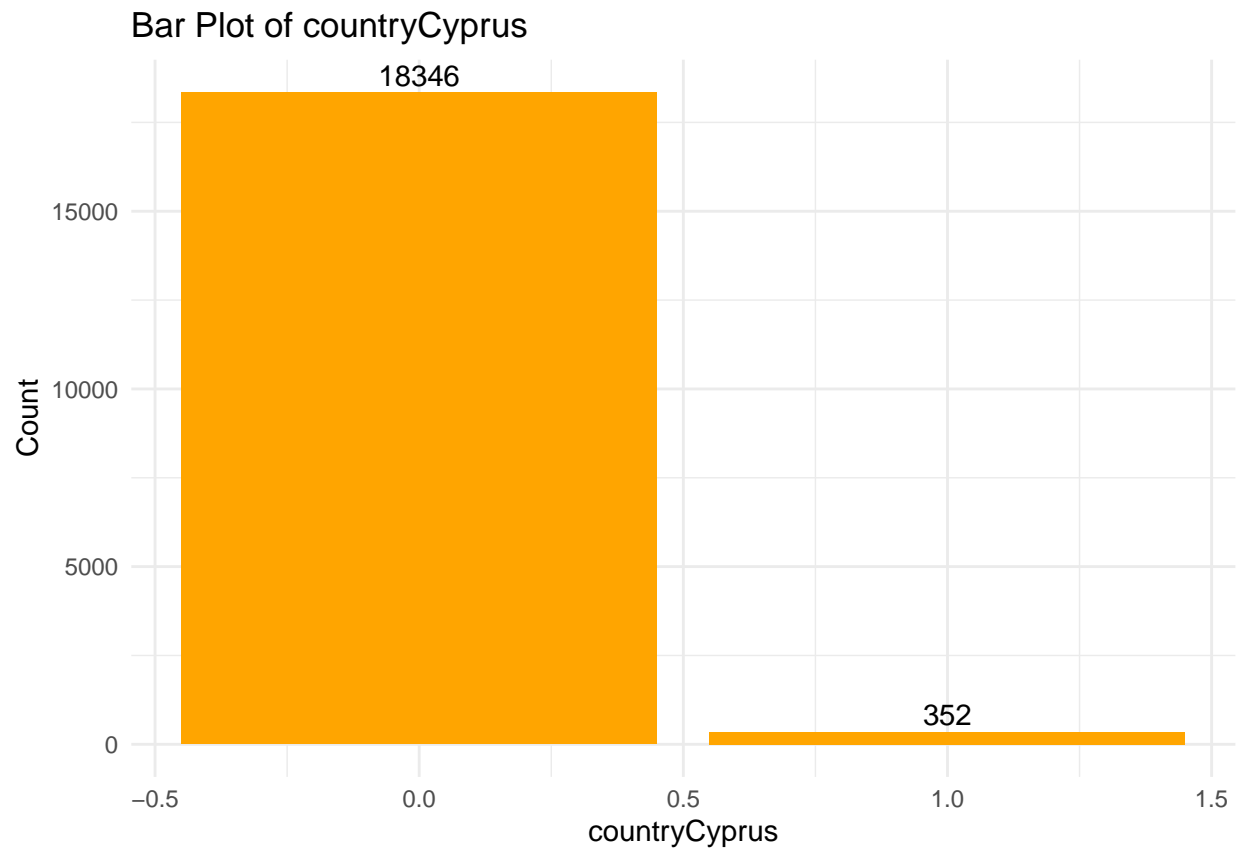


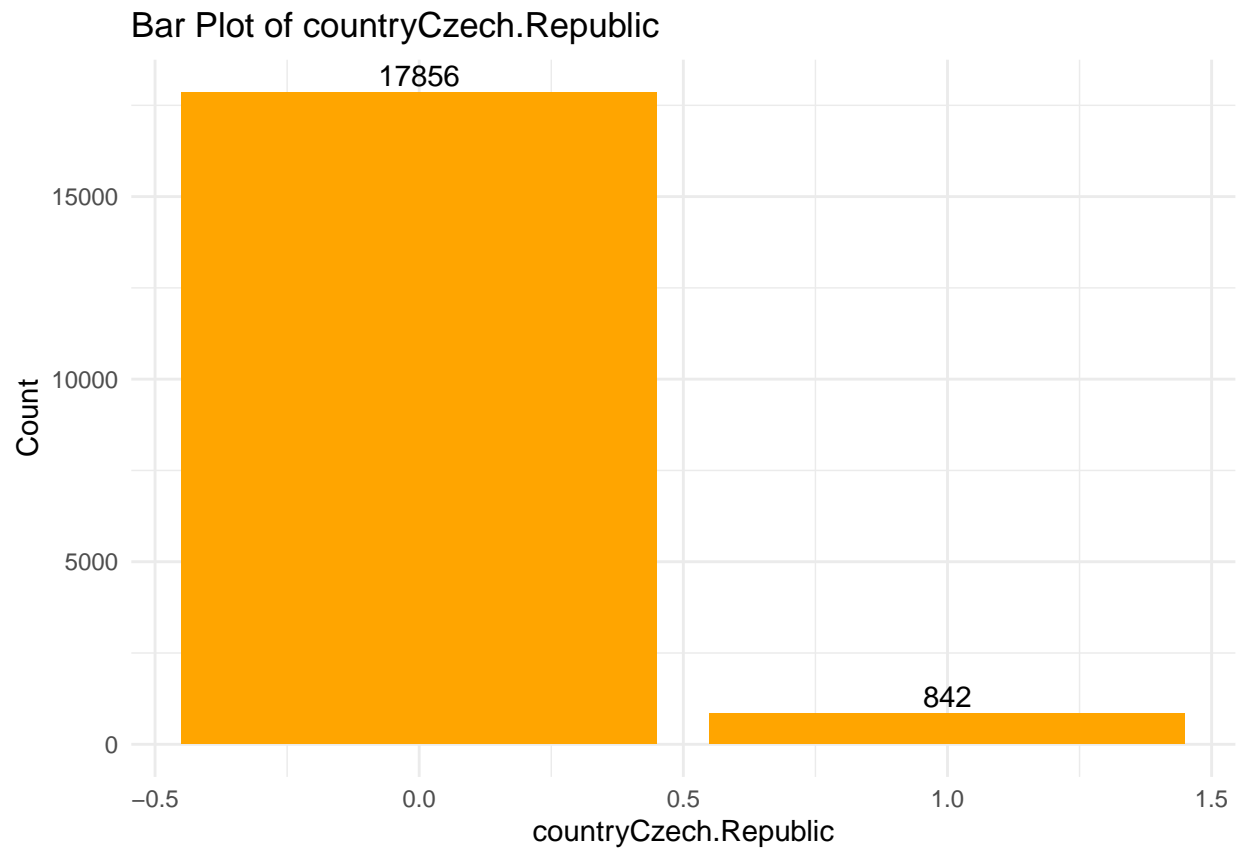


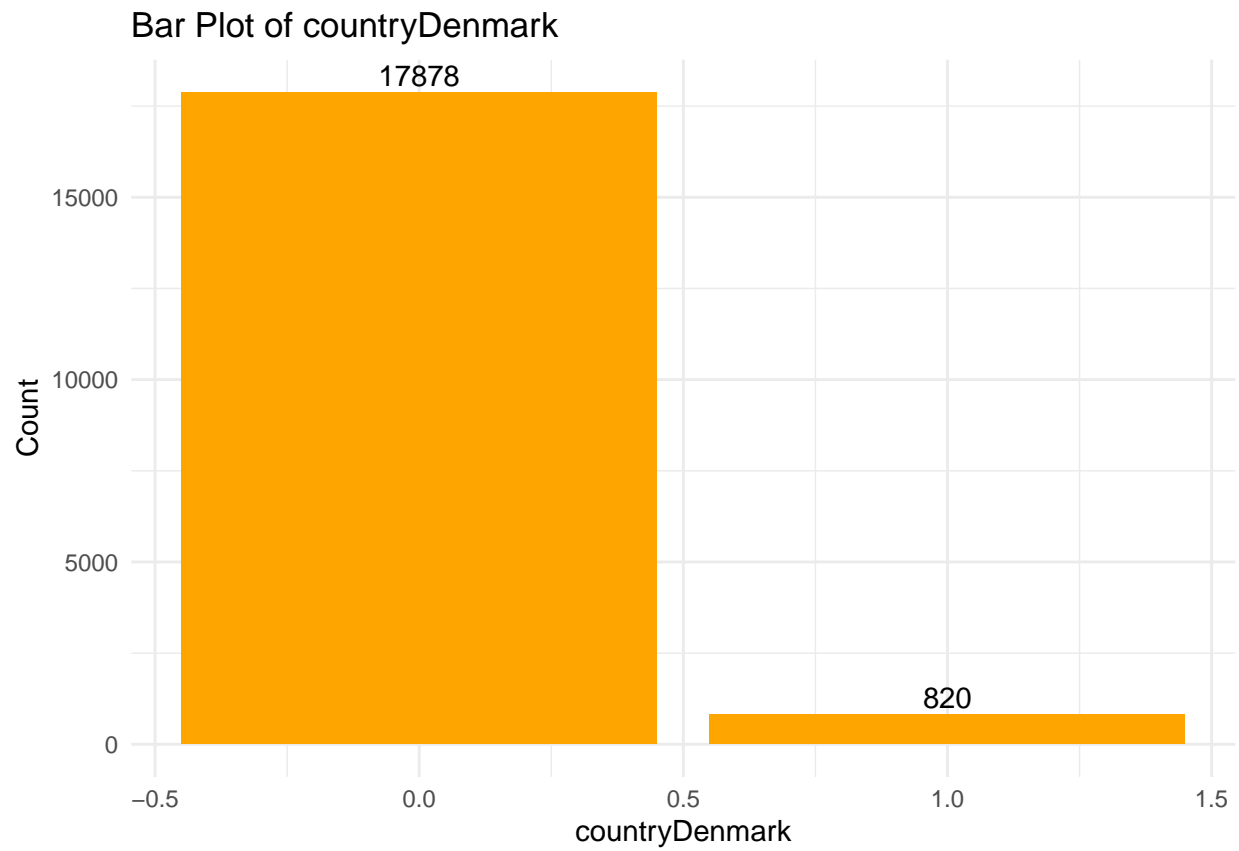


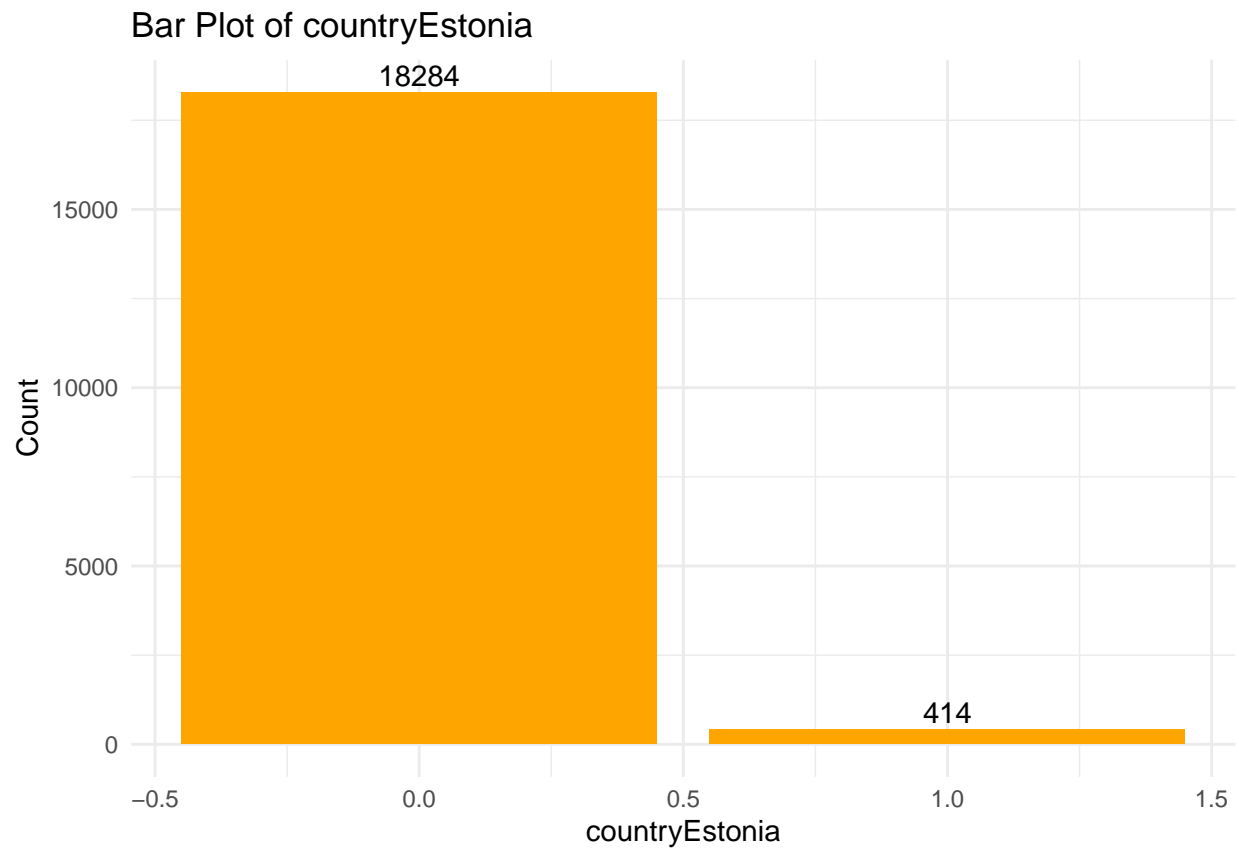


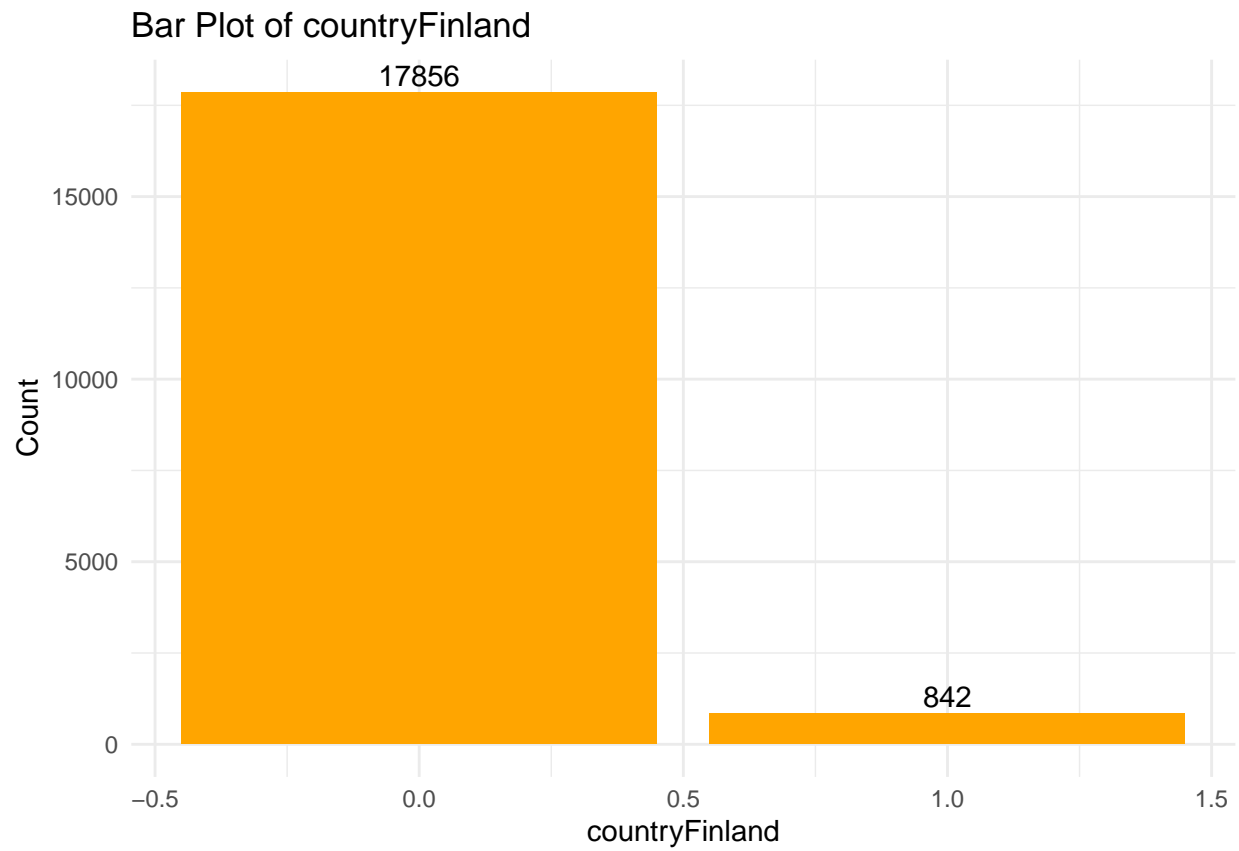


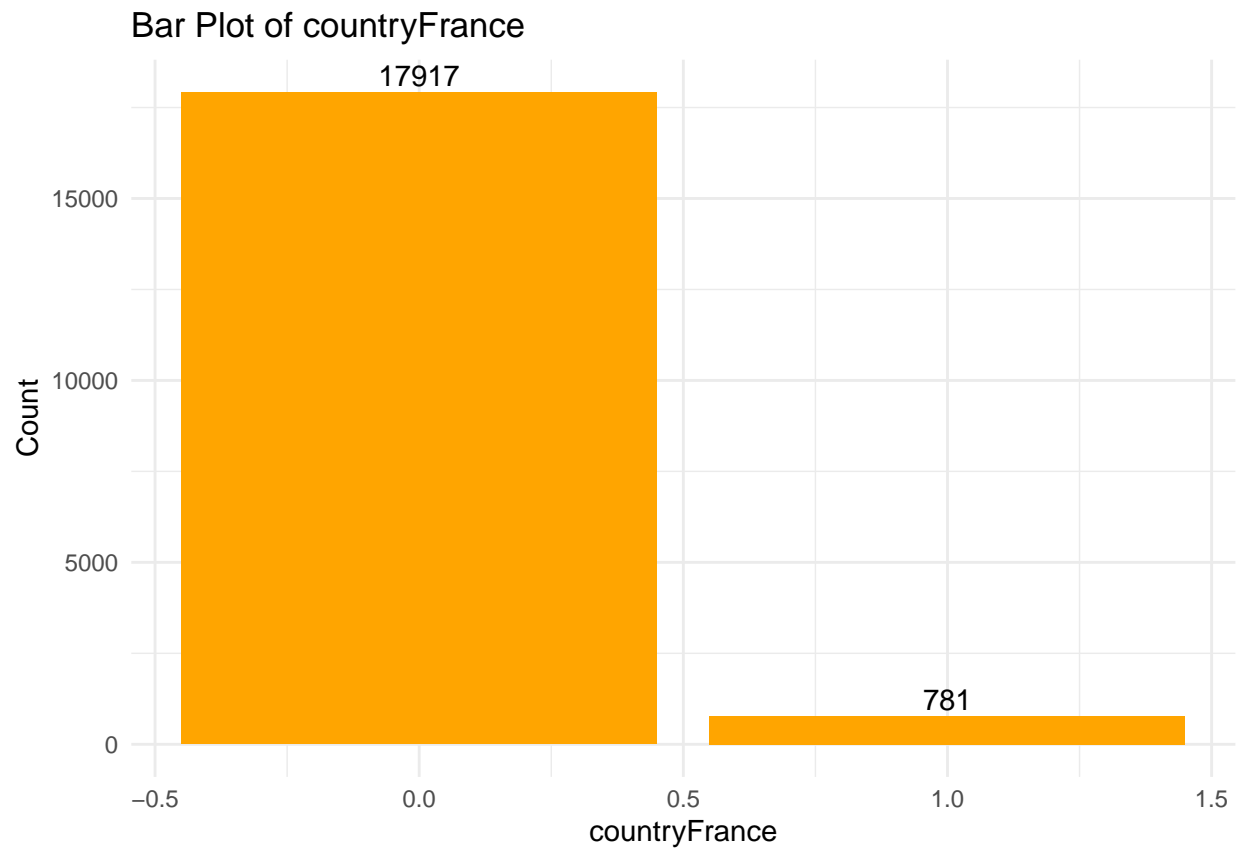


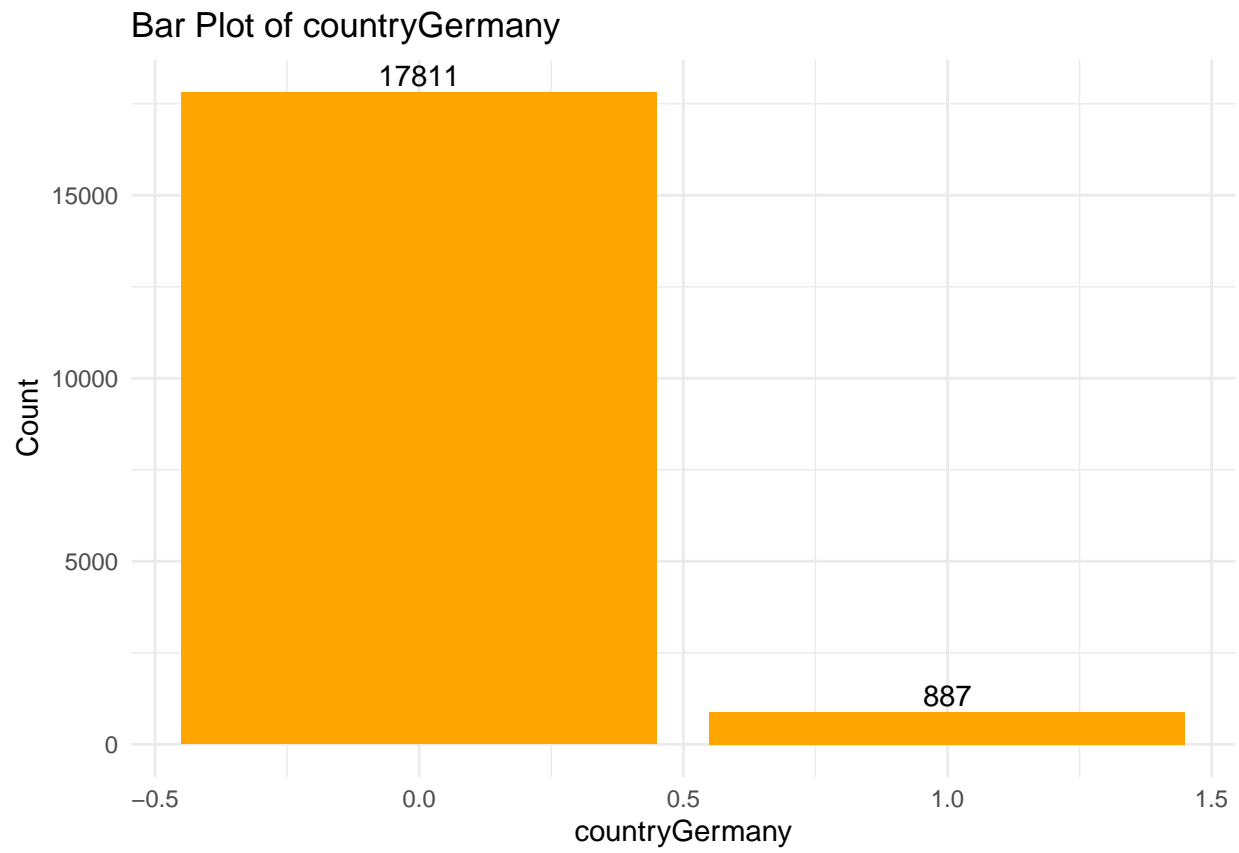


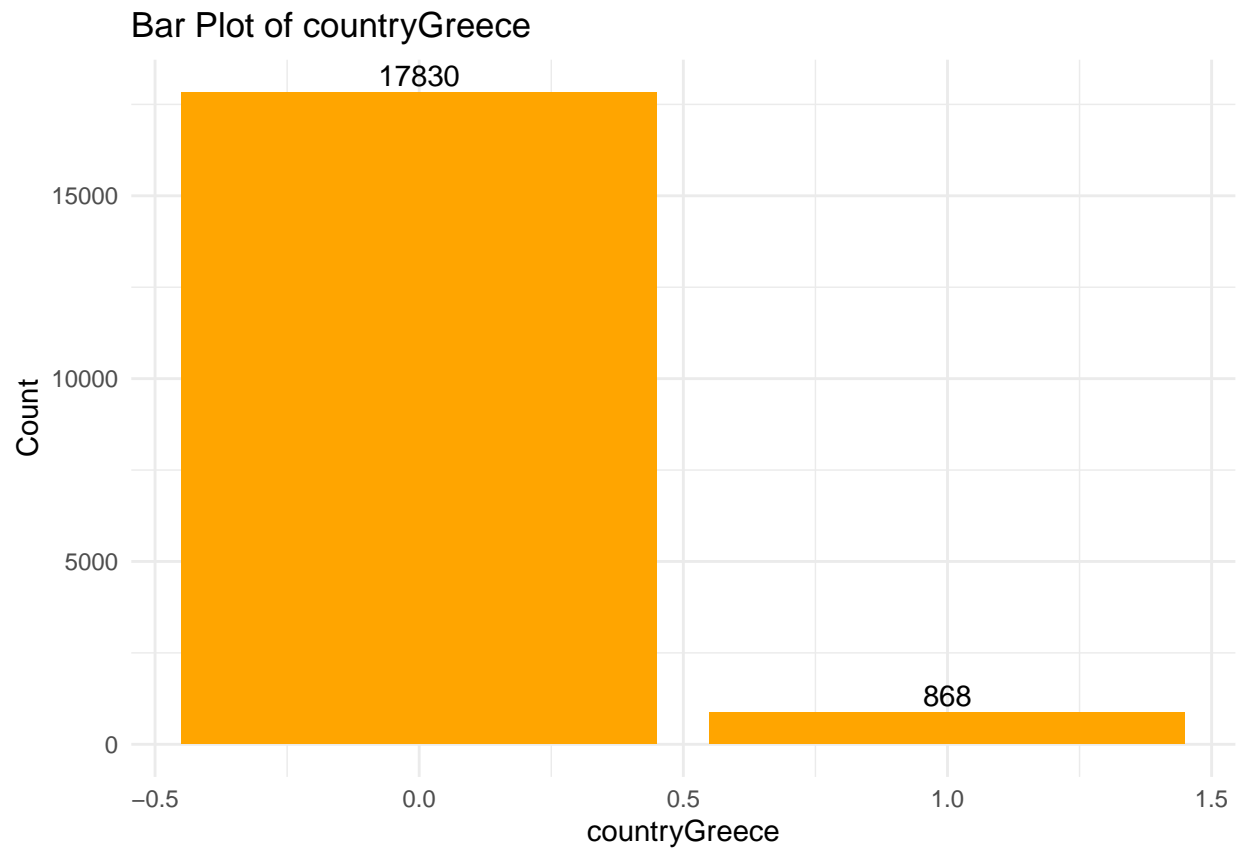


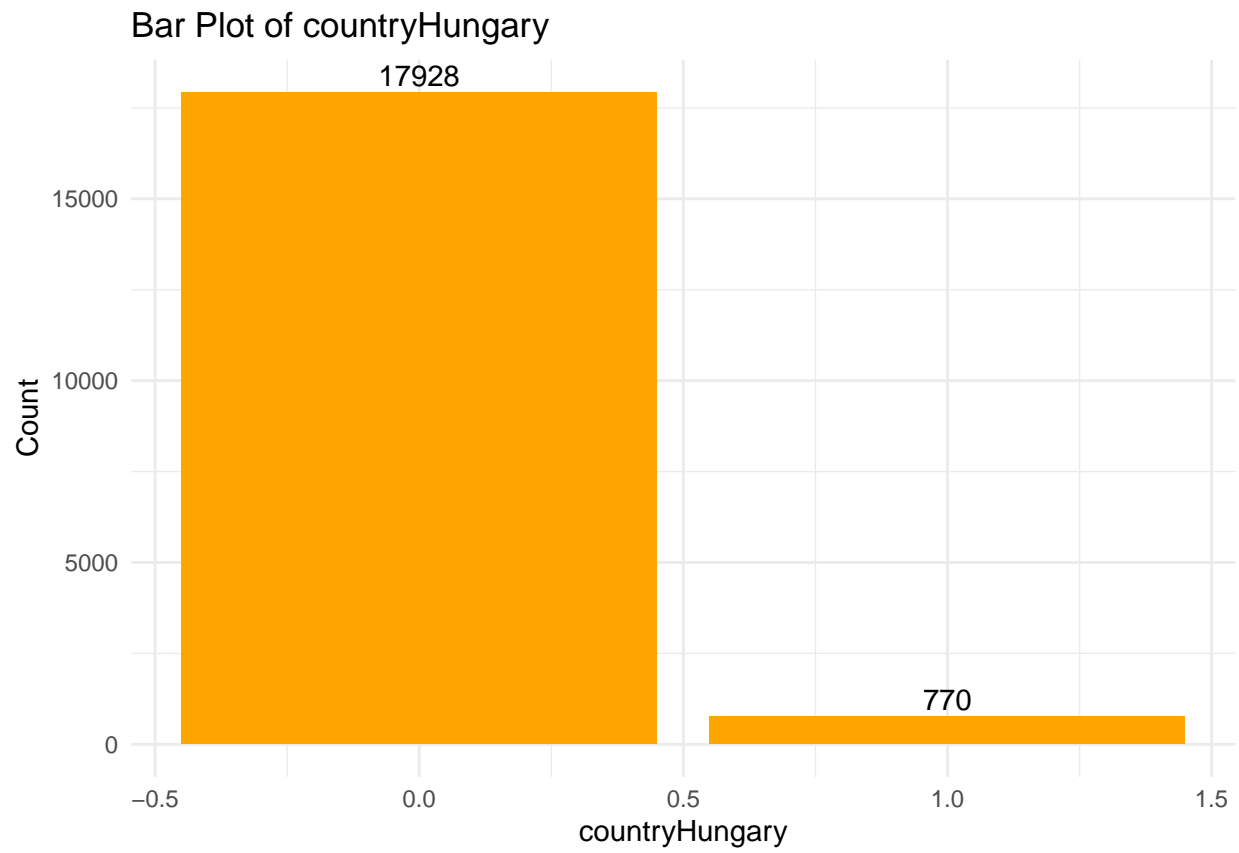


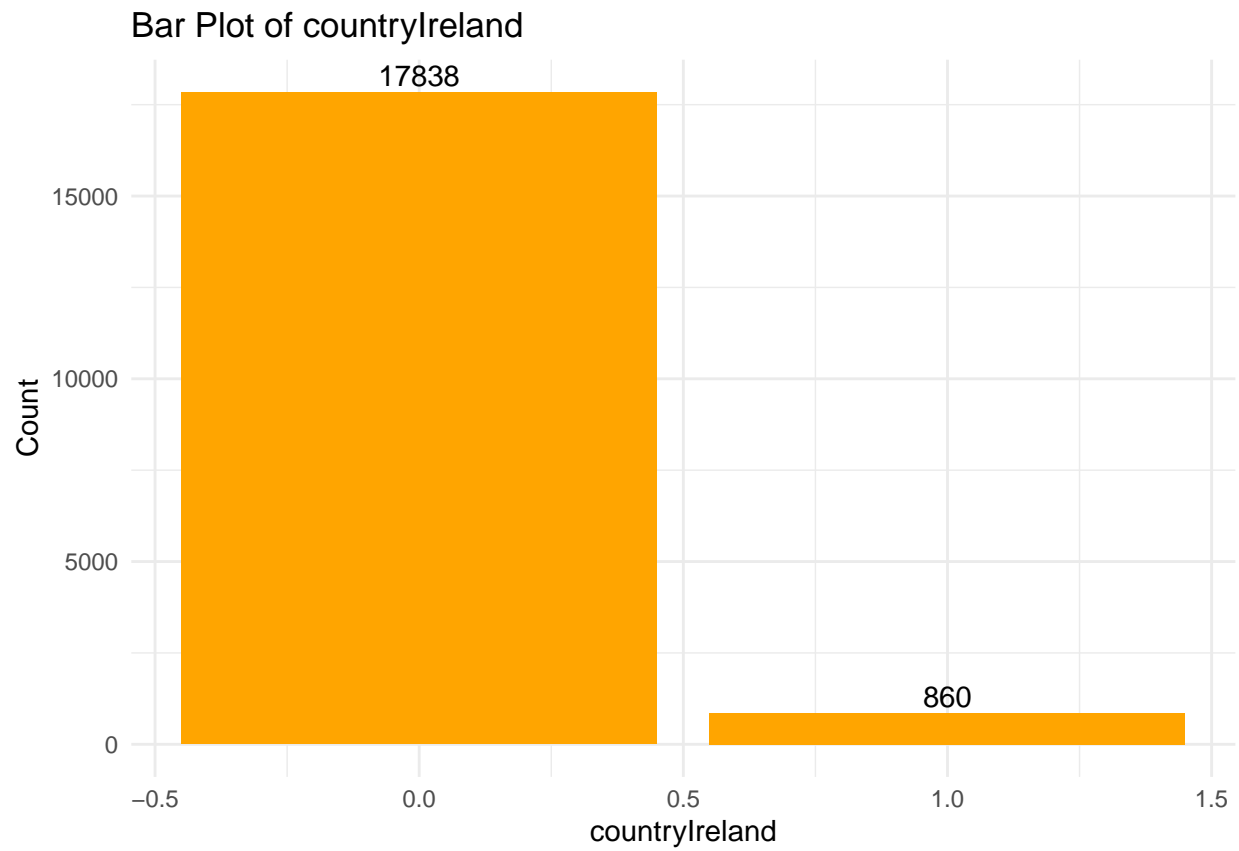


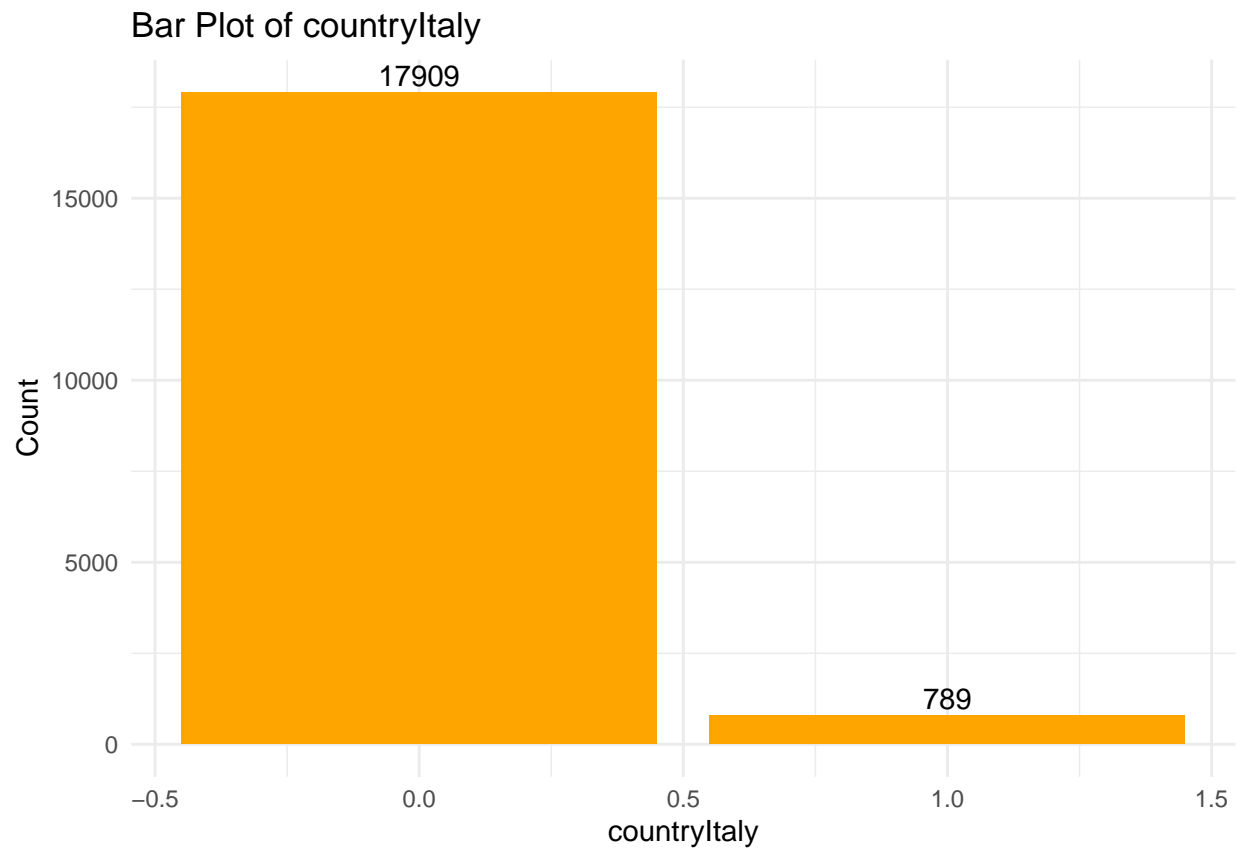


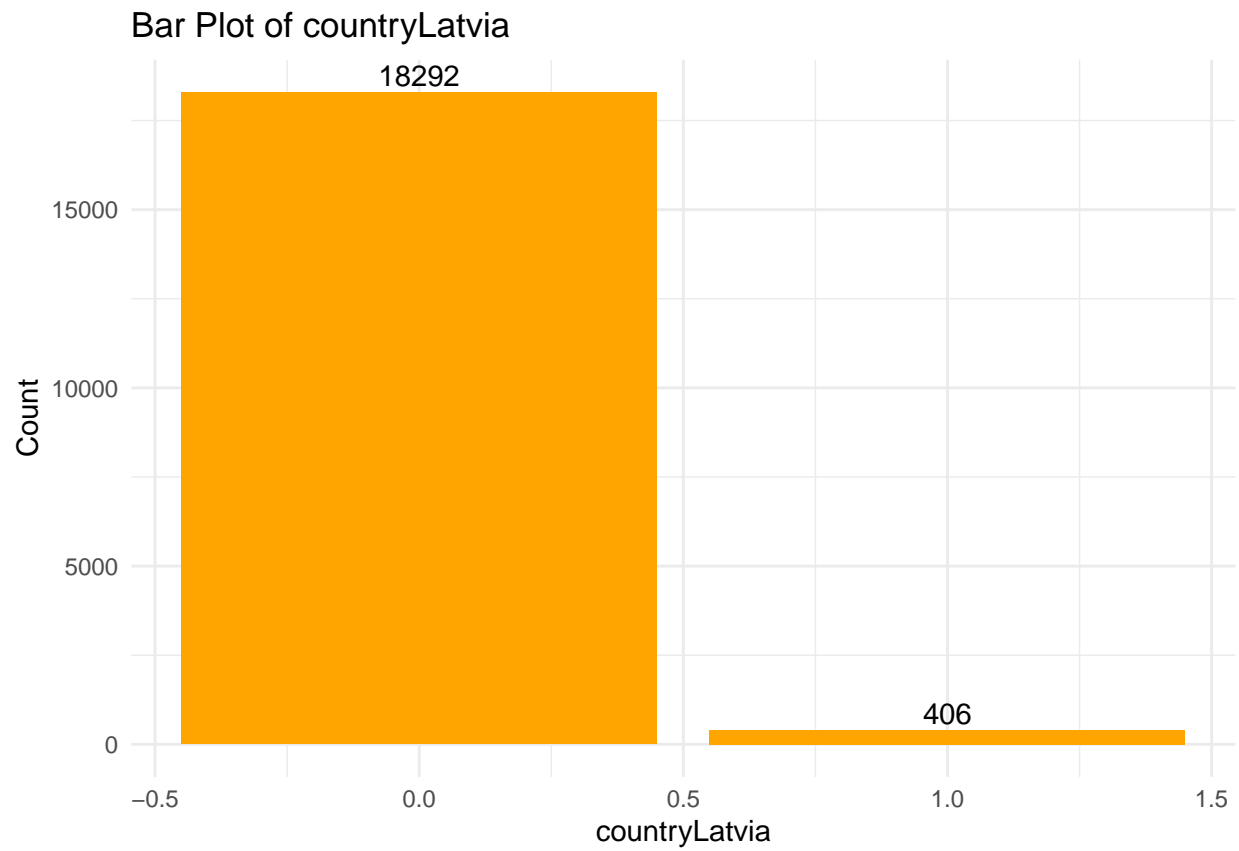


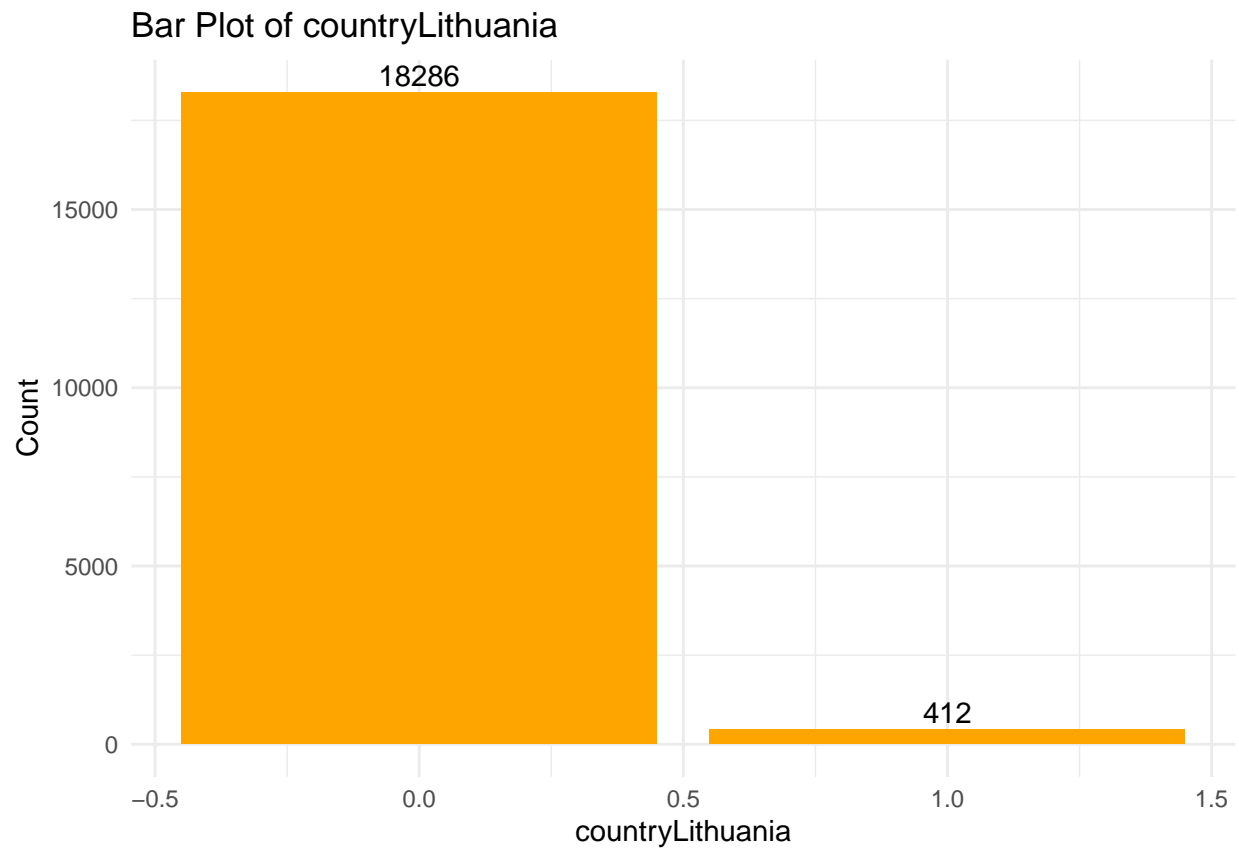


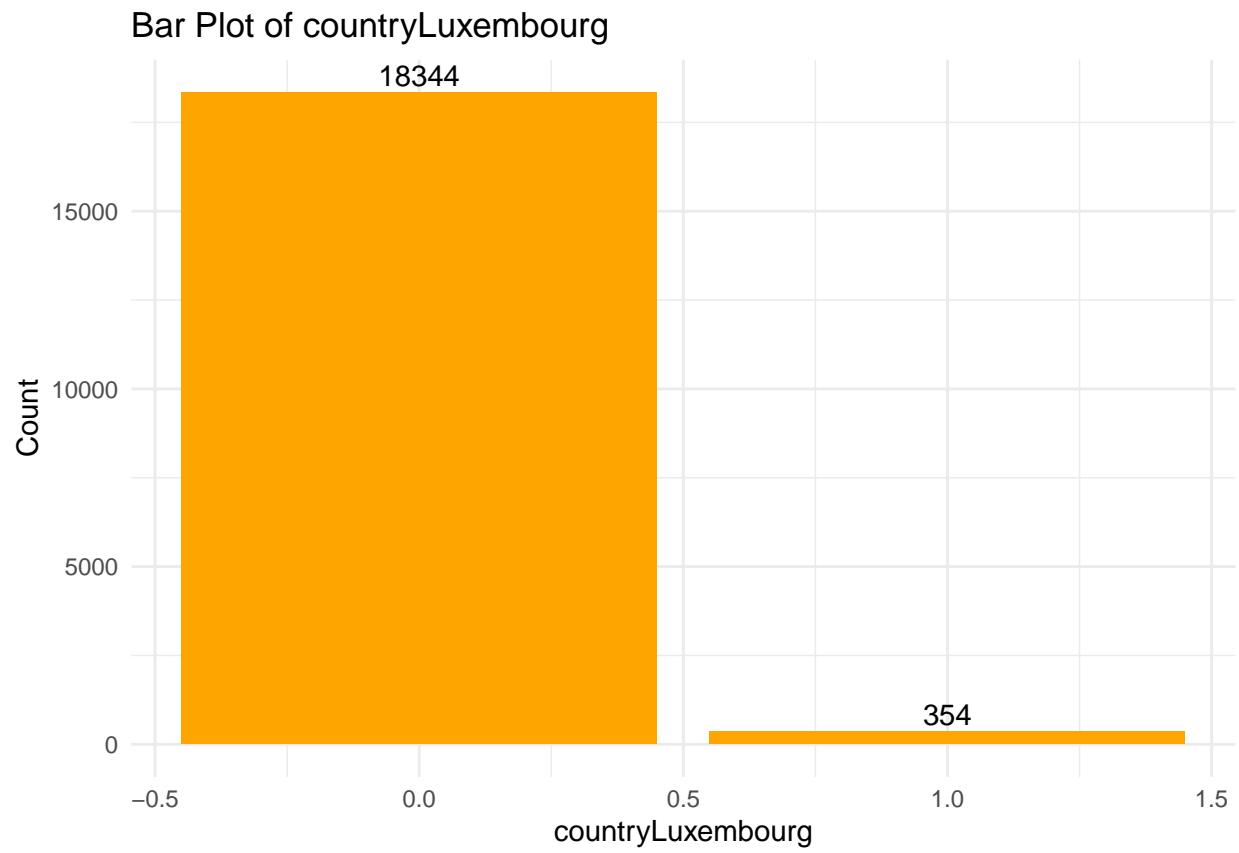


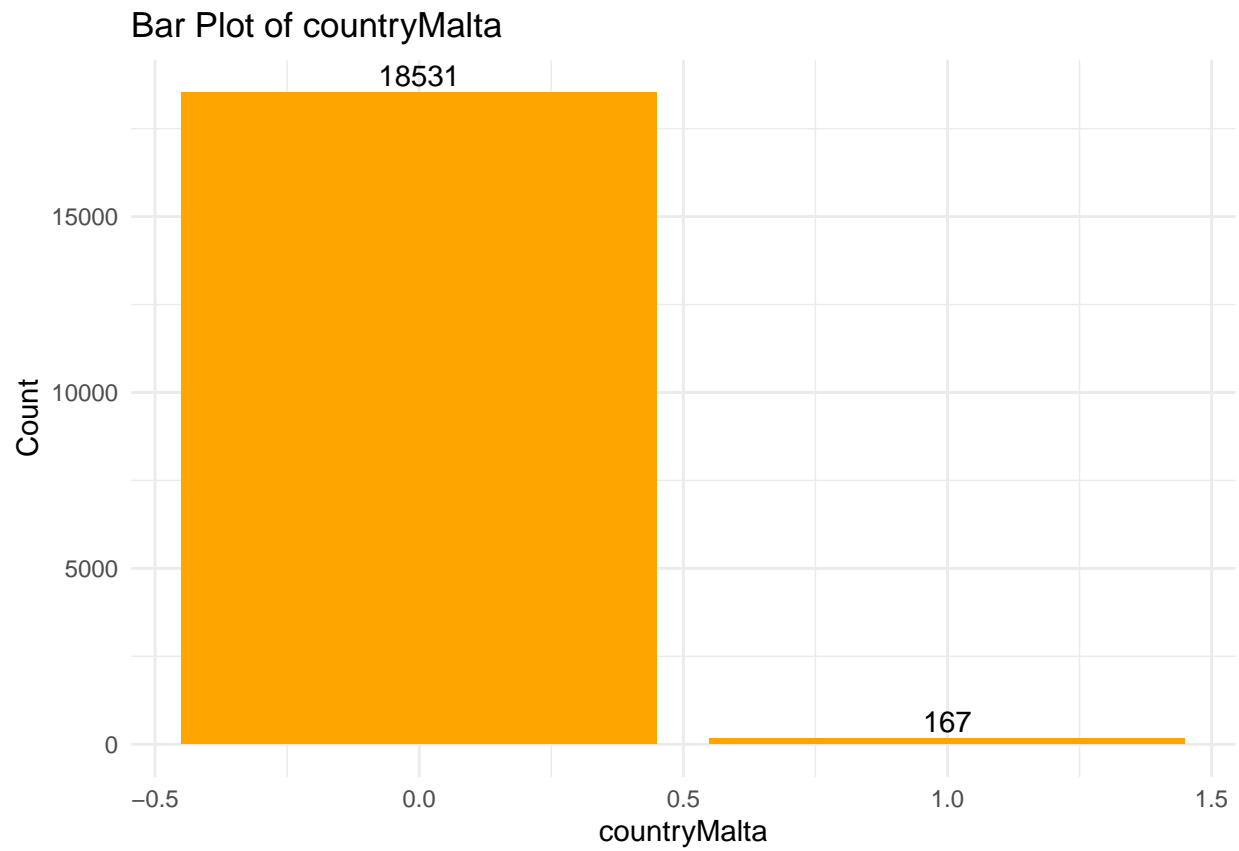


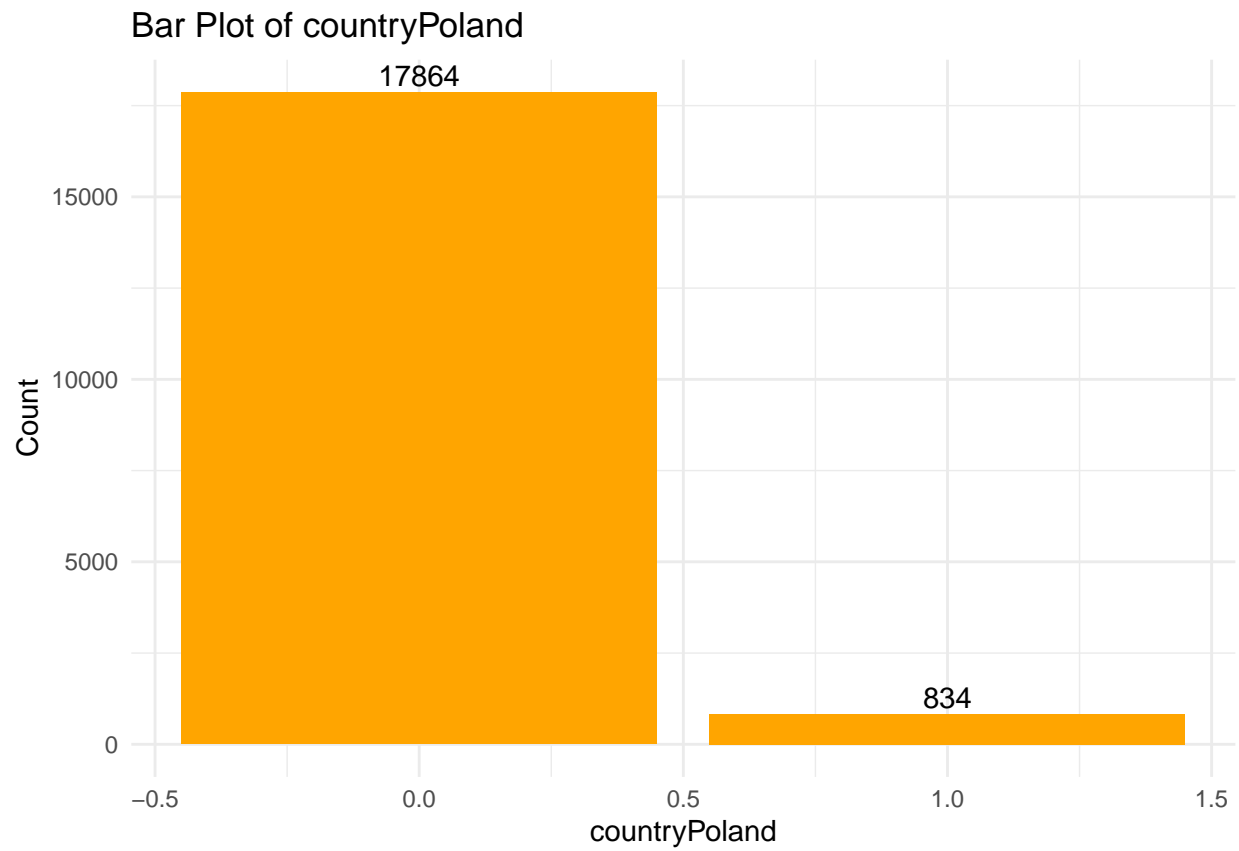


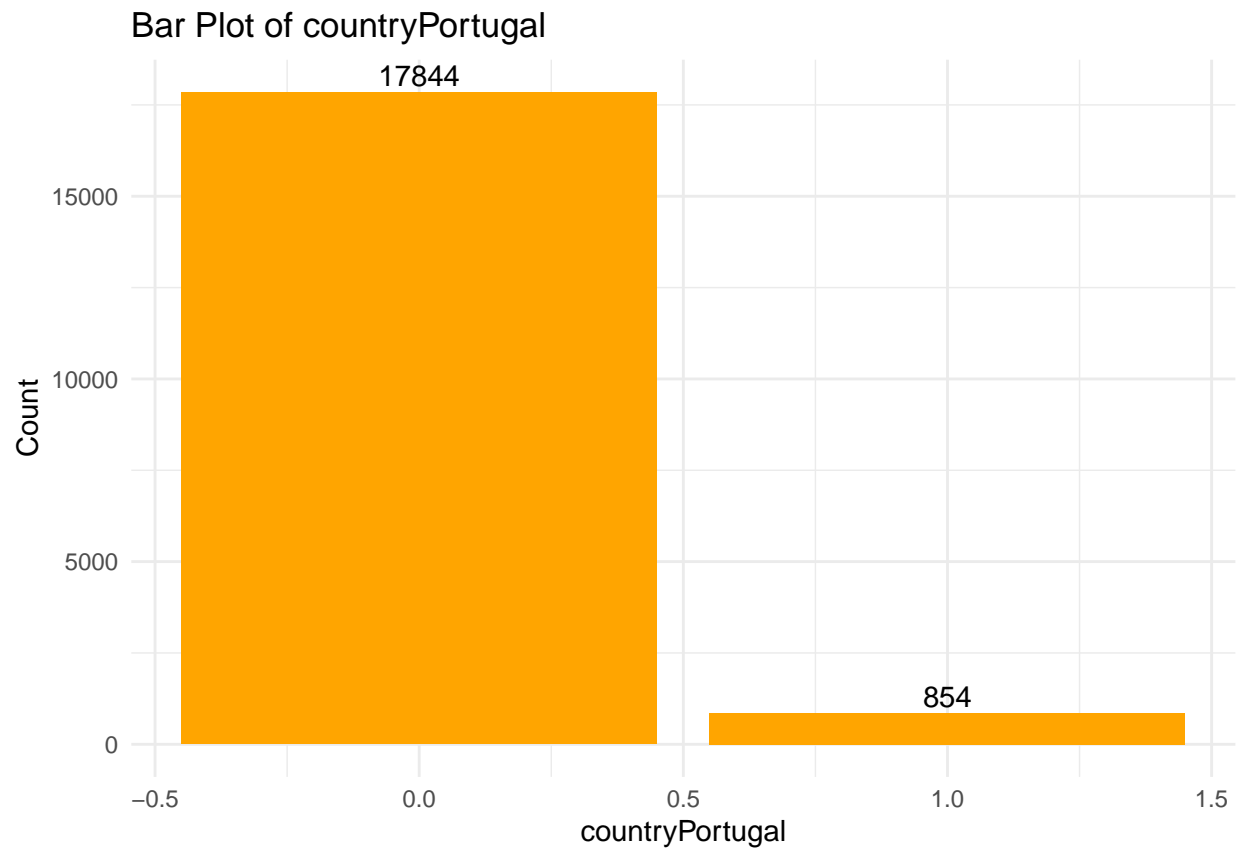


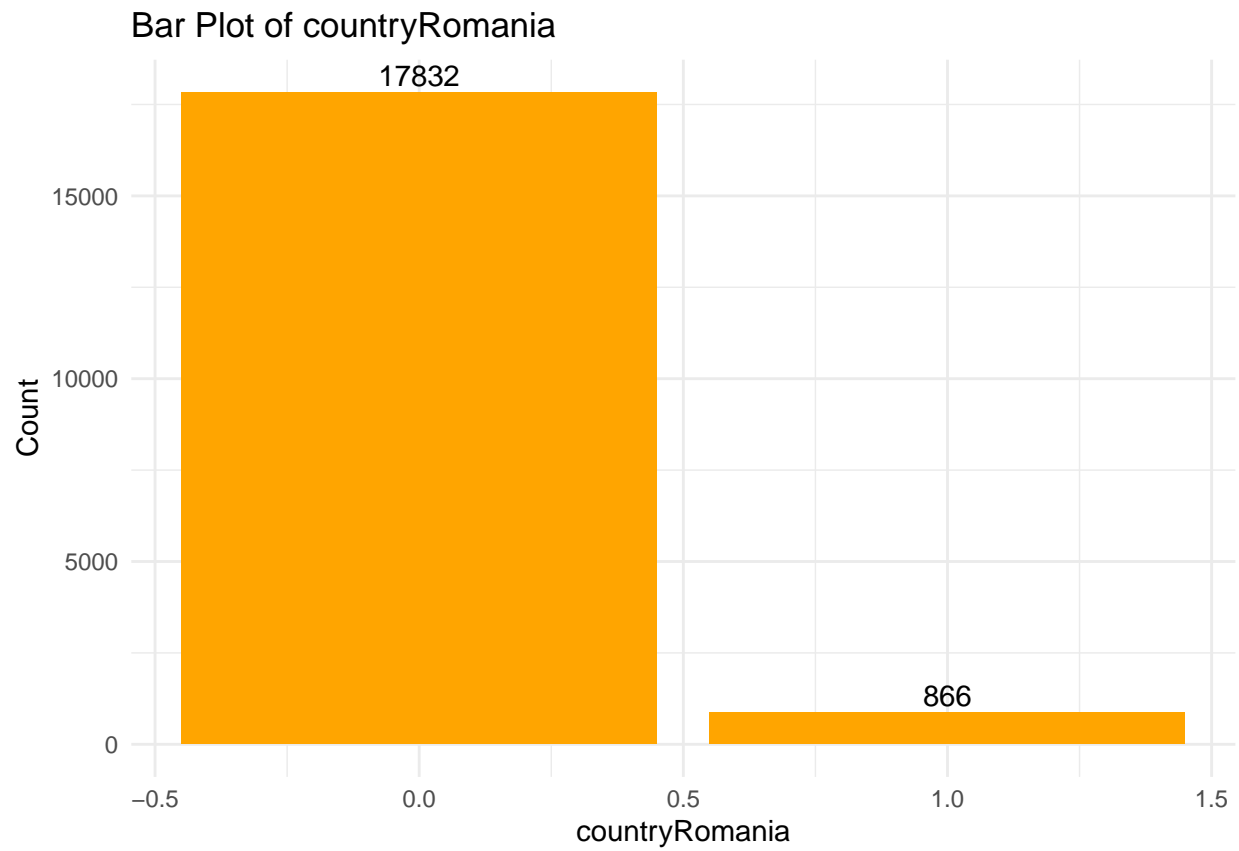


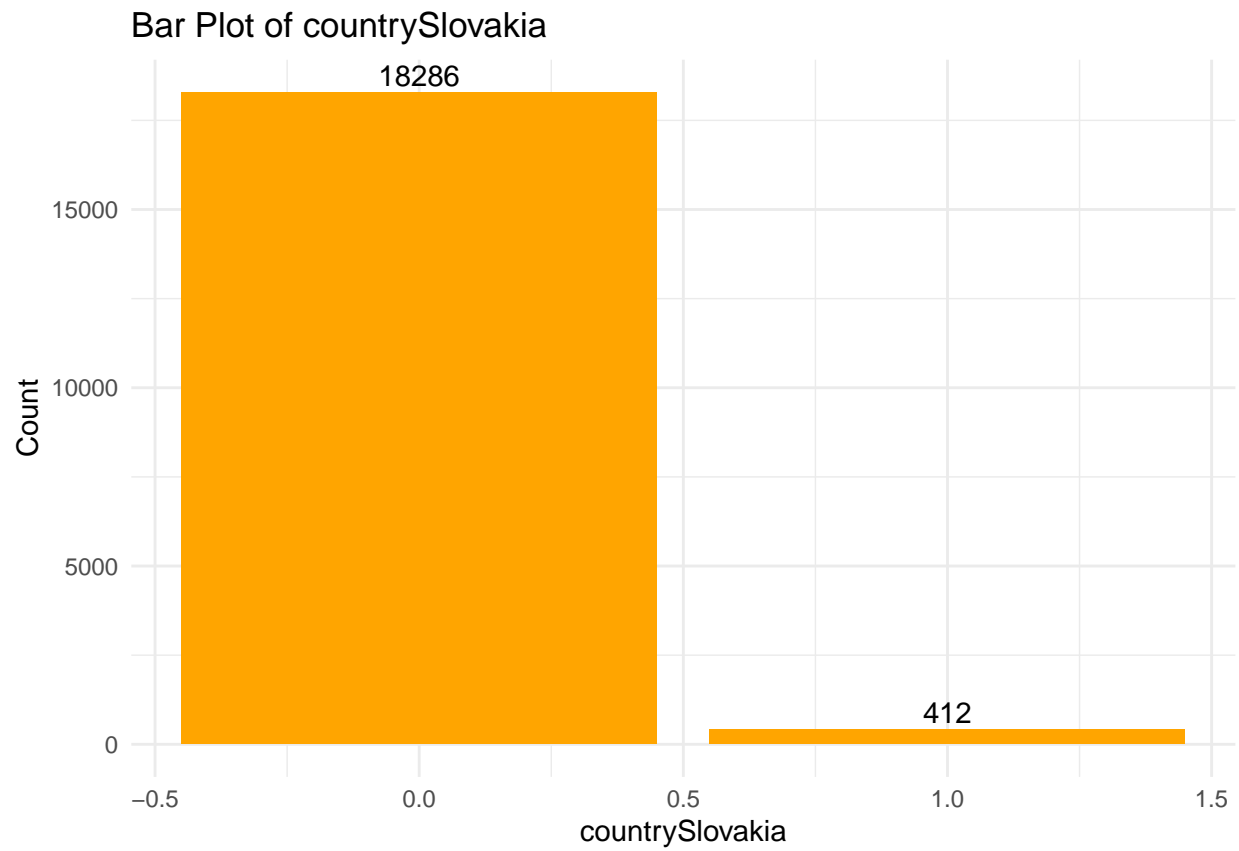


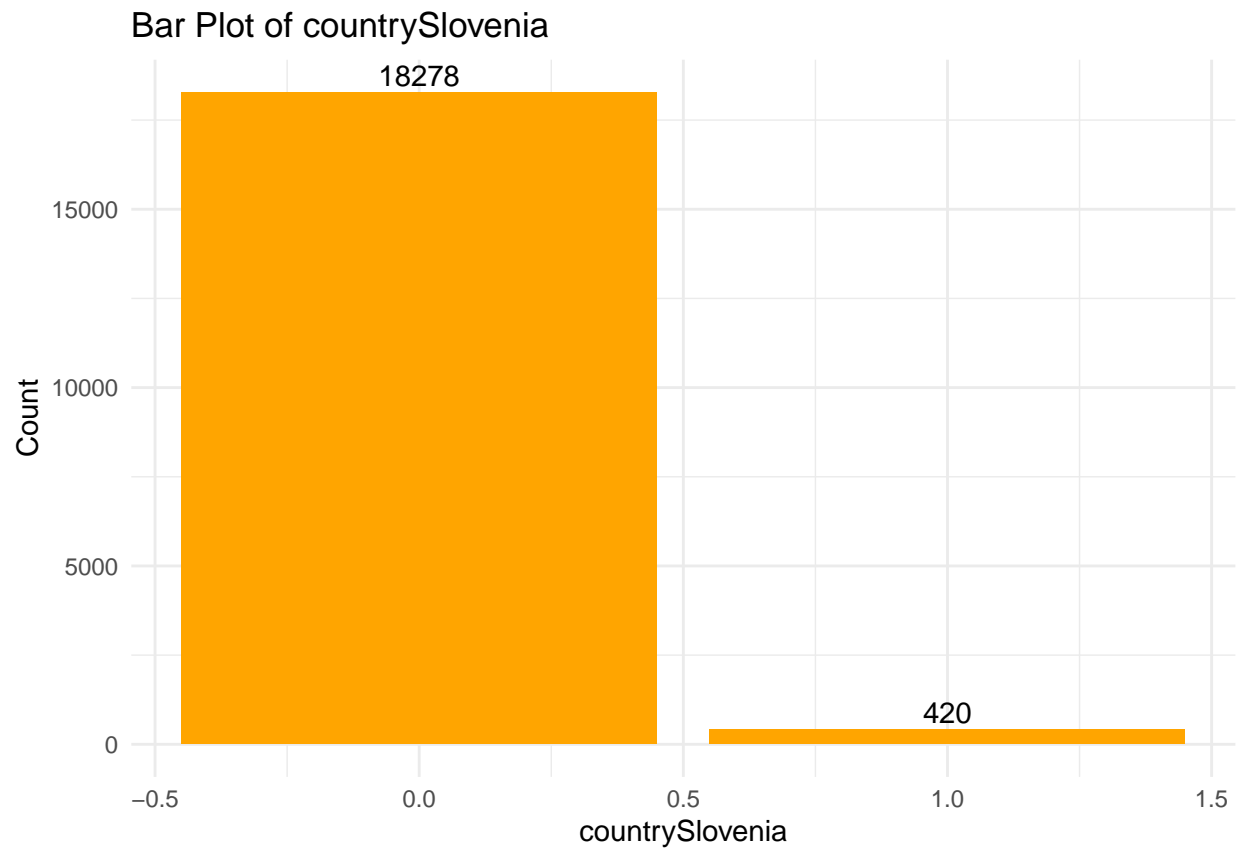


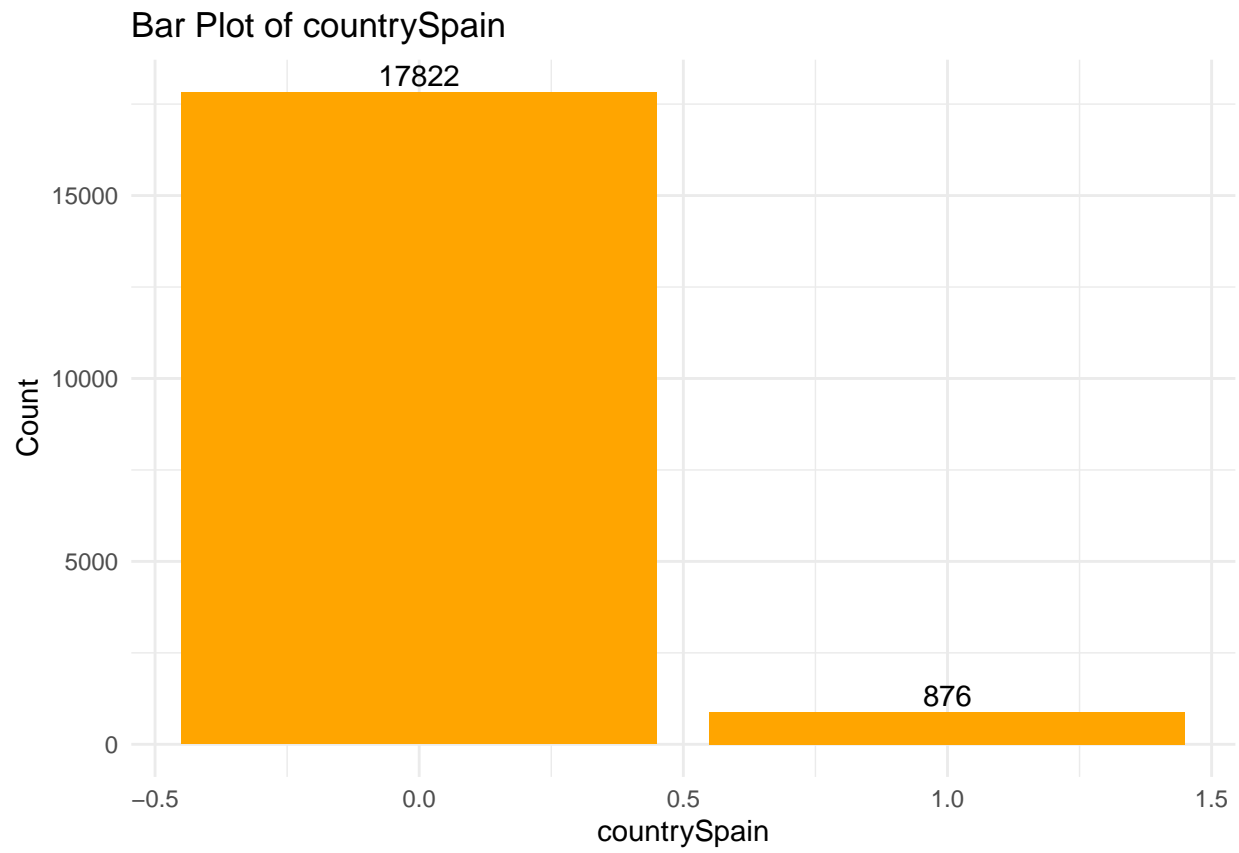


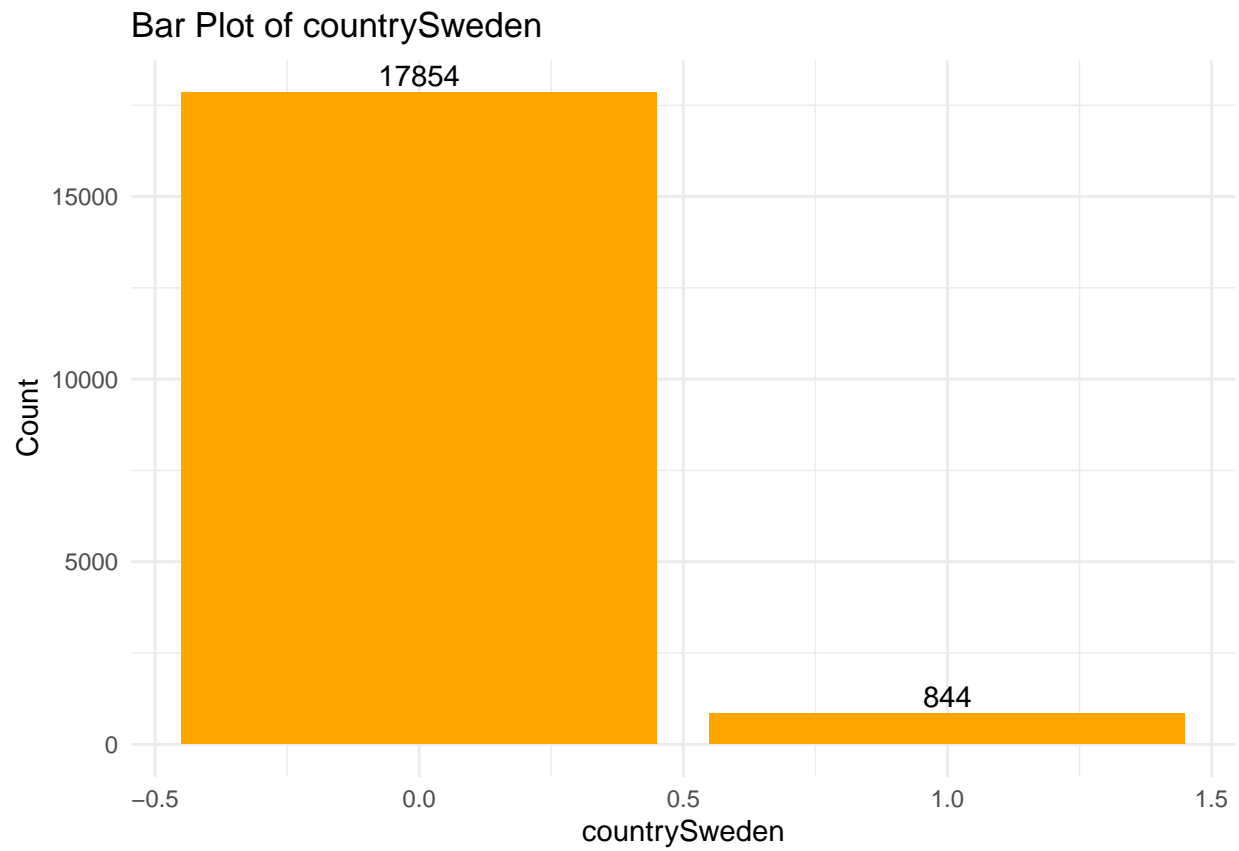


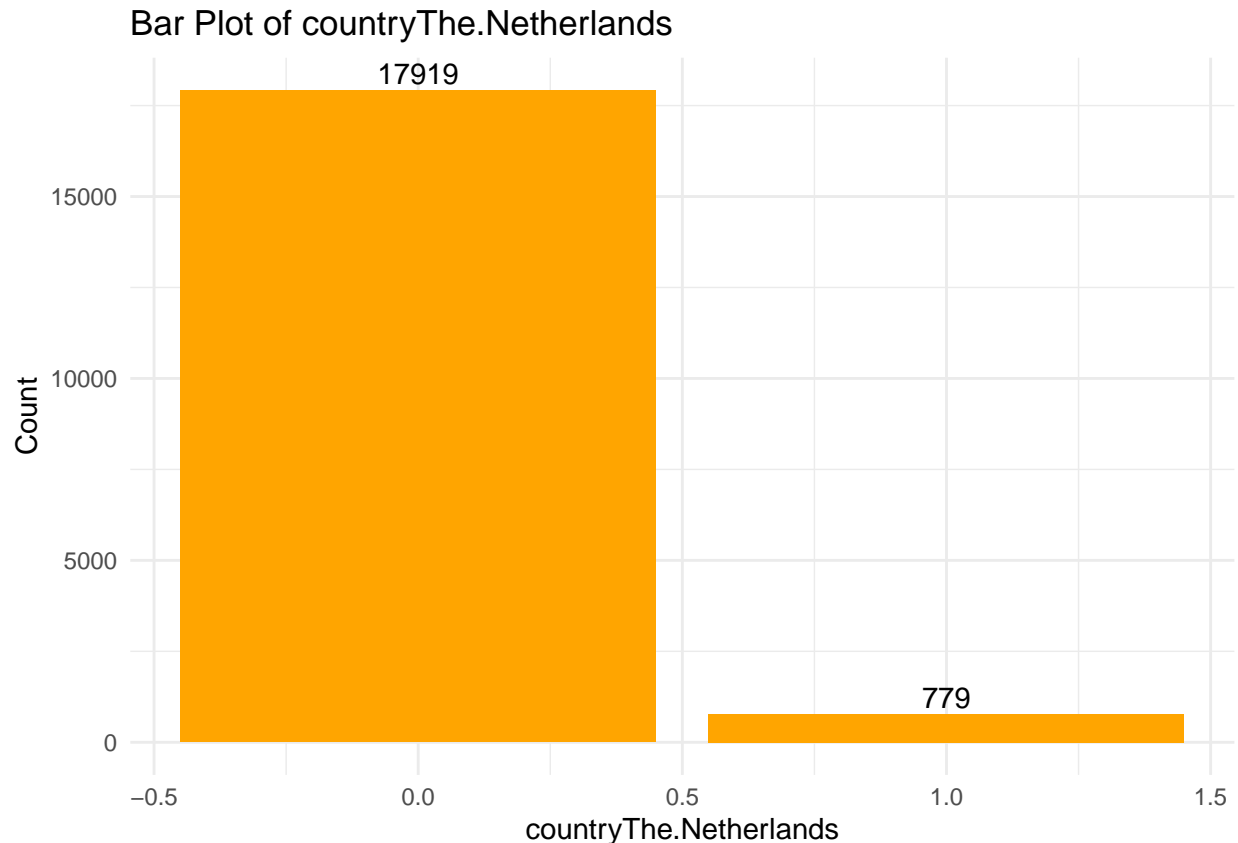












What are the distributions of the independent variables respect to ctax_binary (response) ?

```
# Function to create and save bar plots of ctax_binary against binary/categorical predictors
response_plots <- function(data, response = "ctax_binary") {
  # Ensure response is a factor
  data[[response]] <- as.factor(data[[response]])

  # Loop through all variables except the response
  for (v in setdiff(names(data), response)) {

    # Skip if all NA or constant
    if (all(is.na(data[[v]])) || length(unique(data[[v]])) <= 1) next

    # Determine plot type
    if (is.numeric(data[[v]]) && length(unique(data[[v]])) > 2) {
      # Continuous: use boxplot
      p <- ggplot(data, aes_string(x = response, y = v, fill = response)) +
        geom_boxplot() +
        labs(title = paste(v, "by", response),
             x = response, y = v) +
        theme_minimal()
    } else {
      # Categorical/binary: use bar plot
      p <- ggplot(data, aes_string(x = v, fill = response)) +
        geom_bar(position = "dodge") +
```

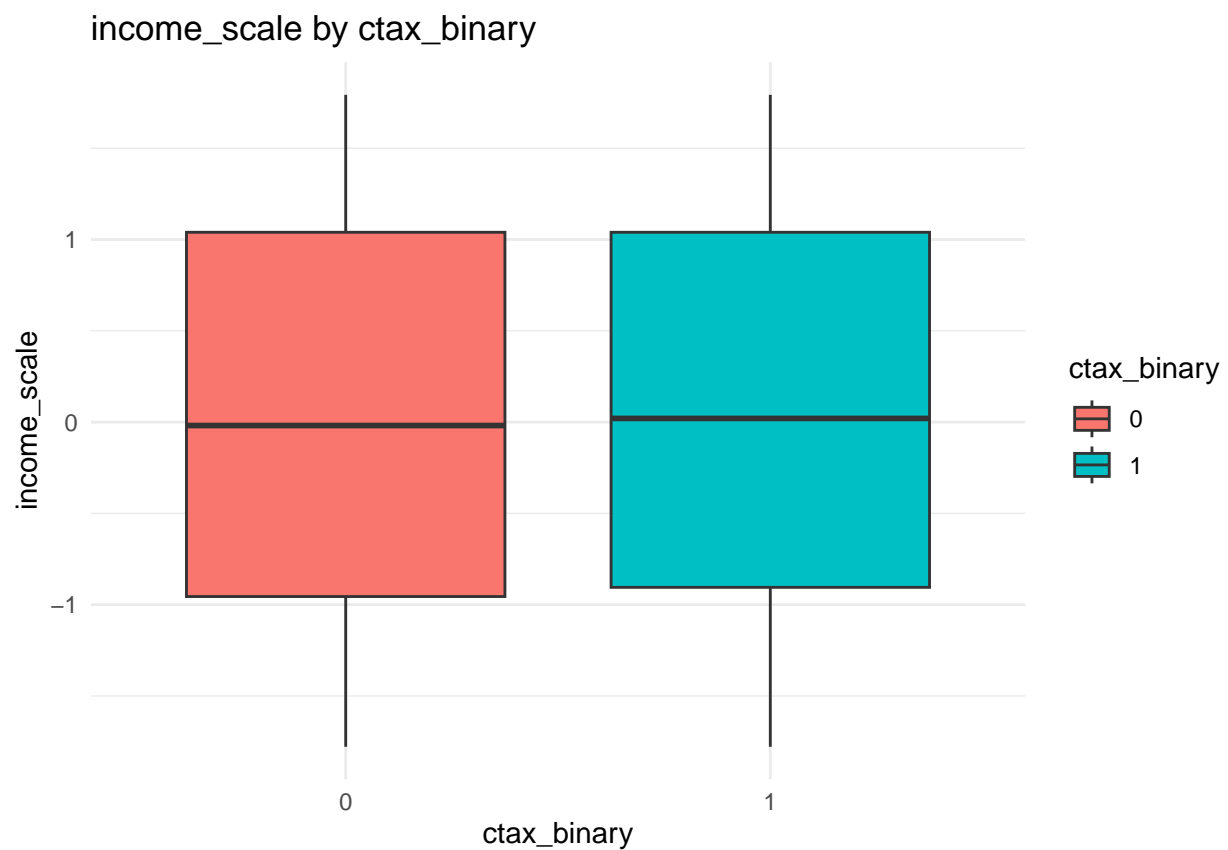
```

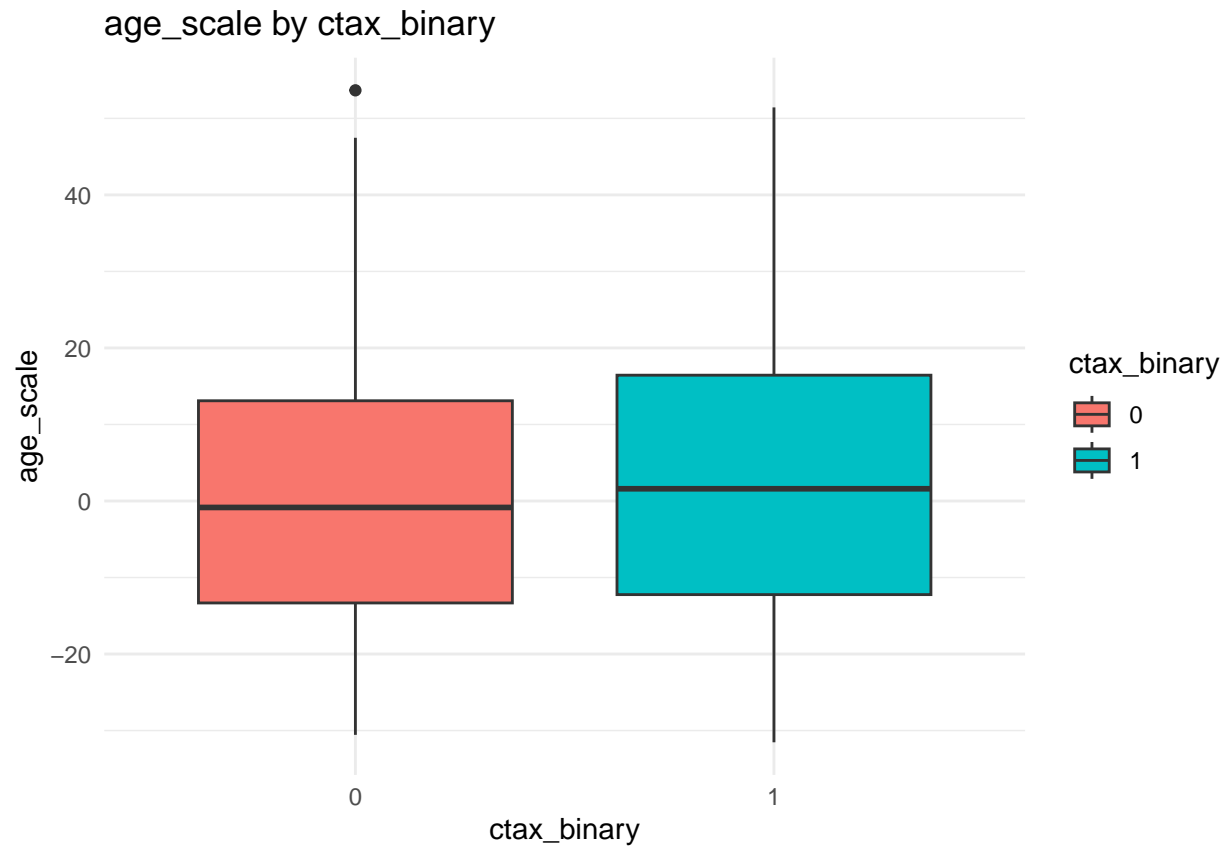
    geom_text(stat = "count", aes(label = ..count..),
              position = position_dodge(width = 0.9), vjust = -0.3) +
    labs(title = paste("ctax_binary by", v),
          x = v, y = "Count", fill = response) +
    theme_minimal()
  }

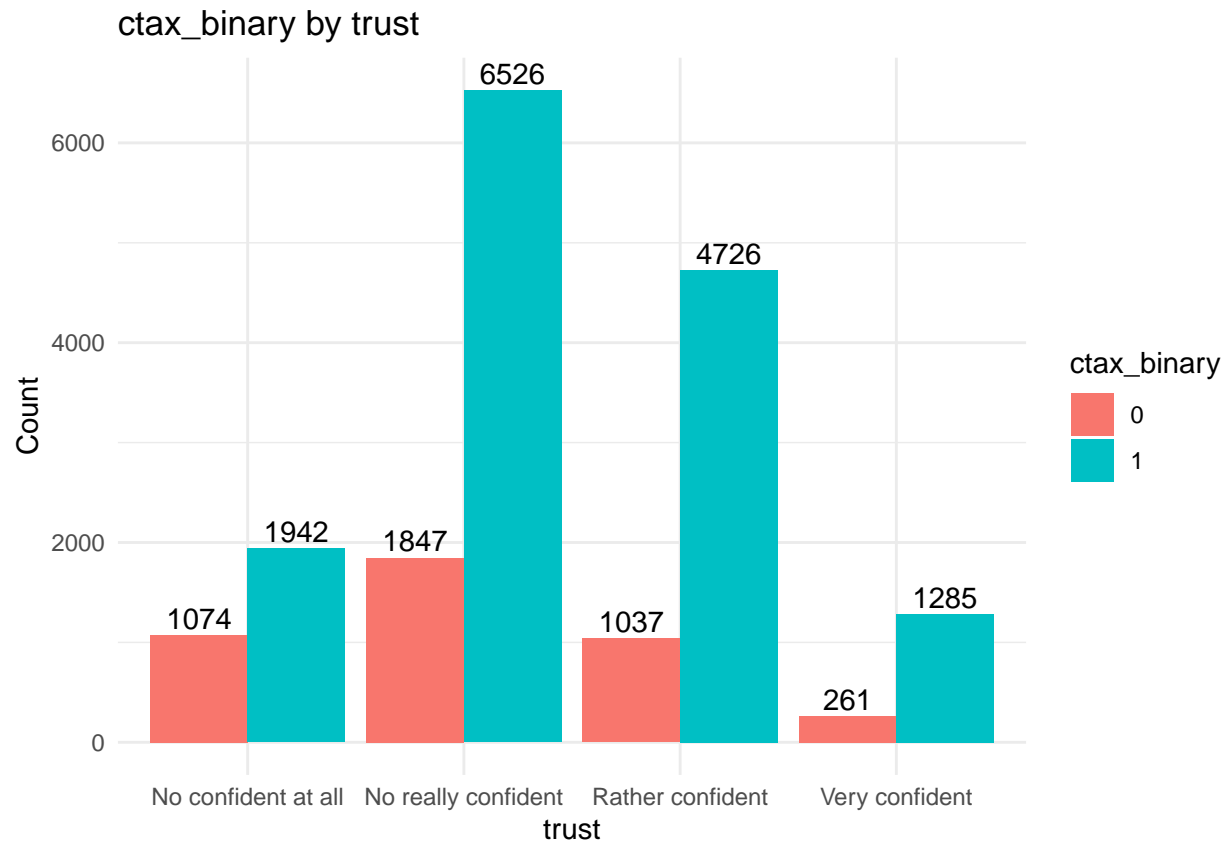
  print(p)
}
}

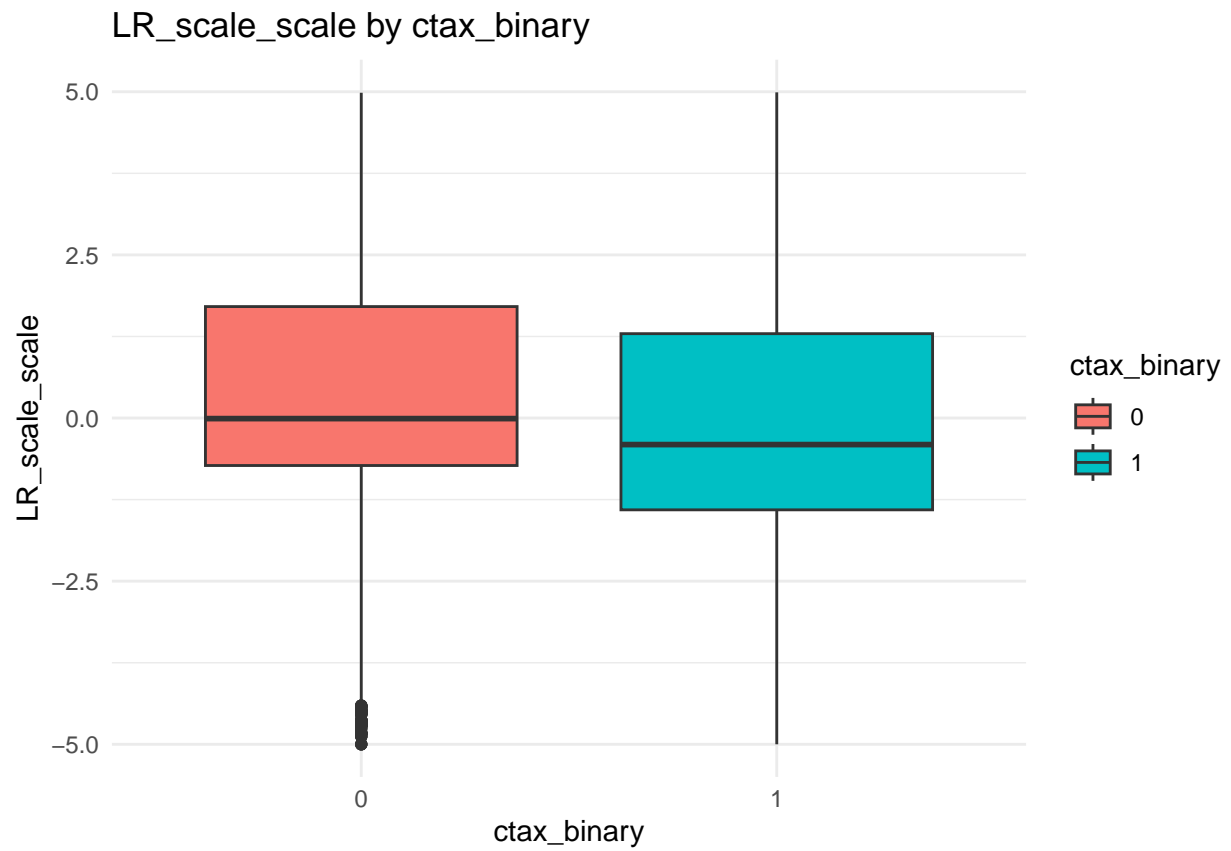
```

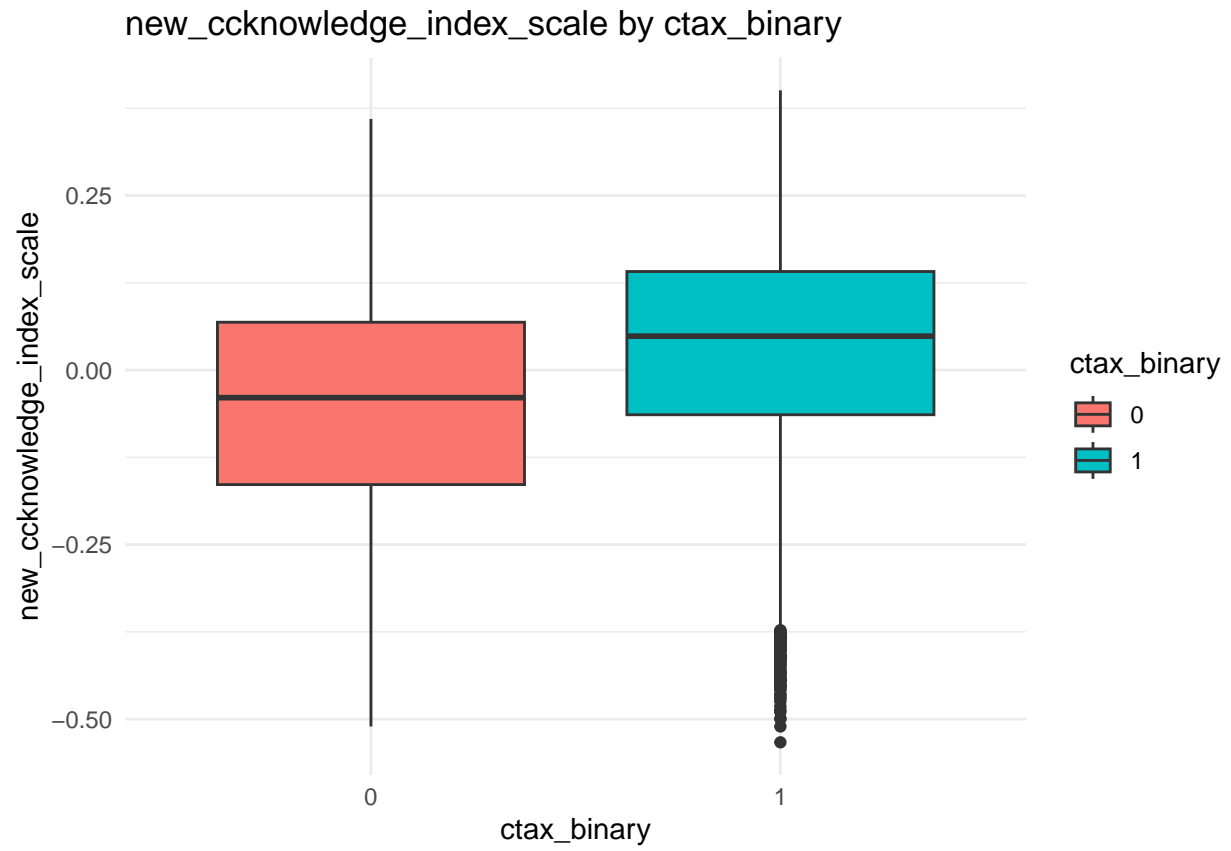
```
response_plots(df_clean, response="ctax_binary")
```

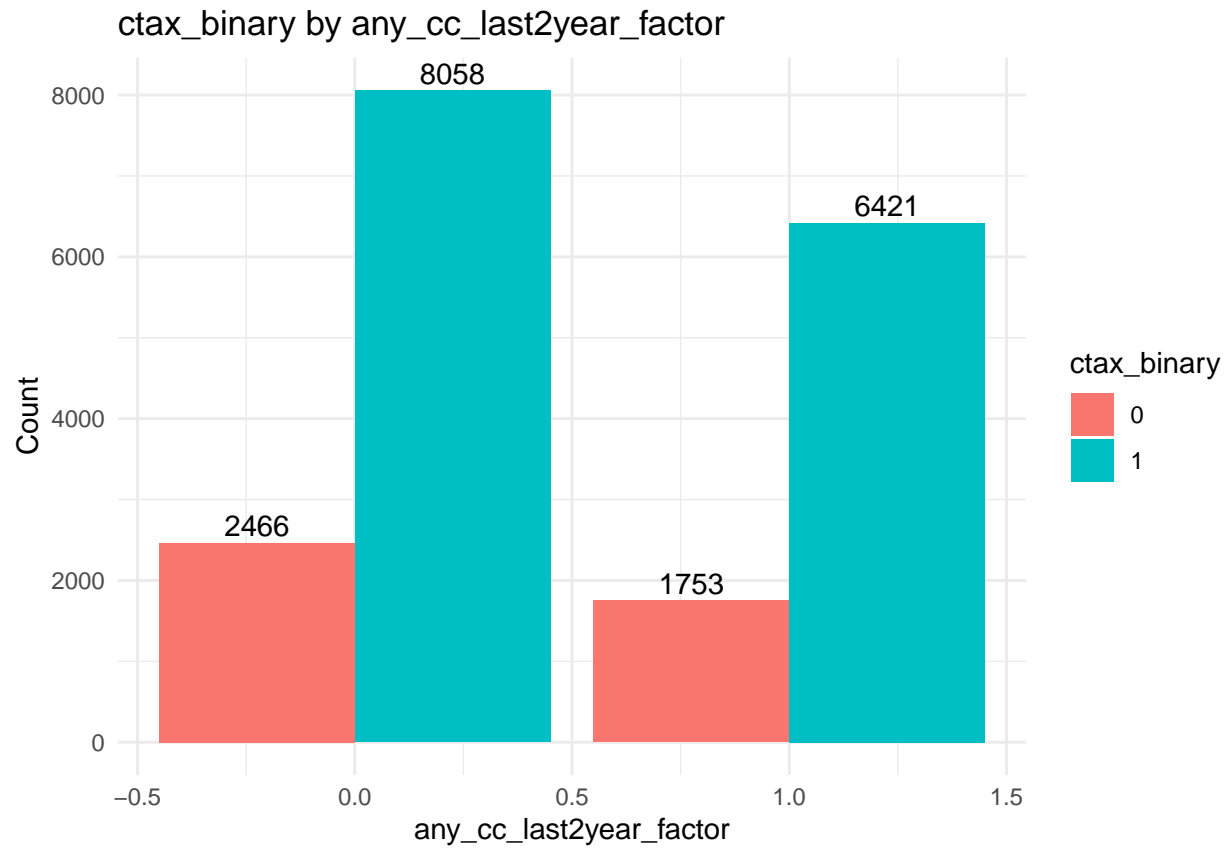


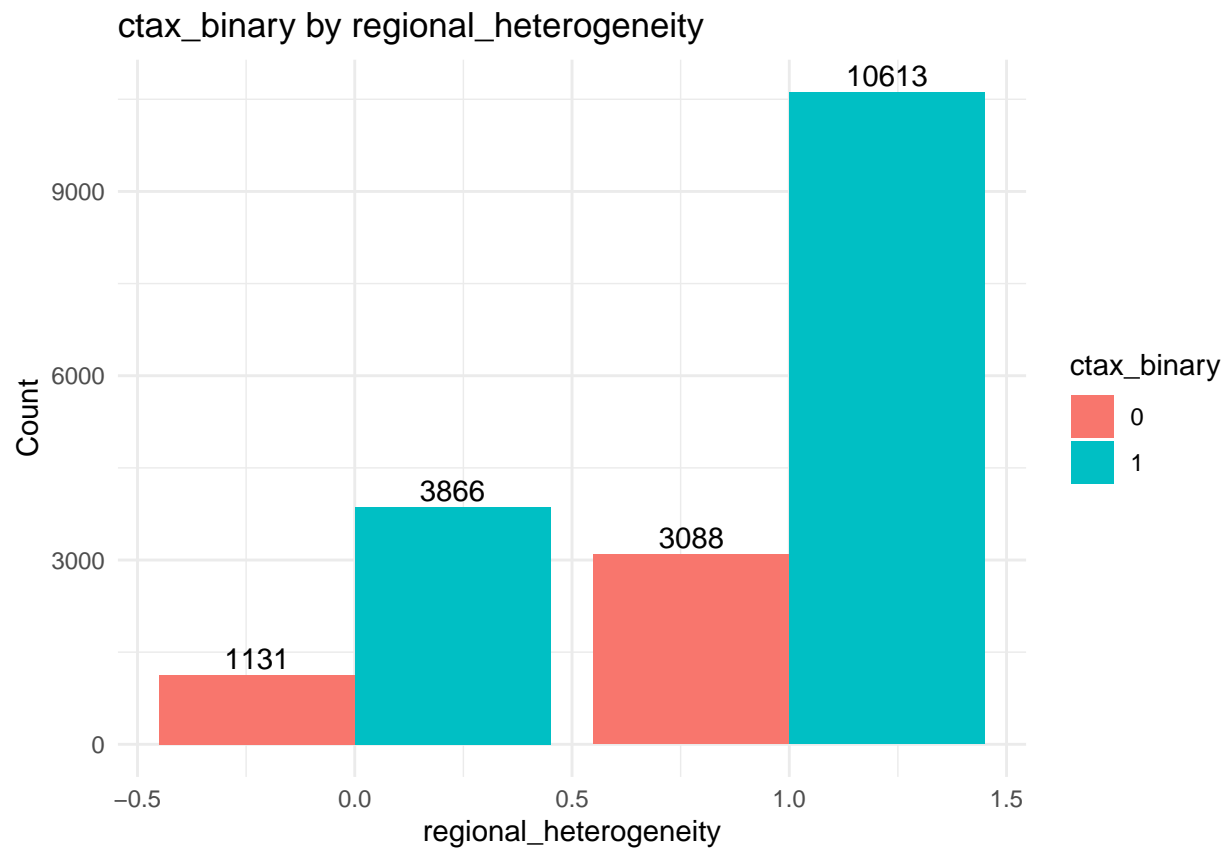


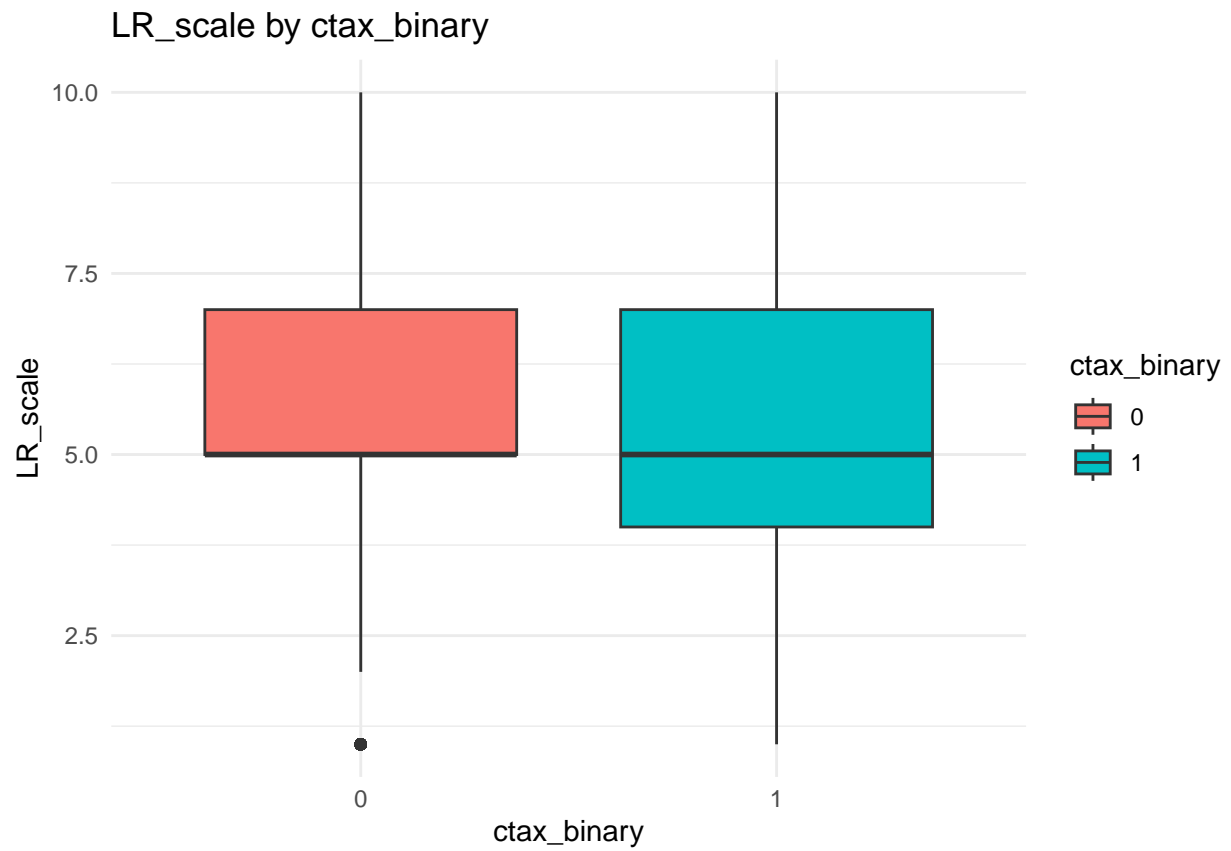


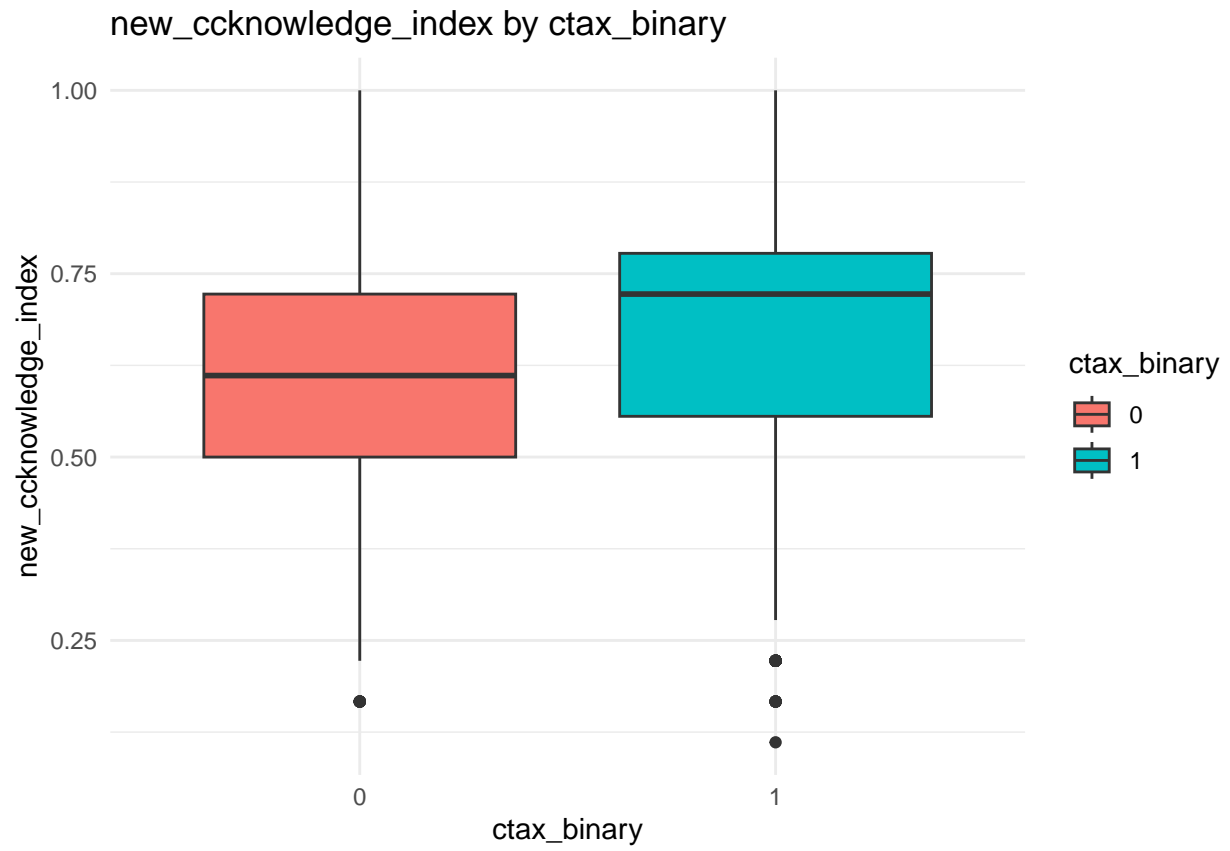


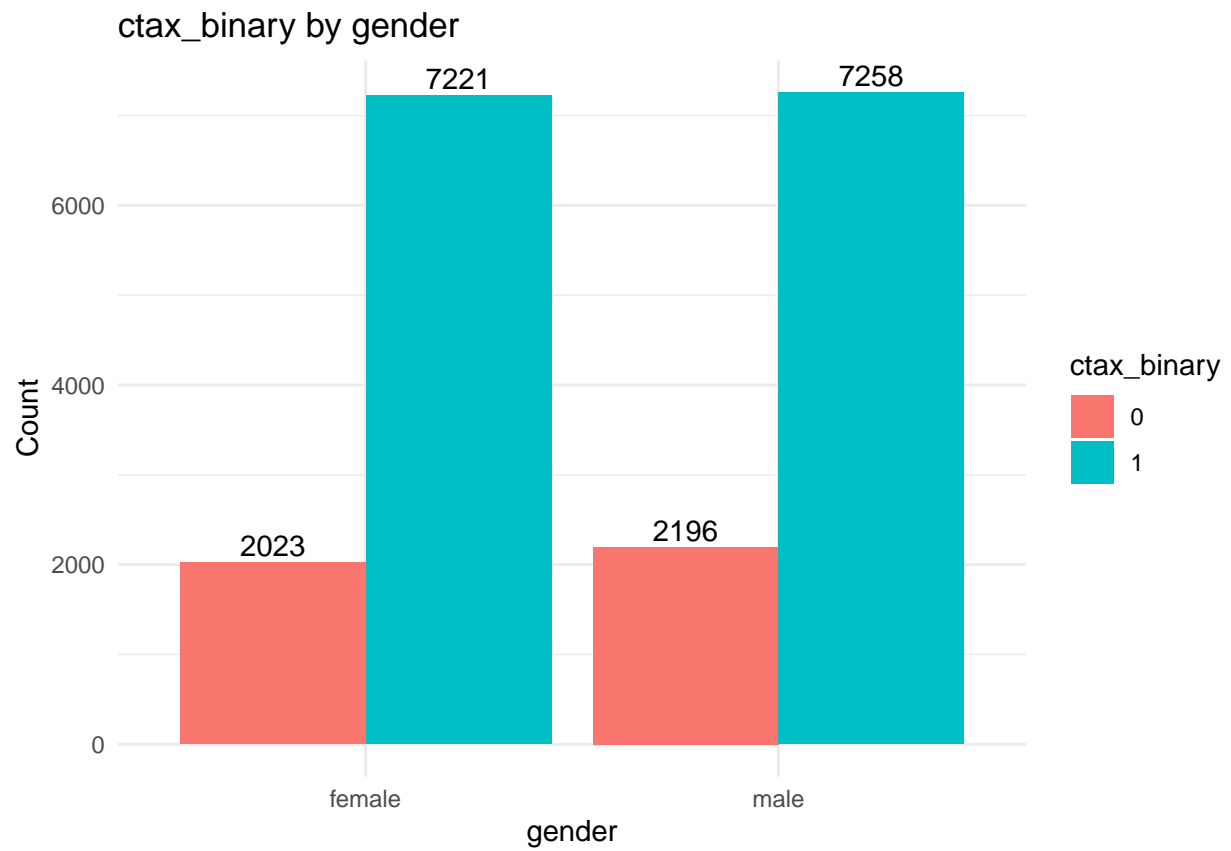


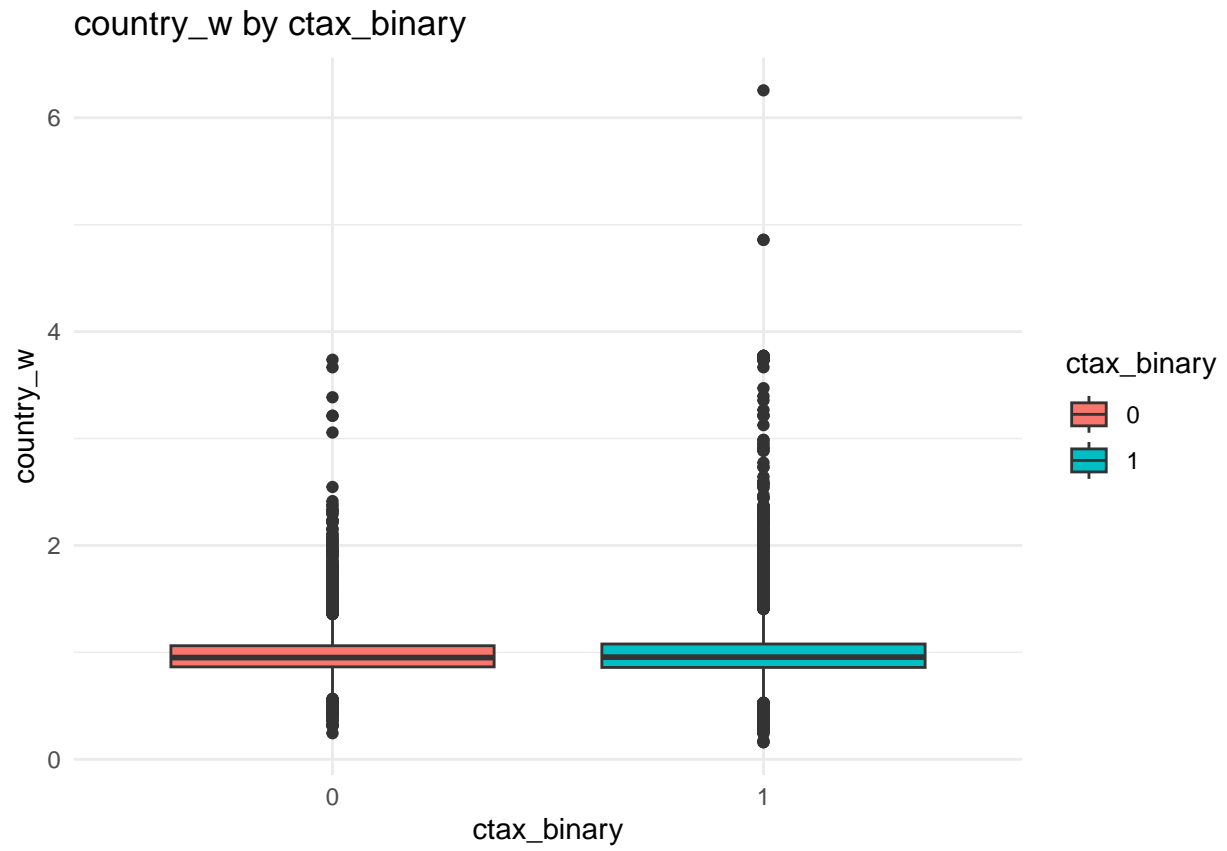


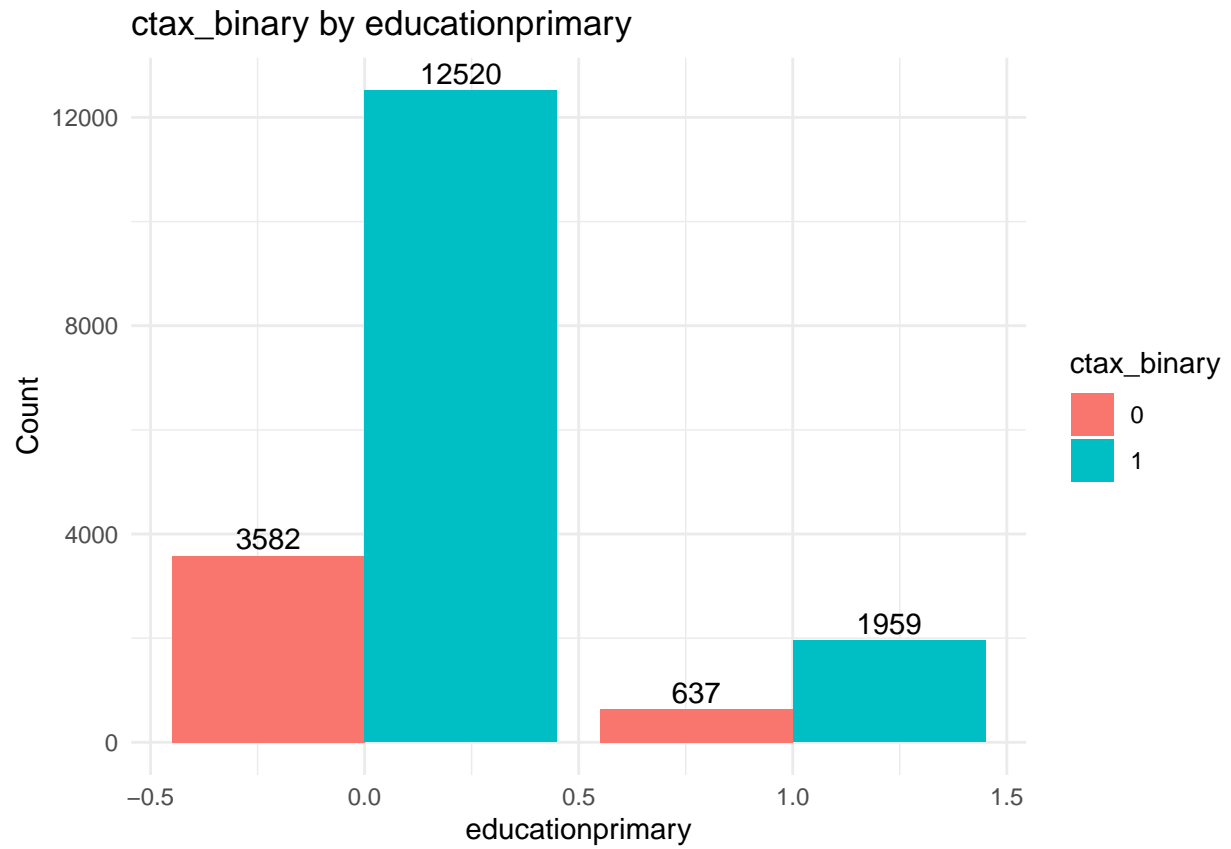


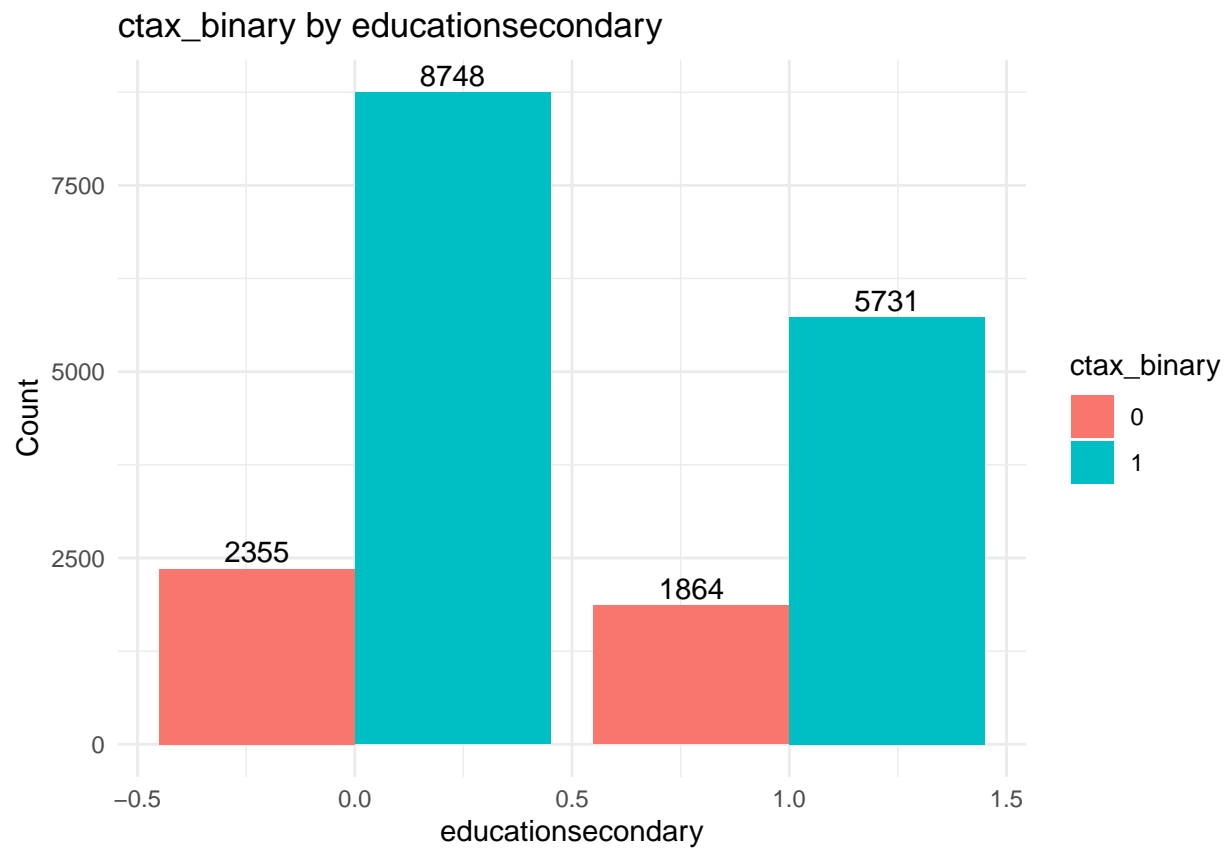


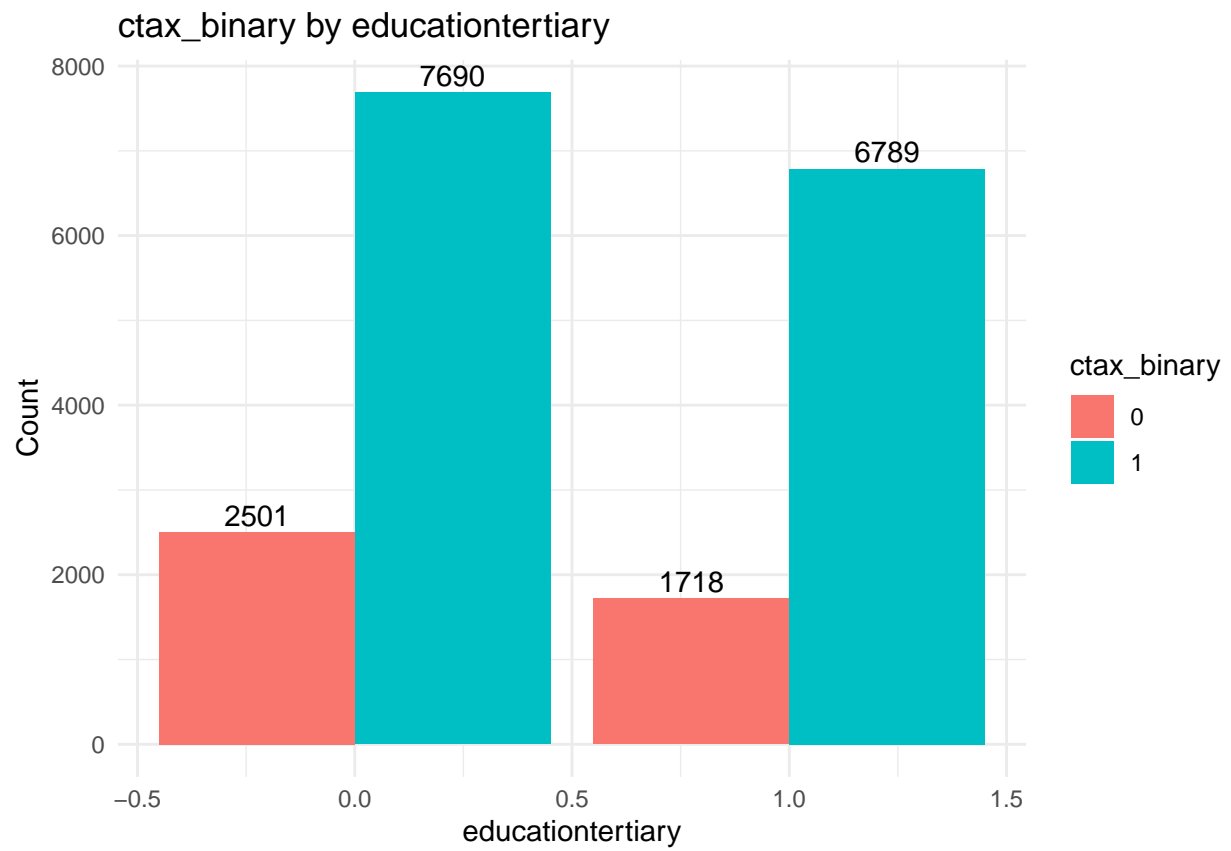


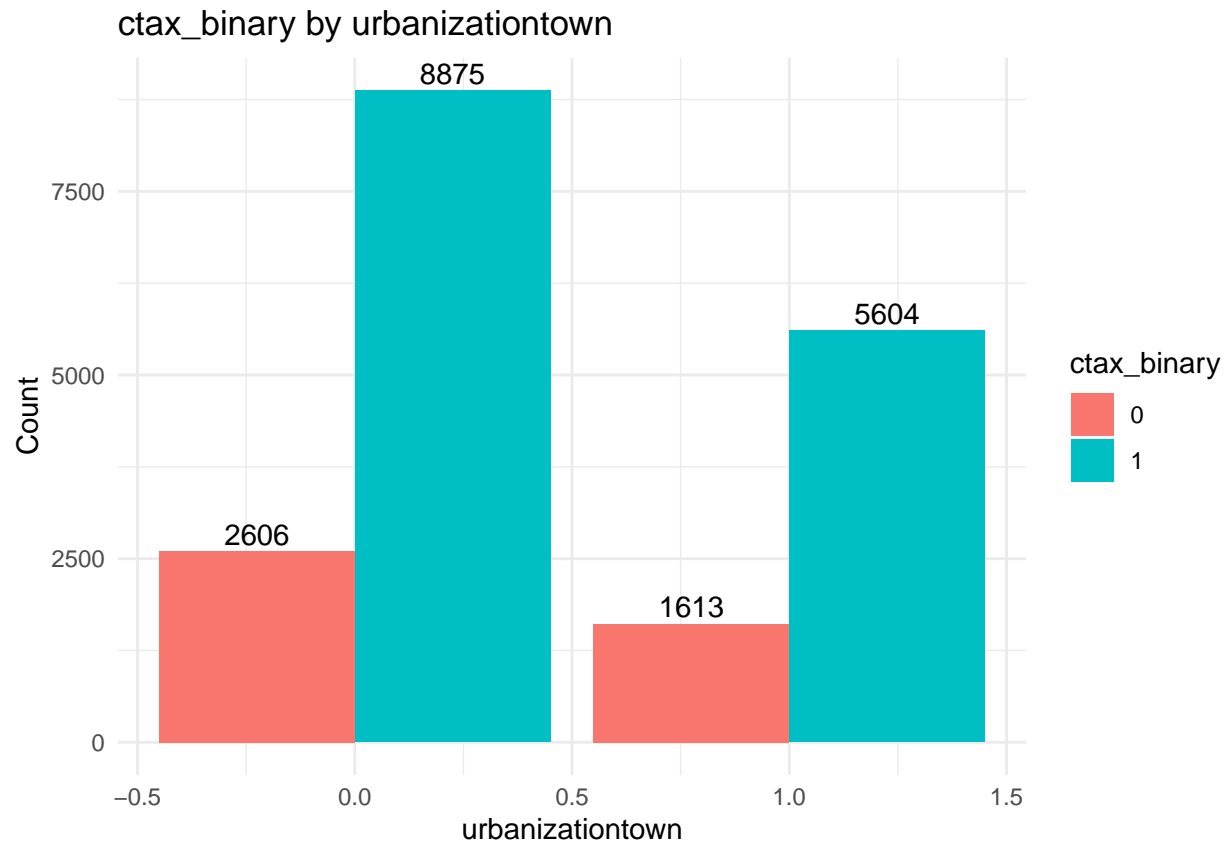


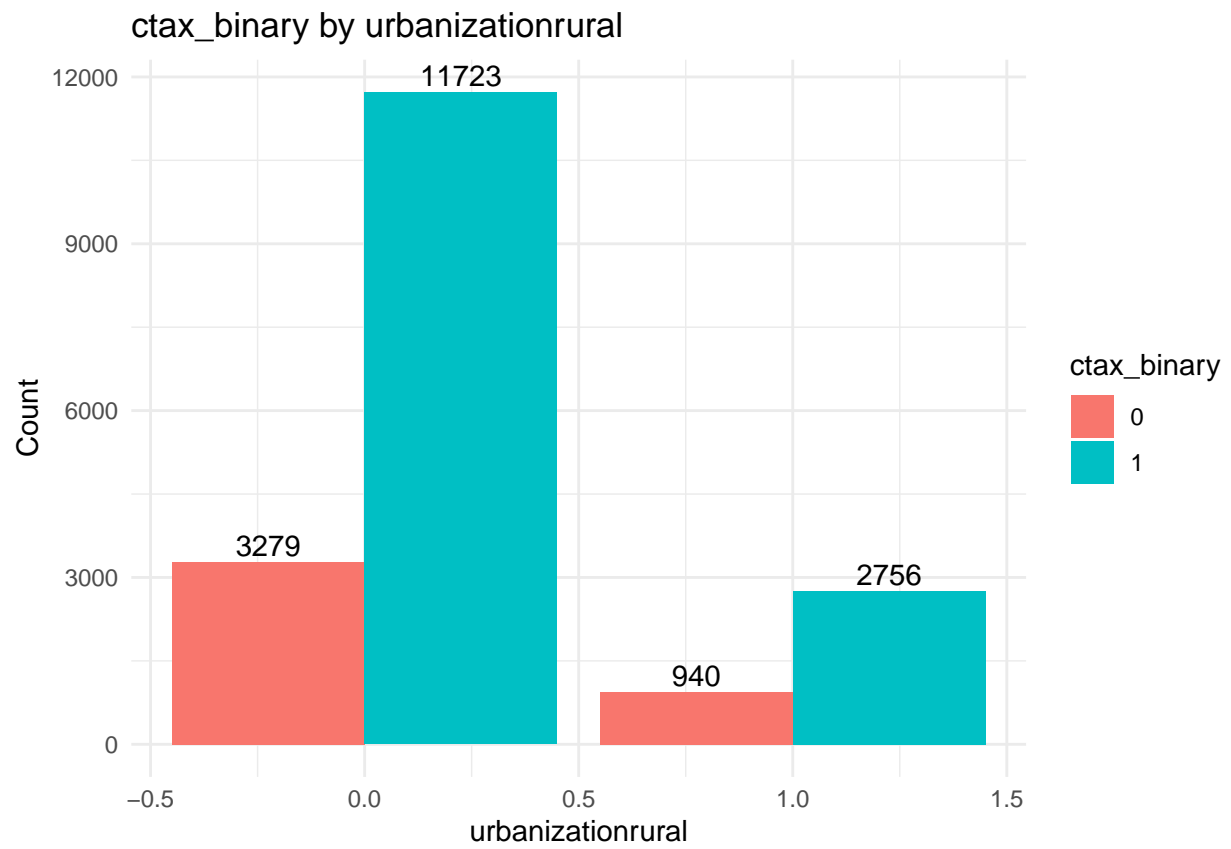


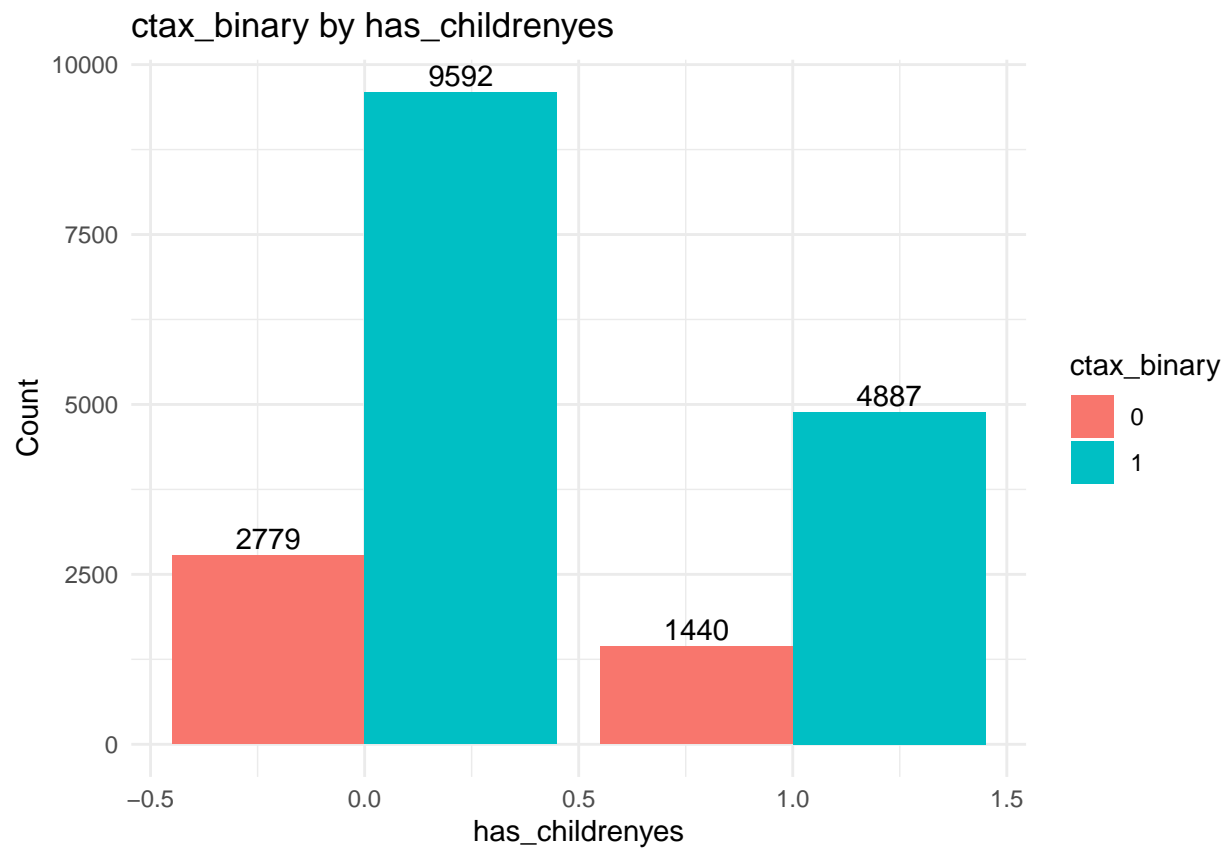


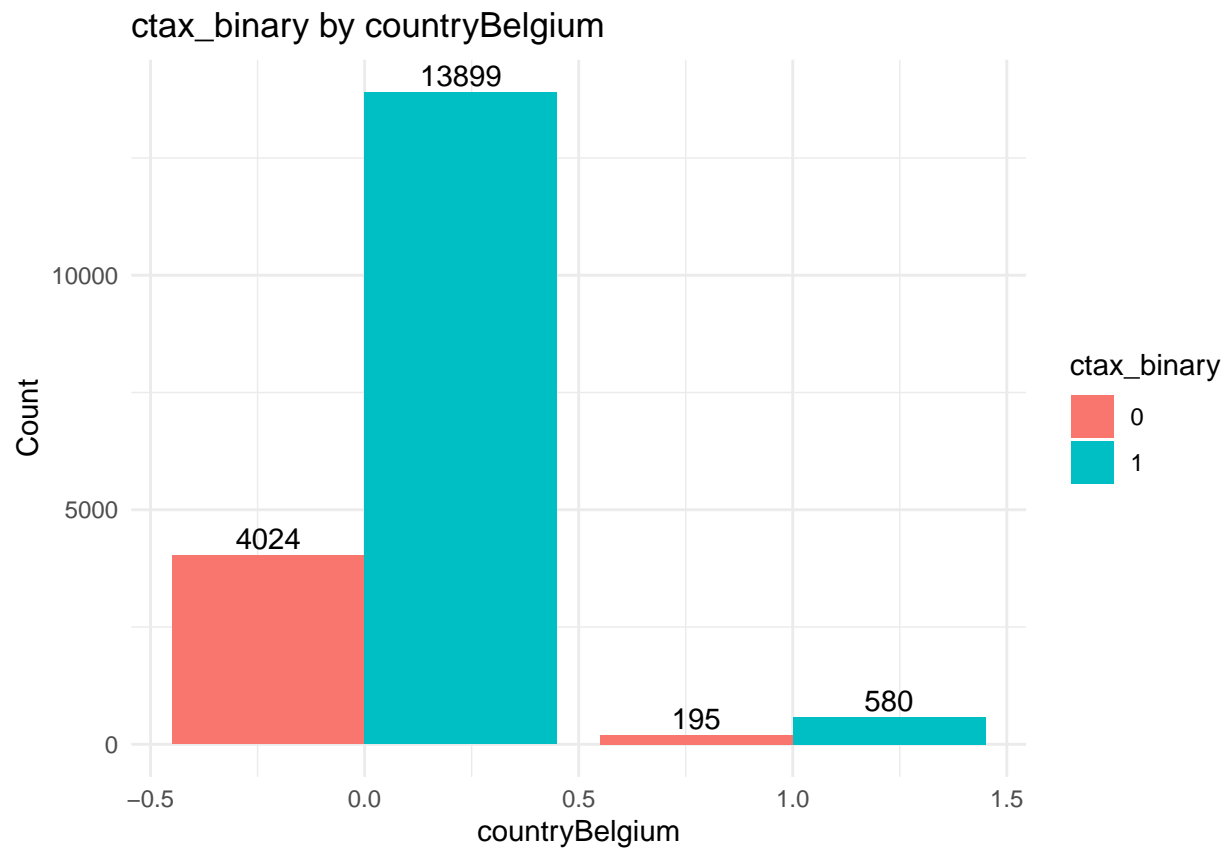


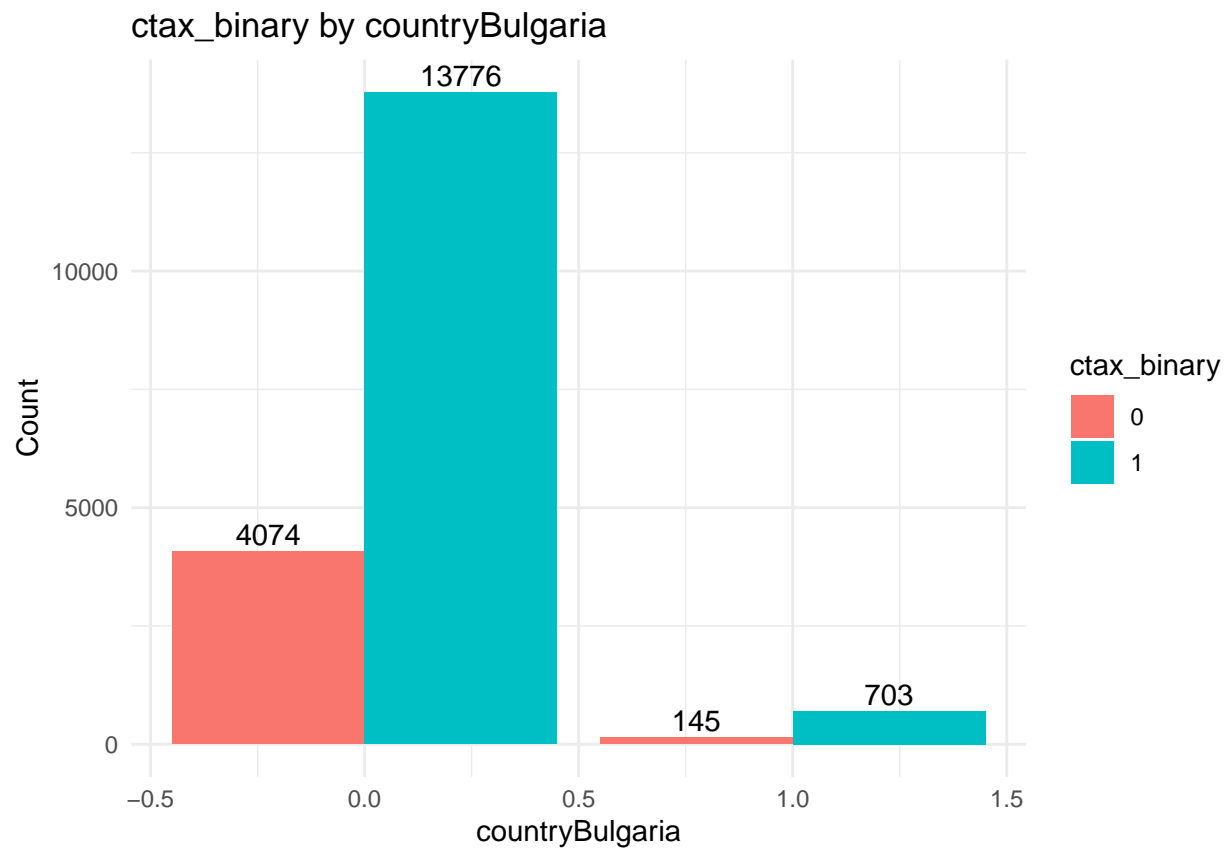


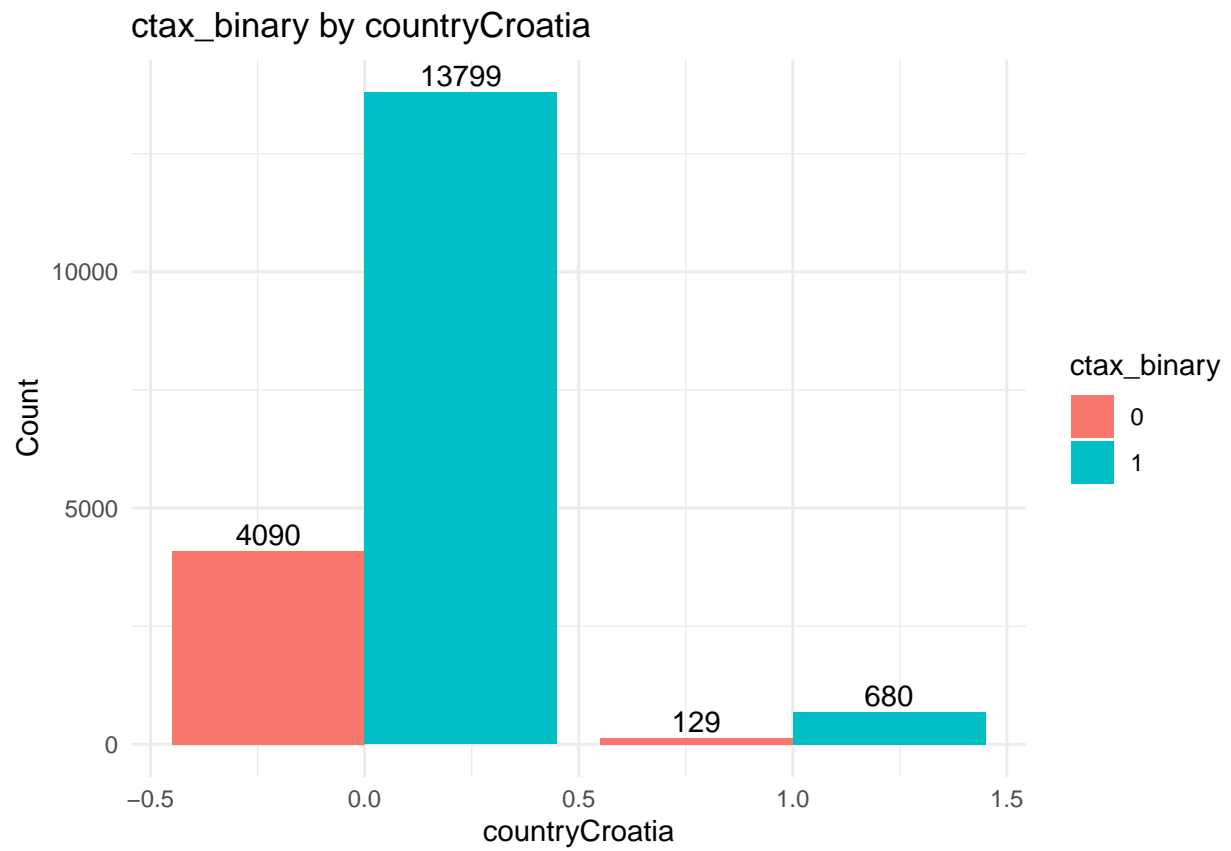


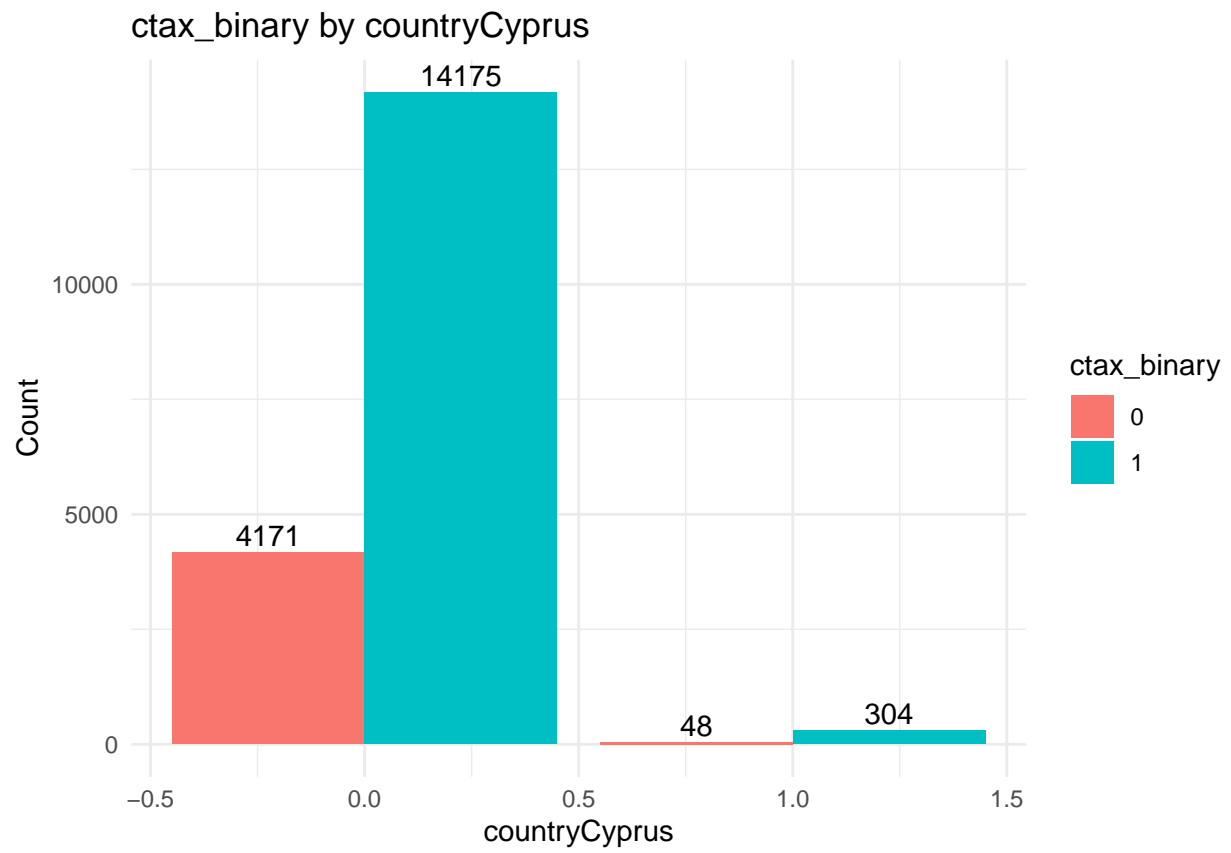


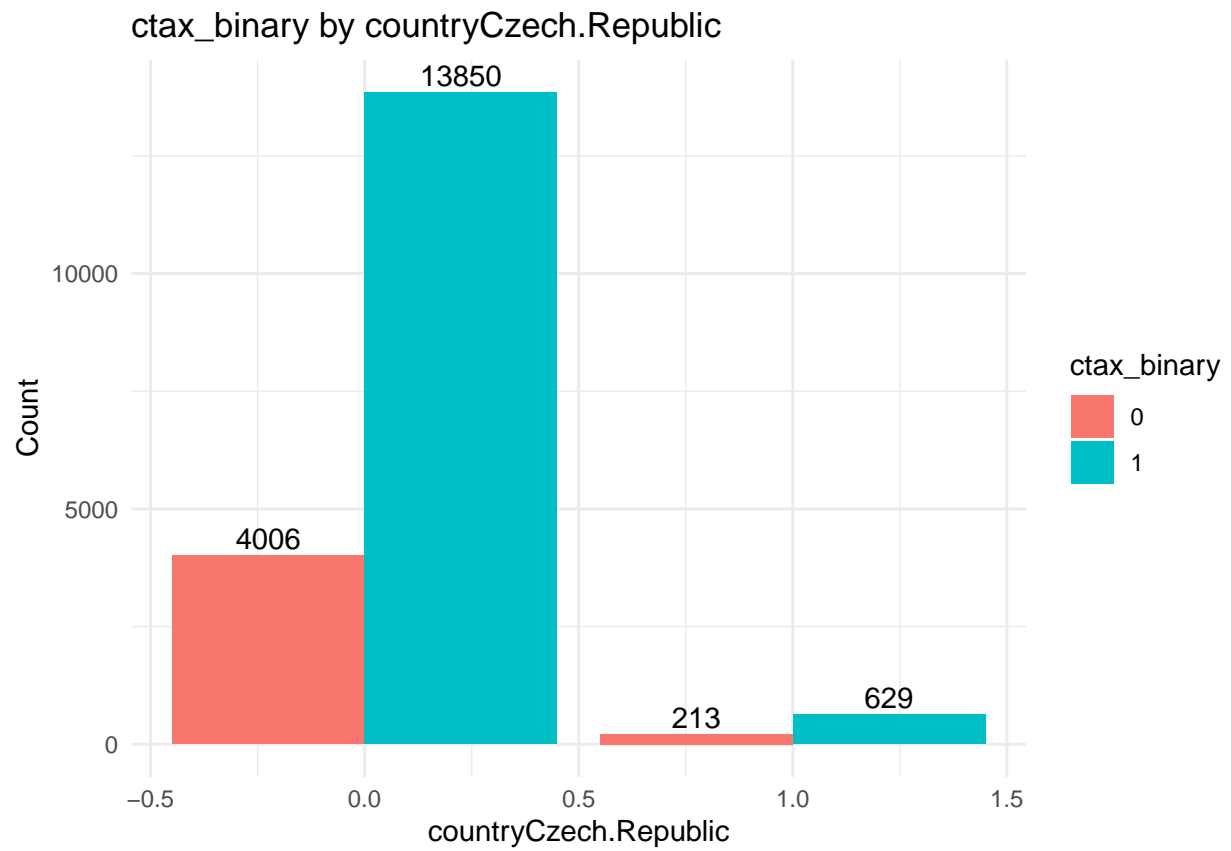


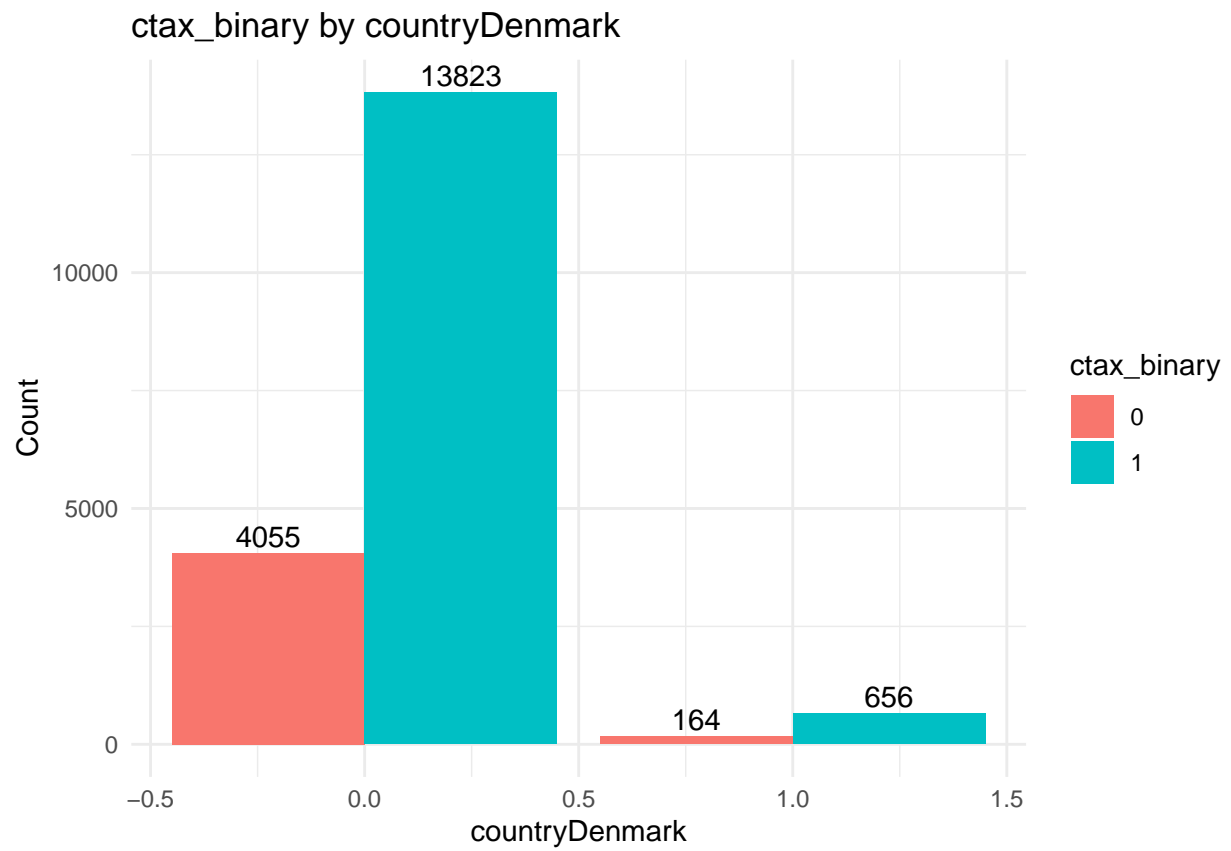


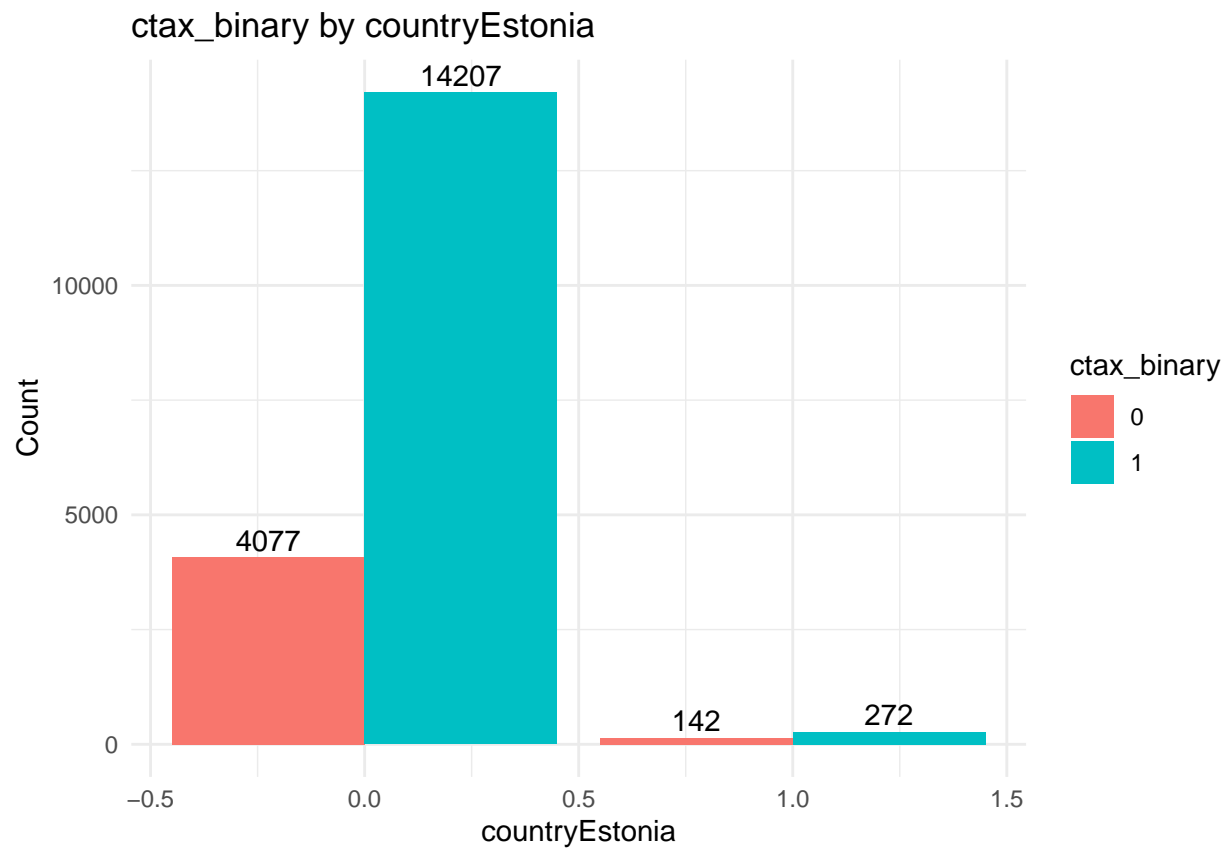


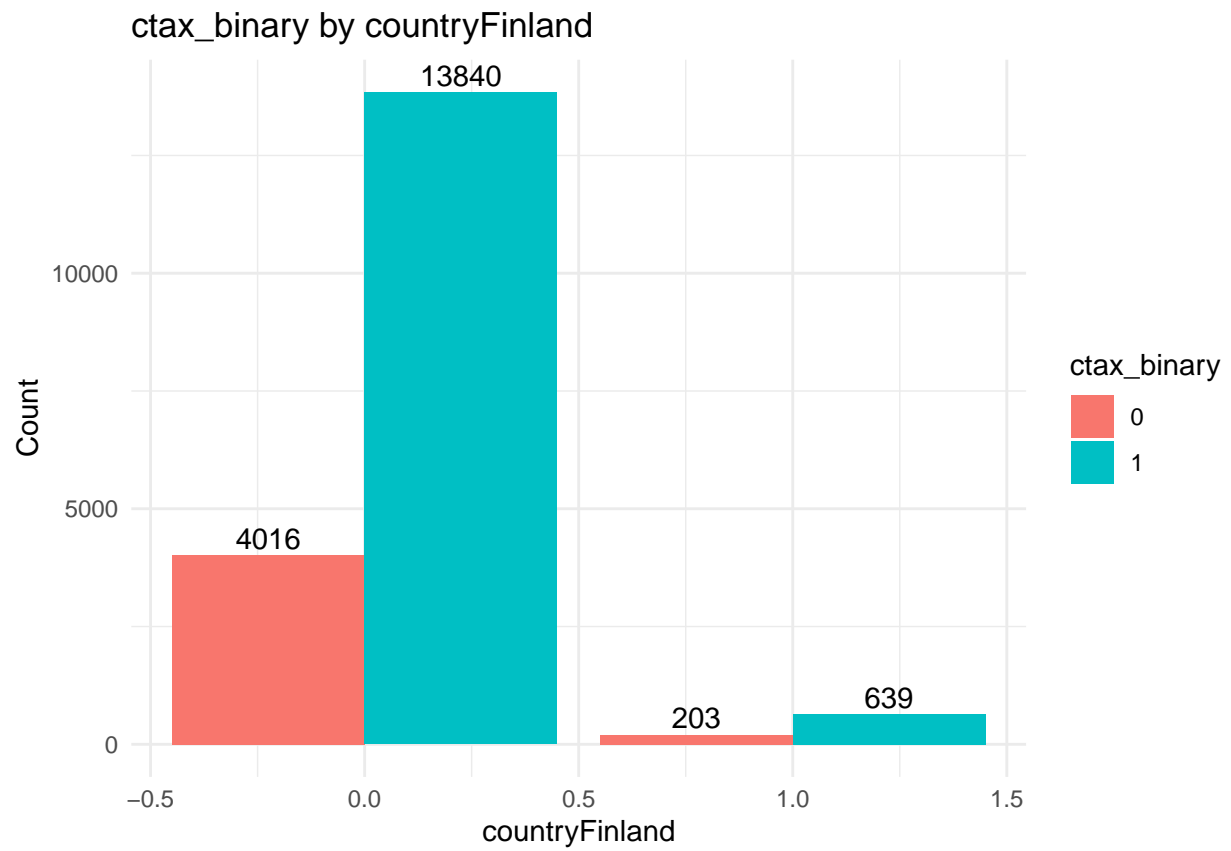


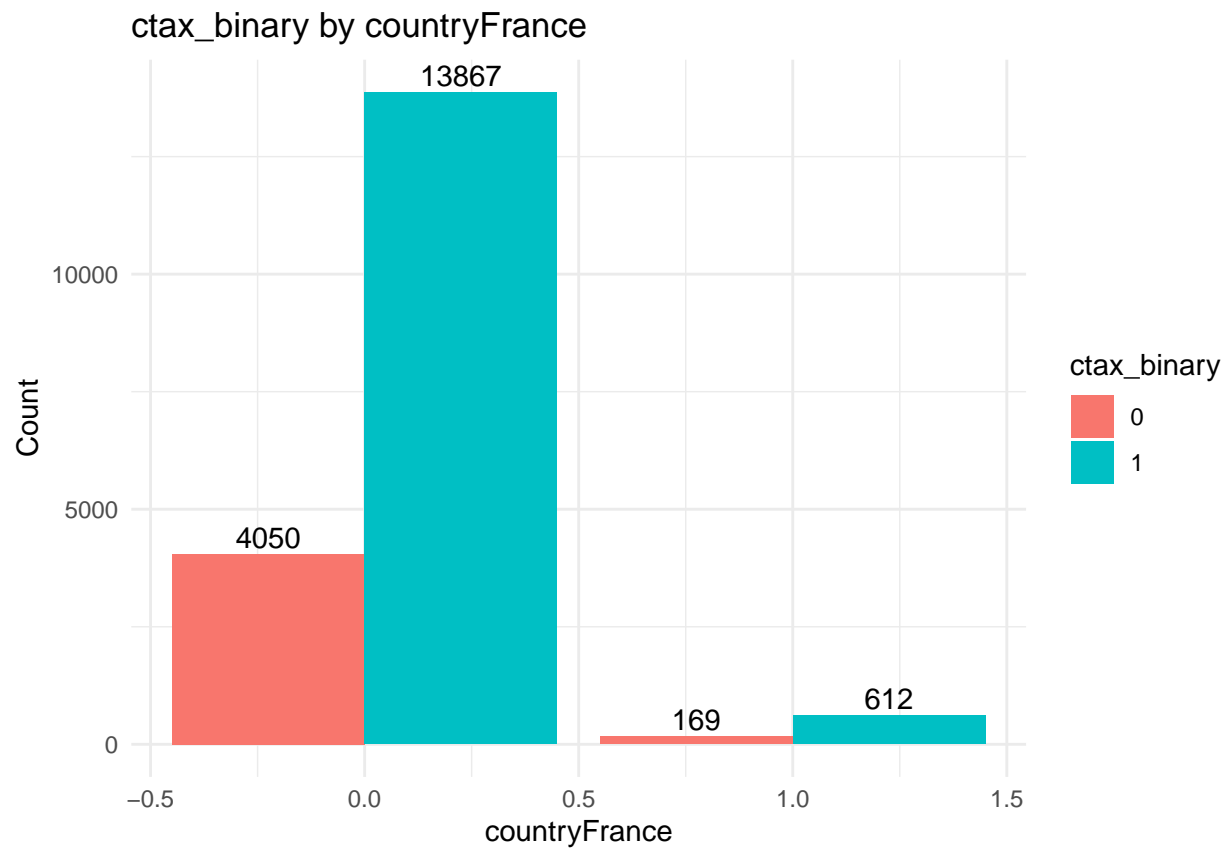


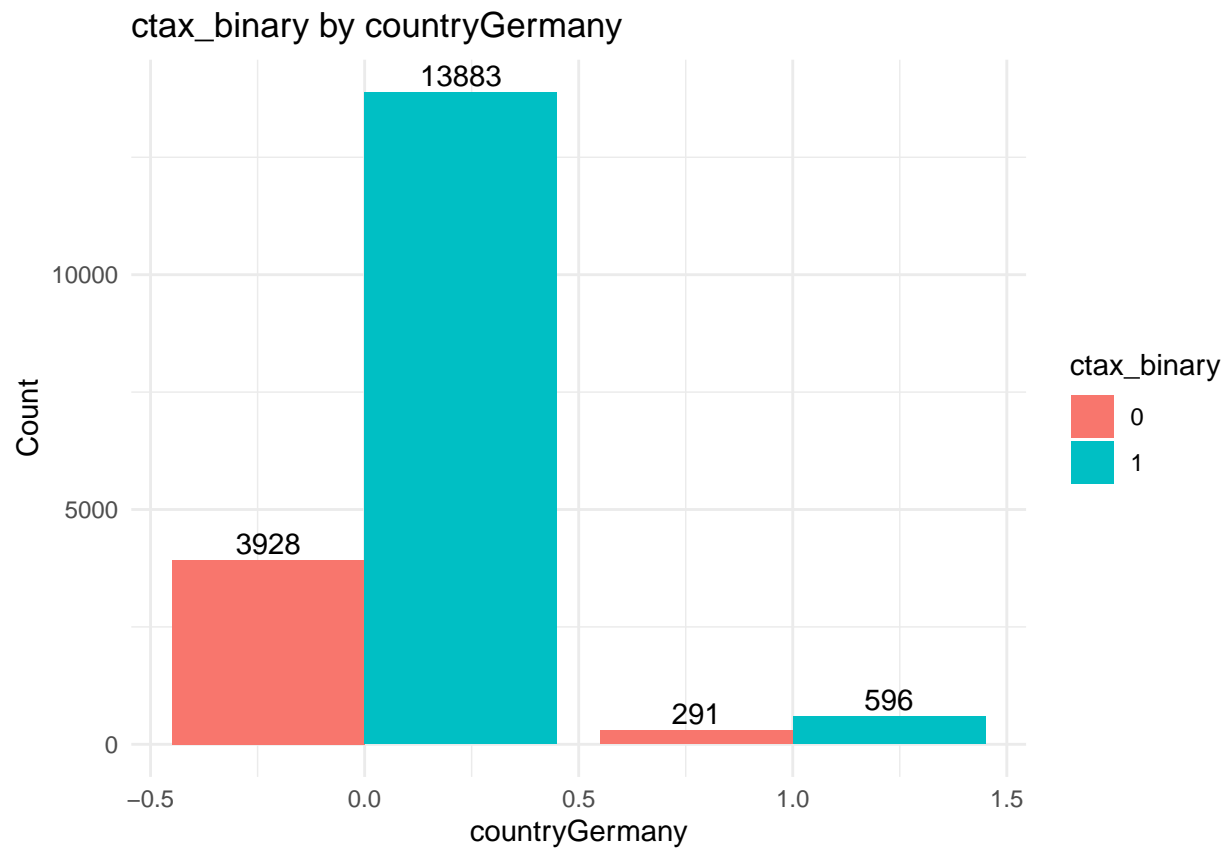


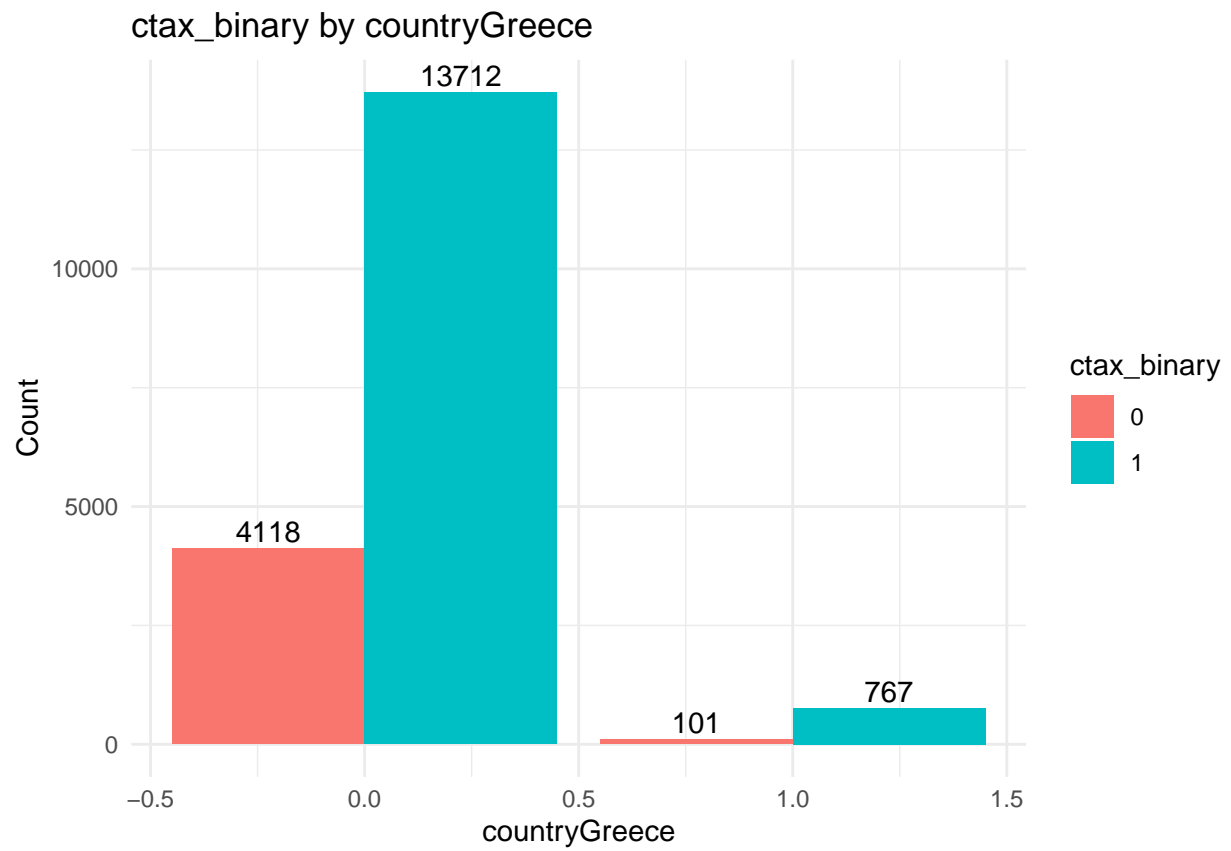


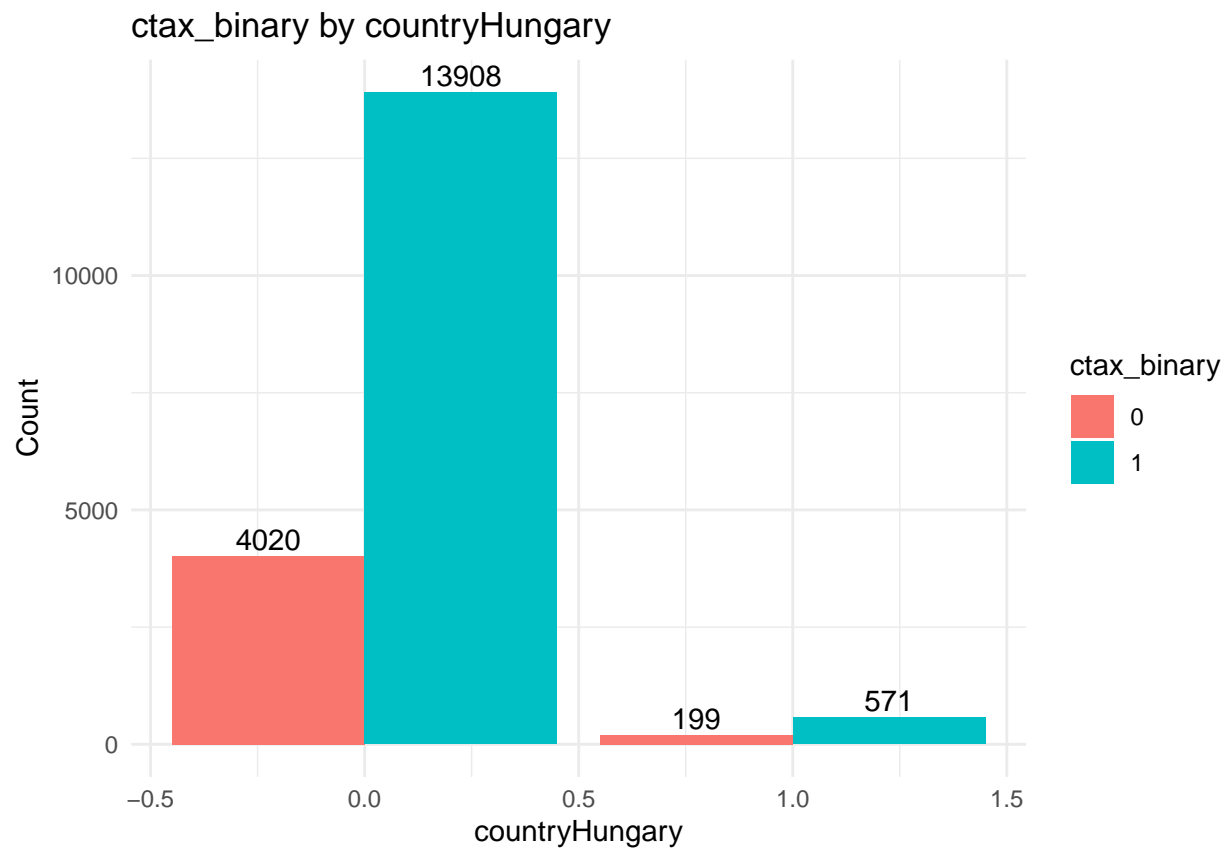


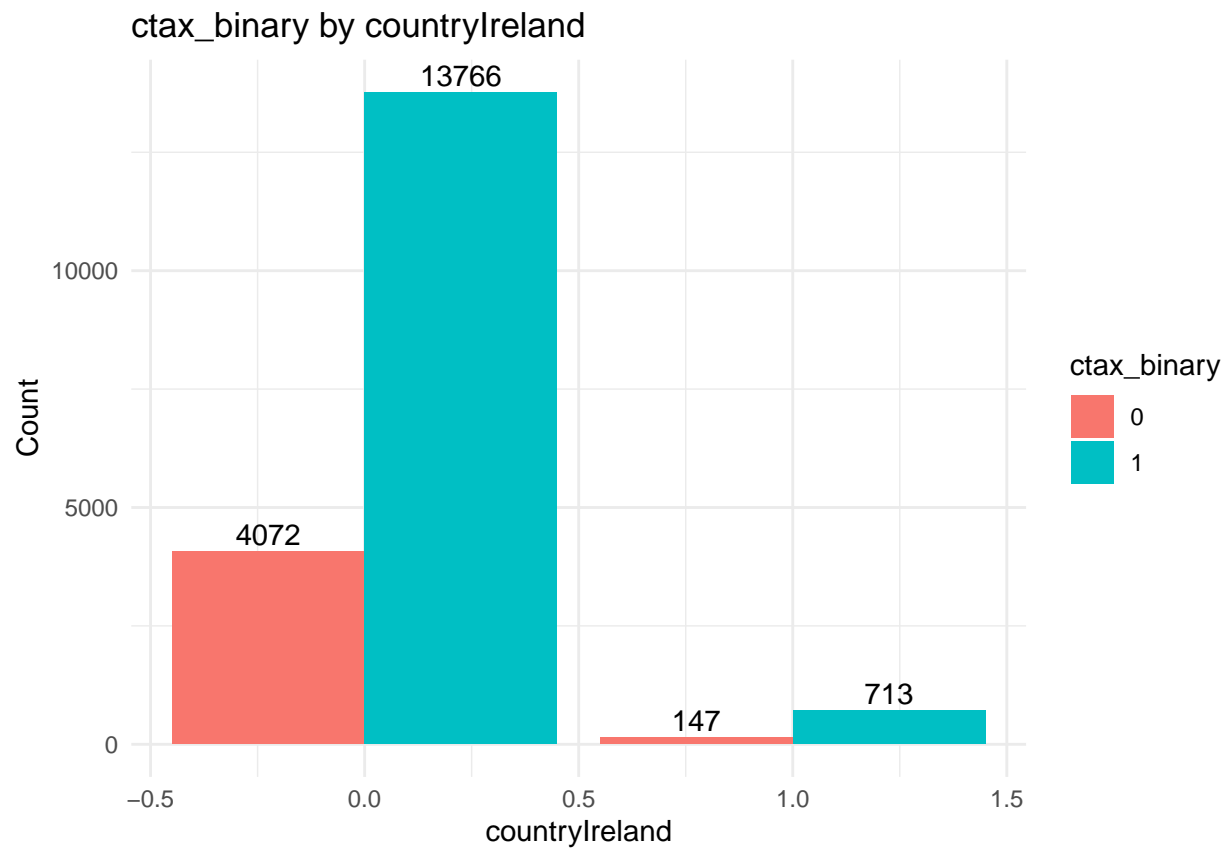


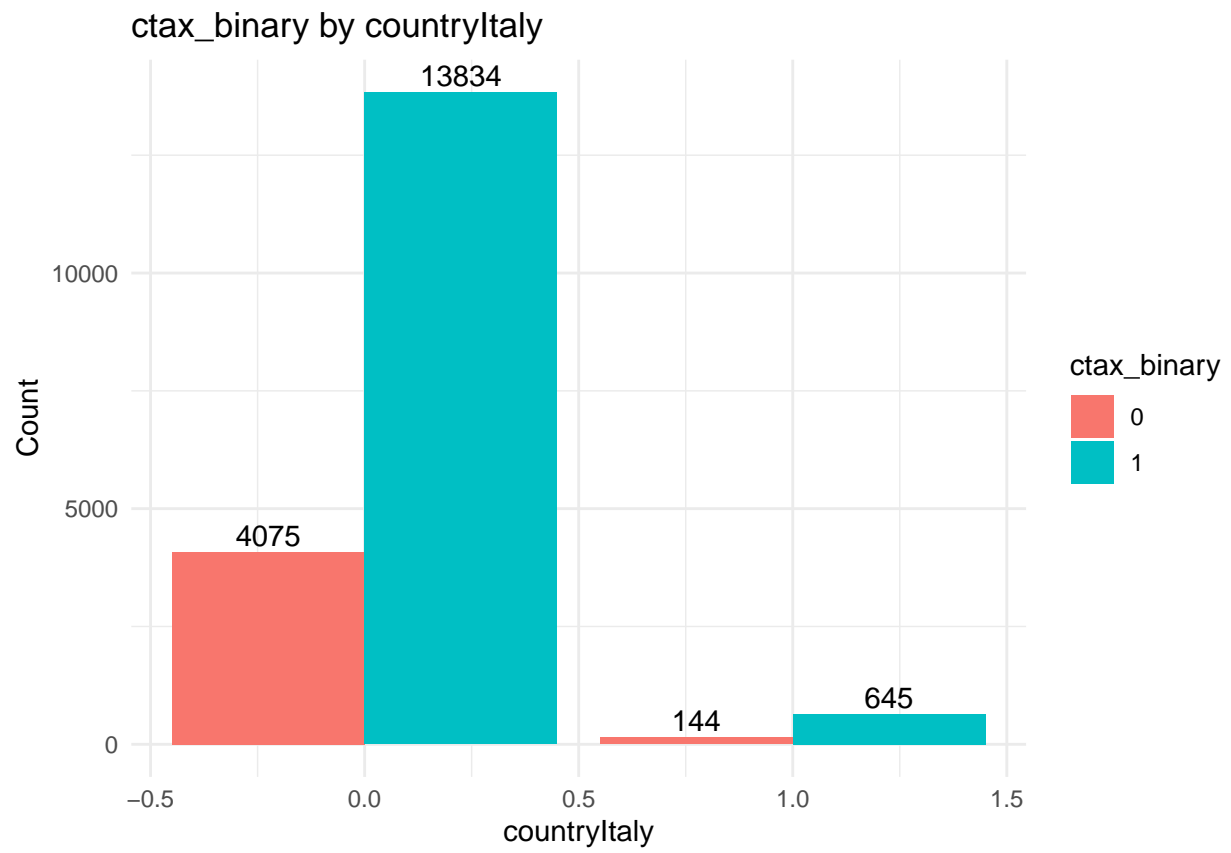


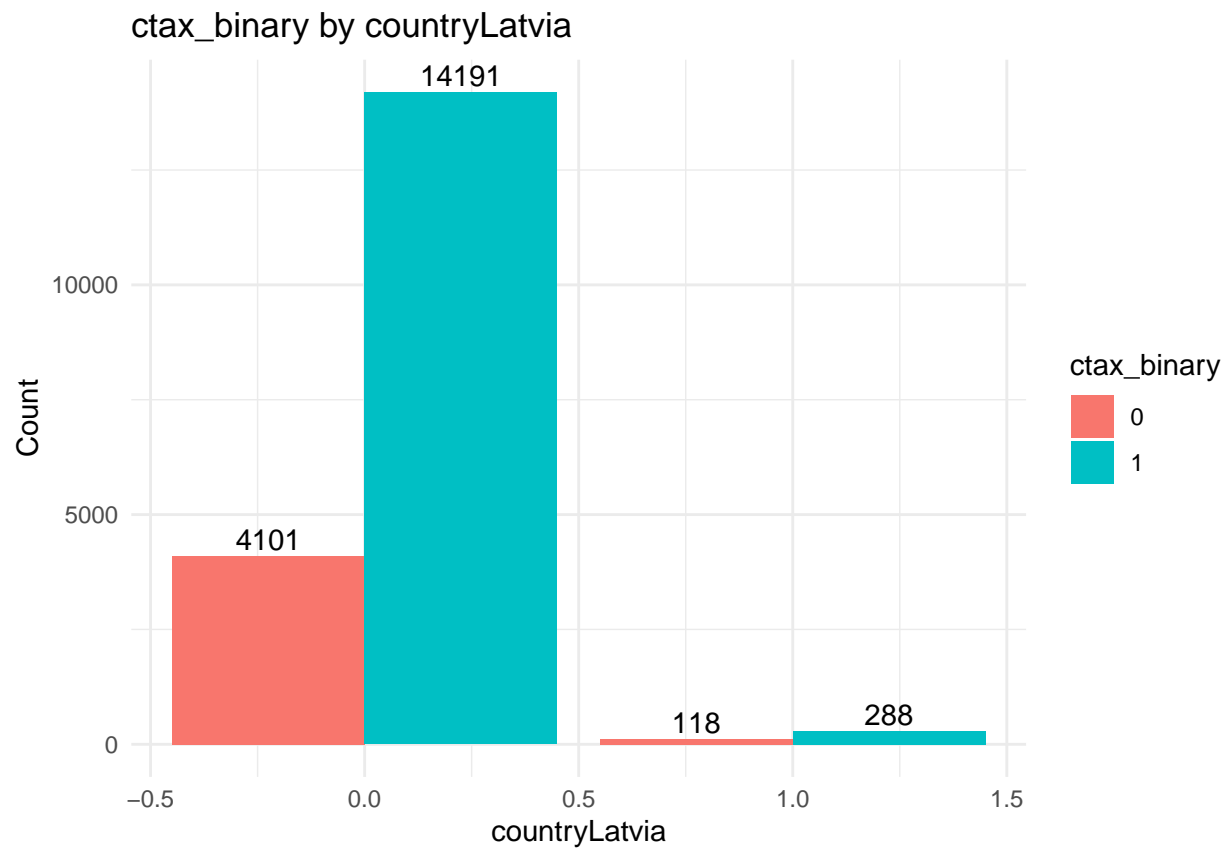


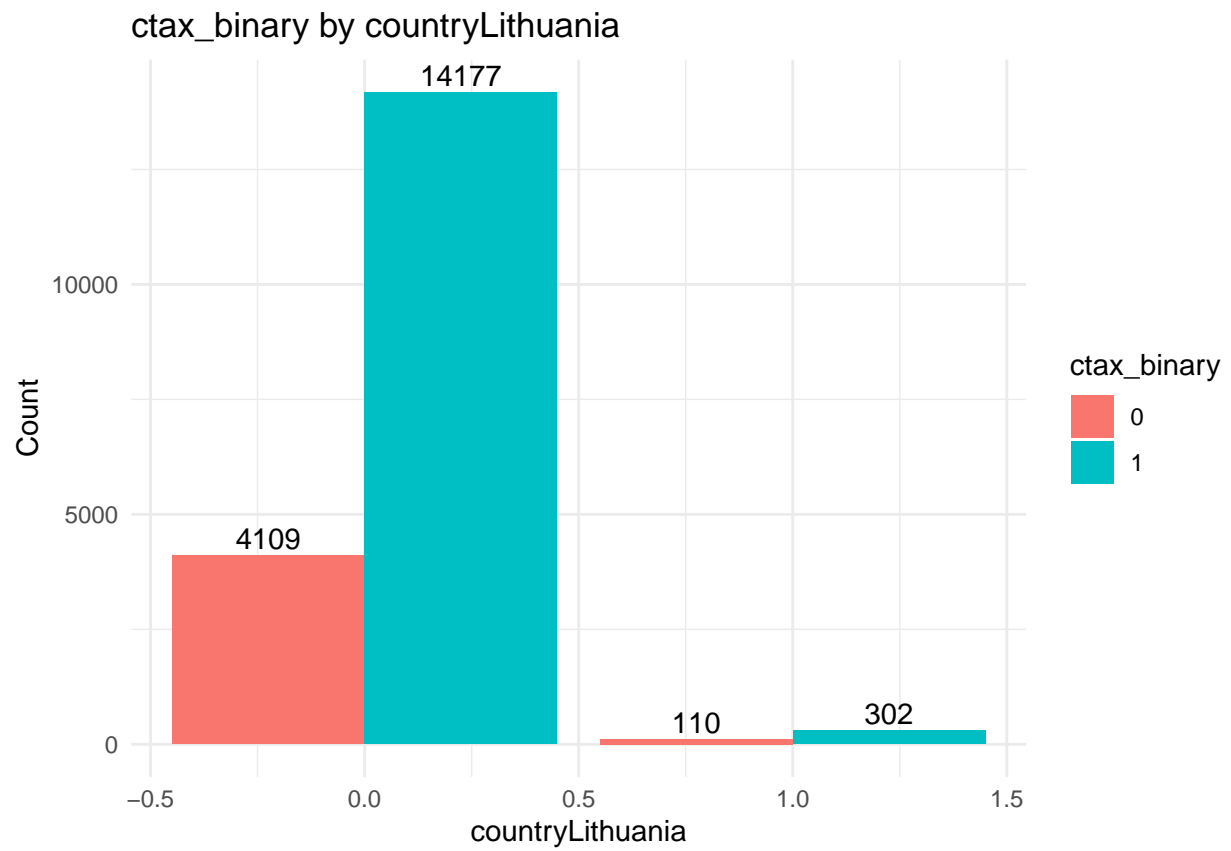


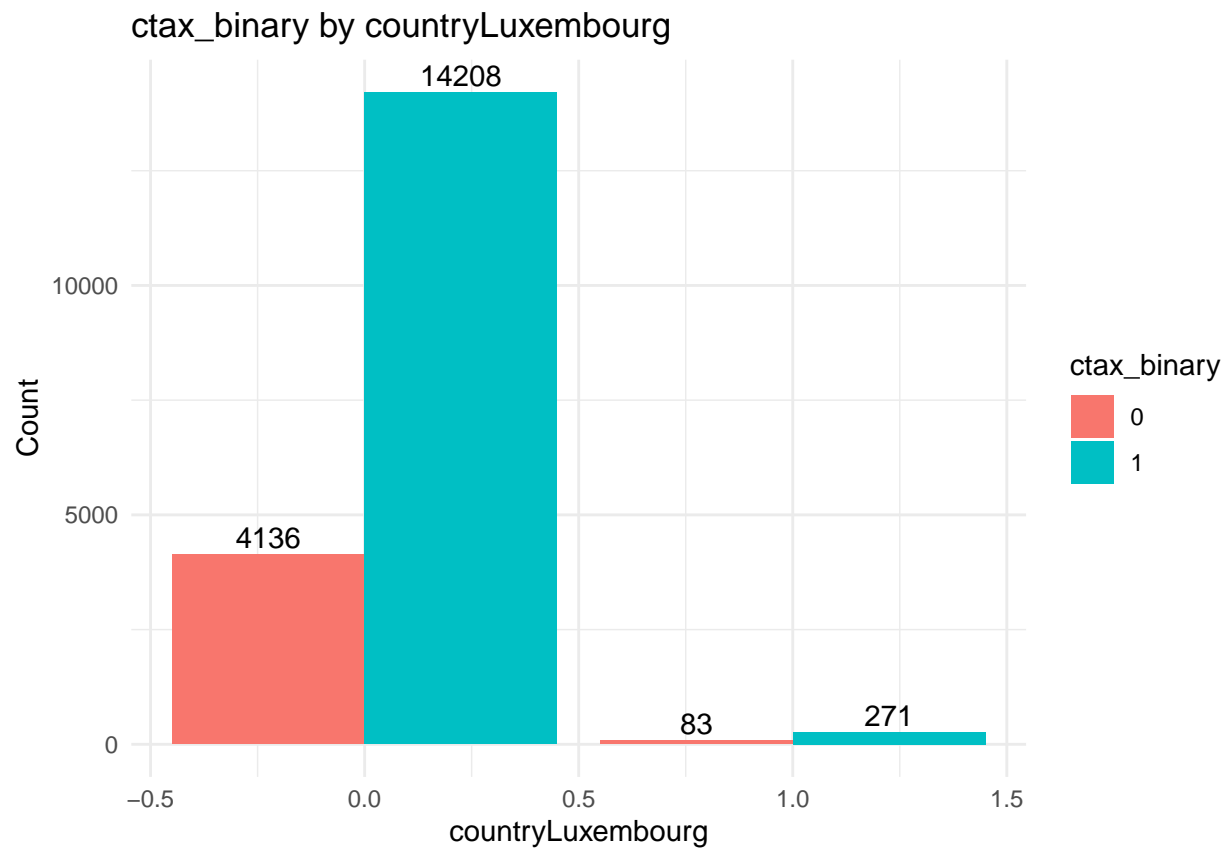


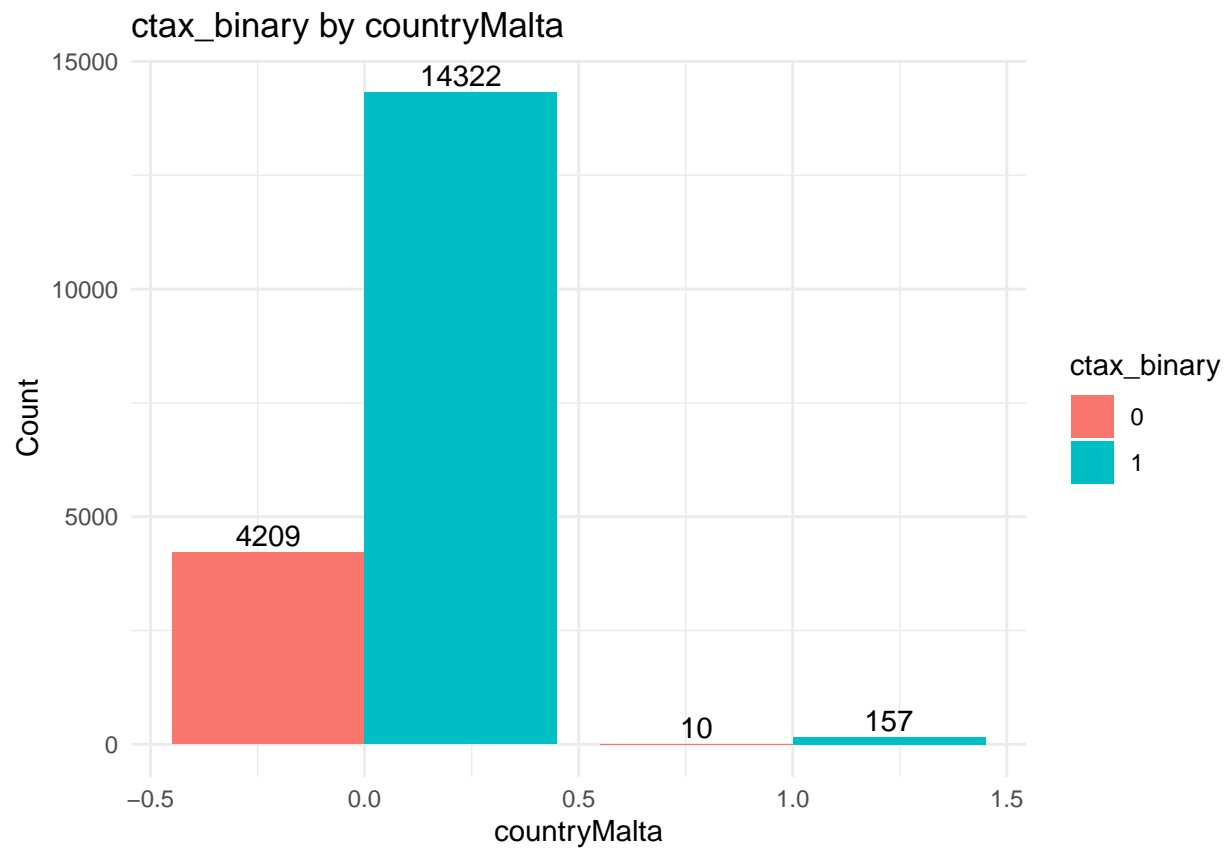


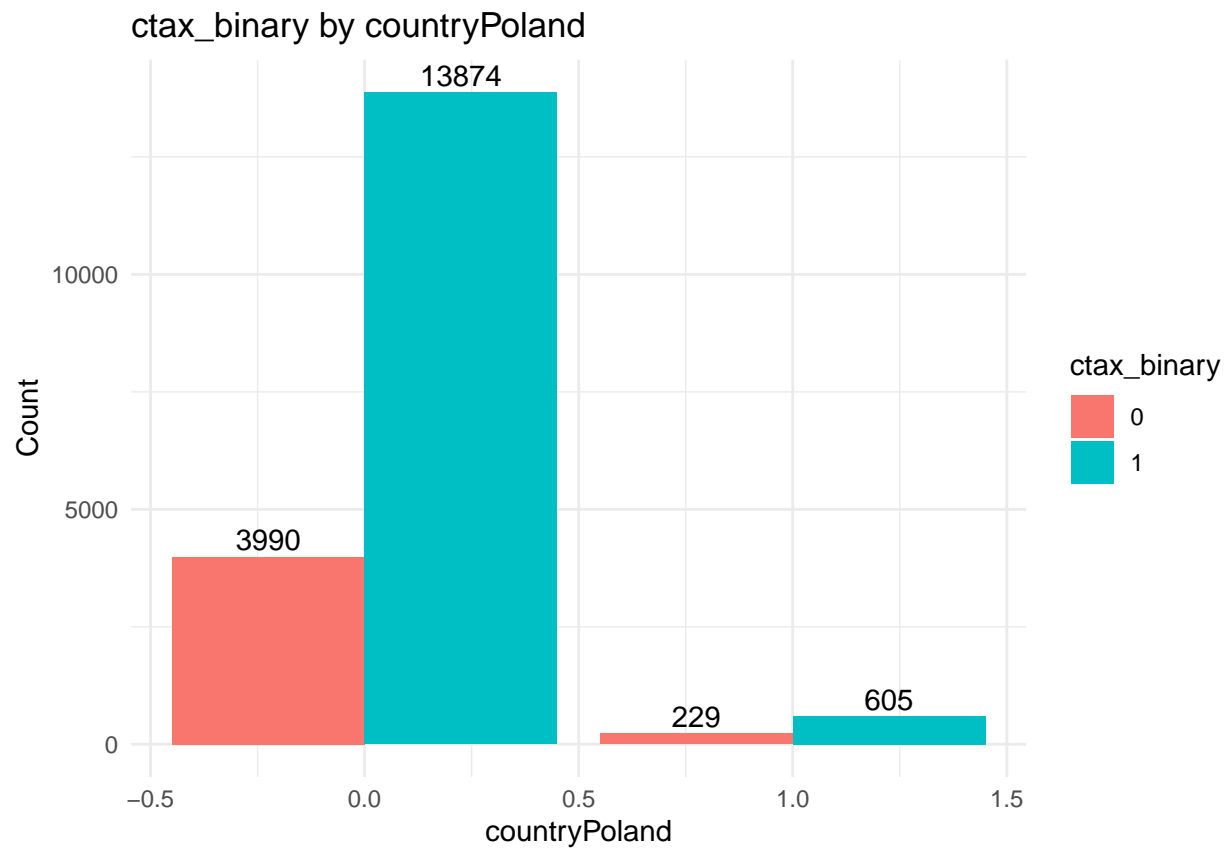


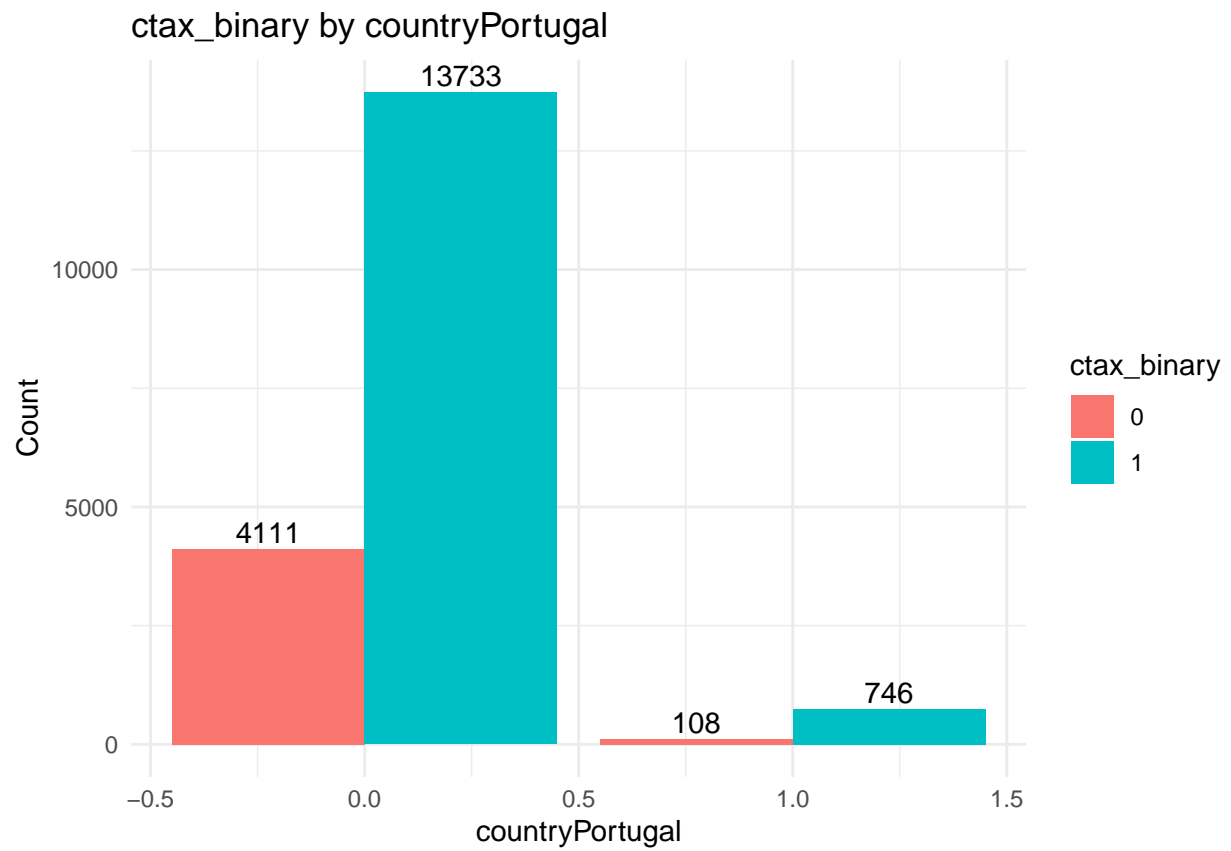


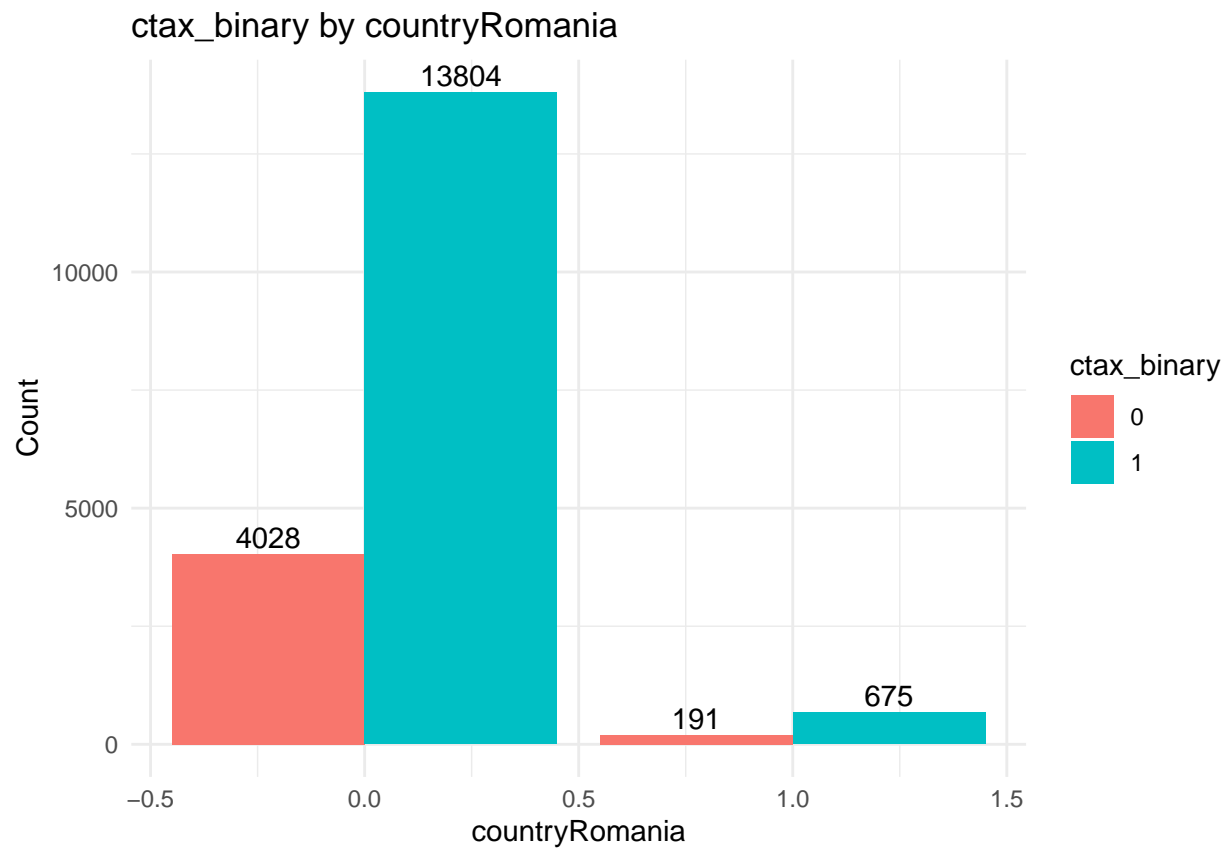


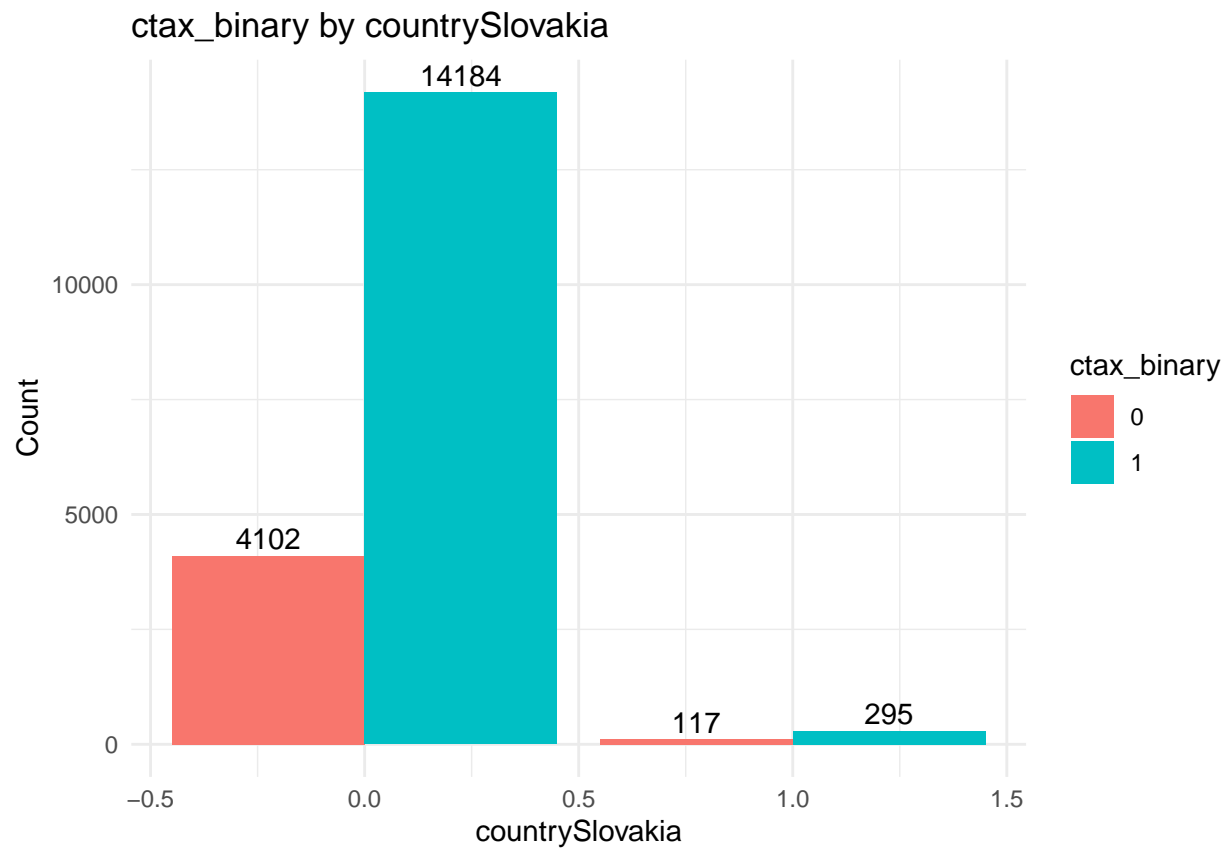


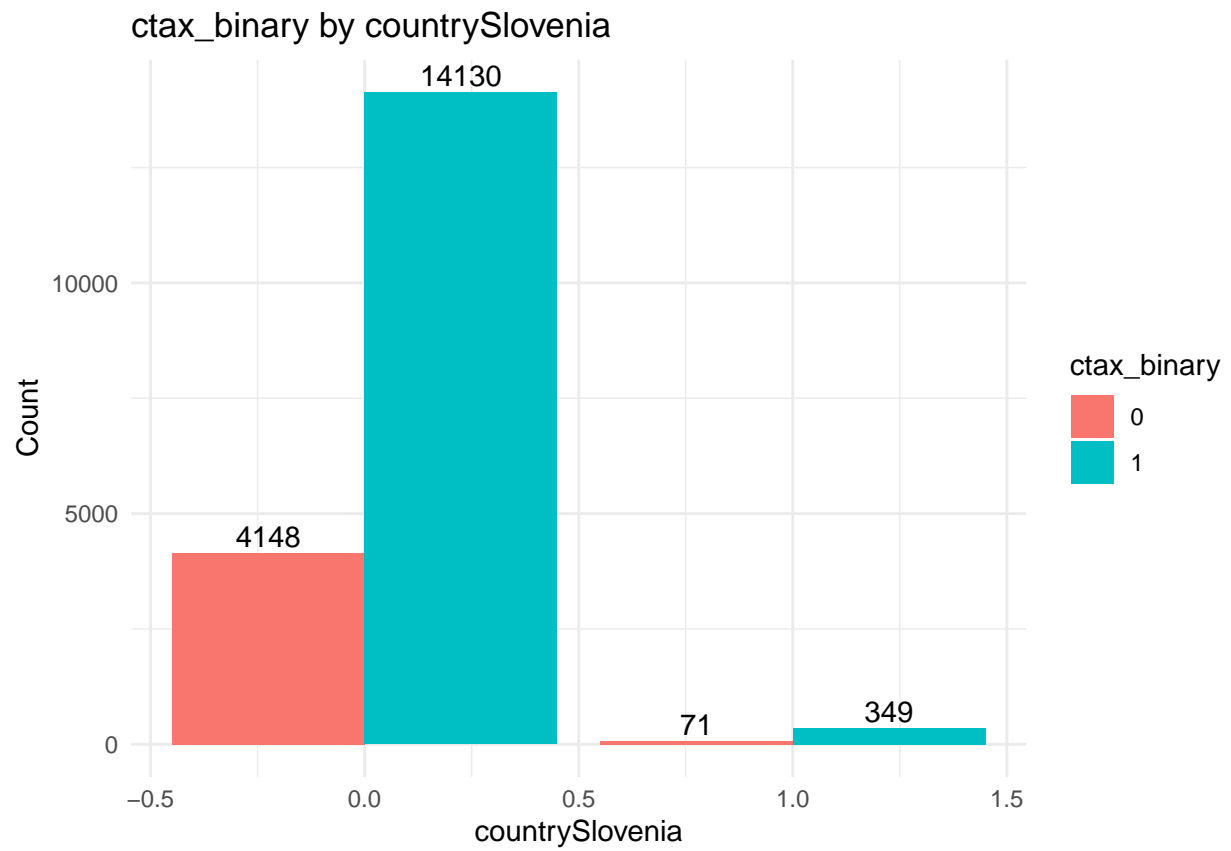


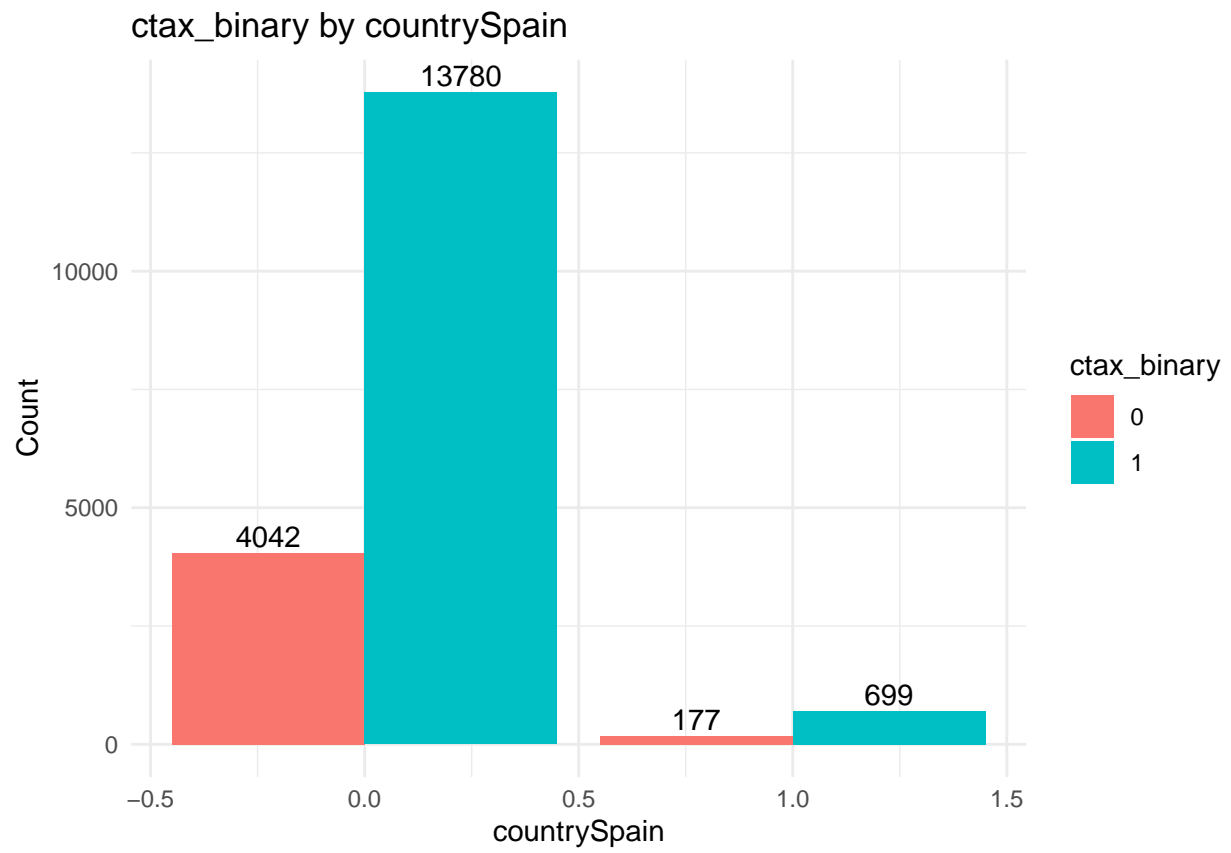


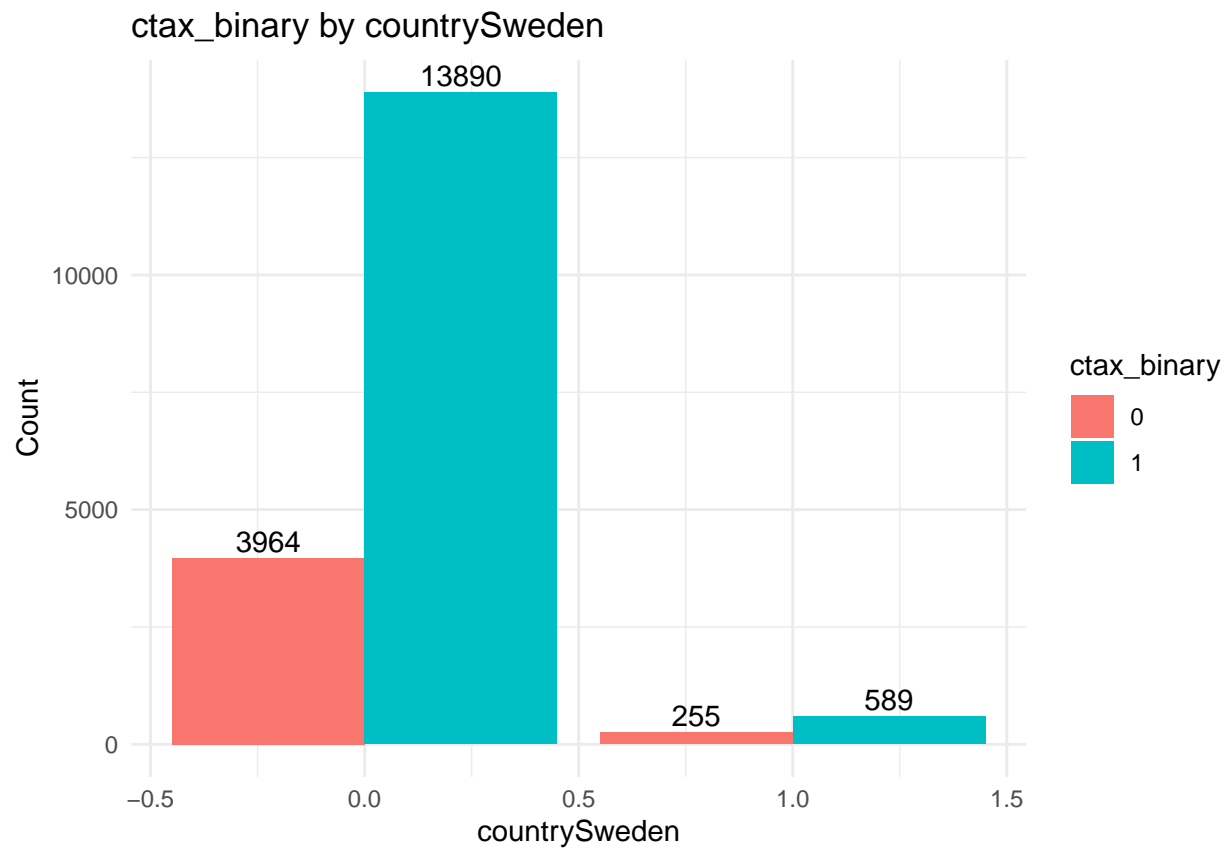


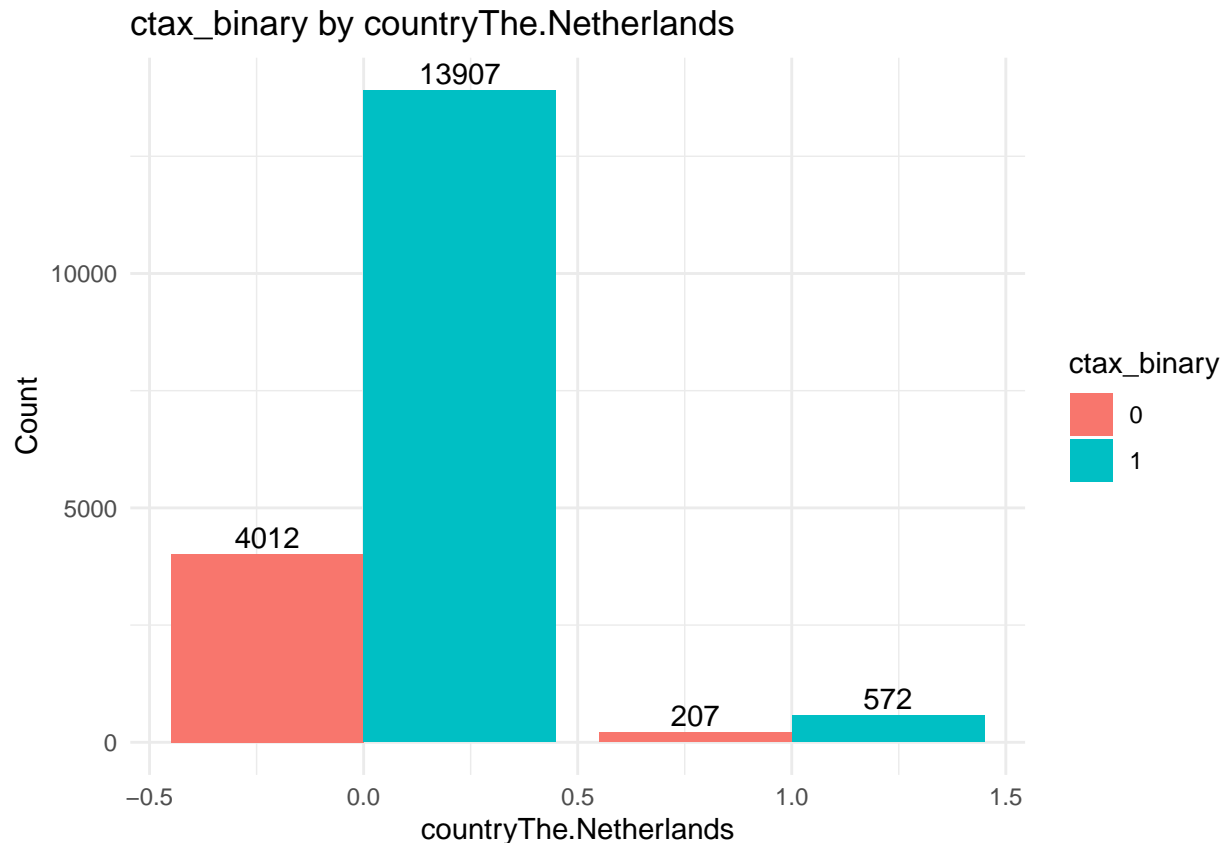












Creating the Education Variable, Residence variable respect to education dummies and residence dummies:

```
df_clean$education <- with(df_clean, ifelse(df_clean$educationprimary == 1, "primary",
                                           ifelse(df_clean$educationtertiary == 1, "tertiary", "secondary")))
df_clean$education <- factor(df_clean$education, levels = c("secondary", "primary", "tertiary"))
df_clean$residence <- with(df_clean, ifelse(urbanizationrural == 1, "rural",
                                           ifelse(urbanizationtown == 1, "town", "city")))
df_clean$residence <- factor(df_clean$residence, levels = c("city", "town", "rural"))
df_clean = df_clean[, -c(13:17)]
```

Creating Validation Sets:

We will create 80/20 Train and Test data set.

```
set.seed(123) # For the reproducibility
train_index <- createDataPartition(df_clean$ctax_binary, p = 0.8, list = FALSE)
train_data <- df_clean[train_index, ]
test_data <- df_clean[-train_index, ]
```

Creating the Baseline Logistic Regression Model:

Create two models : One with and one without any_cc_last2year_factor variable.

Without any_cc_last2year_factor:

```
country_vars <- names(df)[grepl("^country", names(df))]  
country_vars=country_vars[-1]  
train_data$ctax_binary=as.factor(train_data$ctax_binary)
```

```
paste(country_vars, collapse = " + ")
```

```
## [1] "countryBelgium + countryBulgaria + countryCroatia + countryCyprus + countryCzech.Republic + coun
```

```
full_formula=as.formula(ctax_binary~income_scale+trust+LR_scale_scale+new_ccknowledge_index_scale+resid
```

With any_cc_last2year_factor :

```
full_formula_with_cc=as.formula(ctax_binary~any_cc_last2year_factor+income_scale+trust+LR_scale_scale+n
```

```
set.seed(1)  
baseline_model_with_cc=glm(formula = full_formula_with_cc,  
  data = train_data, weights = country_w,family = binomial(link = "logit"),subset = regional_het
```

```
## Warning in eval(family$initialize): non-integer #successes in a binomial glm!
```

```
summary(baseline_model_with_cc)
```

```
##  
## Call:  
## glm(formula = full_formula_with_cc, family = binomial(link = "logit"),  
##      data = train_data, weights = country_w, subset = regional_heterogeneity ==  
##          1)  
##  
## Coefficients: (10 not defined because of singularities)  
##              Estimate Std. Error z value Pr(>|z|)  
## (Intercept)      0.124650   0.115978   1.075  0.28248  
## any_cc_last2year_factor -0.015918   0.054307  -0.293  0.76944  
## income_scale      -0.042643   0.026540  -1.607  0.10811  
## trustNo really confident  0.522837   0.065005   8.043 8.77e-16 ***  
## trustRather confident   1.028047   0.073736  13.942 < 2e-16 ***  
## trustVery confident     1.565422   0.116714  13.412 < 2e-16 ***  
## LR_scale_scale      -0.083681   0.011405  -7.337 2.18e-13 ***  
## new_ccknowledge_index_scale 3.465530   0.162580  21.316 < 2e-16 ***  
## residencetown      -0.057702   0.056009  -1.030  0.30290  
## residencerural     -0.157699   0.068027  -2.318  0.02044 *  
## age_scale           0.001520   0.001599   0.951  0.34185
```

```
## gendermale -0.024335 0.050121 -0.486 0.62730
## educationprimary 0.113683 0.070946 1.602 0.10907
## educationtertiary 0.084285 0.055204 1.527 0.12681
## has_childrenyes 0.037041 0.054953 0.674 0.50028
## countryBelgium 0.368496 0.131468 2.803 0.00506 **
## countryBulgaria 0.963051 0.141415 6.810 9.75e-12 ***
## countryCroatia 1.101087 0.143966 7.648 2.04e-14 ***
## countryCyprus NA NA NA NA
## countryCzech.Republic 0.359325 0.128755 2.791 0.00526 **
## countryDenmark 0.588647 0.136647 4.308 1.65e-05 ***
## countryEstonia NA NA NA NA
## countryFinland NA NA NA NA
## countryFrance 0.568806 0.138433 4.109 3.98e-05 ***
## countryGermany 0.008216 0.124181 0.066 0.94725
## countryGreece 1.347111 0.156113 8.629 < 2e-16 ***
## countryHungary NA NA NA NA
## countryIreland NA NA NA NA
## countryItaly 0.818755 0.140729 5.818 5.96e-09 ***
## countryLatvia NA NA NA NA
## countryLithuania NA NA NA NA
## countryLuxembourg NA NA NA NA
## countryMalta NA NA NA NA
## countryPoland 0.207412 0.128444 1.615 0.10635
## countryPortugal 1.364291 0.152059 8.972 < 2e-16 ***
## countryRomania 0.576754 0.129256 4.462 8.12e-06 ***
## countrySlovakia 0.208981 0.154742 1.351 0.17685
## countrySlovenia NA NA NA NA
## countrySpain 0.626163 0.134919 4.641 3.47e-06 ***
## countrySweden 0.089265 0.126395 0.706 0.48004
## countryThe.Netherlands 0.337780 0.130960 2.579 0.00990 **
## ---
## Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## (Dispersion parameter for binomial family taken to be 1)
##
## Null deviance: 11590 on 10967 degrees of freedom
## Residual deviance: 10487 on 10937 degrees of freedom
## AIC: 10838
##
## Number of Fisher Scoring iterations: 4
```

```
yhat_baseline_train_cc=predict(baseline_model_with_cc,newdata=train_data,type="response")
```

```
## Warning in predict.lm(object, newdata, se.fit, scale = 1, type = if (type == :
## prediction from rank-deficient fit; attr(*, "non-estim") has doubtful cases
```

```
yhat_baseline_test_cc=predict(baseline_model_with_cc,newdata=test_data,type="response")
```

```
## Warning in predict.lm(object, newdata, se.fit, scale = 1, type = if (type == :
## prediction from rank-deficient fit; attr(*, "non-estim") has doubtful cases
```

```
class_pred_train_cc <- ifelse(yhat_baseline_train_cc > 0.5, 1, 0)
class_pred_test_cc <- ifelse(yhat_baseline_test_cc > 0.5, 1, 0)
```

Checking the Accuracy Metrics respect to the Train Data for the Baseline Model:

```
confusionMatrix(as.factor(train_data$ctax_binary),as.factor(class_pred_train_cc))
```

```
## Confusion Matrix and Statistics
##
##              Reference
## Prediction      0      1
##              0  536  2833
##              1  421 11169
##
##              Accuracy : 0.7825
##              95% CI : (0.7758, 0.7891)
##              No Information Rate : 0.936
##              P-Value [Acc > NIR] : 1
##
##              Kappa : 0.1646
##
##              Mcnemar's Test P-Value : <2e-16
##
##              Sensitivity : 0.56008
##              Specificity : 0.79767
##              Pos Pred Value : 0.15910
##              Neg Pred Value : 0.96368
##              Prevalence : 0.06397
##              Detection Rate : 0.03583
##              Detection Prevalence : 0.22522
##              Balanced Accuracy : 0.67888
##
##              'Positive' Class : 0
##
```

Checking the Accuracy Metrics respect to the Test Data for the Baseline Model:

```
confusionMatrix(as.factor(test_data$ctax_binary),as.factor(class_pred_test_cc))
```

```
## Confusion Matrix and Statistics
##
##              Reference
## Prediction      0      1
##              0  144   706
```

```
##          1    91 2798
##
##          Accuracy : 0.7868
##          95% CI : (0.7734, 0.7999)
##    No Information Rate : 0.9371
##    P-Value [Acc > NIR] : 1
##
##          Kappa : 0.1852
##
##    McNemar's Test P-Value : <2e-16
##
##          Sensitivity : 0.61277
##          Specificity : 0.79852
##    Pos Pred Value : 0.16941
##    Neg Pred Value : 0.96850
##          Prevalence : 0.06285
##    Detection Rate : 0.03851
##    Detection Prevalence : 0.22733
##    Balanced Accuracy : 0.70564
##
##    'Positive' Class : 0
##
```

```
set.seed(1)
baseline_model=glm(formula = full_formula,
  data = train_data, weights = country_w,family = binomial(link = "logit"))
```

```
## Warning in eval(family$initialize): non-integer #successes in a binomial glm!
```

```
summary(baseline_model)
```

```
##
## Call:
## glm(formula = full_formula, family = binomial(link = "logit"),
##    data = train_data, weights = country_w)
##
## Coefficients:
##              Estimate Std. Error z value Pr(>|z|)
## (Intercept)    0.111298   0.108305   1.028  0.30412
## income_scale   -0.042261   0.022709  -1.861  0.06275 .
## trustNo really confident  0.542622   0.055923   9.703 < 2e-16 ***
## trustRather confident    1.062523   0.062755  16.931 < 2e-16 ***
## trustVery confident     1.463011   0.096809  15.112 < 2e-16 ***
## LR_scale_scale   -0.078879   0.009956  -7.922 2.33e-15 ***
## new_ccknowledge_index_scale 3.337251   0.138800  24.044 < 2e-16 ***
## residencetown    -0.052793   0.048274  -1.094  0.27413
## residencerural   -0.180987   0.057454  -3.150  0.00163 **
## age_scale        0.004192   0.001354   3.096  0.00196 **
## gendermale      -0.031994   0.042786  -0.748  0.45460
## educationprimary  0.098942   0.062898   1.573  0.11570
## educationtertiary 0.089704   0.047123   1.904  0.05696 .
## has_childrenyes   0.059602   0.046833   1.273  0.20314
## countryBelgium    0.360642   0.131103   2.751  0.00594 **
```

```
## countryBulgaria      0.945159  0.140740  6.716 1.87e-11 ***
## countryCroatia      1.081013  0.140910  7.672 1.70e-14 ***
## countryCyprus        1.175089  0.215932  5.442 5.27e-08 ***
## countryCzech.Republic 0.347182  0.127715  2.718 0.00656 **
## countryDenmark      0.585273  0.135923  4.306 1.66e-05 ***
## countryEstonia     -0.030983  0.150034 -0.207 0.83640
## countryFinland      0.342999  0.129891  2.641 0.00827 **
## countryFrance       0.548620  0.136140  4.030 5.58e-05 ***
## countryGermany     -0.001514  0.121591 -0.012 0.99006
## countryGreece       1.322158  0.152765  8.655 < 2e-16 ***
## countryHungary      0.307740  0.130072  2.366 0.01799 *
## countryIreland      0.958000  0.139828  6.851 7.32e-12 ***
## countryItaly        0.801192  0.139195  5.756 8.62e-09 ***
## countryLatvia       0.205191  0.155111  1.323 0.18588
## countryLithuania    0.159447  0.156838  1.017 0.30933
## countryLuxembourg   0.465739  0.175129  2.659 0.00783 **
## countryMalta        2.225869  0.424701  5.241 1.60e-07 ***
## countryPoland       0.197634  0.127959  1.545 0.12247
## countryPortugal     1.346586  0.151540  8.886 < 2e-16 ***
## countryRomania      0.569869  0.128763  4.426 9.61e-06 ***
## countrySlovakia     0.205084  0.154375  1.328 0.18402
## countrySlovenia     0.903655  0.174709  5.172 2.31e-07 ***
## countrySpain        0.602861  0.132552  4.548 5.41e-06 ***
## countrySweden       0.081917  0.125238  0.654 0.51306
## countryThe.Netherlands 0.325172  0.130432  2.493 0.01267 *
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## (Dispersion parameter for binomial family taken to be 1)
##
## Null deviance: 15803  on 14958  degrees of freedom
## Residual deviance: 14314  on 14919  degrees of freedom
## AIC: 14923
##
## Number of Fisher Scoring iterations: 5
```

```
yhat_baseline_train=predict(baseline_model,newdata=train_data,type="response")
yhat_baseline_test=predict(baseline_model,newdata=test_data,type="response")

class_pred_train <- ifelse(yhat_baseline_train > 0.5, 1, 0)
class_pred_test  <- ifelse(yhat_baseline_test > 0.5, 1, 0)
```

Checking the Accuracy Metrics respect to the Train Data for the Baseline Model:

```
confusionMatrix(as.factor(train_data$ctax_binary),as.factor(class_pred_train))

## Confusion Matrix and Statistics
##
## Reference
```

```

## Prediction      0      1
##              0  432  2937
##              1   259 11331
##
##              Accuracy : 0.7863
##              95% CI : (0.7797, 0.7929)
##      No Information Rate : 0.9538
##      P-Value [Acc > NIR] : 1
##
##              Kappa : 0.1474
##
##      McNemar's Test P-Value : <2e-16
##
##              Sensitivity : 0.62518
##              Specificity : 0.79415
##      Pos Pred Value : 0.12823
##      Neg Pred Value : 0.97765
##              Prevalence : 0.04619
##      Detection Rate : 0.02888
##      Detection Prevalence : 0.22522
##      Balanced Accuracy : 0.70967
##
##      'Positive' Class : 0
##

```

Checking the Accuracy Metrics respect to the Test Data for the Baseline Model:

```
confusionMatrix(as.factor(test_data$ctax_binary),as.factor(class_pred_test))
```

```

## Confusion Matrix and Statistics
##
##              Reference
## Prediction      0      1
##              0  117  733
##              1   53 2836
##
##              Accuracy : 0.7898
##              95% CI : (0.7764, 0.8027)
##      No Information Rate : 0.9545
##      P-Value [Acc > NIR] : 1
##
##              Kappa : 0.1662
##
##      McNemar's Test P-Value : <2e-16
##
##              Sensitivity : 0.68824
##              Specificity : 0.79462
##      Pos Pred Value : 0.13765
##      Neg Pred Value : 0.98165

```

```
##           Prevalence : 0.04547
##           Detection Rate : 0.03129
##           Detection Prevalence : 0.22733
##           Balanced Accuracy : 0.74143
##
##           'Positive' Class : 0
##
```

Creating the ML Models:

XGBoosting Model:

Creating the Grid Search for the XGBoosting Model:

```
xgb_grid <- expand.grid(
  nrounds = c(50,100,200),
  max_depth = c(500,700,1000),
  eta = c(0.01),
  gamma = c(0,1),
  colsample_bytree = 1,
  min_child_weight = 1,
  subsample = 1
)
```

Creating the model without cc__

Training the Model respect to Tuning Parameters

```
set.seed(123)
xgb_model <- train(
  full_formula,
  data = train_data,
  method = "xgbTree",
  trControl = trainControl(method = "cv", number = 5),
  tuneGrid = xgb_grid,
  metric = "Accuracy"
)
```

```
## [13:06:31] WARNING: src/c_api/c_api.cc:935: 'ntree_limit' is deprecated, use 'iteration_range' instead
## [13:06:31] WARNING: src/c_api/c_api.cc:935: 'ntree_limit' is deprecated, use 'iteration_range' instead
## [13:06:47] WARNING: src/c_api/c_api.cc:935: 'ntree_limit' is deprecated, use 'iteration_range' instead
## [13:06:47] WARNING: src/c_api/c_api.cc:935: 'ntree_limit' is deprecated, use 'iteration_range' instead
## [13:07:06] WARNING: src/c_api/c_api.cc:935: 'ntree_limit' is deprecated, use 'iteration_range' instead
## [13:07:06] WARNING: src/c_api/c_api.cc:935: 'ntree_limit' is deprecated, use 'iteration_range' instead
## [13:07:24] WARNING: src/c_api/c_api.cc:935: 'ntree_limit' is deprecated, use 'iteration_range' instead
## [13:07:24] WARNING: src/c_api/c_api.cc:935: 'ntree_limit' is deprecated, use 'iteration_range' instead
## [13:07:43] WARNING: src/c_api/c_api.cc:935: 'ntree_limit' is deprecated, use 'iteration_range' instead
## [13:07:43] WARNING: src/c_api/c_api.cc:935: 'ntree_limit' is deprecated, use 'iteration_range' instead
## [13:08:01] WARNING: src/c_api/c_api.cc:935: 'ntree_limit' is deprecated, use 'iteration_range' instead
## [13:08:01] WARNING: src/c_api/c_api.cc:935: 'ntree_limit' is deprecated, use 'iteration_range' instead
```



```

## xgboost::xgb.train(params = list(eta = param$eta, max_depth = param$max_depth,
##   gamma = param$gamma, colsample_bytree = param$colsample_bytree,
##   min_child_weight = param$min_child_weight, subsample = param$subsample),
##   data = x, nrounds = param$nrounds, objective = "binary:logistic")
## params (as set within xgb.train):
##   eta = "0.01", max_depth = "500", gamma = "1", colsample_bytree = "1", min_child_weight = "1", subsample = "1"
## xgb.attributes:
##   niter
## callbacks:
##   cb.print.evaluation(period = print_every_n)
## # of features: 39
## niter: 200
## nfeatures : 39
## xNames : income_scale trustNo really confident trustRather confident trustVery confident LR_scale_score
## problemType : Classification
## tuneValue :
##   nrounds max_depth eta gamma colsample_bytree min_child_weight subsample
## 6      200      500 0.01      1              1              1              1
## obsLevels : 0 1
## param :
##   list()

```

Creating The Accuracy Metrics for the XgBoost Classification model for the Train data:

```

yhat_xgboost_train<- predict(xgb_model, newdata = train_data, type = "prob")[, "1"]
yhat_xgboost_train_binary=ifelse(yhat_xgboost_train > 0.5, 1, 0)
confusionMatrix(as.factor(train_data$ctax_binary),as.factor(yhat_xgboost_train_binary))

```

```

## Confusion Matrix and Statistics
##
##               Reference
## Prediction      0      1
##      0  2448   921
##      1    29 11561
##
##               Accuracy : 0.9365
##               95% CI : (0.9325, 0.9403)
##      No Information Rate : 0.8344
##      P-Value [Acc > NIR] : < 2.2e-16
##
##               Kappa : 0.7992
##
##      McNemar's Test P-Value : < 2.2e-16
##
##               Sensitivity : 0.9883
##               Specificity : 0.9262
##               Pos Pred Value : 0.7266
##               Neg Pred Value : 0.9975
##               Prevalence : 0.1656
##               Detection Rate : 0.1636
##      Detection Prevalence : 0.2252
##               Balanced Accuracy : 0.9573

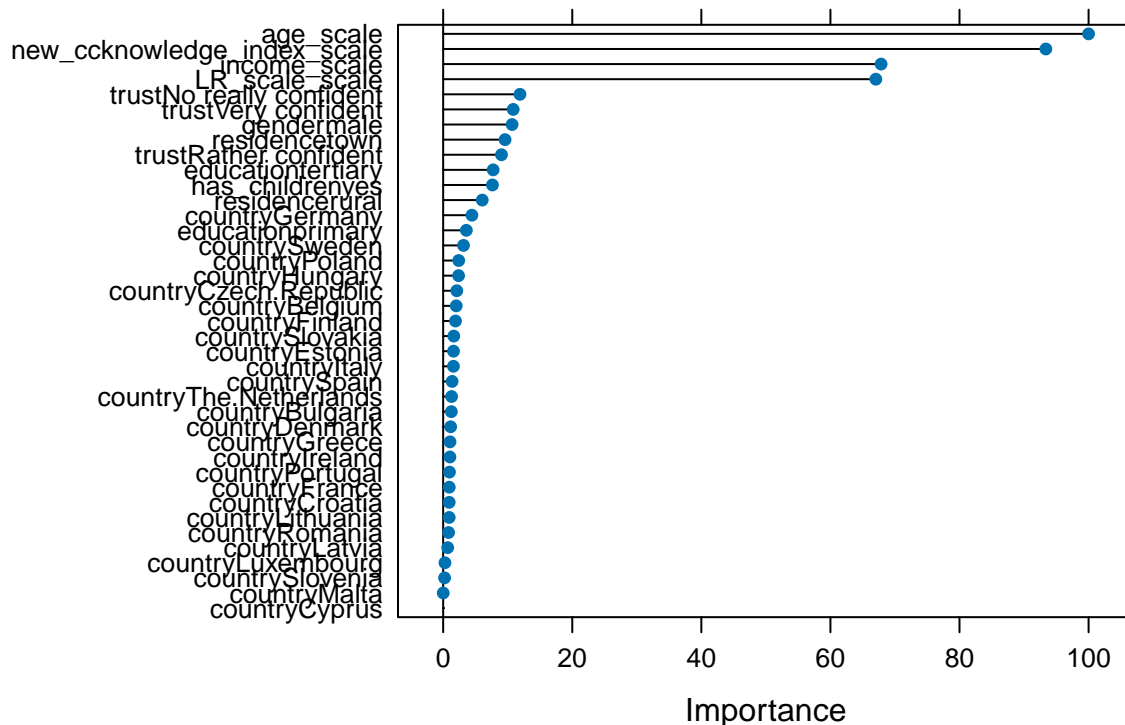
```

```
##
##      'Positive' Class : 0
##
```

Creating Variable Importance to see which features are important for our XGBoosting model:

```
xgb_varimp <- varImp(xgb_model)
plot(xgb_varimp, main = "Variable Importance - XGBoost")
```

Variable Importance – XGBoost



Creating The Accuracy Metrics for the XgBoost Classification model for the Test data:

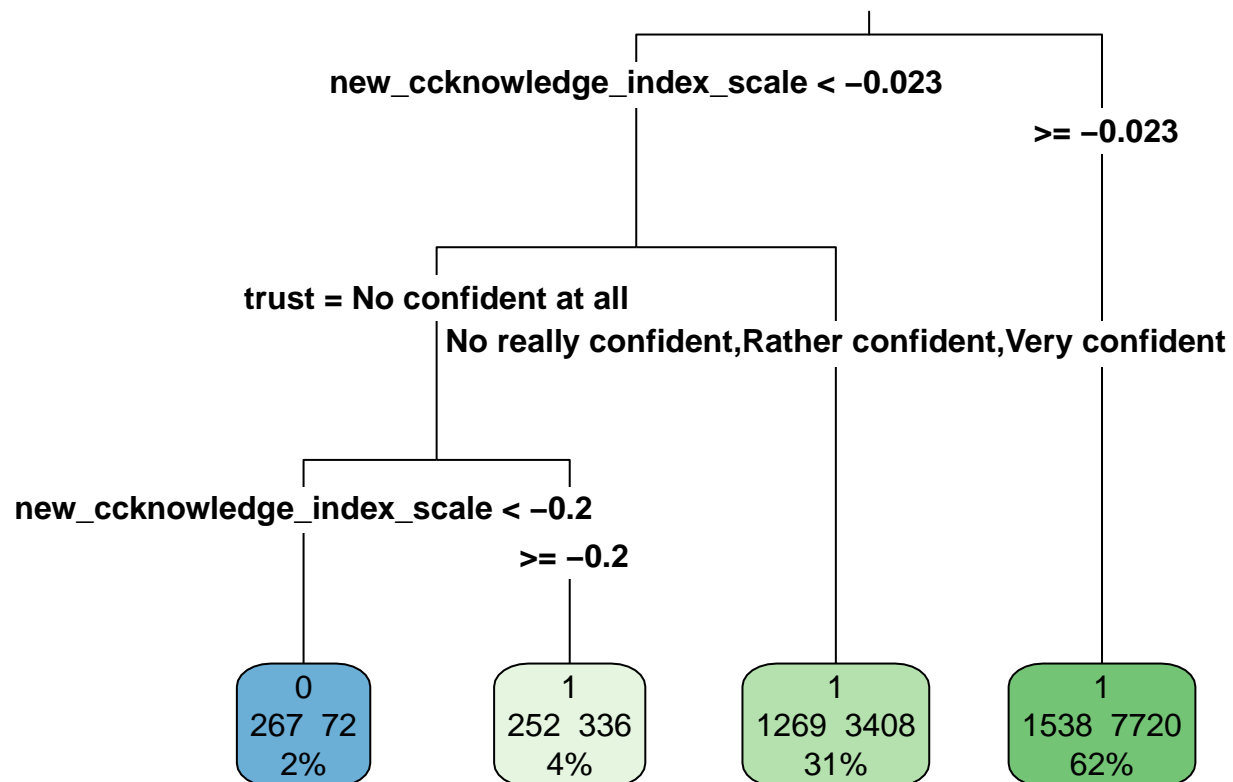
```
yhat_xgboost_test<- predict(xgb_model, newdata = test_data, type = "prob")[, "1"]
yhat_xgboost_test_binary=ifelse(yhat_xgboost_test > 0.5, 1, 0)
confusionMatrix(as.factor(test_data$ctax_binary),as.factor(yhat_xgboost_test_binary))
```

```
## Confusion Matrix and Statistics
##
##      Reference
## Prediction  0    1
##      0  160  690
##      1  156 2733
```

```
##
##           Accuracy : 0.7737
##           95% CI   : (0.76, 0.7871)
##    No Information Rate : 0.9155
##    P-Value [Acc > NIR] : 1
##
##           Kappa : 0.1725
##
##    McNemar's Test P-Value : <2e-16
##
##           Sensitivity : 0.50633
##           Specificity : 0.79842
##           Pos Pred Value : 0.18824
##           Neg Pred Value : 0.94600
##           Prevalence : 0.08451
##           Detection Rate : 0.04279
##           Detection Prevalence : 0.22733
##           Balanced Accuracy : 0.65238
##
##           'Positive' Class : 0
##
```

Classification Tree Model:

```
set.seed(123)
tree_model <- rpart(formula = full_formula, data = train_data, method = "class", weights = country_w )
rpart.plot(tree_model, type = 3, extra = 101, fallen.leaves = TRUE)
```

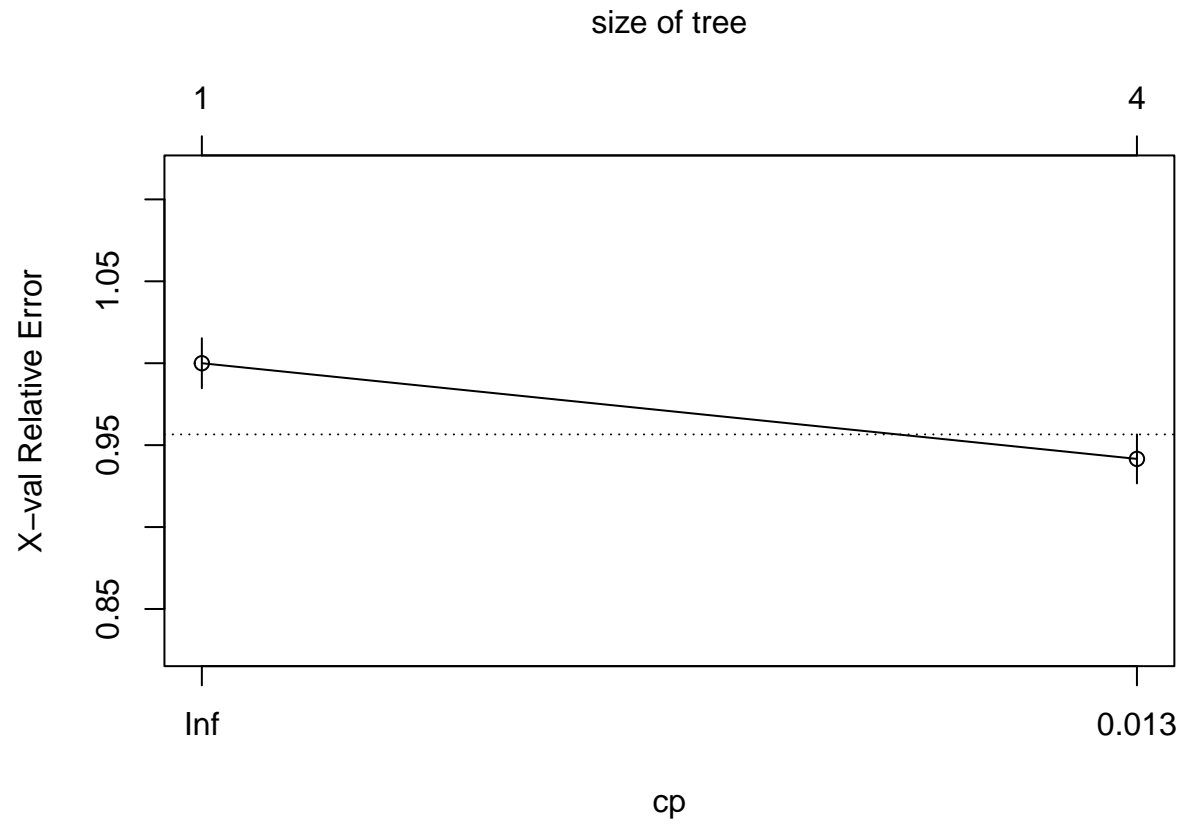


Pruning the classification tree:

```
printcp(tree_model)
```

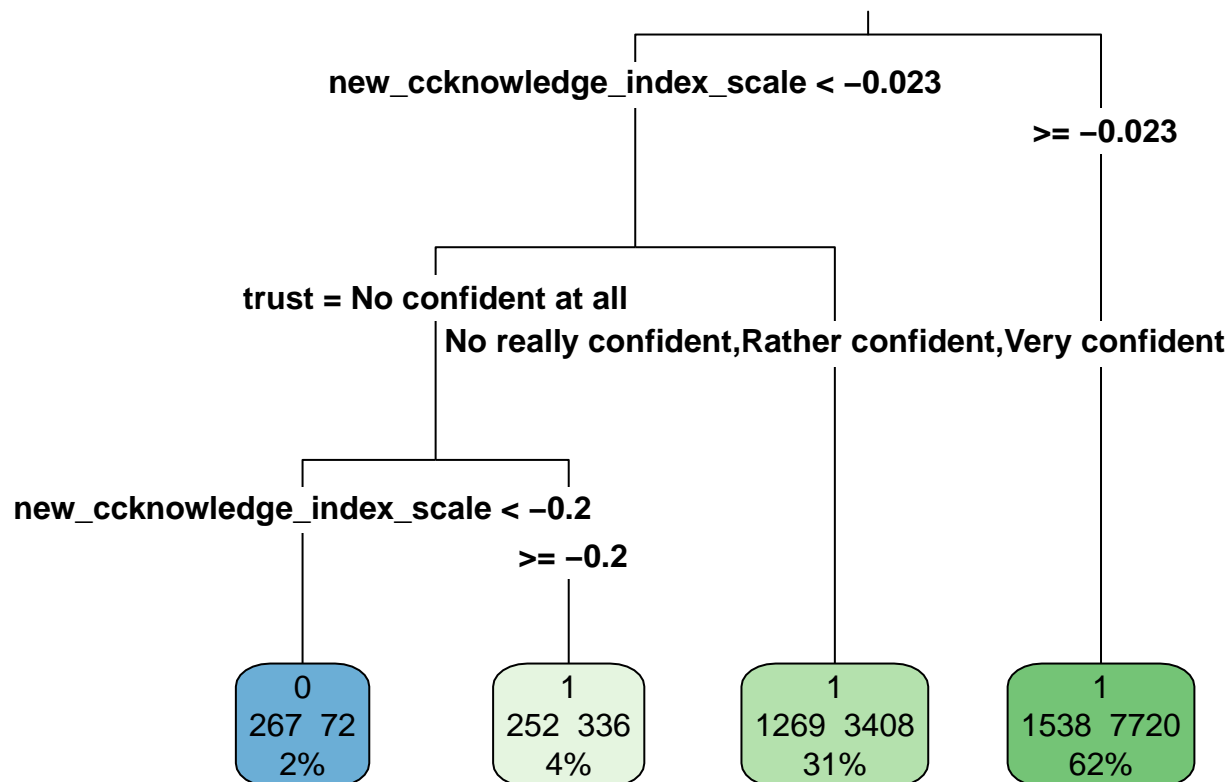
```
##
## Classification tree:
## rpart(formula = full_formula, data = train_data, weights = country_w,
##       method = "class")
##
## Variables actually used in tree construction:
## [1] new_ccknowledge_index_scale trust
##
## Root node error: 3325.8/14959 = 0.22233
##
## n= 14959
##
##      CP nsplit rel error  xerror    xstd
## 1 0.016588      0  1.00000 1.00000 0.015277
## 2 0.010000      3  0.94161 0.94161 0.014949
```

```
plotcp(tree_model)
```



```
best_cp <- tree_model$cptable[which.min(tree_model$cptable[, "xerror"]), "CP"] #Finding the best complex
```

```
pruned_tree <- prune(tree_model, cp = best_cp)  
rpart.plot(pruned_tree, type = 3, extra = 101, fallen.leaves = TRUE)
```



Grid Search for hyperparameter tuning in classification tree:

```

minsplit_vals <- c(10, 20, 30)
maxdepth_vals <- c(2, 4, 6)

best_model <- NULL
best_error <- Inf

for (minsplit in minsplit_vals) {
  for (maxdepth in maxdepth_vals) {
    model <- rpart(formula=full_formula, data = train_data, method = "class", weights=country_w,
                    control = rpart.control(minsplit = minsplit, maxdepth = maxdepth, cp = 0))

    # xerror from cross-validated tree
    err <- model$cptable[which.min(model$cptable[, "xerror"]), "xerror"]

    cat("minsplit =", minsplit, ", maxdepth =", maxdepth, " -> xerror =", err, "\n")

    if (err < best_error) {
      best_error <- err
      best_model <- model
    }
  }
}

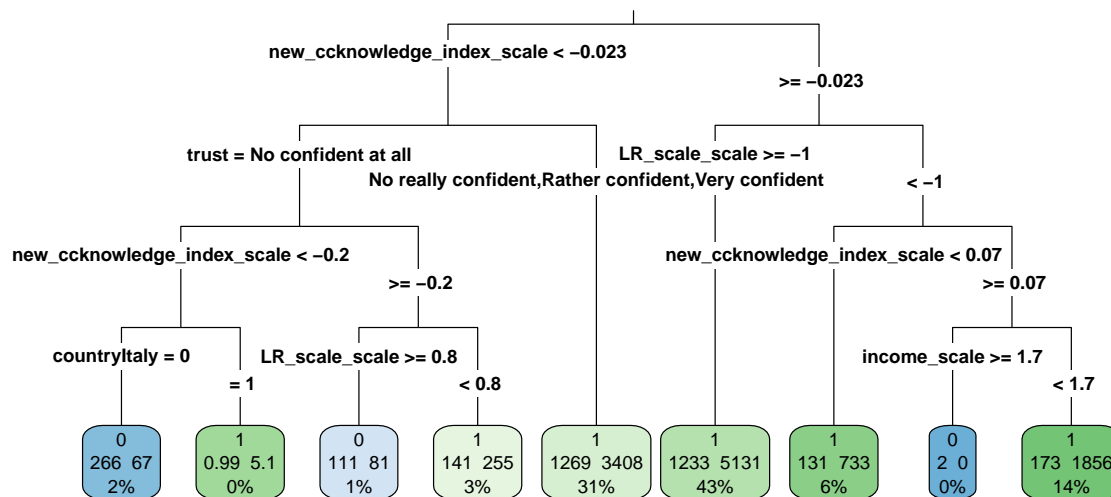
```

```
## minsplit = 10 , maxdepth = 2 -> xerror = 0.9727628
```

```
## minsplit = 10 , maxdepth = 4 -> xerror = 0.9355584
## minsplit = 10 , maxdepth = 6 -> xerror = 0.9359682
## minsplit = 20 , maxdepth = 2 -> xerror = 0.9715462
## minsplit = 20 , maxdepth = 4 -> xerror = 0.9416051
## minsplit = 20 , maxdepth = 6 -> xerror = 0.9360847
## minsplit = 30 , maxdepth = 2 -> xerror = 0.9681962
## minsplit = 30 , maxdepth = 4 -> xerror = 0.9420926
## minsplit = 30 , maxdepth = 6 -> xerror = 0.9366461
```

```
# Plot best tree
```

```
rpart.plot(best_model, type = 3, extra = 101, fallen.leaves = TRUE)
```



Creating The Accuracy Metrics for the Pruned Classification Tree model for the Train data:

```
yhat_tree_pruned_train <- predict(pruned_tree, newdata = train_data, type = "class")
confusionMatrix(as.factor(train_data$ctax_binary), yhat_tree_pruned_train)
```

```
## Confusion Matrix and Statistics
##
##           Reference
## Prediction    0    1
##           0  270 3099
##           1   73 11517
```



```
##
##           Accuracy : 0.788
##           95% CI : (0.7813, 0.7945)
##      No Information Rate : 0.9771
##      P-Value [Acc > NIR] : 1
##
##           Kappa : 0.1084
##
##      McNemar's Test P-Value : <2e-16
##
##           Sensitivity : 0.78717
##           Specificity : 0.78797
##      Pos Pred Value : 0.08014
##      Neg Pred Value : 0.99370
##           Prevalence : 0.02293
##      Detection Rate : 0.01805
##      Detection Prevalence : 0.22522
##      Balanced Accuracy : 0.78757
##
##      'Positive' Class : 0
##
```

Creating The Accuracy Metrics for the Pruned Classification Tree model for the Test data:

```
yhat_tree_pruned_test <- predict(pruned_tree, newdata = test_data, type = "class")
confusionMatrix(as.factor(test_data$ctax_binary), yhat_tree_pruned_test)
```

```
## Confusion Matrix and Statistics
##
##           Reference
## Prediction    0    1
##           0   77  773
##           1   23 2866
##
##           Accuracy : 0.7871
##           95% CI : (0.7736, 0.8001)
##      No Information Rate : 0.9733
##      P-Value [Acc > NIR] : 1
##
##           Kappa : 0.12
##
##      McNemar's Test P-Value : <2e-16
##
##           Sensitivity : 0.77000
##           Specificity : 0.78758
##      Pos Pred Value : 0.09059
##      Neg Pred Value : 0.99204
##           Prevalence : 0.02675
##      Detection Rate : 0.02059
##      Detection Prevalence : 0.22733
##      Balanced Accuracy : 0.77879
##
```

```
##          'Positive' Class : 0
##
```

Creating The Accuracy Metrics for the Tuned Classification Tree model for the Train data:

```
yhat_tree_tuned_train <- predict(best_model, newdata = train_data, type = "class")
confusionMatrix(as.factor(train_data$ctax_binary),yhat_tree_tuned_train)
```

```
## Confusion Matrix and Statistics
##
##              Reference
## Prediction      0      1
##              0  385 2984
##              1  150 11440
##
##              Accuracy : 0.7905
##              95% CI : (0.7839, 0.797)
##              No Information Rate : 0.9642
##              P-Value [Acc > NIR] : 1
##
##              Kappa : 0.1444
##
## Mcnemar's Test P-Value : <2e-16
##
##              Sensitivity : 0.71963
##              Specificity : 0.79312
##              Pos Pred Value : 0.11428
##              Neg Pred Value : 0.98706
##              Prevalence : 0.03576
##              Detection Rate : 0.02574
##              Detection Prevalence : 0.22522
##              Balanced Accuracy : 0.75637
##
##          'Positive' Class : 0
##
```

Creating The Accuracy Metrics for the Tuned Classification Tree model for the Test data:

```
yhat_tree_tuned_test <- predict(best_model, newdata = test_data, type = "class")
confusionMatrix(as.factor(test_data$ctax_binary),yhat_tree_tuned_test)
```

```
## Confusion Matrix and Statistics
##
##              Reference
## Prediction      0      1
##              0   98   752
##              1   42  2847
##
##              Accuracy : 0.7876
```

```
##          95% CI : (0.7742, 0.8007)
##    No Information Rate : 0.9626
##    P-Value [Acc > NIR] : 1
##
##          Kappa : 0.1429
##
##    McNemar's Test P-Value : <2e-16
##
##          Sensitivity : 0.70000
##          Specificity : 0.79105
##    Pos Pred Value : 0.11529
##    Neg Pred Value : 0.98546
##          Prevalence : 0.03744
##    Detection Rate : 0.02621
##    Detection Prevalence : 0.22733
##    Balanced Accuracy : 0.74553
##
##    'Positive' Class : 0
##
```

RandomForest Model:

Creating Grid Search Parameters for the Random Forest Model

```
mtry_grid <- c(2, 4, 6)
ntree_grid <- c(100, 300, 500)
nodesize_grid <- c(1, 5)

best_oob <- Inf
best_model <- NULL
results <- data.frame()
```

Training the Model respect to Grid Search

```
set.seed(123)

# Grid search loop
for (m in mtry_grid) {
  for (n in ntree_grid) {
    for (node in nodesize_grid) {

      model <- randomForest(
        full_formula,
        data = train_data,
        mtry = m,
        ntree = n,
        nodesize = node,
      )

      oob_err <- tail(model$err.rate[, "OOB"], 1)
```

```

# Store results
results <- rbind(results, data.frame(mtry = m, ntree = n, nodesize = node, OOB_Error = oob_err))

cat("mtry =", m, "| ntree =", n, "| nodesize =", node, "| OOB error =", oob_err, "\n")

if (oob_err < best_oob) {
  best_oob <- oob_err
  best_model <- model
}
}
}
}

```

```

## mtry = 2 | ntree = 100 | nodesize = 1 | OOB error = 0.2236781
## mtry = 2 | ntree = 100 | nodesize = 5 | OOB error = 0.2228759
## mtry = 2 | ntree = 300 | nodesize = 1 | OOB error = 0.2245471
## mtry = 2 | ntree = 300 | nodesize = 5 | OOB error = 0.2246139
## mtry = 2 | ntree = 500 | nodesize = 1 | OOB error = 0.2246139
## mtry = 2 | ntree = 500 | nodesize = 5 | OOB error = 0.2250819
## mtry = 4 | ntree = 100 | nodesize = 1 | OOB error = 0.2132495
## mtry = 4 | ntree = 100 | nodesize = 5 | OOB error = 0.2133164
## mtry = 4 | ntree = 300 | nodesize = 1 | OOB error = 0.2119126
## mtry = 4 | ntree = 300 | nodesize = 5 | OOB error = 0.2120463
## mtry = 4 | ntree = 500 | nodesize = 1 | OOB error = 0.2115783
## mtry = 4 | ntree = 500 | nodesize = 5 | OOB error = 0.2111772
## mtry = 6 | ntree = 100 | nodesize = 1 | OOB error = 0.2133164
## mtry = 6 | ntree = 100 | nodesize = 5 | OOB error = 0.2133832
## mtry = 6 | ntree = 300 | nodesize = 1 | OOB error = 0.2101745
## mtry = 6 | ntree = 300 | nodesize = 5 | OOB error = 0.2109098
## mtry = 6 | ntree = 500 | nodesize = 1 | OOB error = 0.2113778
## mtry = 6 | ntree = 500 | nodesize = 5 | OOB error = 0.2119794

```

```
best_model
```

```

##
## Call:
## randomForest(formula = full_formula, data = train_data, mtry = m, ntree = n, nodesize = node,
##               Type of random forest: classification
##               Number of trees: 300
## No. of variables tried at each split: 6
##
## OOB estimate of error rate: 21.02%
## Confusion matrix:
##      0      1 class.error
## 0 475  2894  0.85900861
## 1 250 11340  0.02157032

```

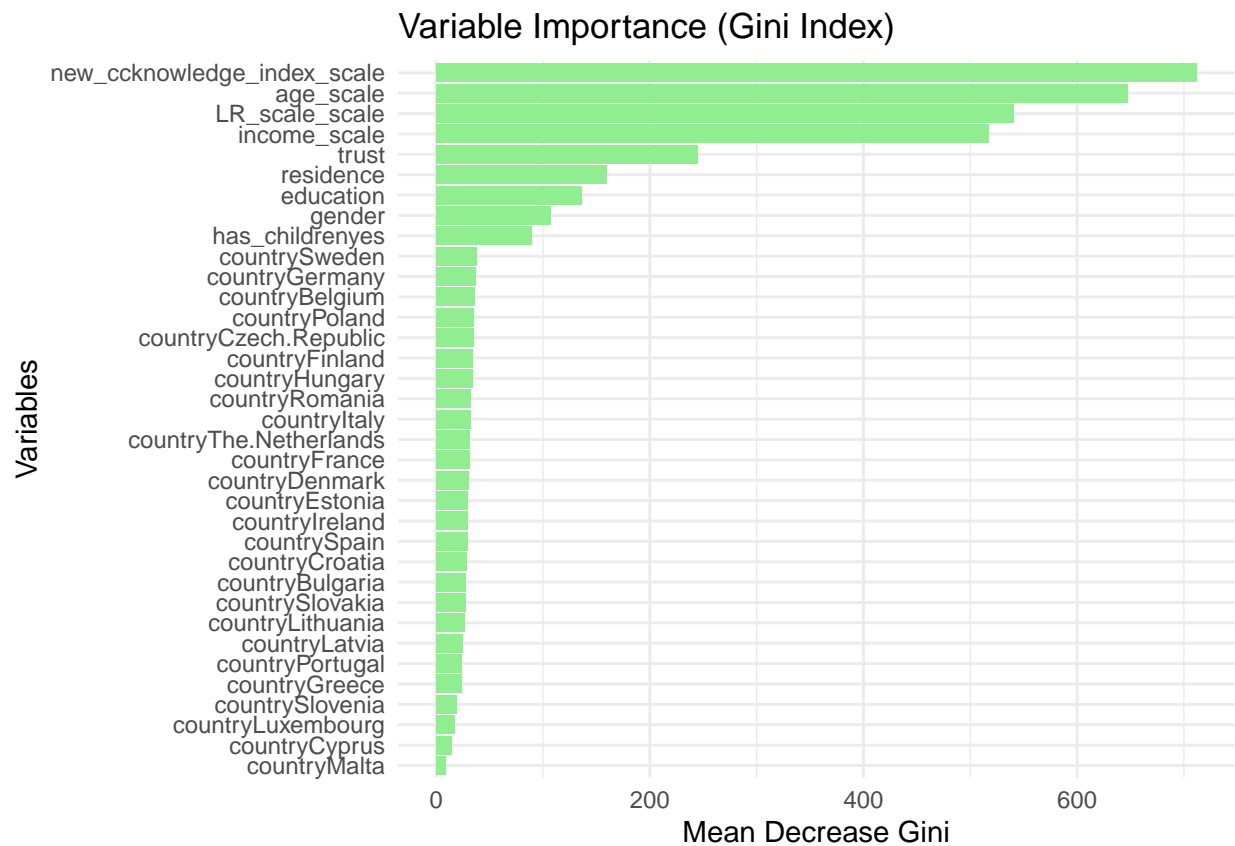
Creating Variable Importance to see which features are important for our randomForest model:

```

imp <- importance(best_model, type = 2) # Getting the feature importance of the best model respect to Me
imp_df <- data.frame(Variable = rownames(imp), Importance = imp[, 1])

```

```
# Plot using ggplot2
ggplot(imp_df, aes(x = reorder(Variable, Importance), y = Importance)) +
  geom_col(fill = "lightgreen") +
  coord_flip() +
  labs(title = "Variable Importance (Gini Index)",
       x = "Variables", y = "Mean Decrease Gini") +
  theme_minimal()
```



Creating The Accuracy Metrics for the RandomForest model for the Train data:

```
yhat_rf_train=predict(best_model, newdata = train_data, type = "response")
confusionMatrix(as.factor(train_data$ctax_binary),yhat_rf_train)
```

```
## Confusion Matrix and Statistics
##
##           Reference
## Prediction      0      1
##           0 2625   744
##           1    2 11588
##
##           Accuracy : 0.9501
##           95% CI : (0.9465, 0.9536)
```

```
##      No Information Rate : 0.8244
##      P-Value [Acc > NIR] : < 2.2e-16
##
##              Kappa : 0.845
##
##  Mcnemar's Test P-Value : < 2.2e-16
##
##      Sensitivity : 0.9992
##      Specificity : 0.9397
##      Pos Pred Value : 0.7792
##      Neg Pred Value : 0.9998
##      Prevalence : 0.1756
##      Detection Rate : 0.1755
##      Detection Prevalence : 0.2252
##      Balanced Accuracy : 0.9695
##
##      'Positive' Class : 0
##
```

Creating The Accuracy Metrics for the RandomForest model for the Test data:

```
yhat_rf_test= predict(best_model, newdata = test_data, type = "response")
confusionMatrix(as.factor(test_data$ctax_binary),yhat_rf_test)
```

```
## Confusion Matrix and Statistics
##
##      Reference
## Prediction    0    1
##      0  123  727
##      1   52 2837
##
##      Accuracy : 0.7917
##      95% CI : (0.7783, 0.8046)
##      No Information Rate : 0.9532
##      P-Value [Acc > NIR] : 1
##
##      Kappa : 0.176
##
##  Mcnemar's Test P-Value : <2e-16
##
##      Sensitivity : 0.7029
##      Specificity : 0.7960
##      Pos Pred Value : 0.1447
##      Neg Pred Value : 0.9820
##      Prevalence : 0.0468
##      Detection Rate : 0.0329
##      Detection Prevalence : 0.2273
##      Balanced Accuracy : 0.7494
##
##      'Positive' Class : 0
##
```

```
table(test_data$ctax_binary,yhat_rf_test)
```

```
##      yhat_rf_test
##           0      1
##    0  123   727
##    1   52 2837
```

Starting SRM:

Defining the smoothed residual function:

```
compute_smoothed_residuals=function(yhat_ml,yhat_baseline,y_true){
  sm_res=((y_true-yhat_baseline)^2) +
    (2 * (y_true - yhat_baseline) * (yhat_ml - y_true))+
    ((yhat_ml - y_true)^2)
  return(sm_res)
}
```

```
yhat_baseline=predict(baseline_model,newdata=test_data,type="response")
yhat_rf= predict(best_model, newdata = test_data, type = "prob")[, "1"]
```

Creating the Smoothed Residuals:

```
srm_df=test_data
srm_df["Smoothed_Residuals"]=compute_smoothed_residuals(yhat_ml=yhat_rf,yhat_baseline,yhat_baseline,test_data$ctax_binary)
```

Finding top 100 smoothed residuals respect to the RandomForest and baseline Logit model:

```
srm_df_sorted <- srm_df %>% arrange(desc(Smoothed_Residuals))
srm_final=head(srm_df_sorted,100)
head(srm_final)
```

```
##      ctax_binary income_scale age_scale      trust LR_scale_scale
## 1             1  1.49539445   9.664 No really confident -0.01511879
## 2             0  1.33867734  -5.537   Very confident   3.94488189
## 3             1  0.29317216 -13.714 No really confident   3.59303591
## 4             1  0.37555168 -24.474   Rather confident   2.57435345
## 5             1 -1.28288379 -11.778 No really confident   1.66515837
## 6             1 -0.01744122 -21.233   Rather confident   2.36582694
##      new_ccknowledge_index_scale any_cc_last2year_factor regional_heterogeneity
## 1                0.13633333                1                0
## 2                0.05433333                1                1
## 3               -0.15372222                0                1
## 4               -0.39711111                1                1
## 5               -0.28488889                1                1
## 6               -0.30327778                0                1
##      LR_scale new_ccknowledge_index gender country_w has_childrenyes
## 1           5                0.7777778  male  0.800470                0
## 2           9                0.7222222  male  1.172674                1
## 3           9                0.5000000  male  0.970778                1
```

```

## 4      8      0.2777778 female 0.971153      1
## 5      7      0.3333333   male 1.152172      0
## 6      8      0.3888889   male 0.747723      1
## countryBelgium countryBulgaria countryCroatia countryCyprus
## 1      0      0      0      0
## 2      0      0      1      0
## 3      0      0      0      0
## 4      0      0      0      0
## 5      0      0      0      0
## 6      0      0      0      0
## countryCzech.Republic countryDenmark countryEstonia countryFinland
## 1      0      0      0      0
## 2      0      0      0      0
## 3      0      0      0      0
## 4      0      0      0      0
## 5      0      0      0      0
## 6      0      0      0      0
## countryFrance countryGermany countryGreece countryHungary countryIreland
## 1      0      0      0      0      0
## 2      0      0      0      0      0
## 3      0      0      0      0      0
## 4      0      1      0      0      0
## 5      0      0      0      0      0
## 6      0      0      0      0      0
## countryItaly countryLatvia countryLithuania countryLuxembourg countryMalta
## 1      0      0      0      0      0
## 2      0      0      0      0      0
## 3      0      0      0      0      0
## 4      0      0      0      0      0
## 5      0      0      0      0      0
## 6      0      0      0      0      0
## countryPoland countryPortugal countryRomania countrySlovakia countrySlovenia
## 1      0      0      0      0      1
## 2      0      0      0      0      0
## 3      0      0      0      0      0
## 4      0      0      0      0      0
## 5      0      0      0      1      0
## 6      0      0      0      0      0
## countrySpain countrySweden countryThe.Netherlands education residence
## 1      0      0      0 secondary town
## 2      0      0      0 tertiary rural
## 3      0      0      0 tertiary city
## 4      0      0      0 tertiary city
## 5      0      0      0 primary rural
## 6      0      1      0 secondary city
## Smoothed_Residuals
## 1      0.2539485
## 2      0.2356578
## 3      0.1630132
## 4      0.1373319
## 5      0.1361468
## 6      0.1185048

```

Creating Plotting Function for the Smoothed Residuals


```

plot_srm_residuals <- function(data, residual_var = "Smoothed_Residuals") {
  predictors <- setdiff(names(data), residual_var)

  for (var in predictors) {
    x <- data[[var]]

    # Determine type
    unique_vals <- unique(na.omit(x))
    is_binary <- length(unique_vals) == 2 && is.numeric(x)
    is_categorical <- is.factor(x) || is.character(x) || is_binary
    is_continuous <- is.numeric(x) && length(unique_vals) > 10

    # Build plot
    if (is_continuous) {
      p <- ggplot(data, aes_string(x = var, y = residual_var)) +
        geom_point(alpha = 0.5) +
        geom_smooth(method = "loess", se = FALSE, color = "blue") +
        labs(title = paste(residual_var, "vs", var),
              x = var, y = residual_var) +
        theme_minimal()

    } else if (is_categorical) {
      # Convert to factor for group-wise plotting if not already
      data[[var]] <- as.factor(data[[var]])

      p <- ggplot(data, aes_string(x = var, y = residual_var)) +
        geom_boxplot(fill = "darkorange", alpha = 0.7) +
        labs(title = paste(residual_var, "by", var),
              x = var, y = residual_var) +
        theme_minimal()
    } else {
      next # Skip variables that don't meet either condition
    }

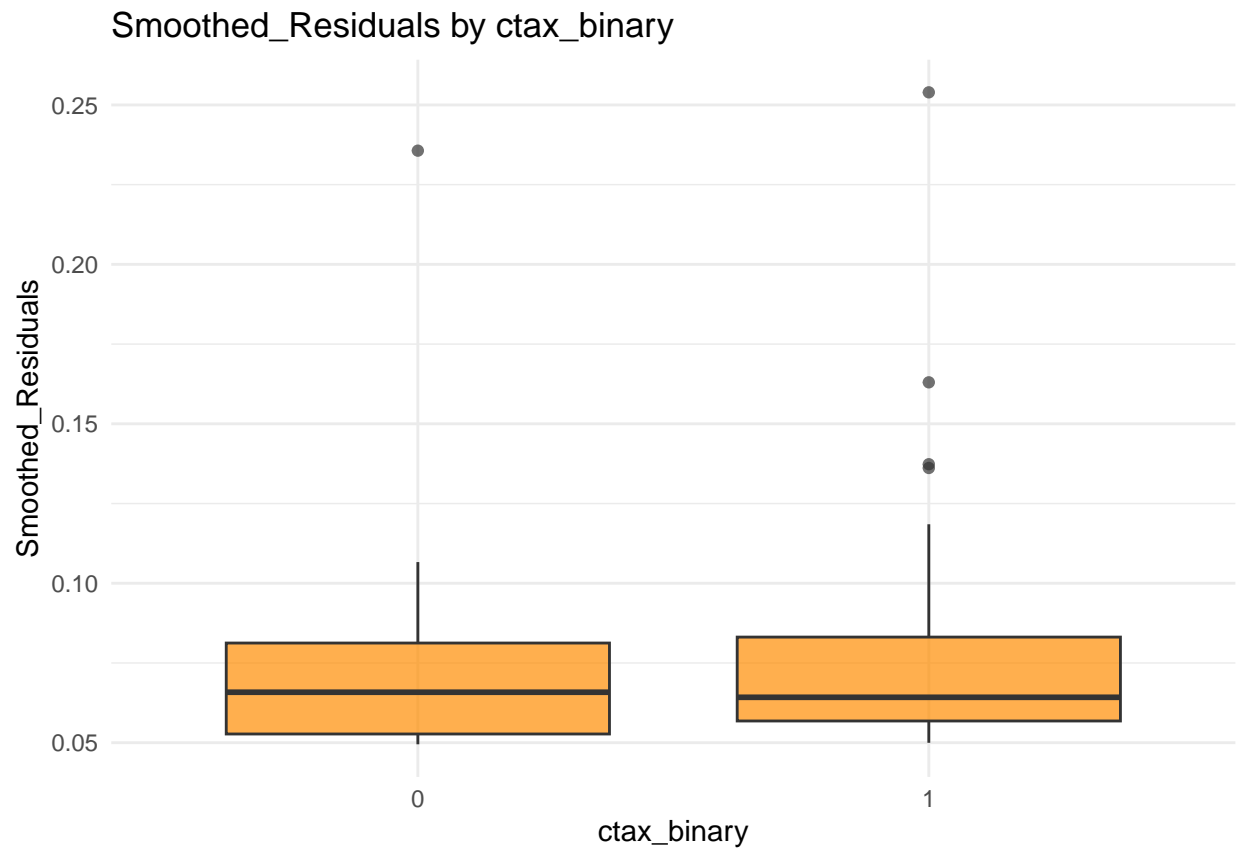
    print(p)
  }
}

```

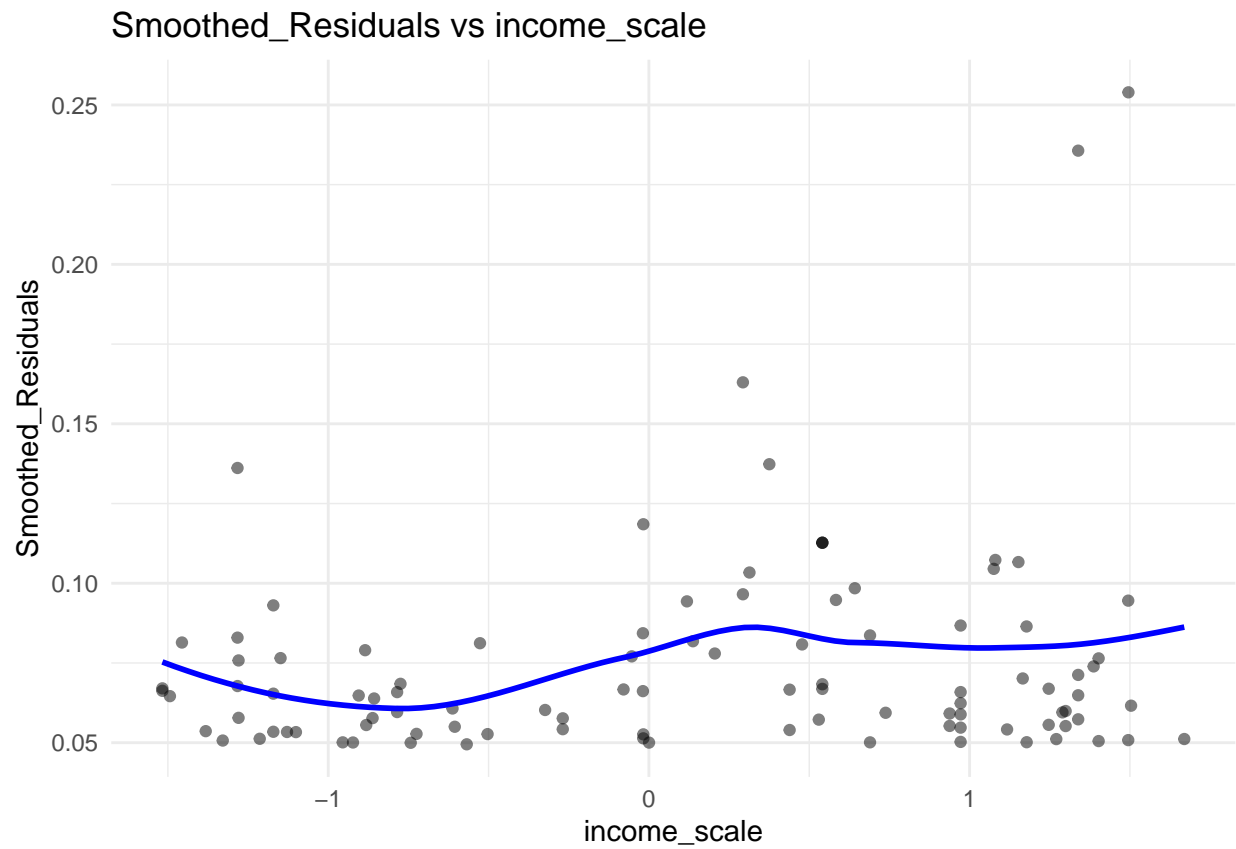
```

plot_srm_residuals(srm_final)

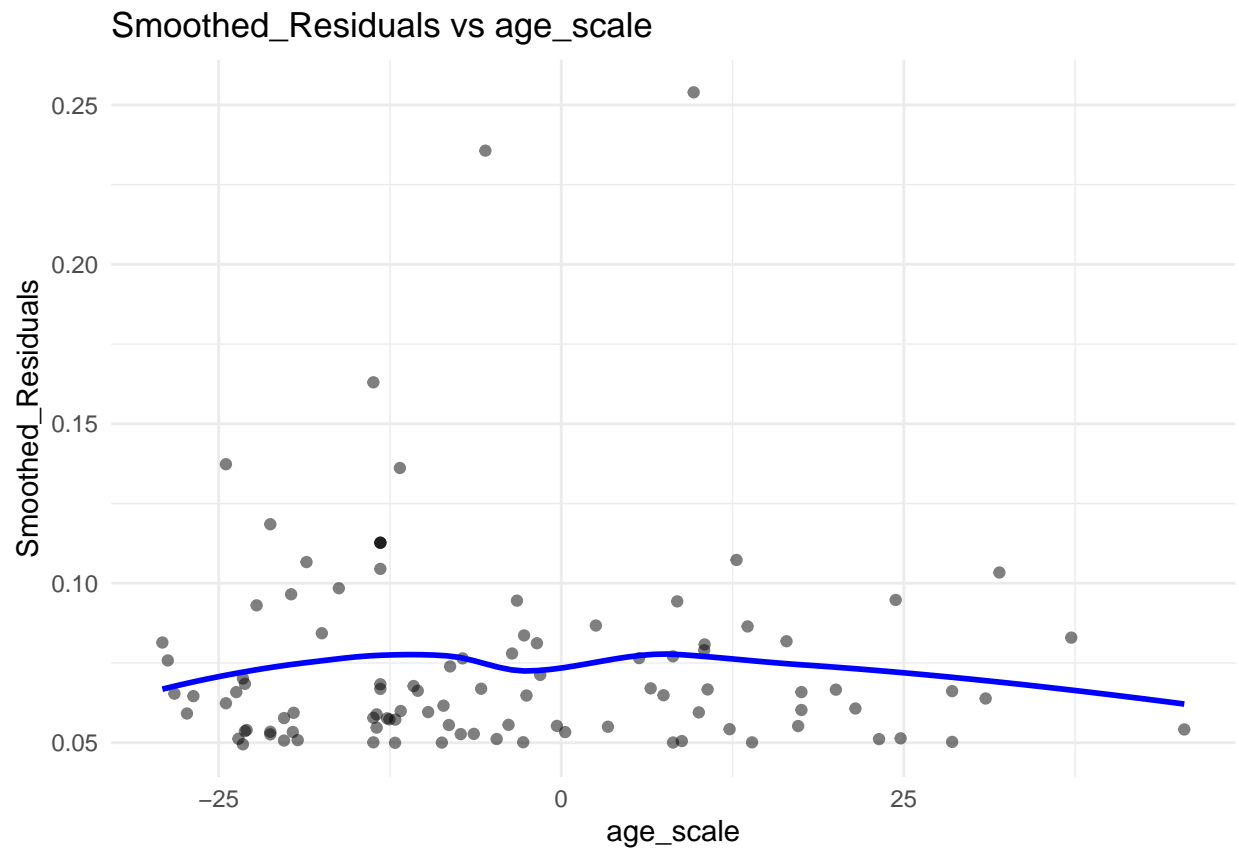
```

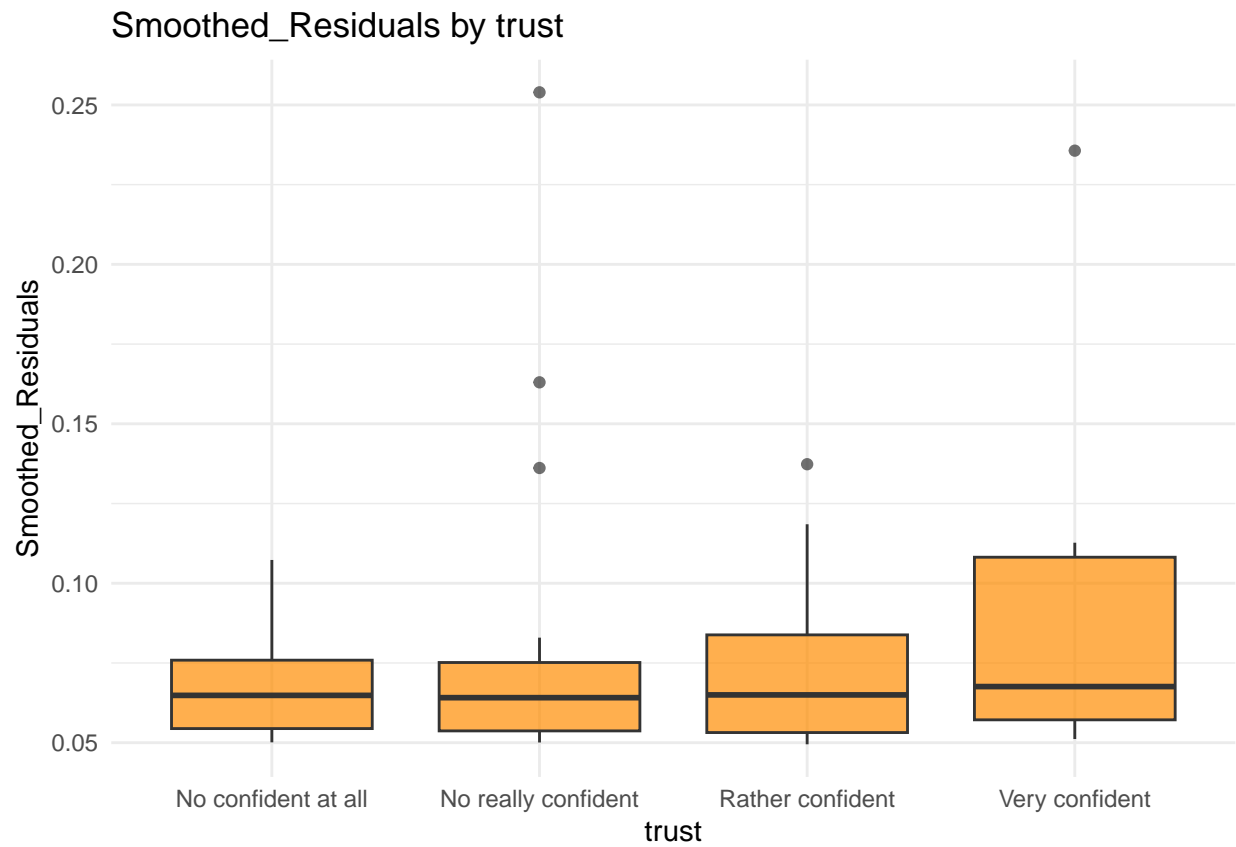


```
## 'geom_smooth()' using formula = 'y ~ x'
```

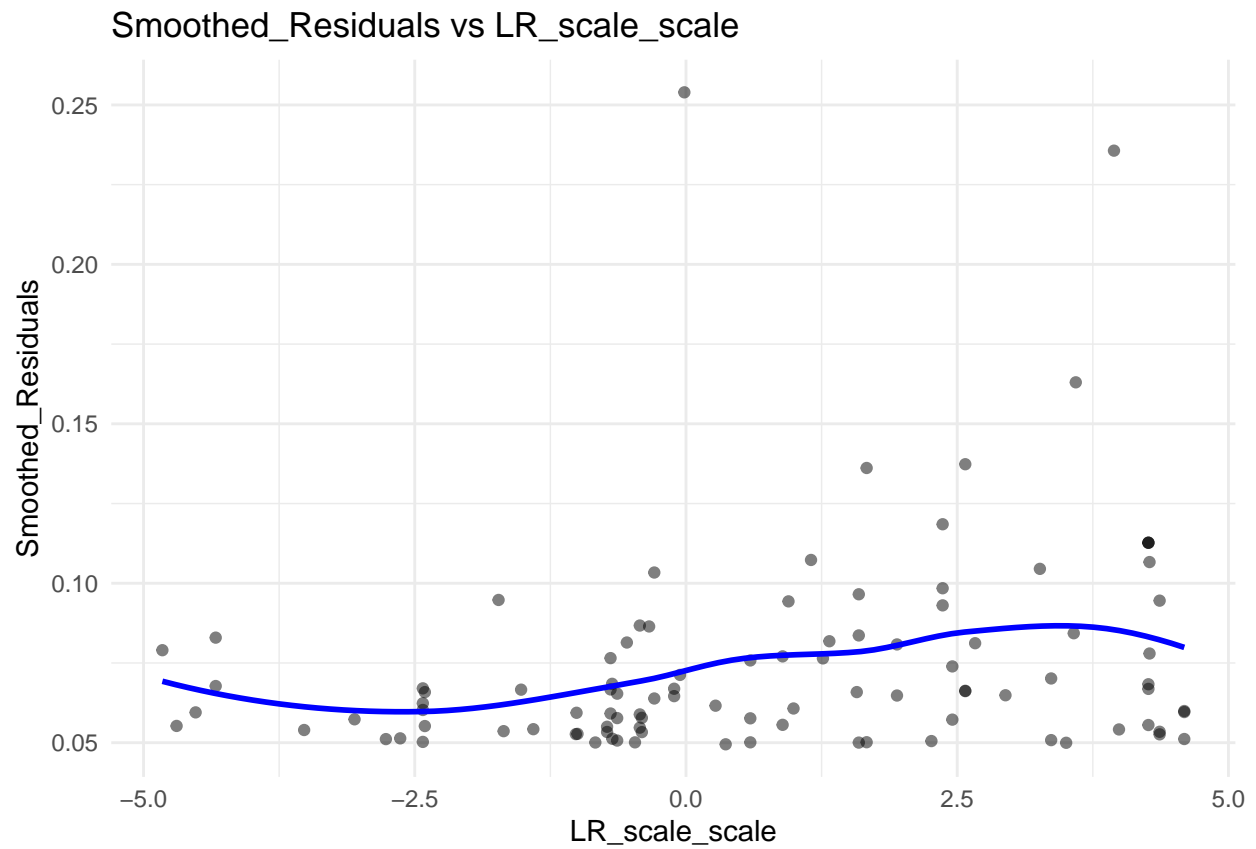


```
## 'geom_smooth()' using formula = 'y ~ x'
```

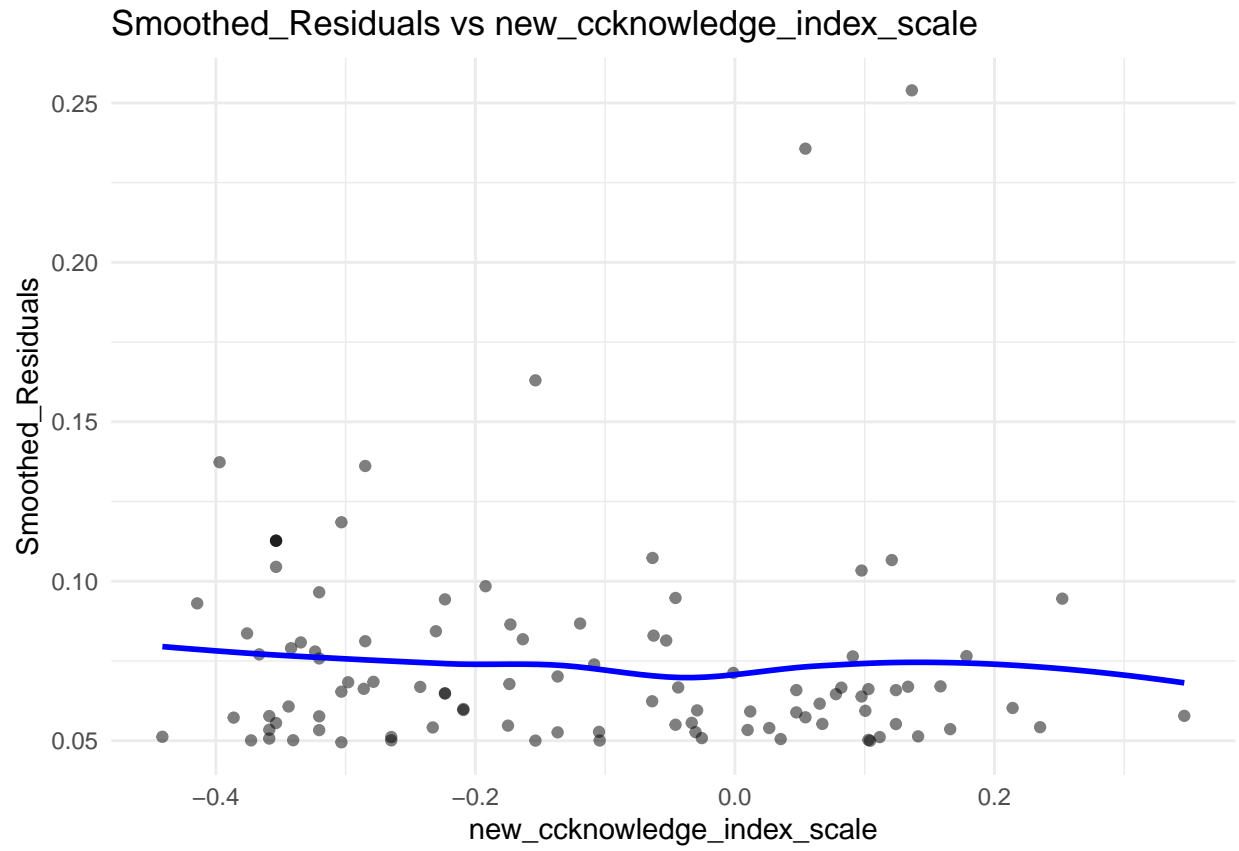


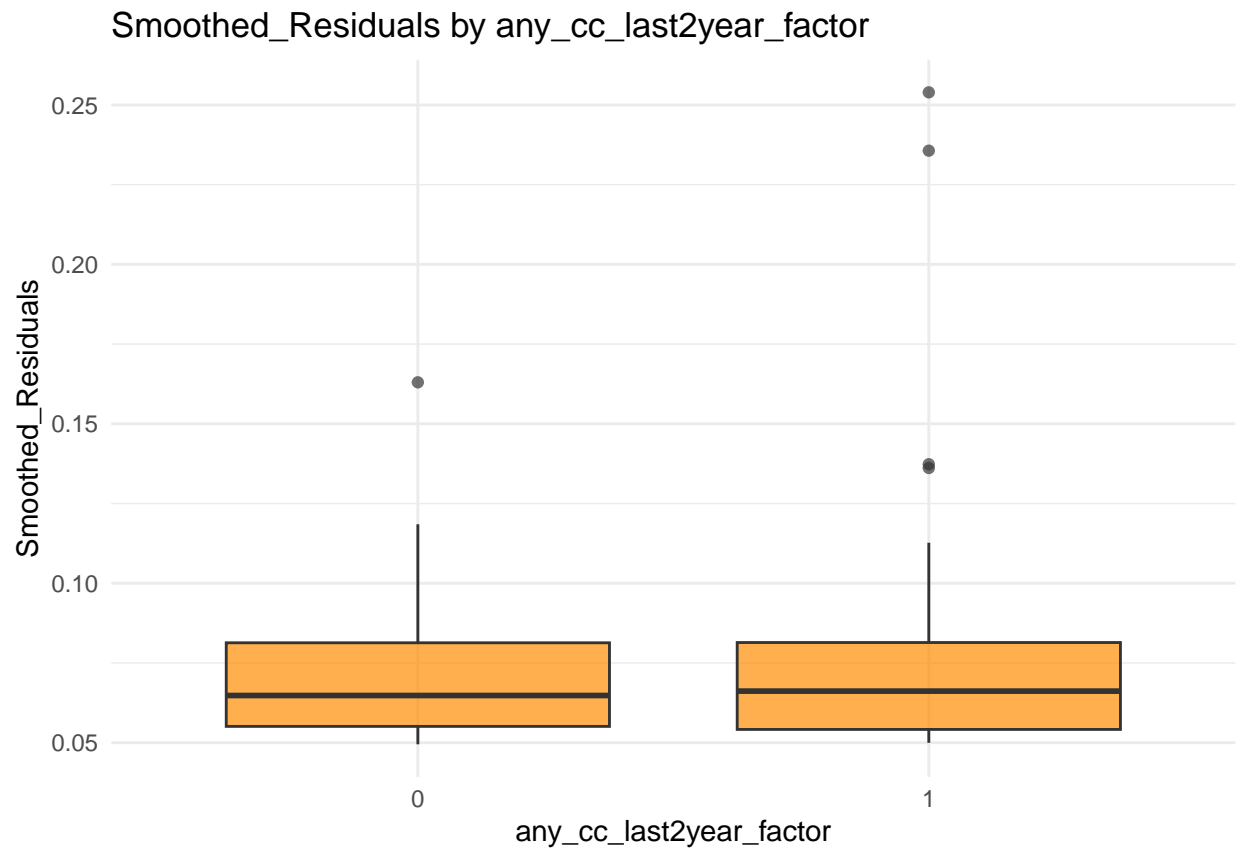


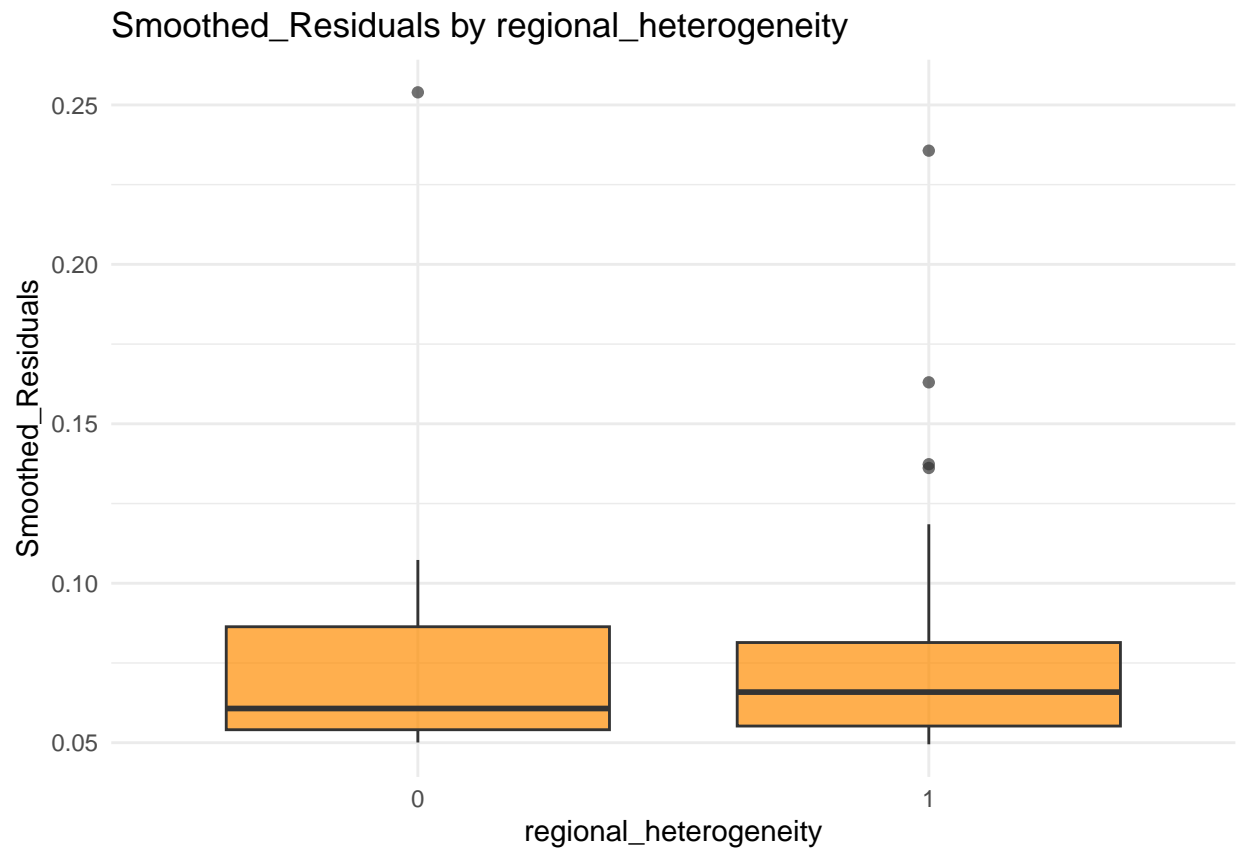
```
## 'geom_smooth()' using formula = 'y ~ x'
```



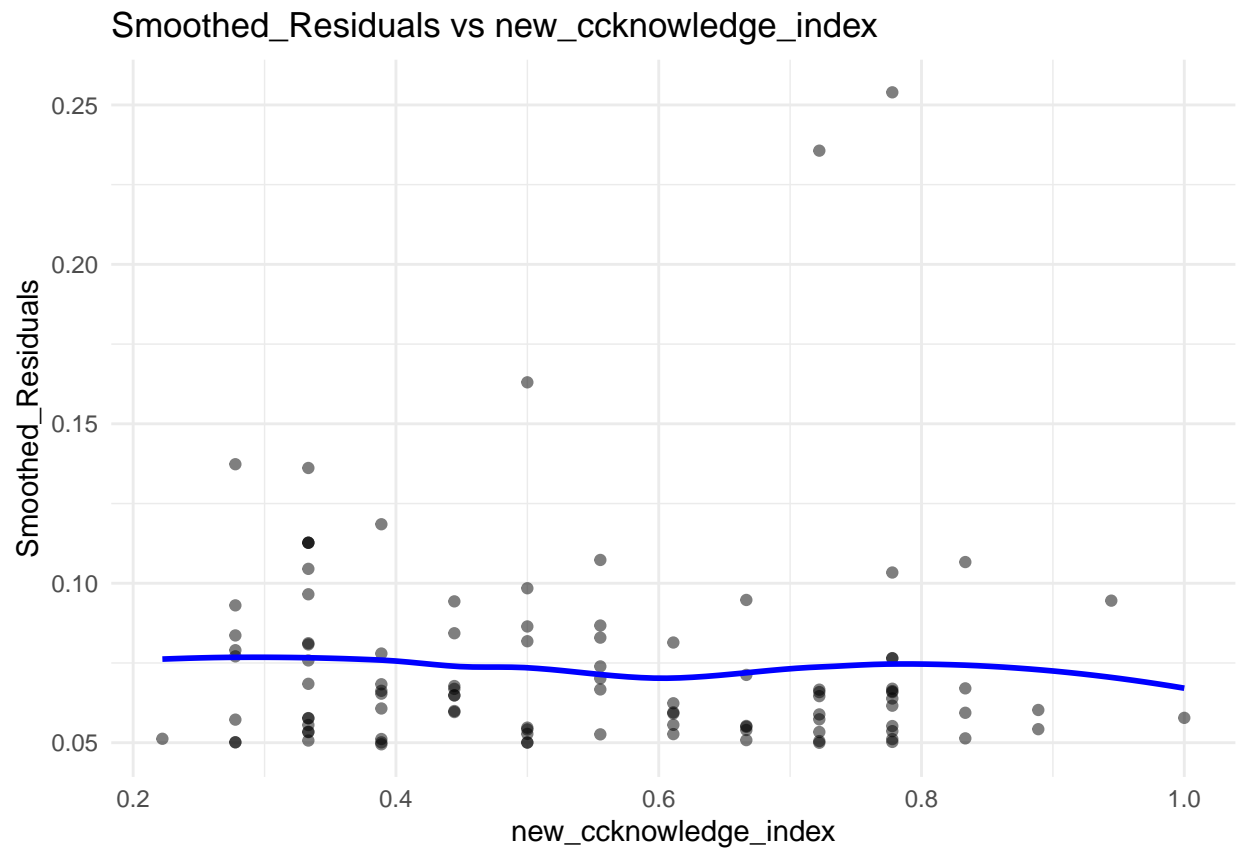
```
## 'geom_smooth()' using formula = 'y ~ x'
```

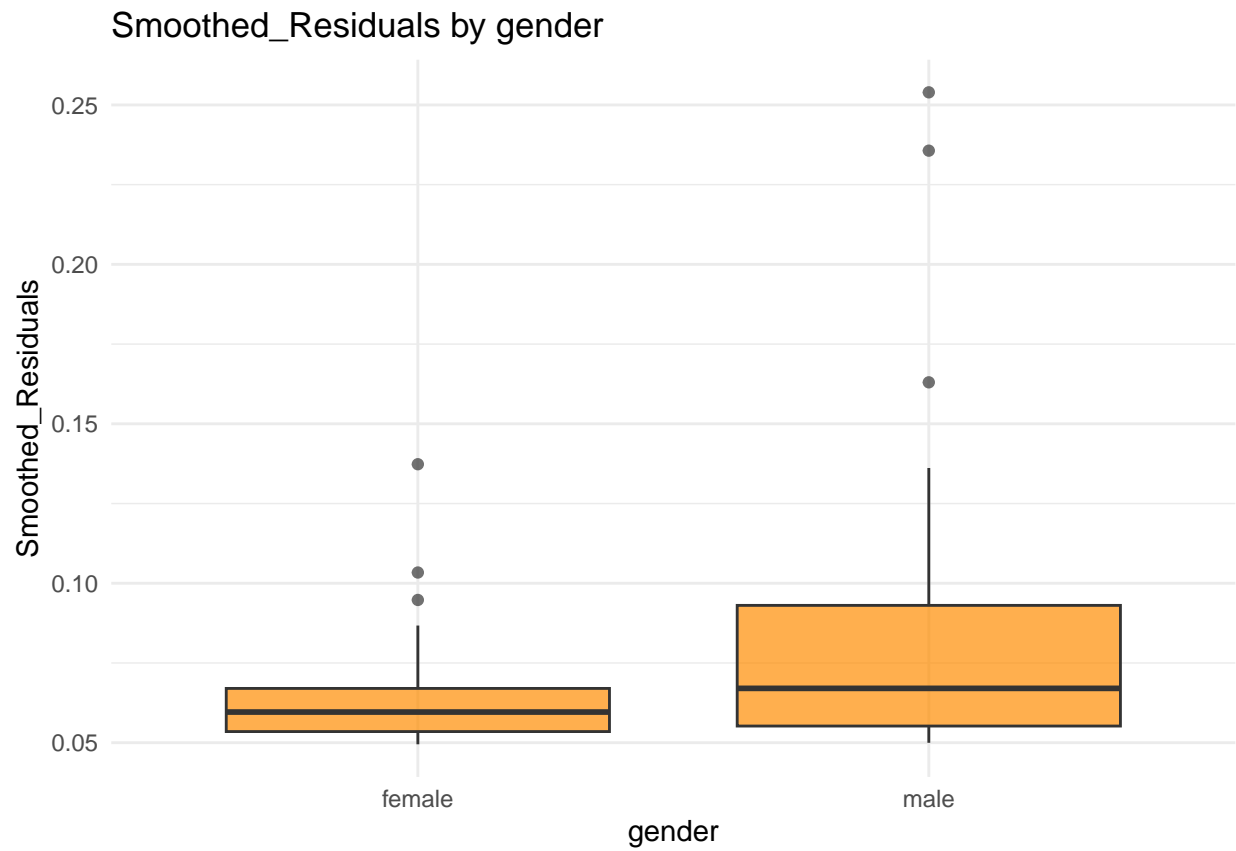




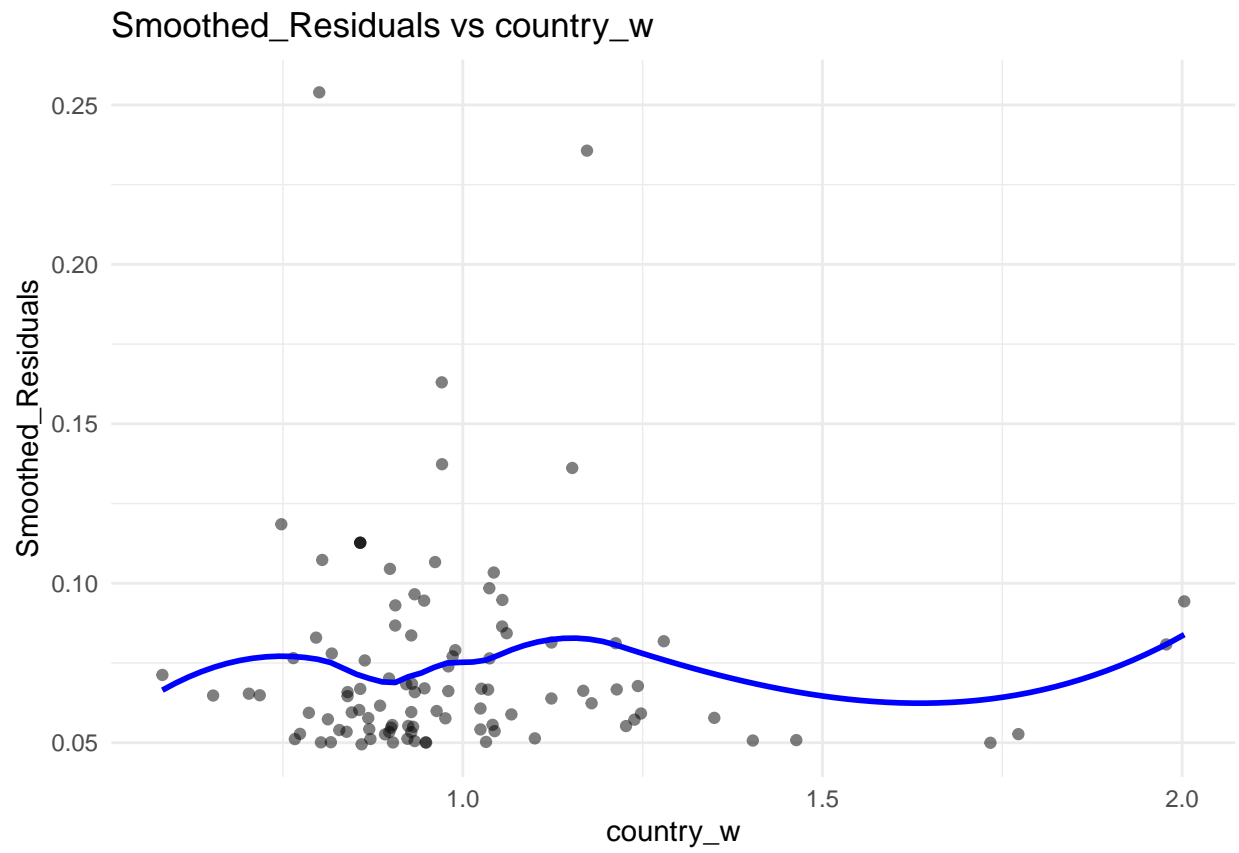


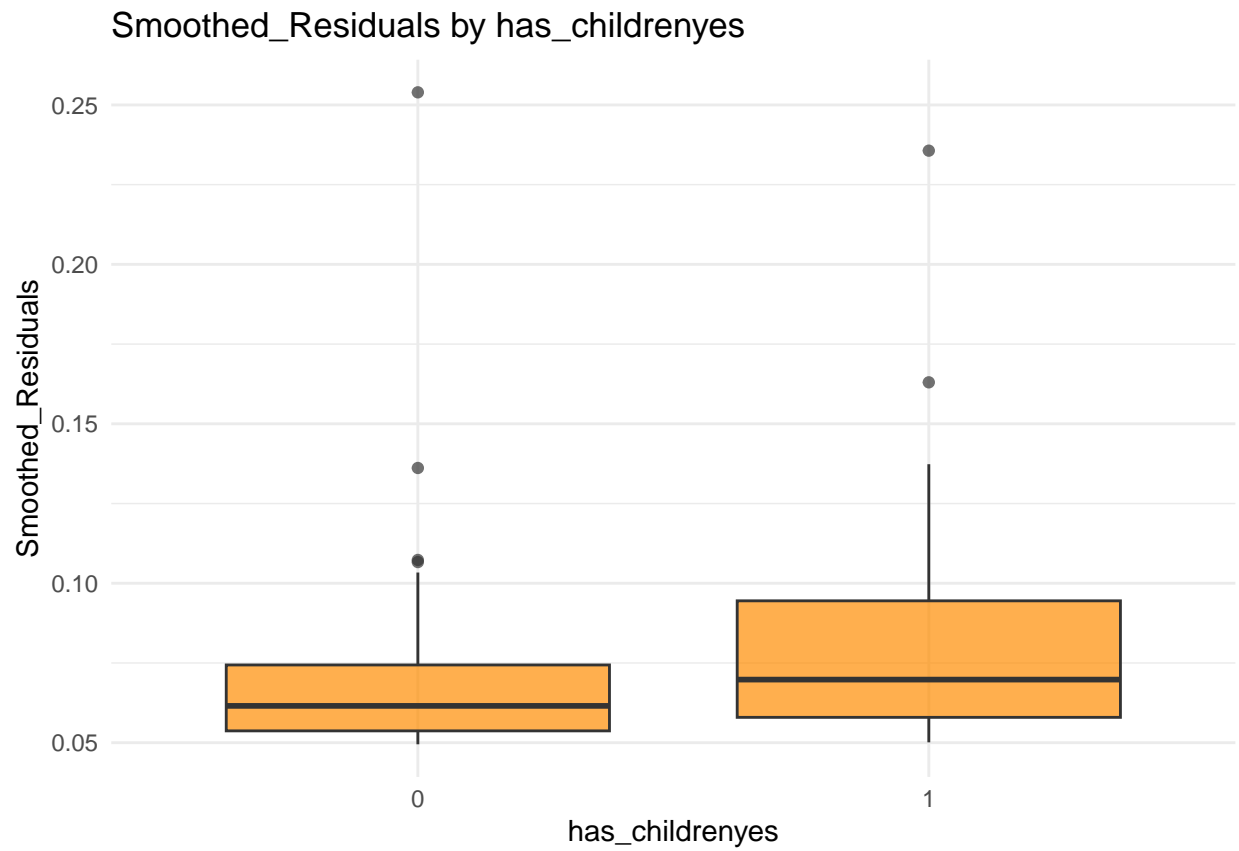
```
## 'geom_smooth()' using formula = 'y ~ x'
```

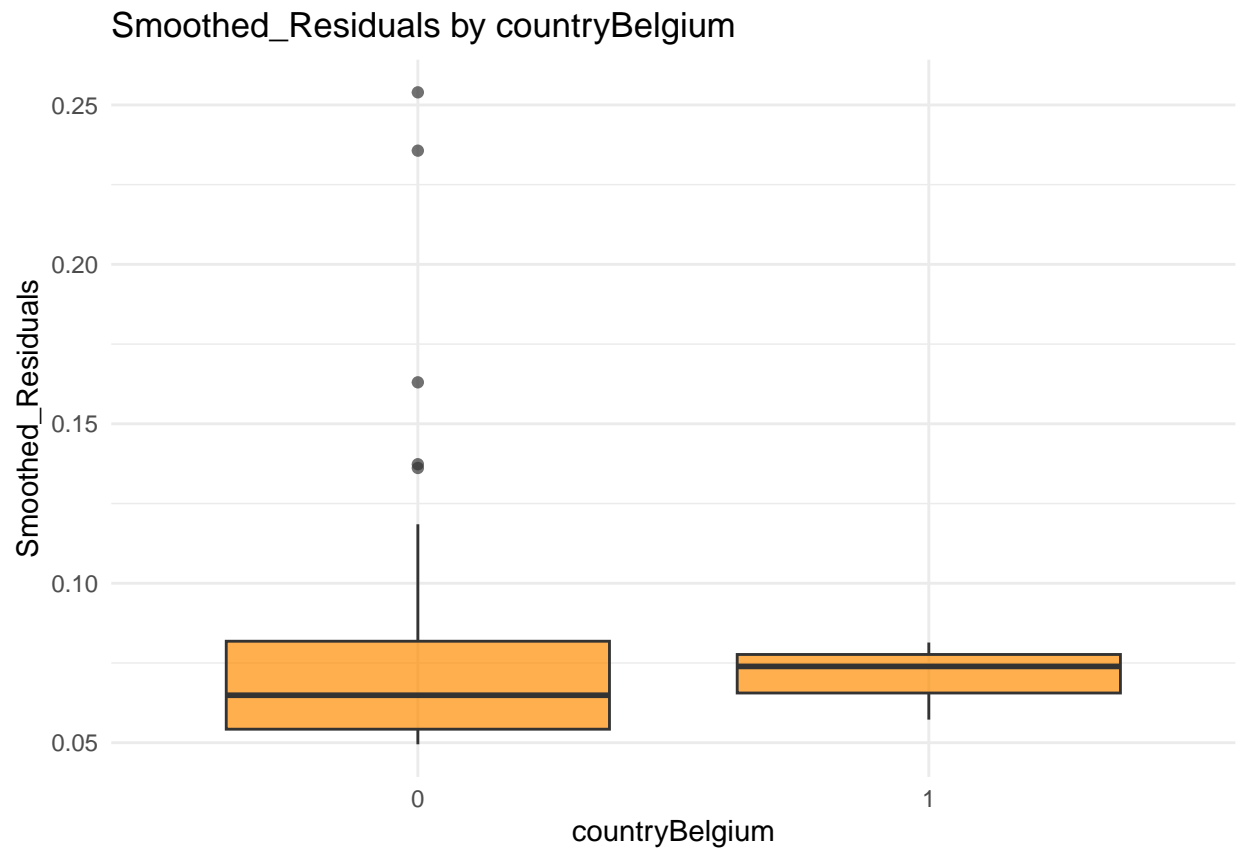


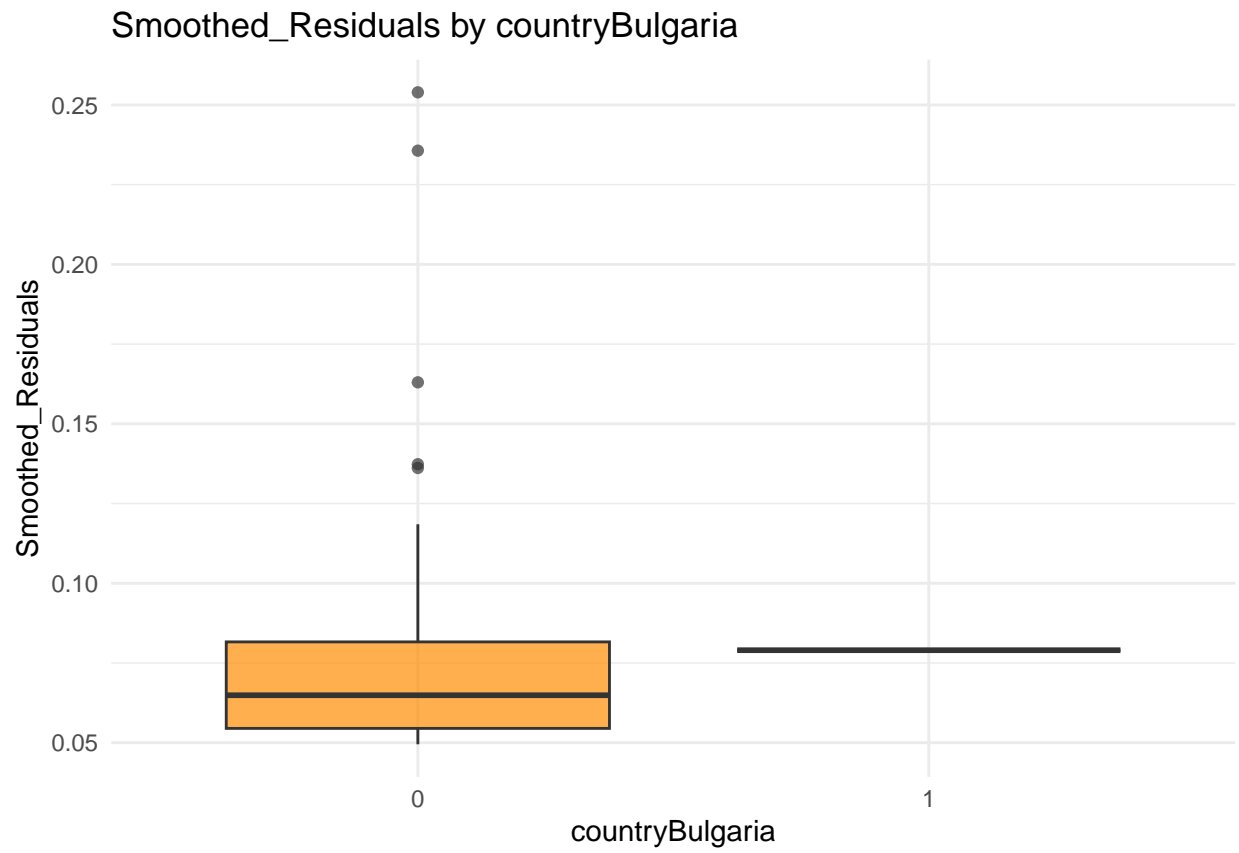


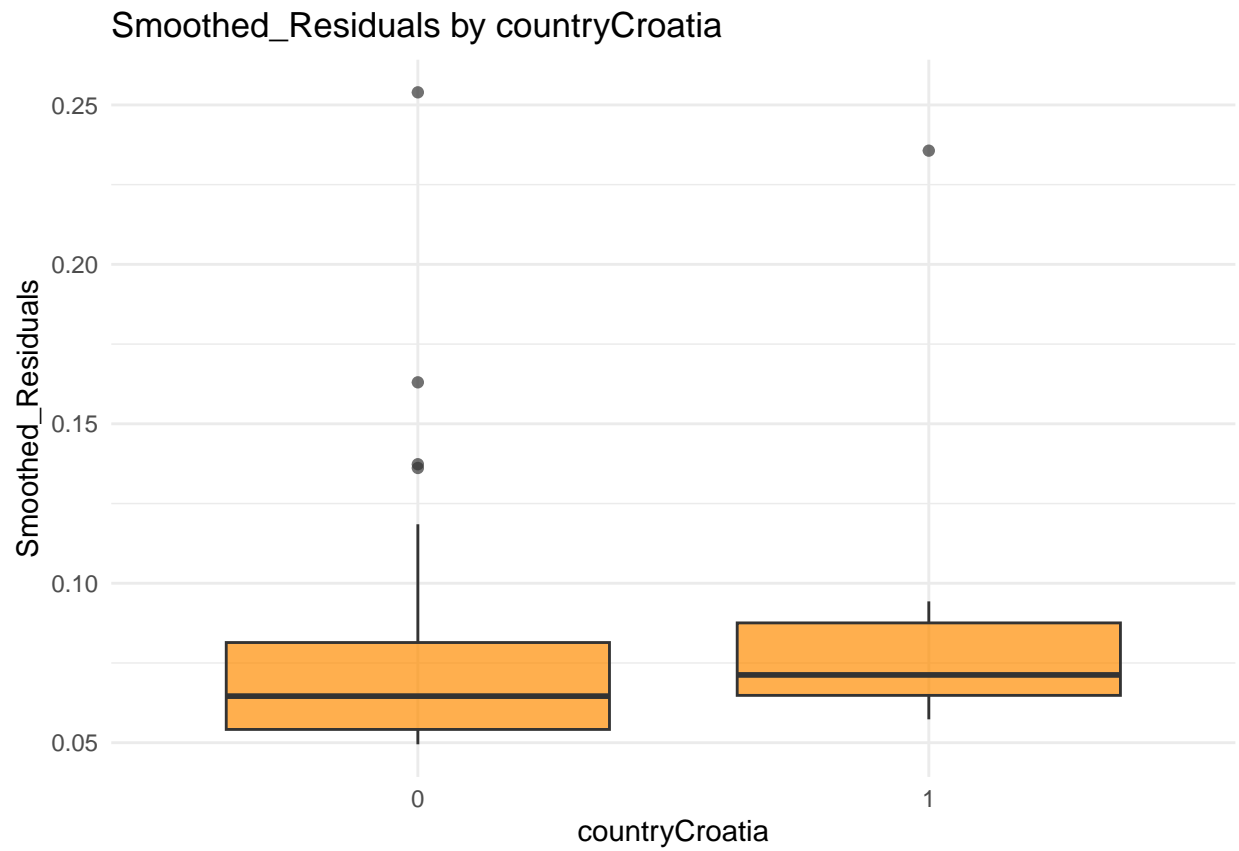
```
## 'geom_smooth()' using formula = 'y ~ x'
```

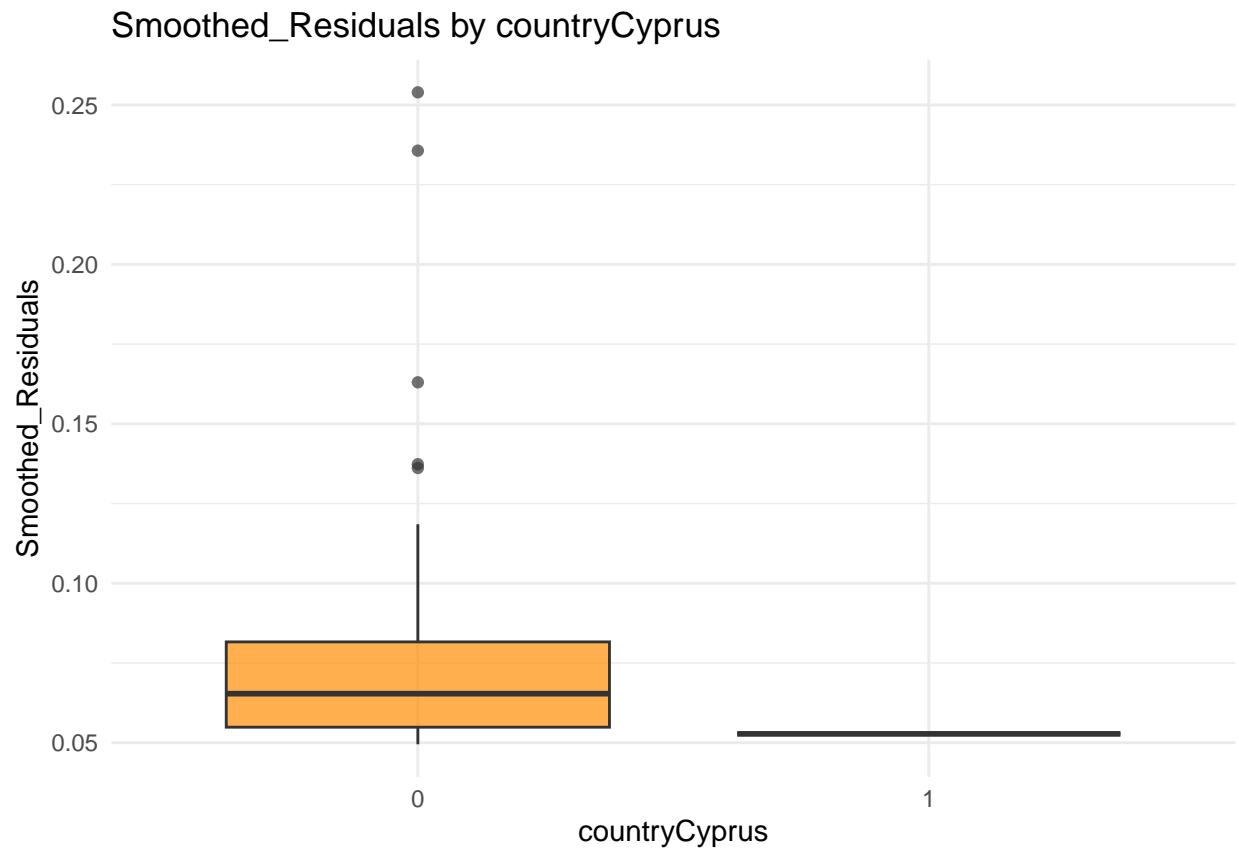


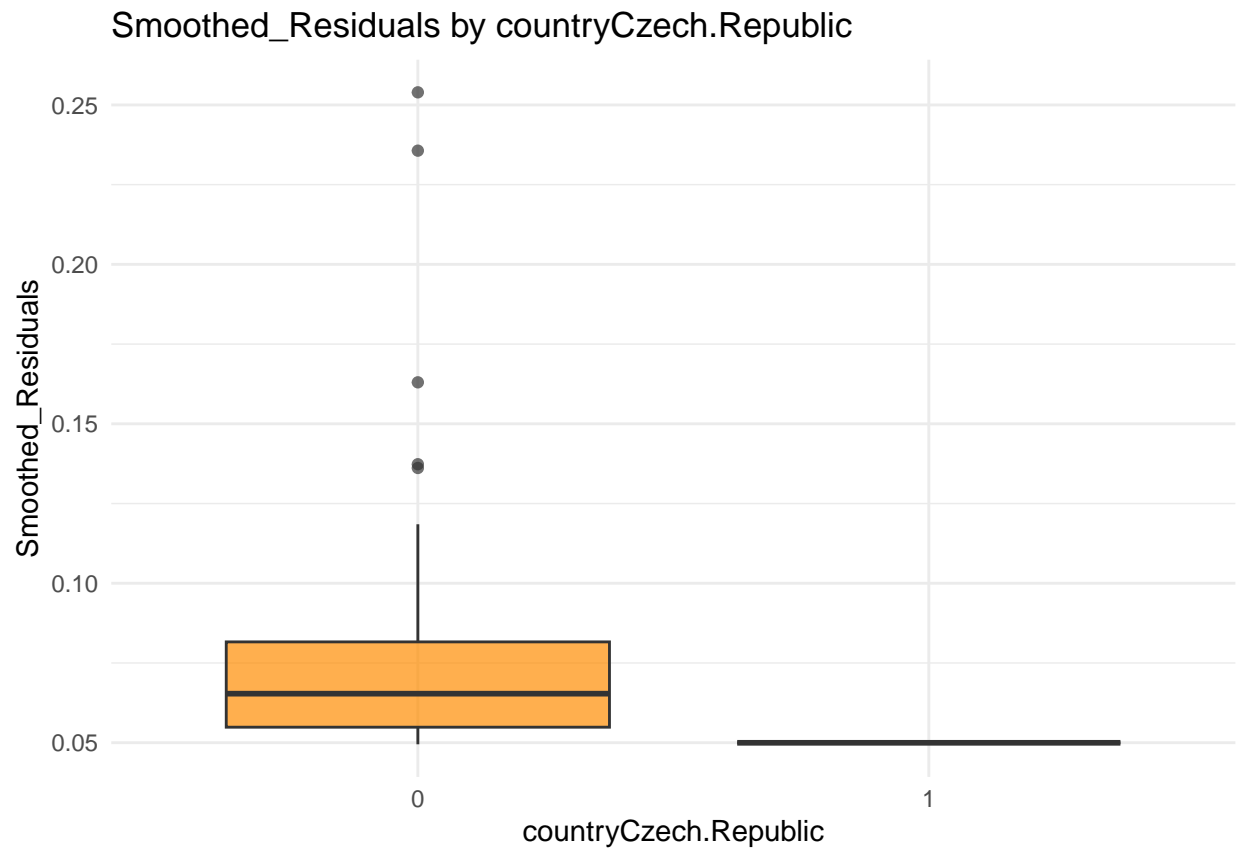


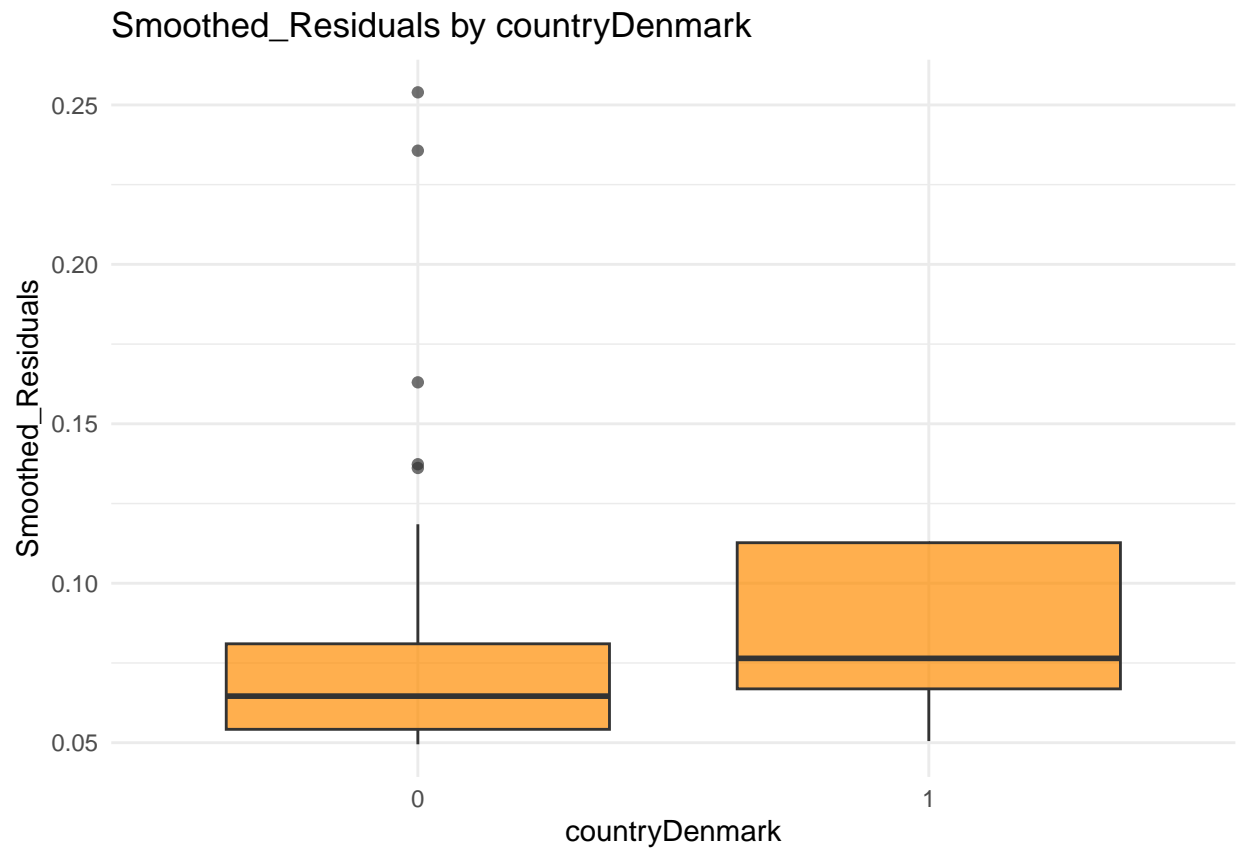


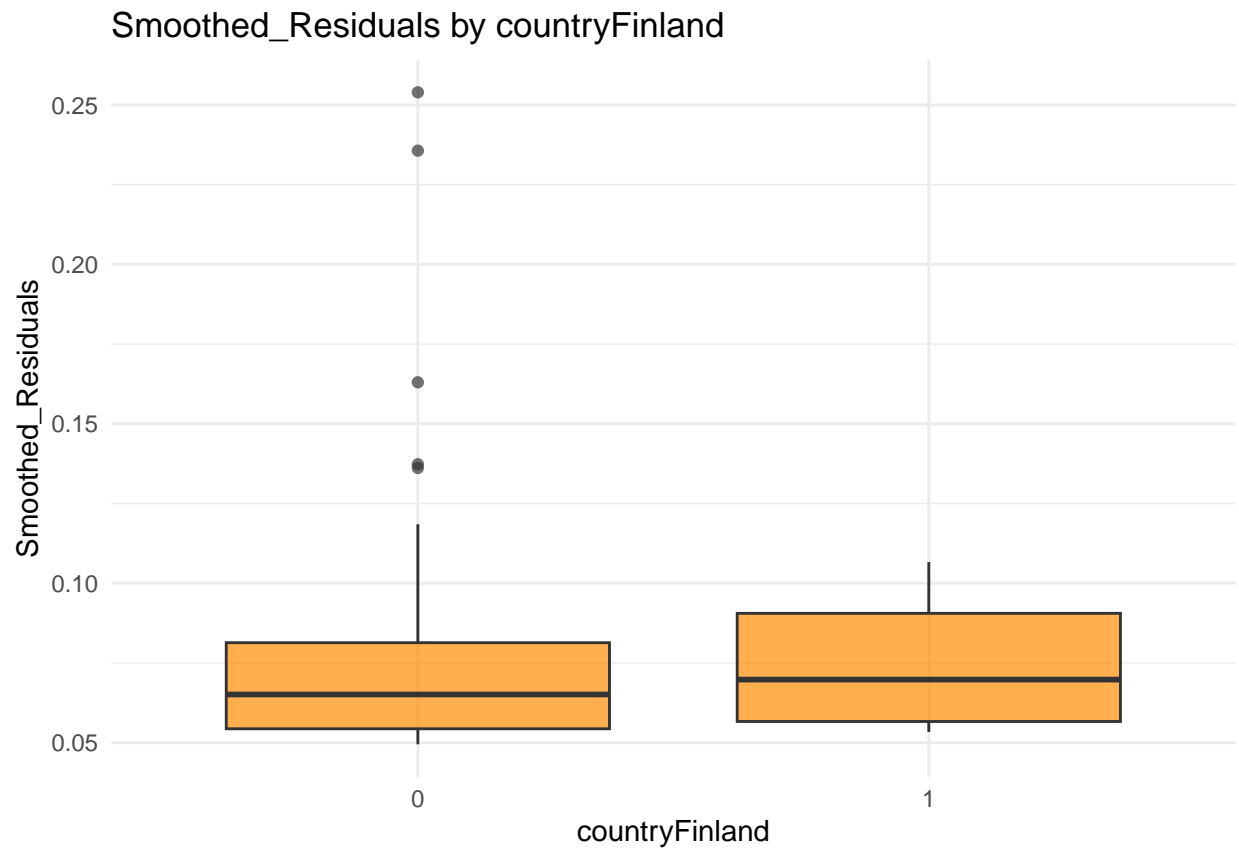


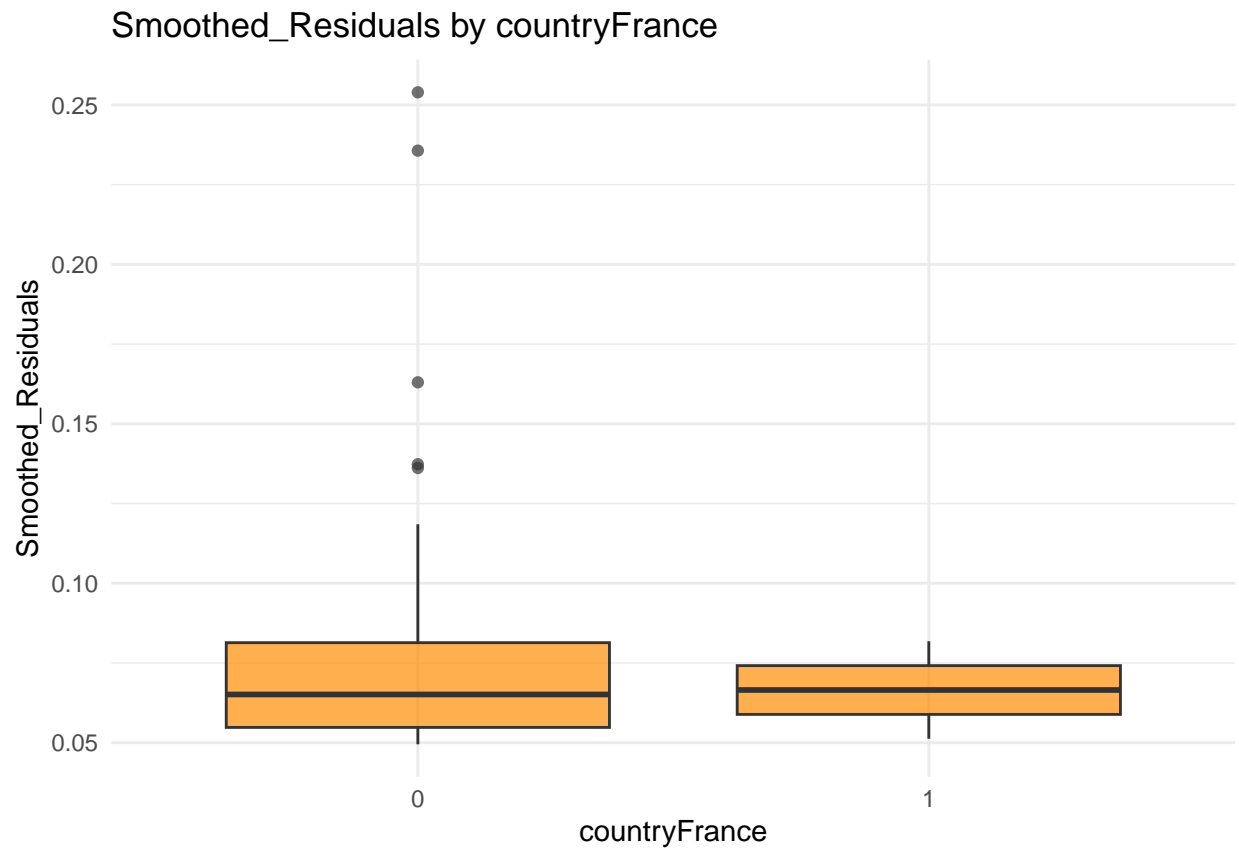


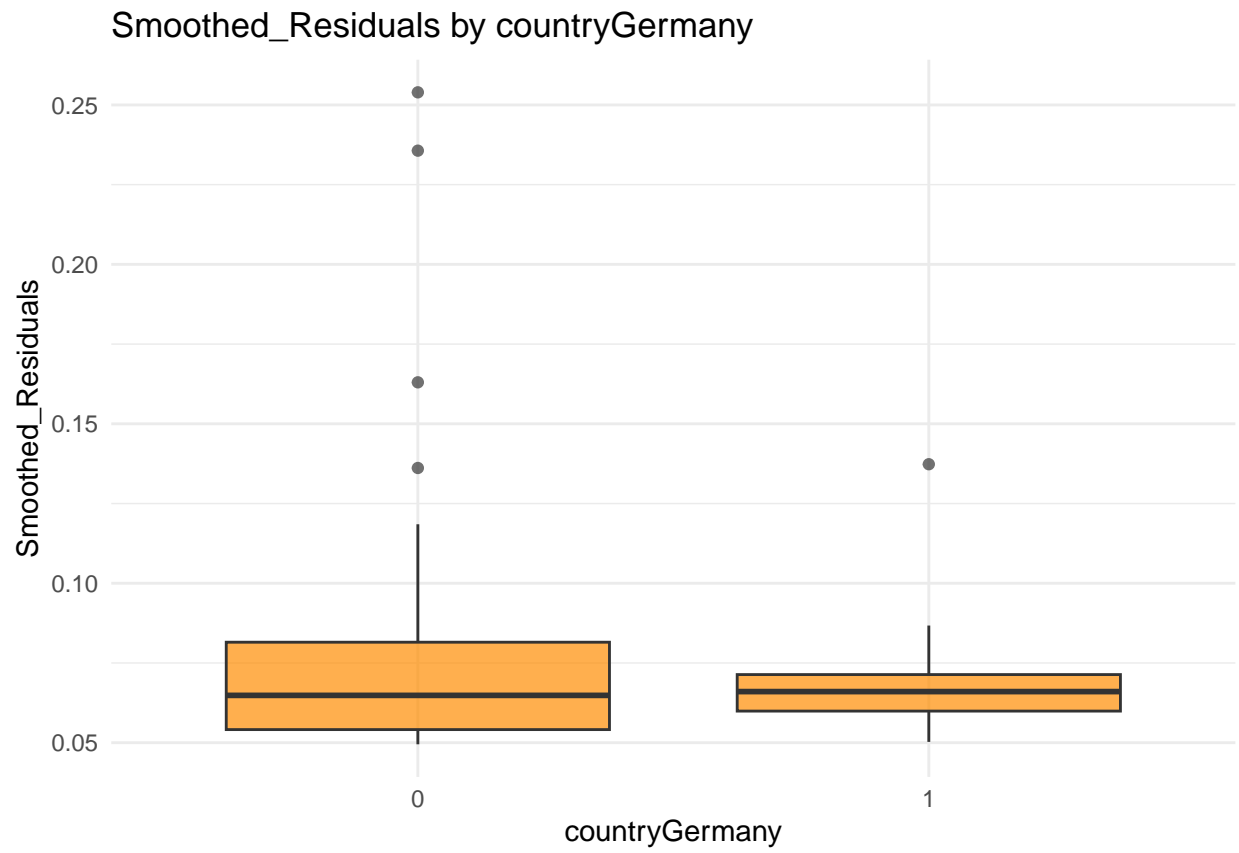


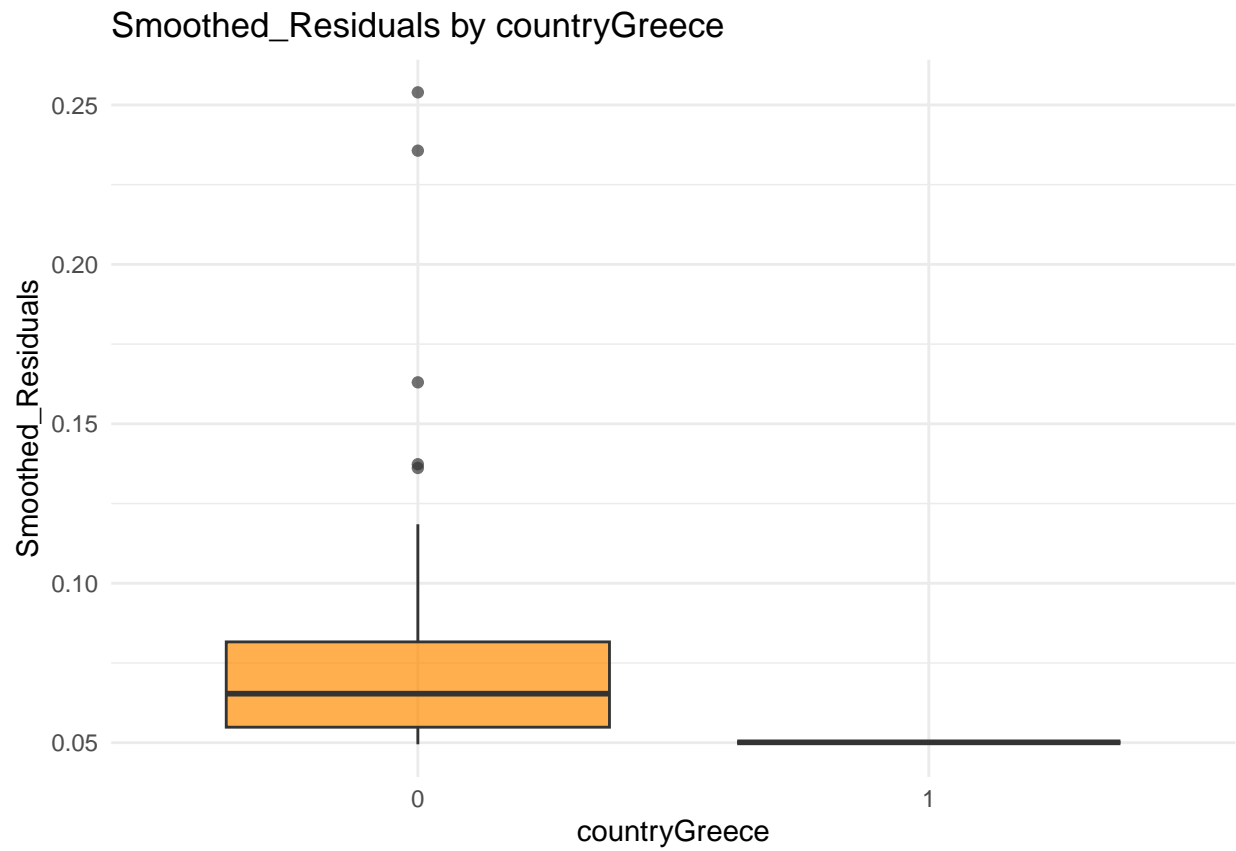


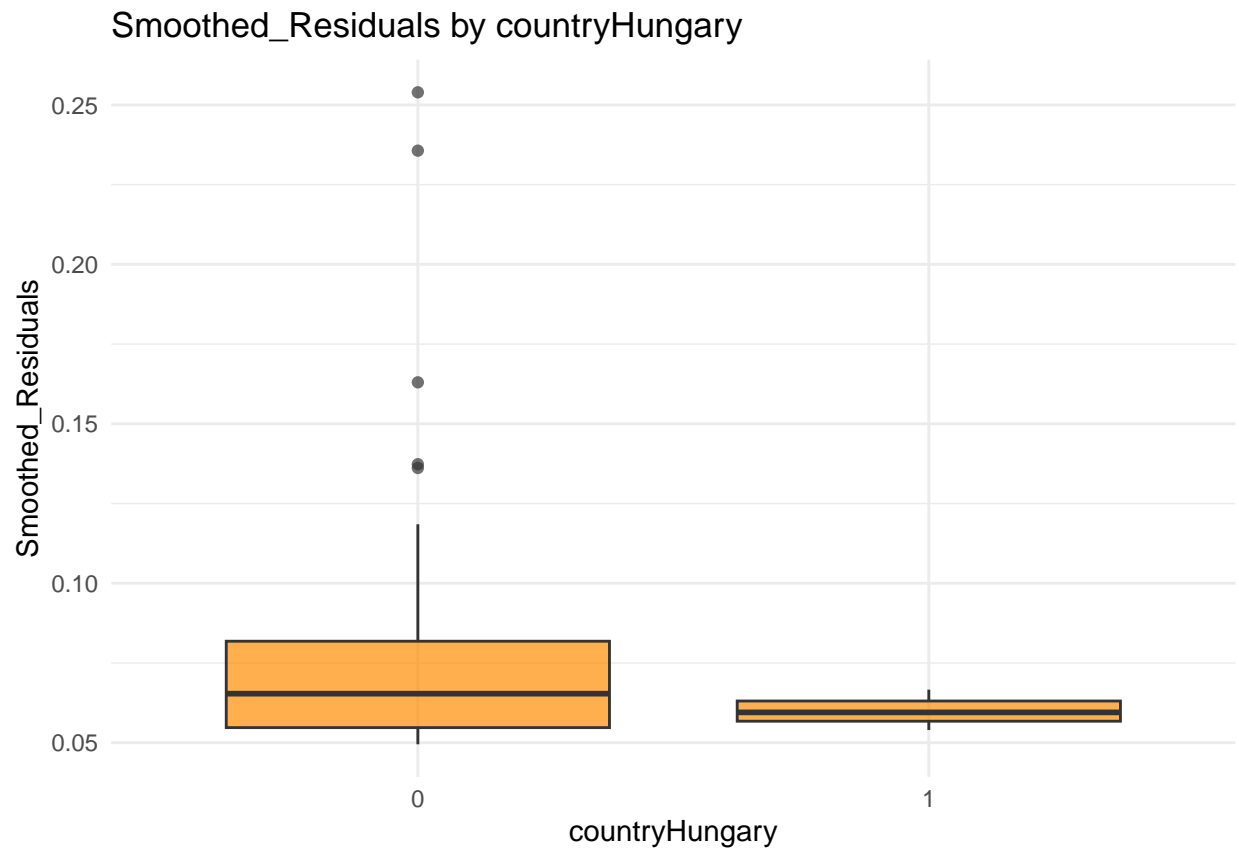


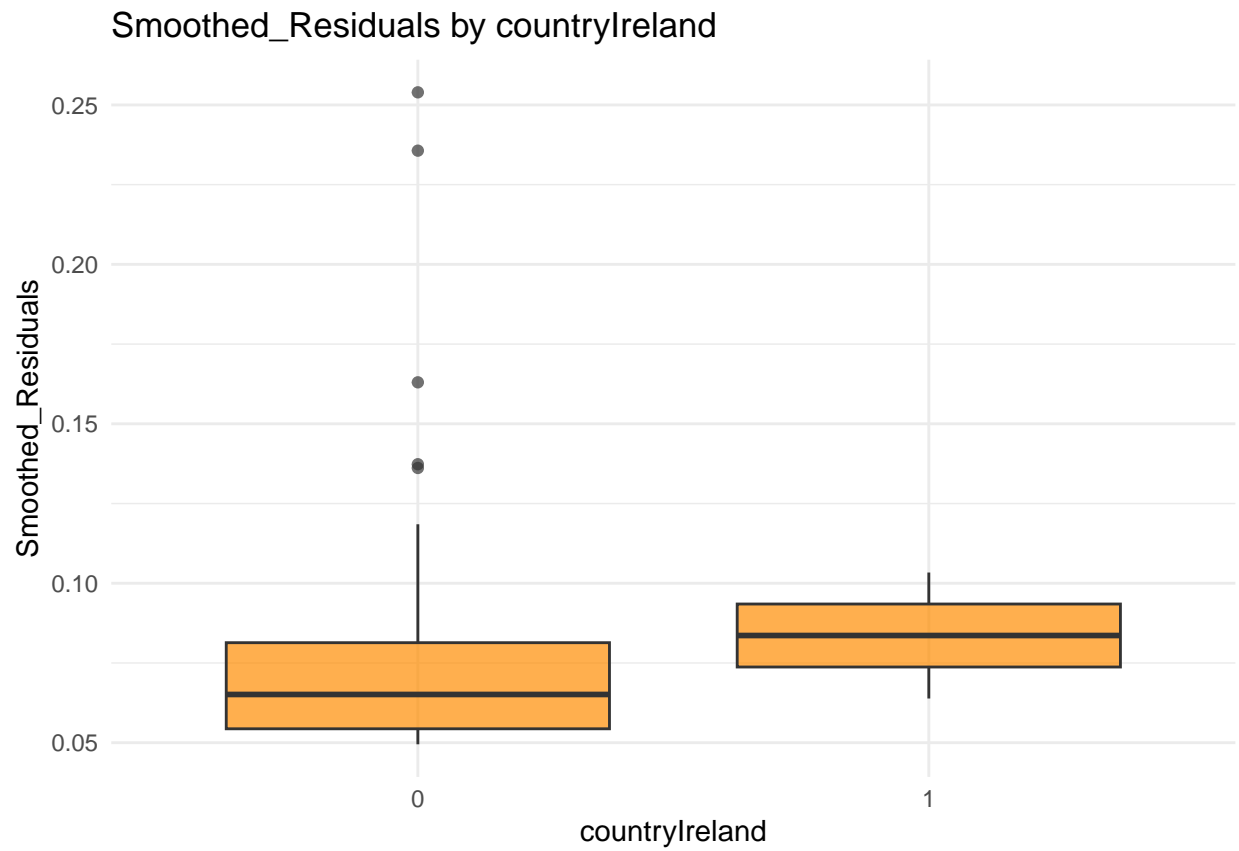


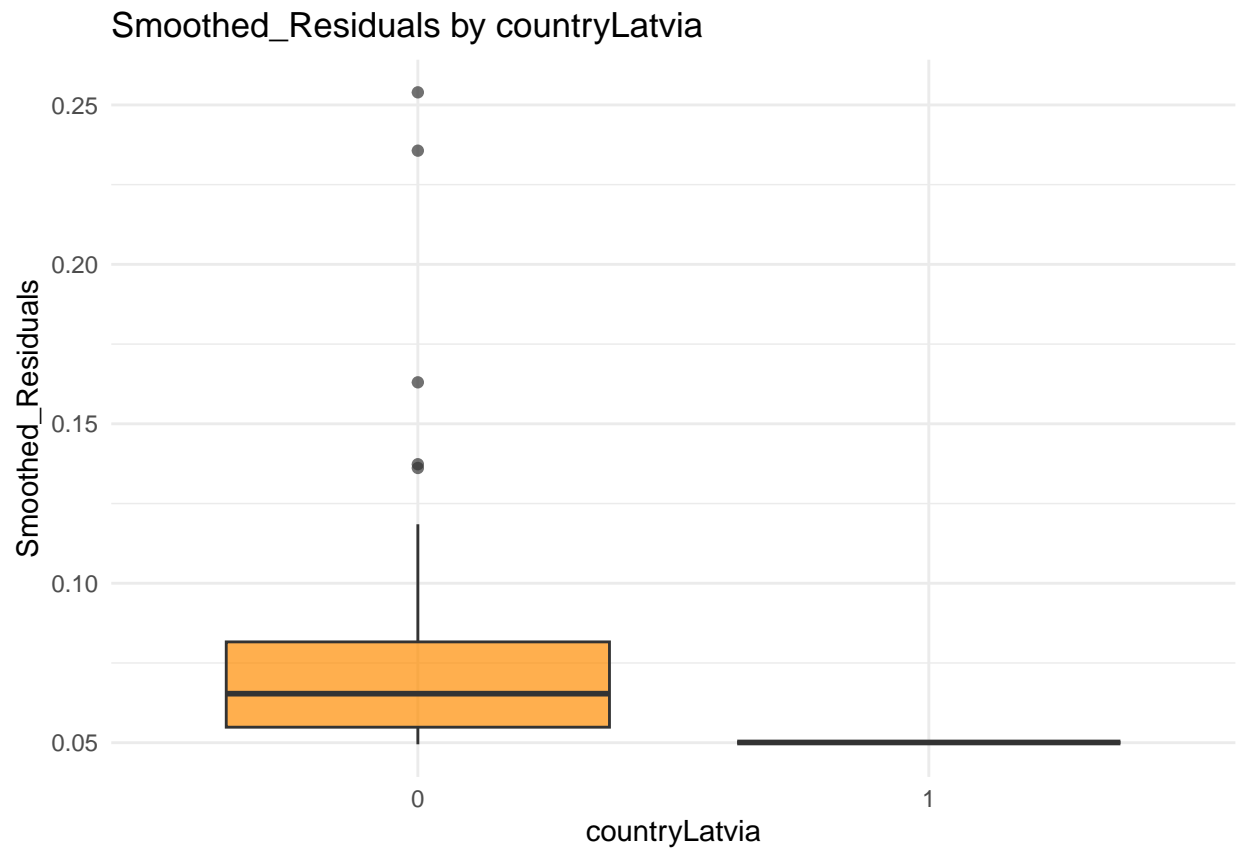


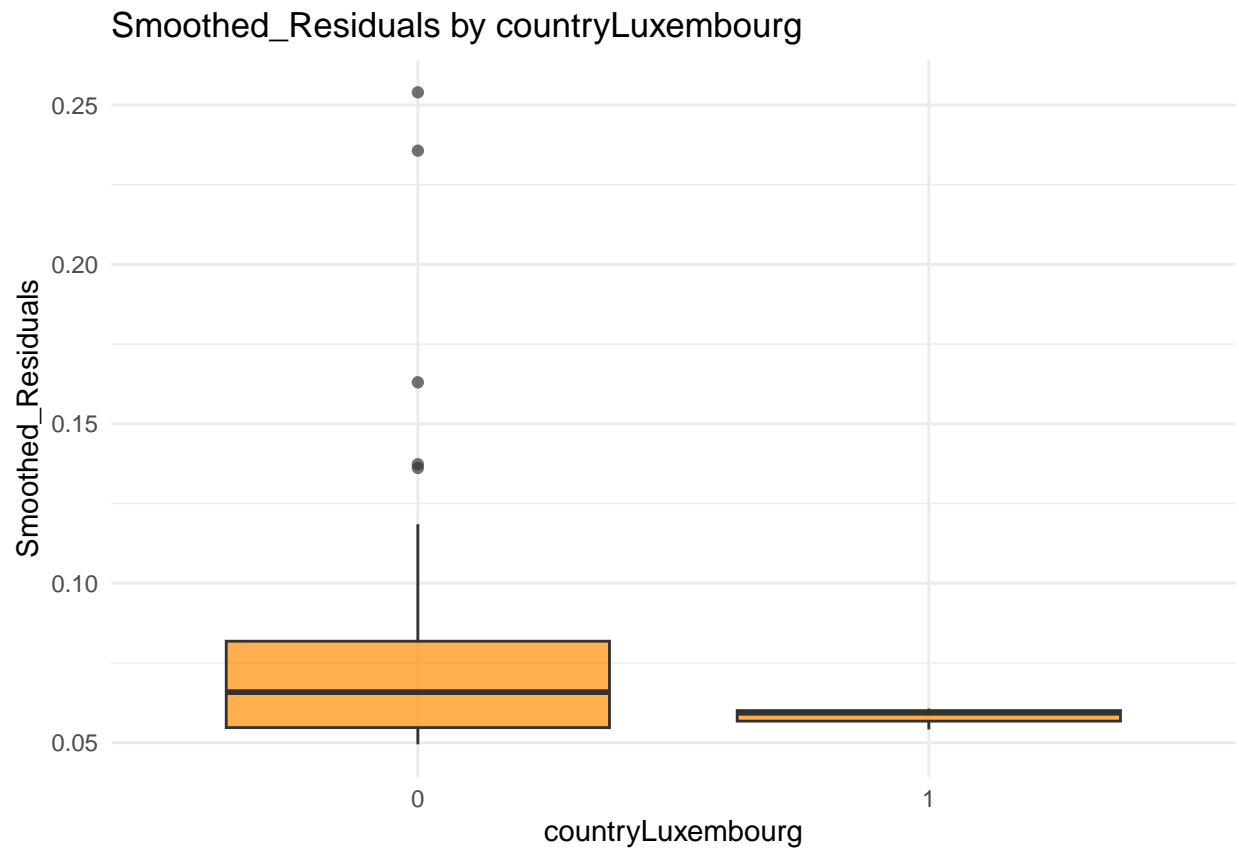


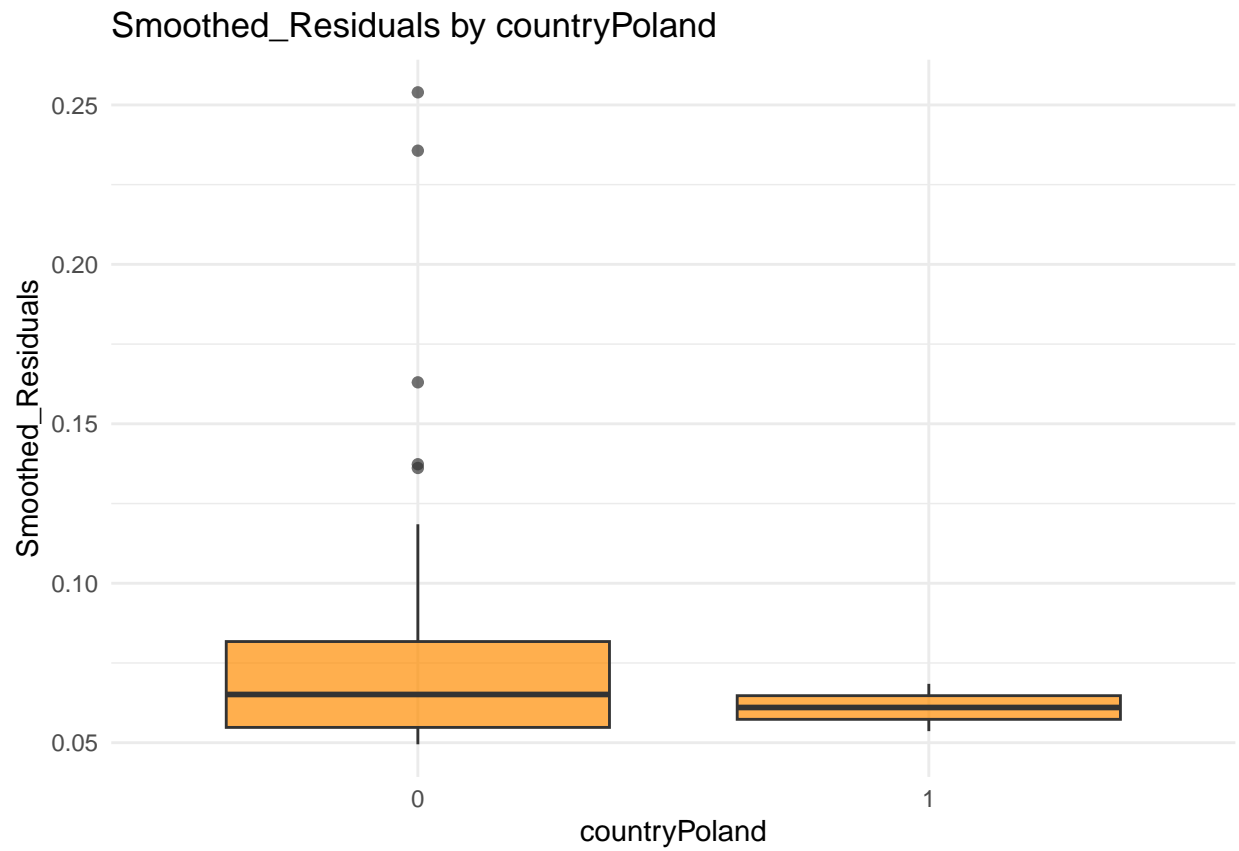


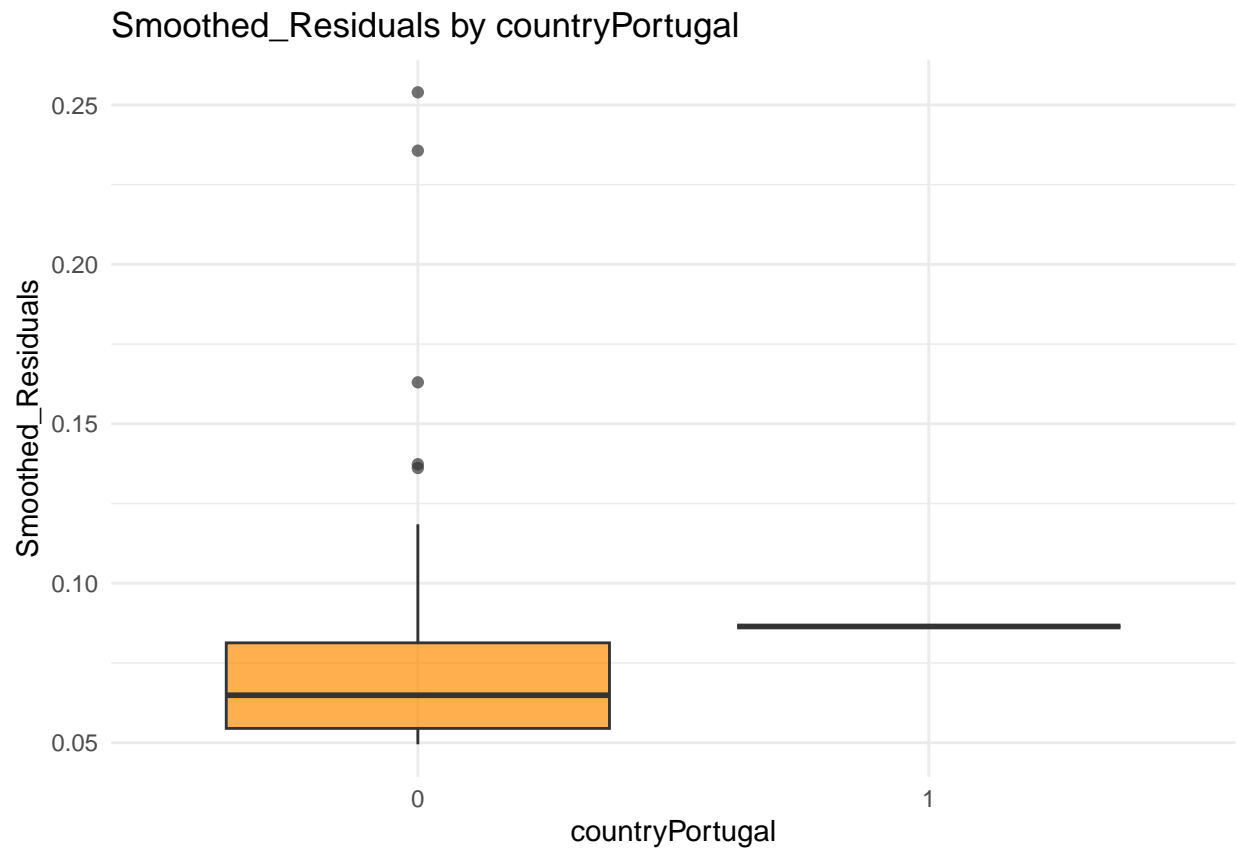


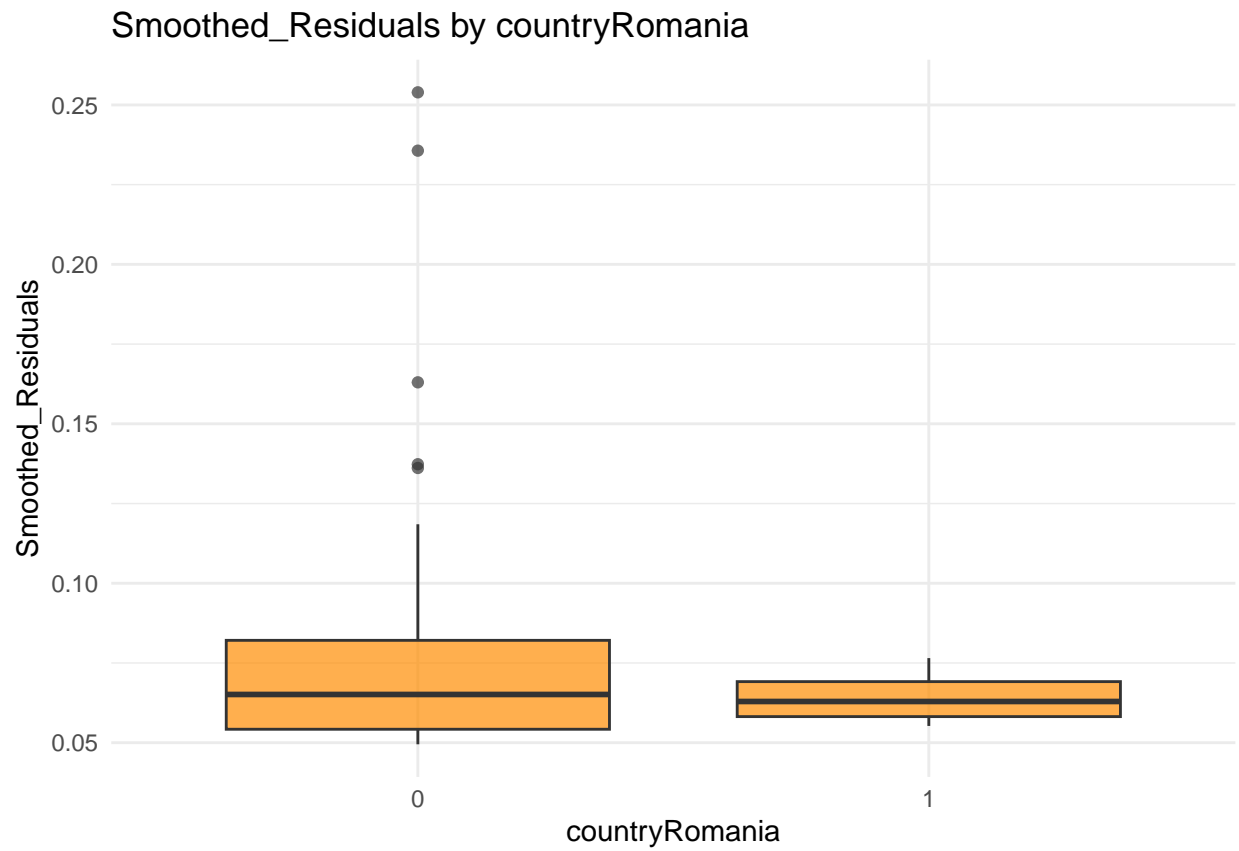


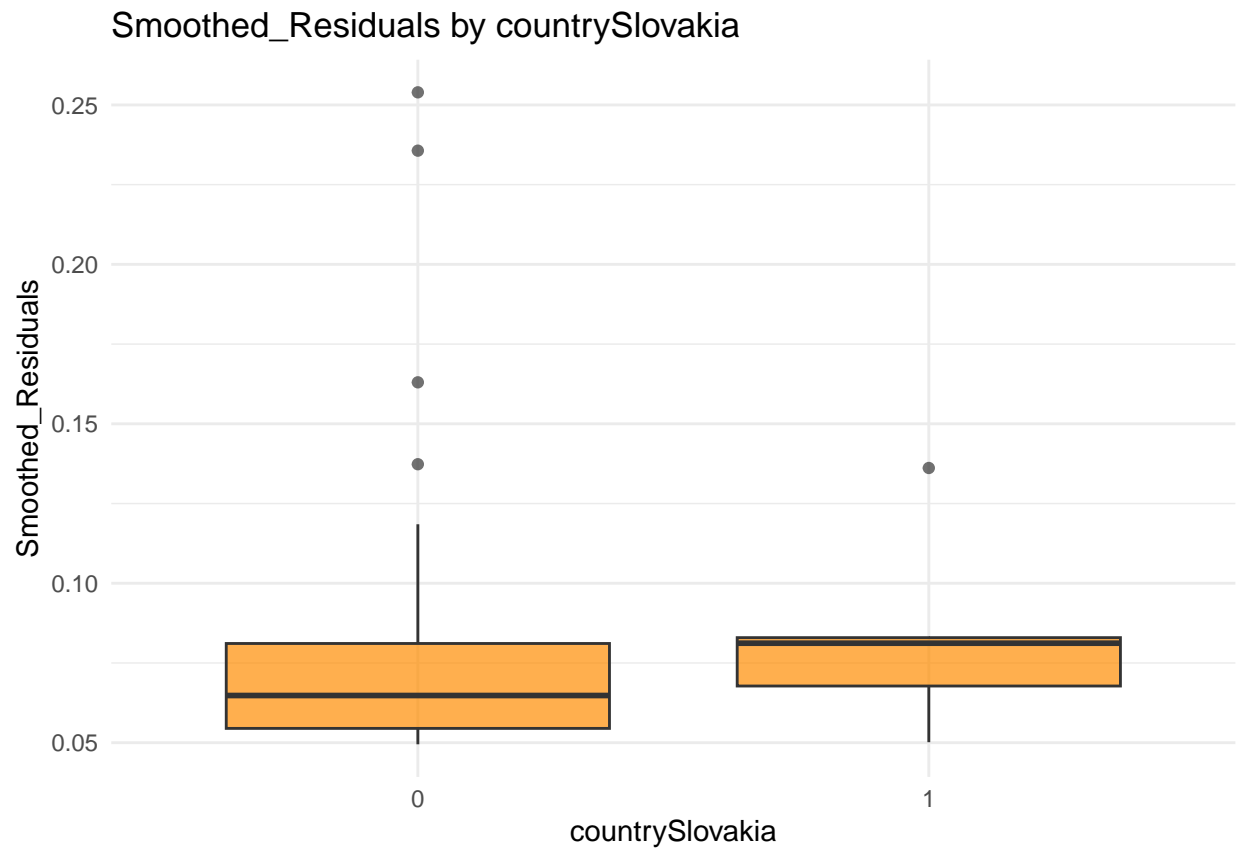


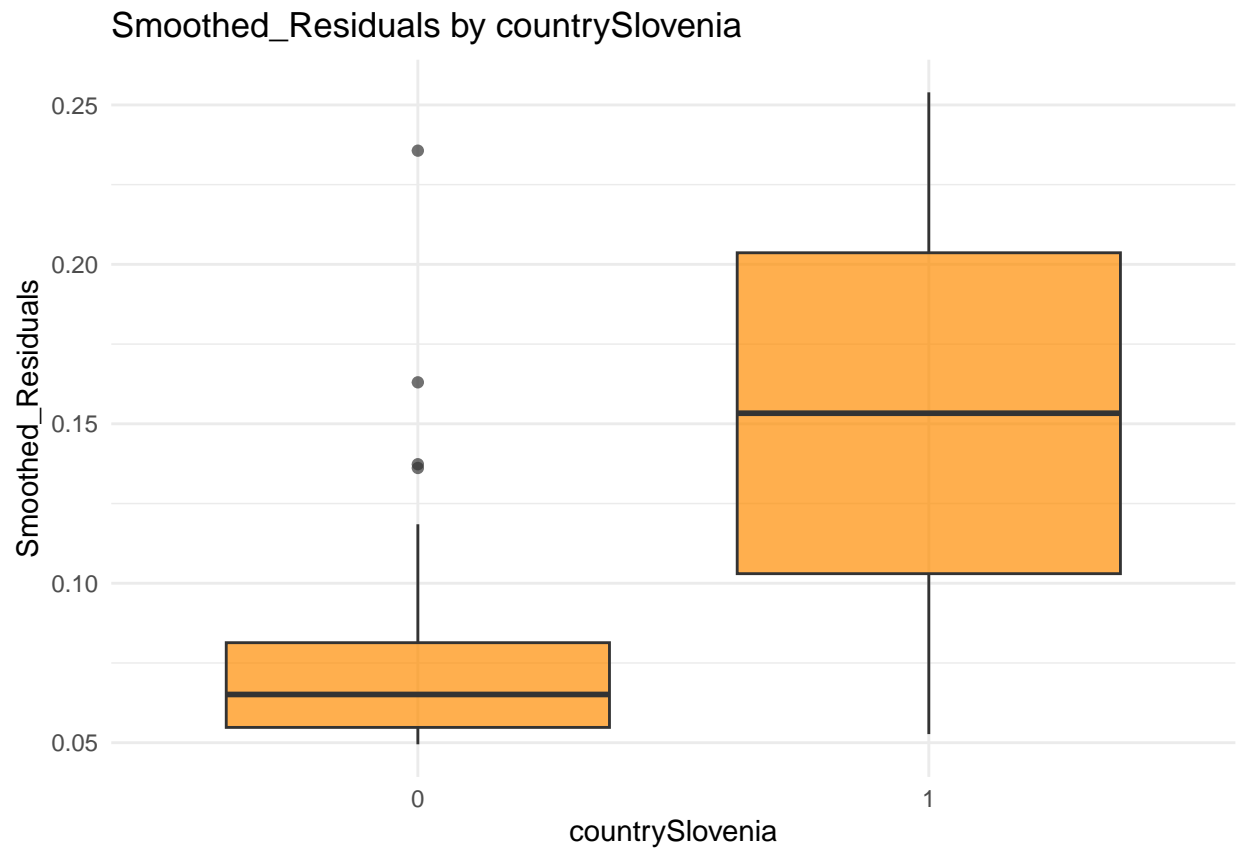


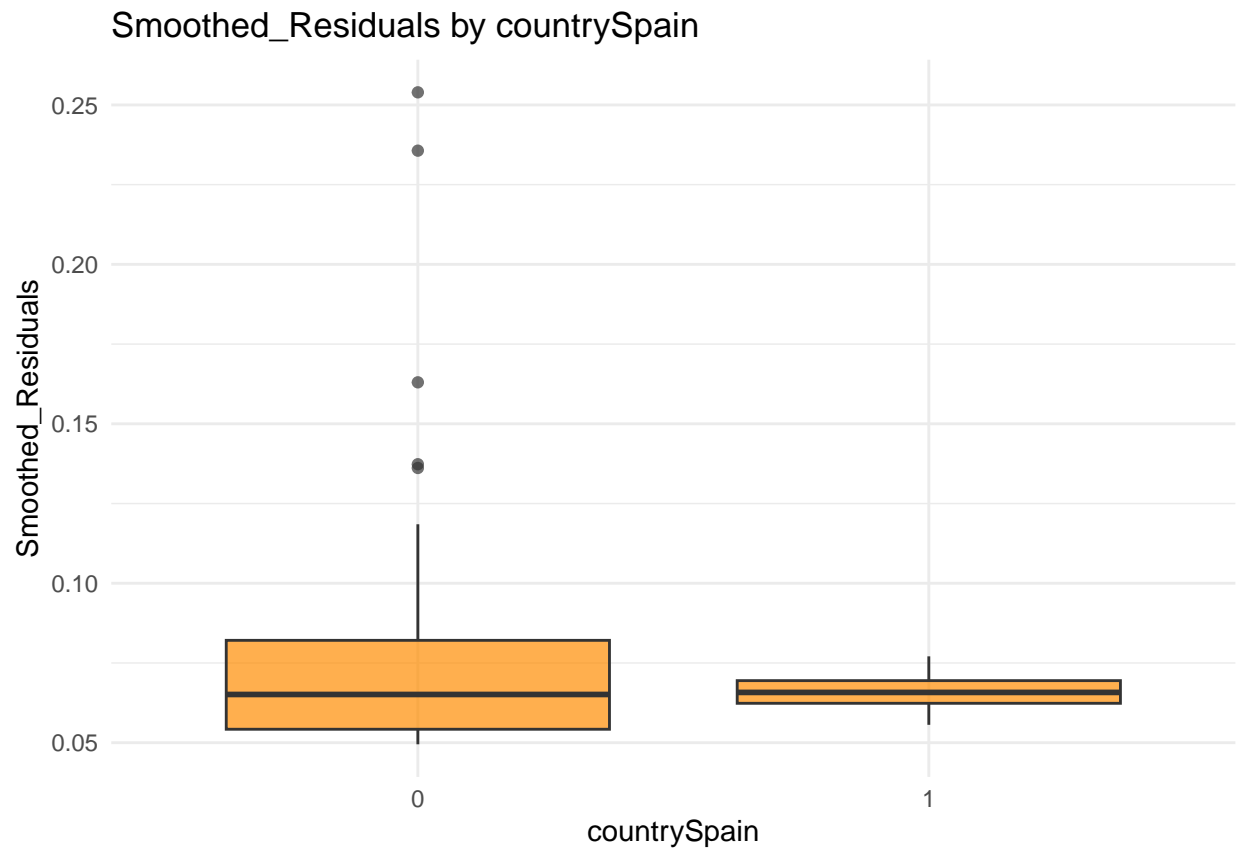


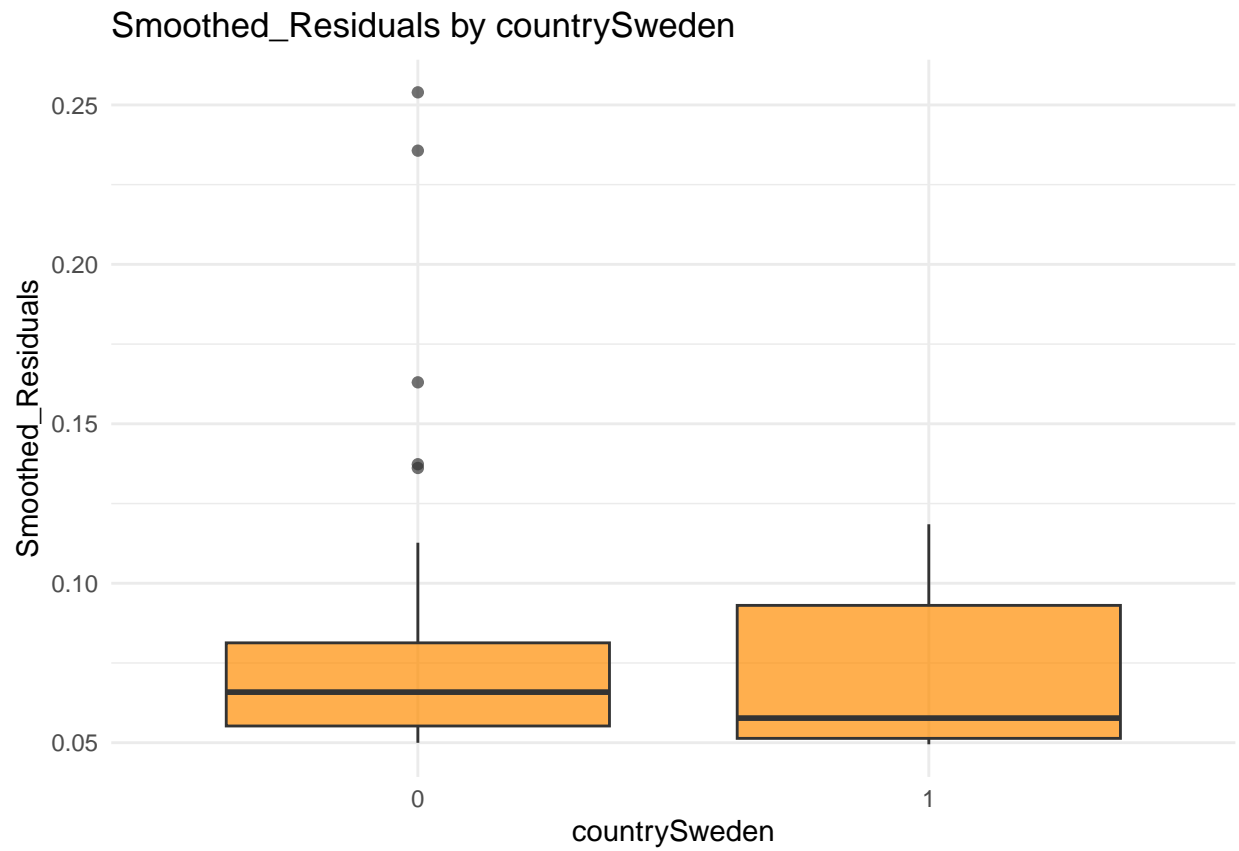


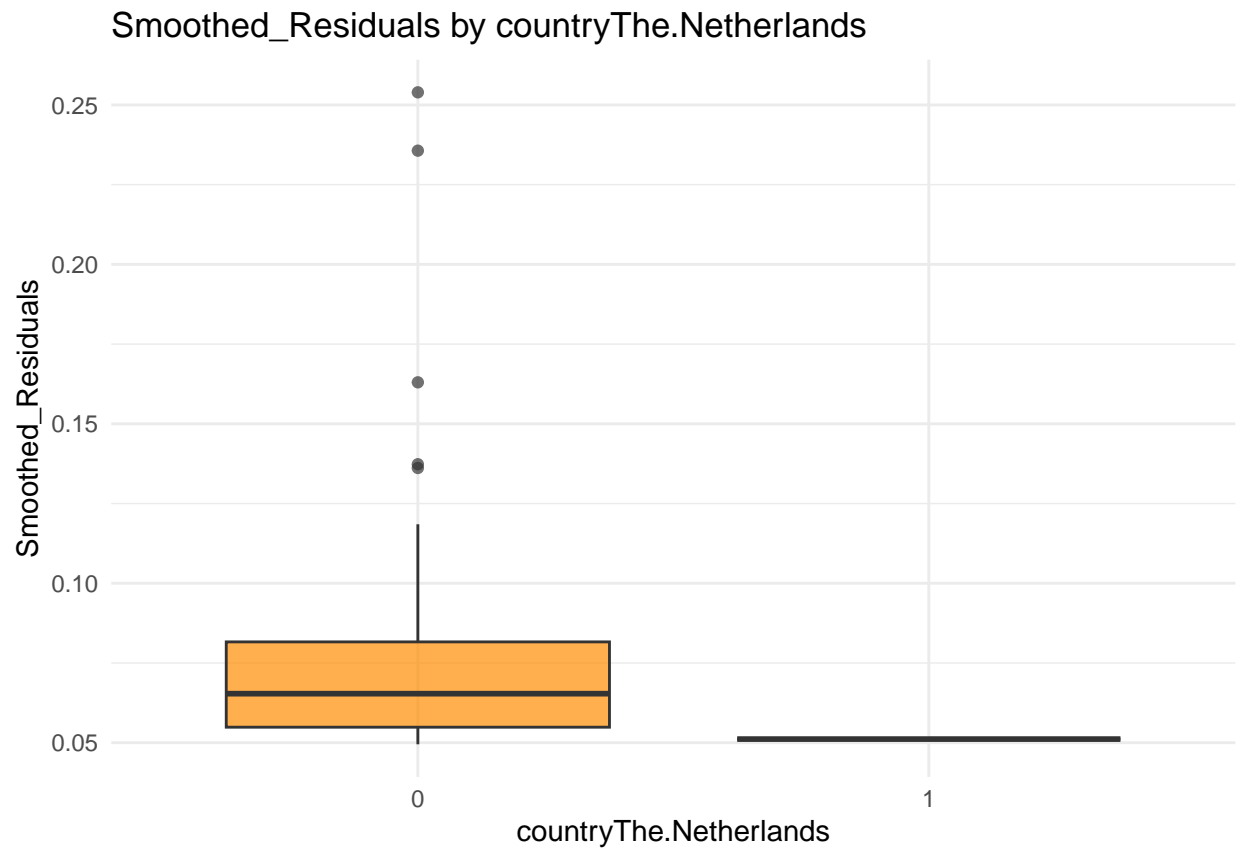


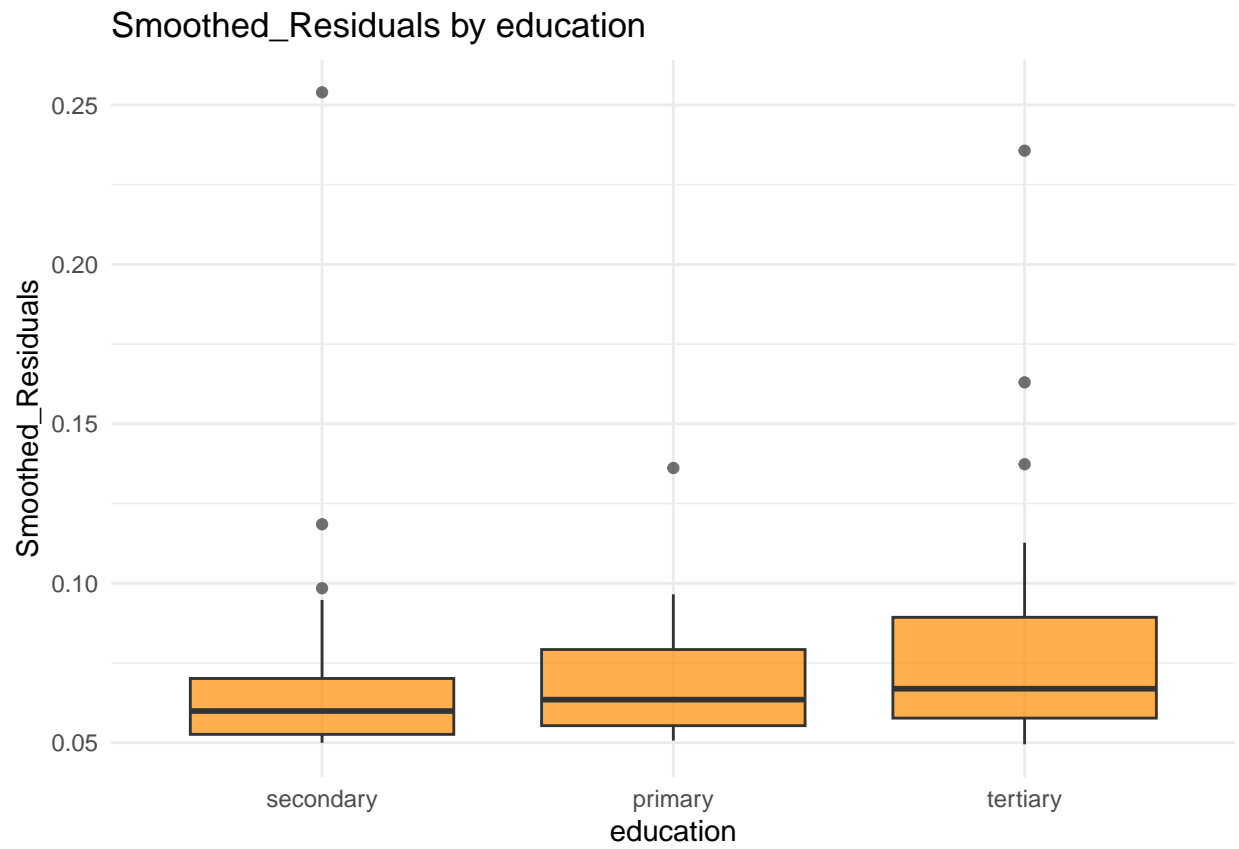


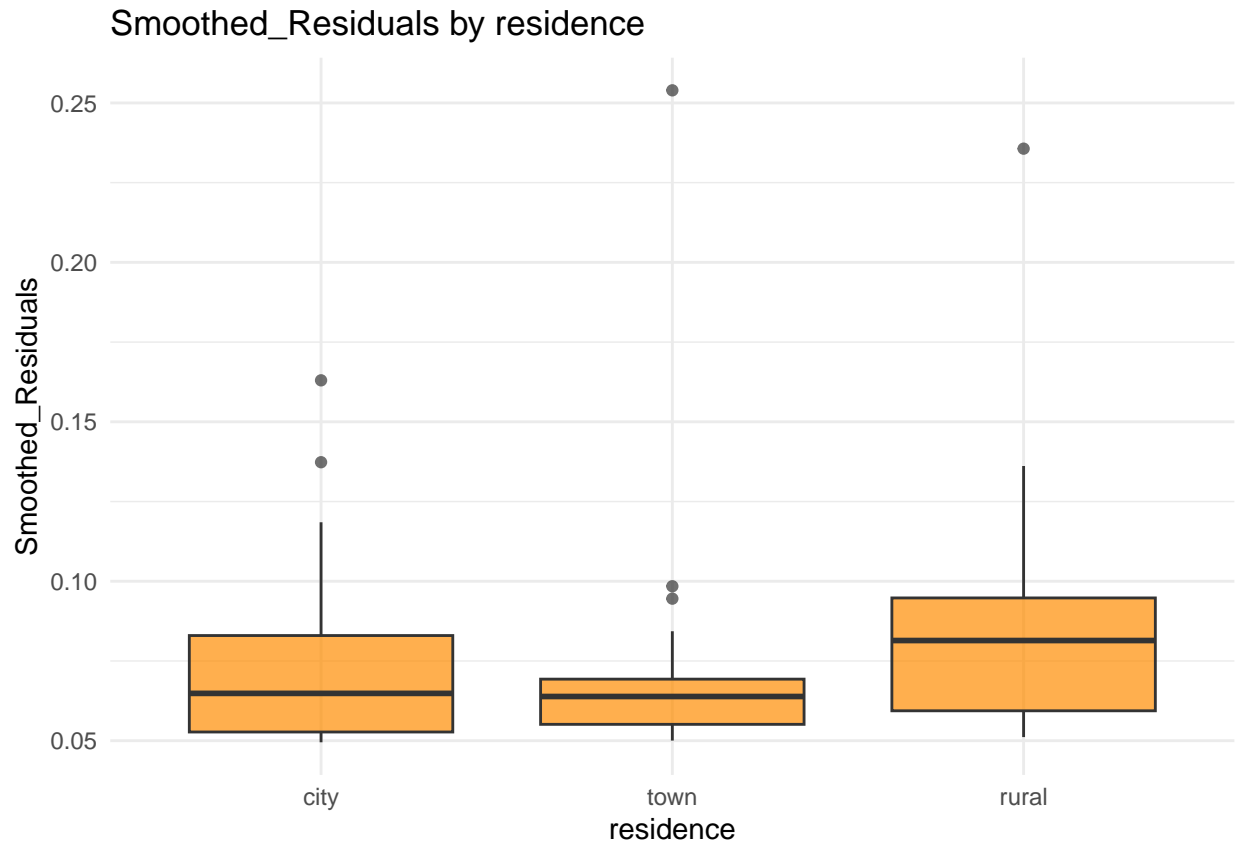












Checking the overrepresentation and underrepresentation of any group of people:

```
count_groups <- function(data) {
  for (var in names(data)) {
    x <- data[[var]]

    # Only apply to categorical, character, or binary numeric variables
    is_binary <- is.numeric(x) && length(unique(na.omit(x))) == 2
    is_categorical <- is.factor(x) || is.character(x) || is_binary

    if (is_categorical) {
      cat("-----", var, "-----\n")
      print(table(x, useNA = "ifany"))
      cat("\n")
    }
  }
}

barplot_srm <- function(data) {
  for (var in names(data)) {
    x <- data[[var]]
```

```

# Only process binary, factor, or character variables
is_binary <- is.numeric(x) && length(unique(na.omit(x))) == 2
is_categorical <- is.factor(x) || is.character(x) || is_binary

if (is_categorical) {
  df <- data.frame(group = as.factor(x))

  p <- ggplot(df, aes(x = group)) +
    geom_bar(fill = "steelblue") +
    geom_text(stat = "count", aes(label = ..count..), vjust = -0.3) +
    labs(title = paste("Counts by", var),
         x = var, y = "Count") +
    theme_minimal()

  print(p)
}
}

```

```
count_groups(srm_final)
```

```

## ---- ctax_binary ----
## x
## 0 1
## 44 56
##
## ---- trust ----
## x
## No confident at all No really confident Rather confident Very confident
## 34 18 32 16
##
## ---- any_cc_last2year_factor ----
## x
## 0 1
## 47 53
##
## ---- regional_heterogeneity ----
## x
## 0 1
## 19 81
##
## ---- gender ----
## x
## female male
## 39 61
##
## ---- has_childrenyes ----
## x
## 0 1
## 58 42
##
## ---- countryBelgium ----
## x

```

```

## 0 1
## 97 3
##
## ---- countryBulgaria ----
## x
## 0 1
## 99 1
##
## ---- countryCroatia ----
## x
## 0 1
## 93 7
##
## ---- countryCyprus ----
## x
## 0 1
## 99 1
##
## ---- countryCzech.Republic ----
## x
## 0 1
## 99 1
##
## ---- countryDenmark ----
## x
## 0 1
## 91 9
##
## ---- countryFinland ----
## x
## 0 1
## 94 6
##
## ---- countryFrance ----
## x
## 0 1
## 98 2
##
## ---- countryGermany ----
## x
## 0 1
## 88 12
##
## ---- countryGreece ----
## x
## 0 1
## 99 1
##
## ---- countryHungary ----
## x
## 0 1
## 97 3
##
## ---- countryIreland ----

```



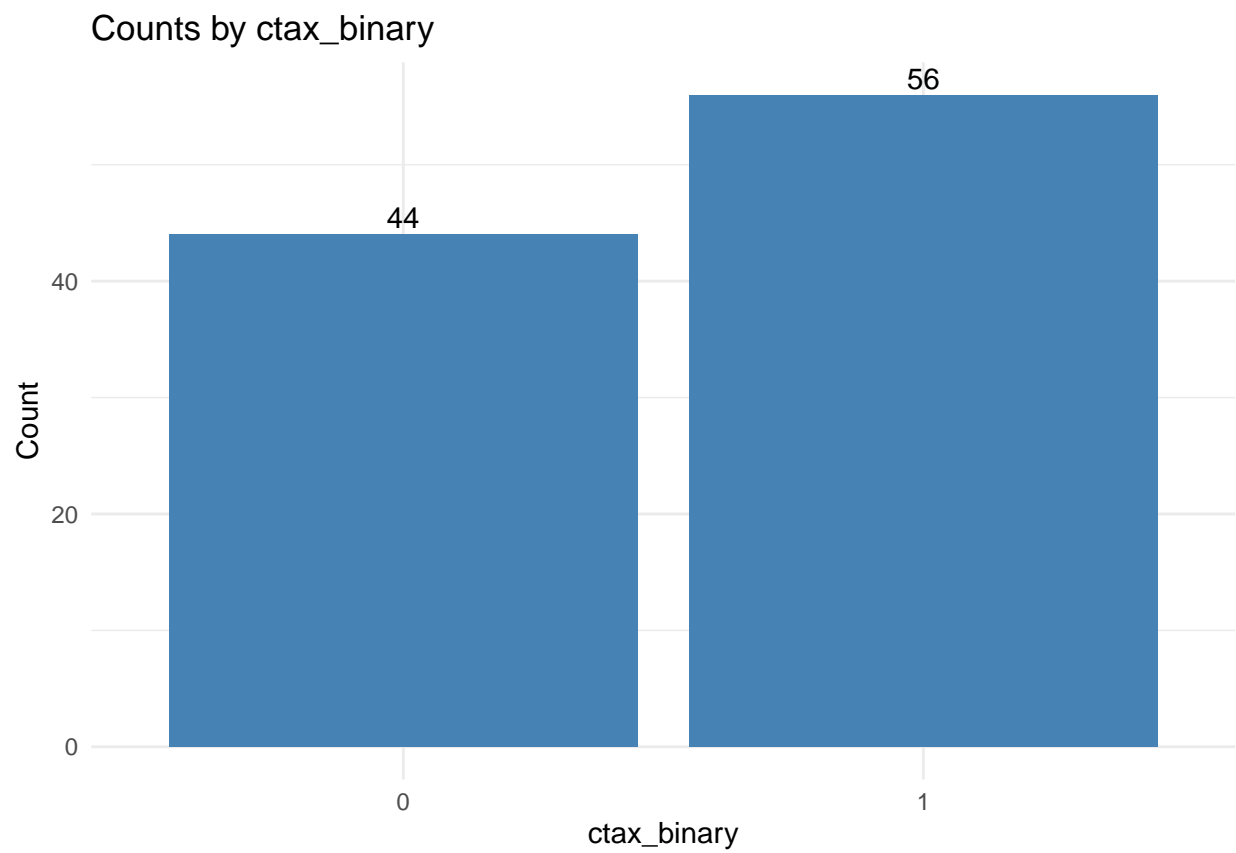
```

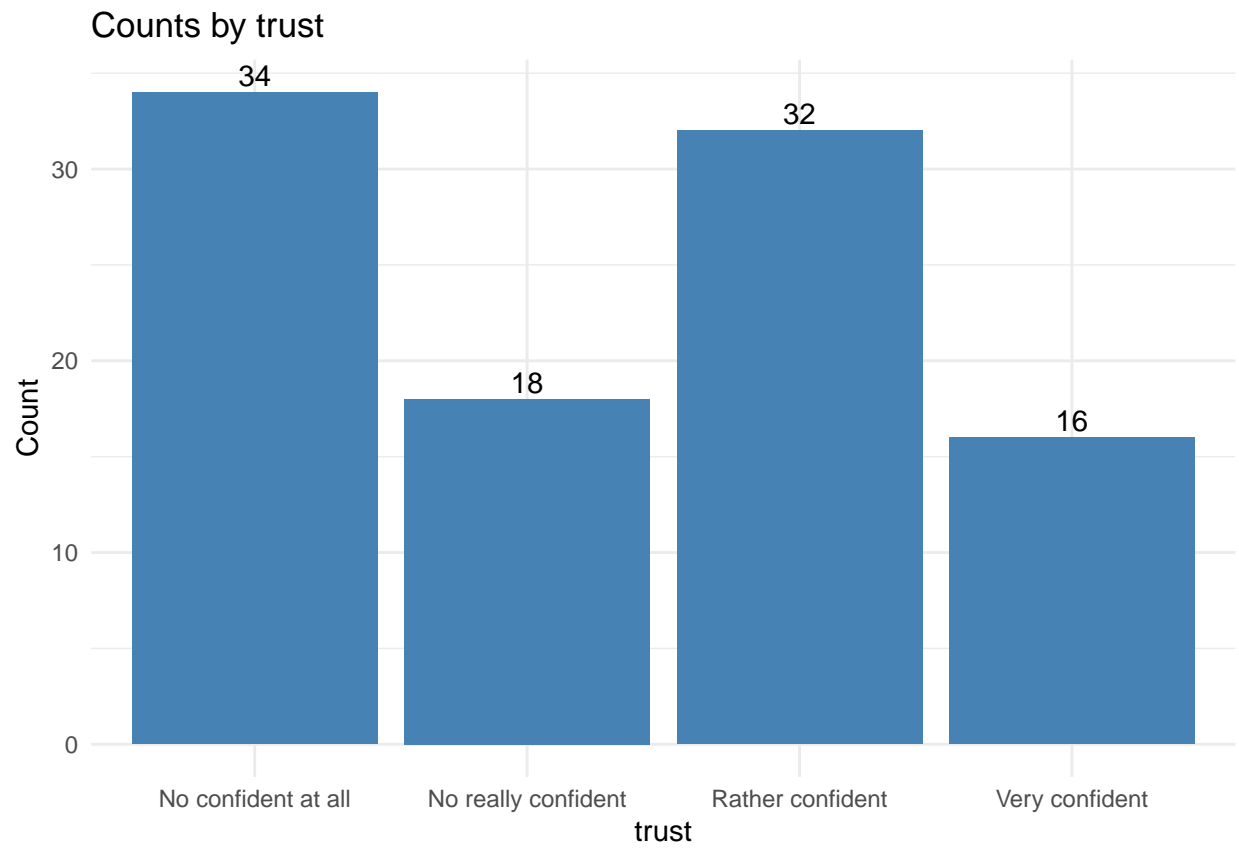
## x
## 0 1
## 98 2
##
## ---- countryLatvia ----
## x
## 0 1
## 99 1
##
## ---- countryLuxembourg ----
## x
## 0 1
## 97 3
##
## ---- countryMalta ----
## x
## 0 1
## 99 1
##
## ---- countryPoland ----
## x
## 0 1
## 98 2
##
## ---- countryPortugal ----
## x
## 0 1
## 99 1
##
## ---- countryRomania ----
## x
## 0 1
## 96 4
##
## ---- countrySlovakia ----
## x
## 0 1
## 95 5
##
## ---- countrySlovenia ----
## x
## 0 1
## 98 2
##
## ---- countrySpain ----
## x
## 0 1
## 96 4
##
## ---- countrySweden ----
## x
## 0 1
## 87 13
##

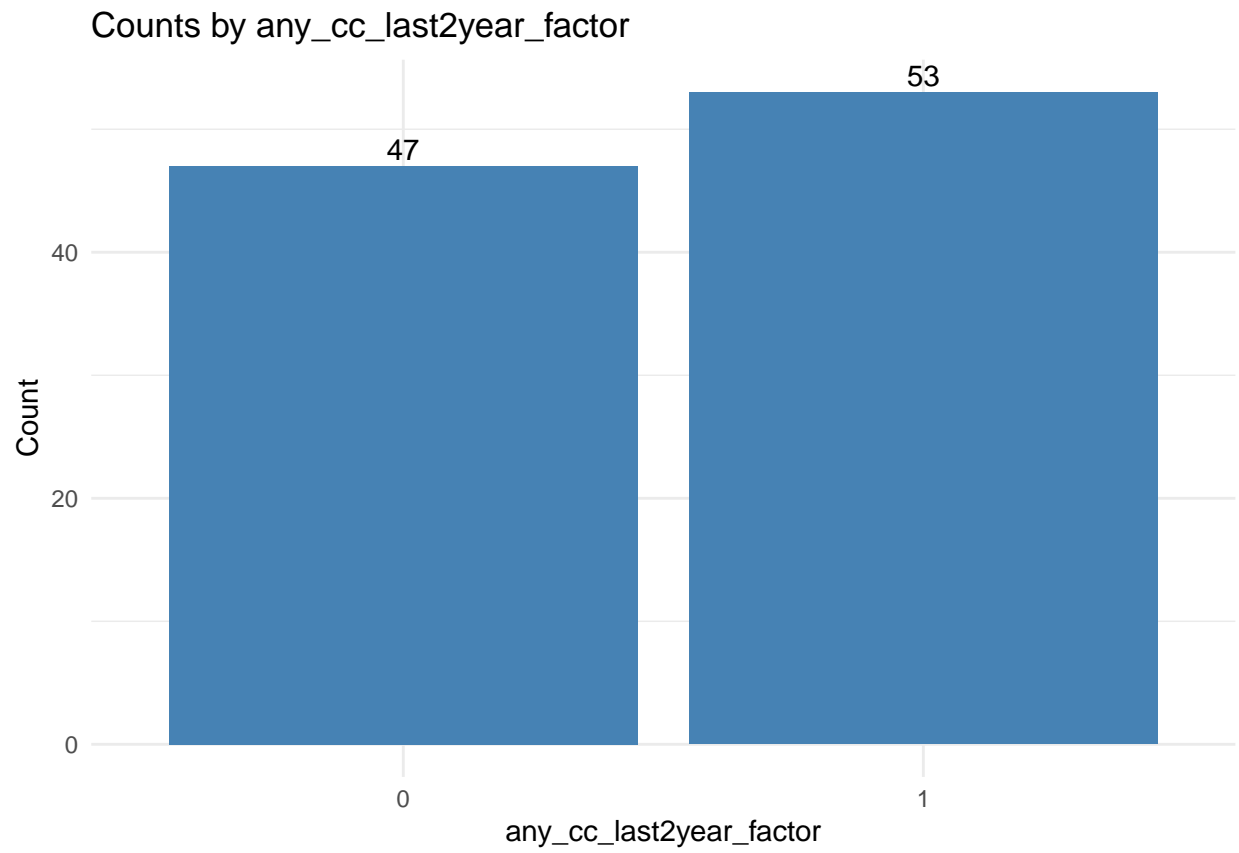
```

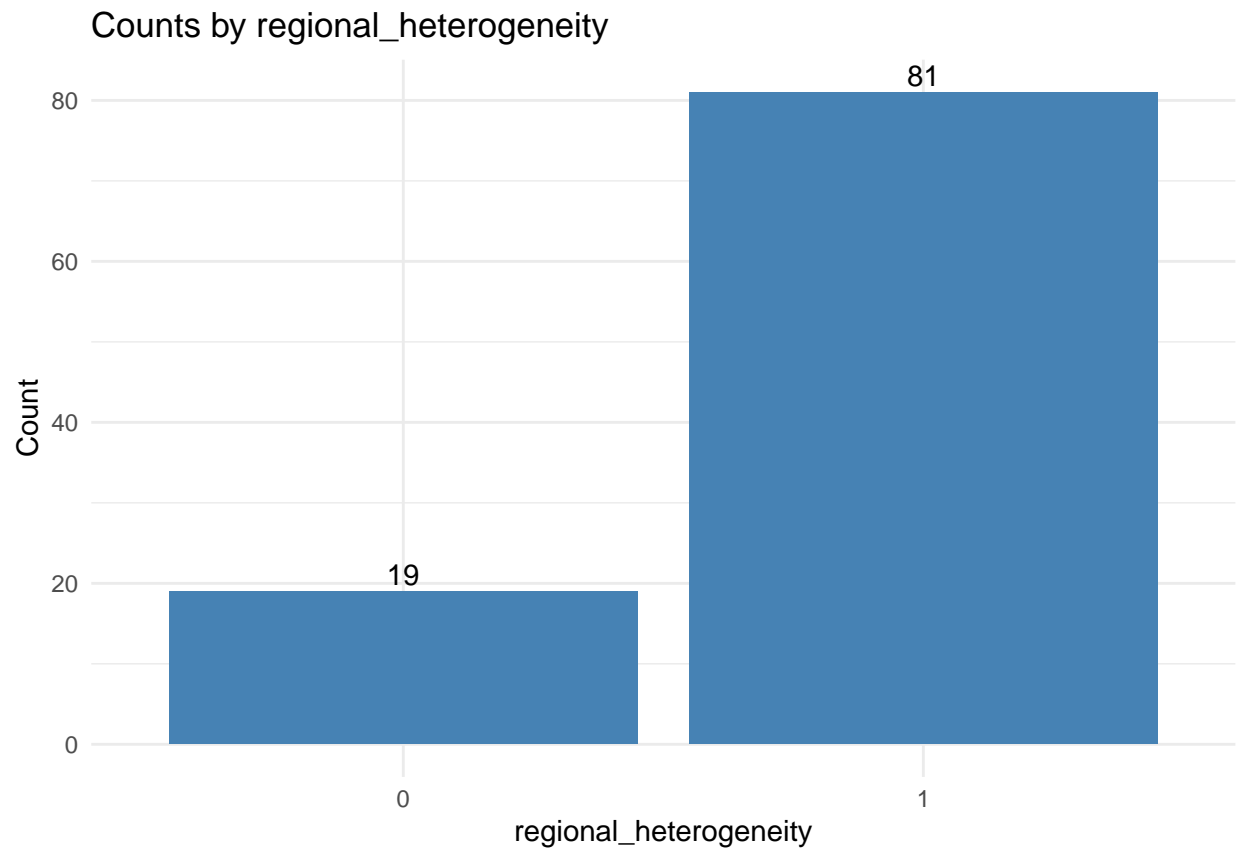
```
## ---- countryThe.Netherlands ----
## x
## 0 1
## 99 1
##
## ---- education ----
## x
## secondary primary tertiary
## 41 12 47
##
## ---- residence ----
## x
## city town rural
## 48 35 17
```

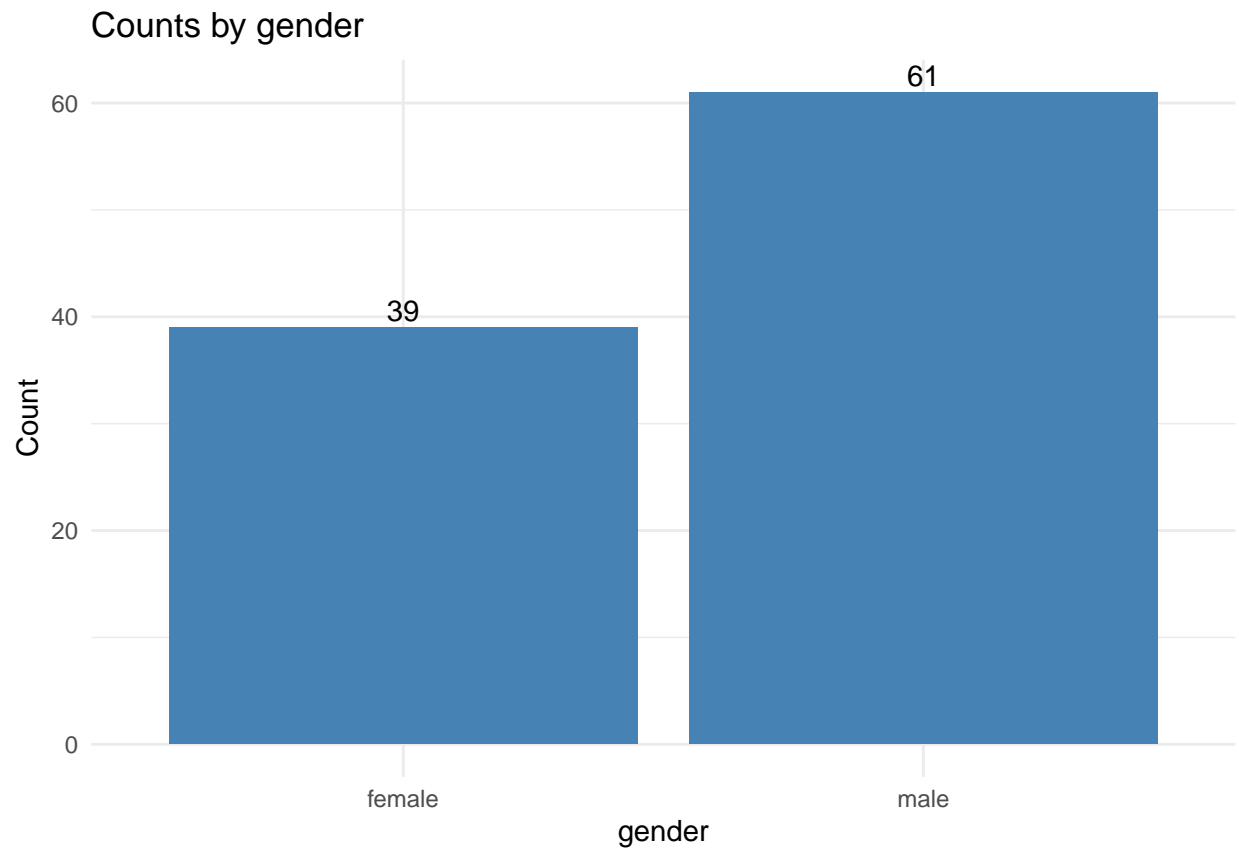
```
barplot_srm(srm_final)
```

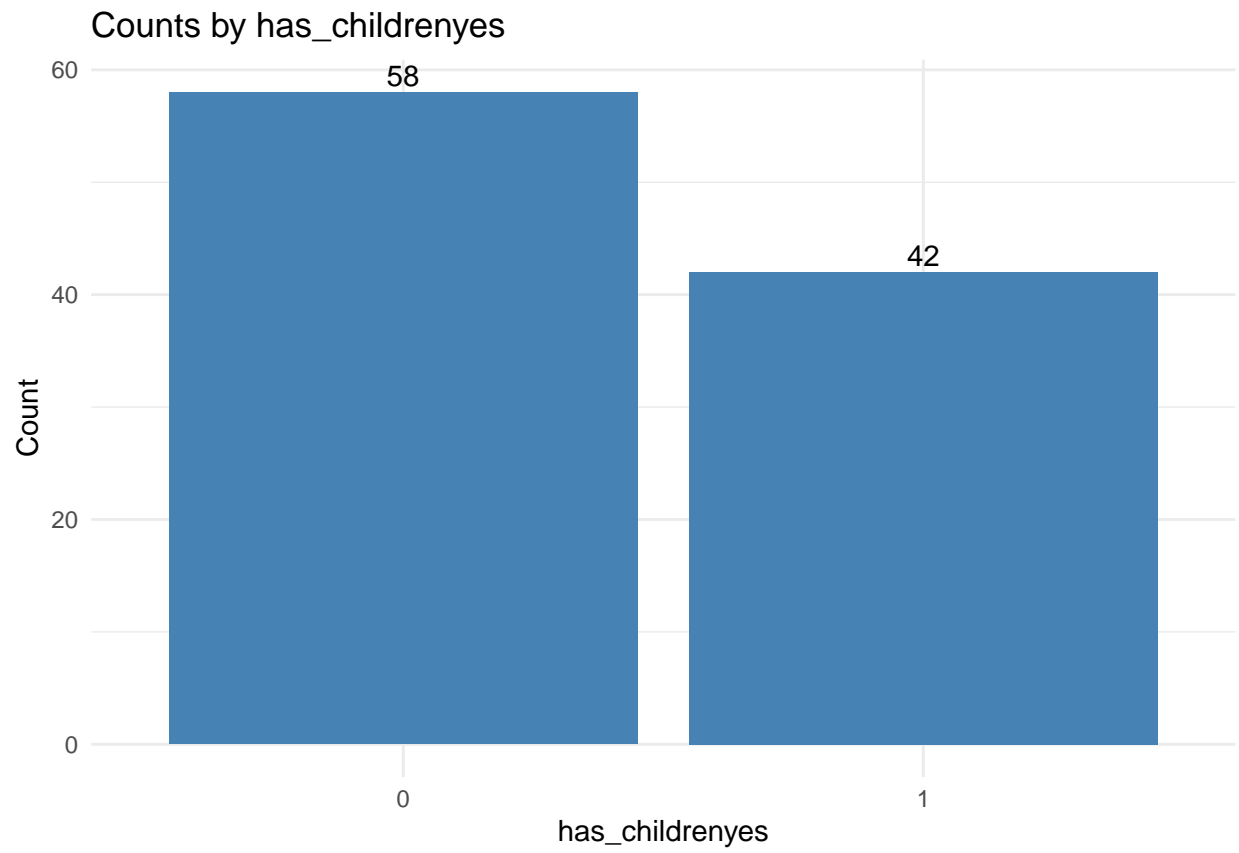


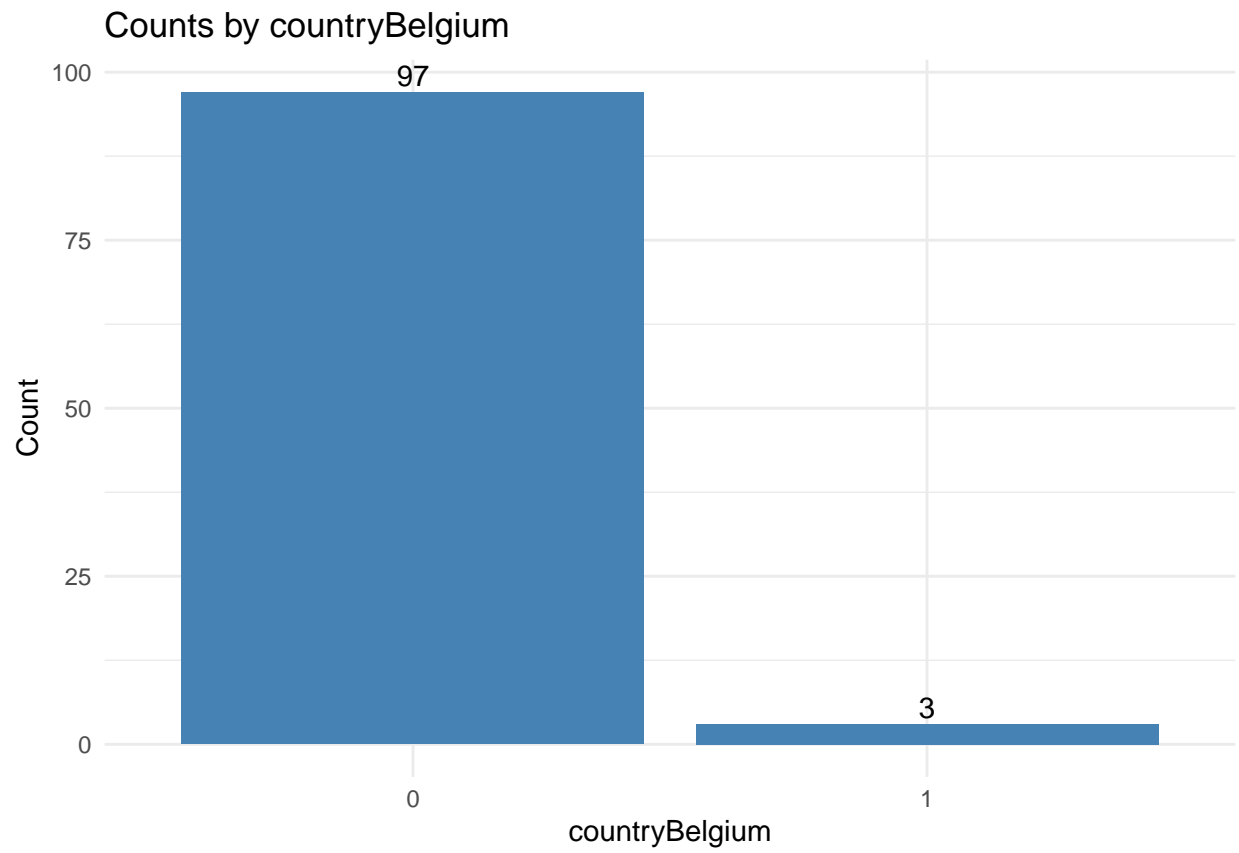


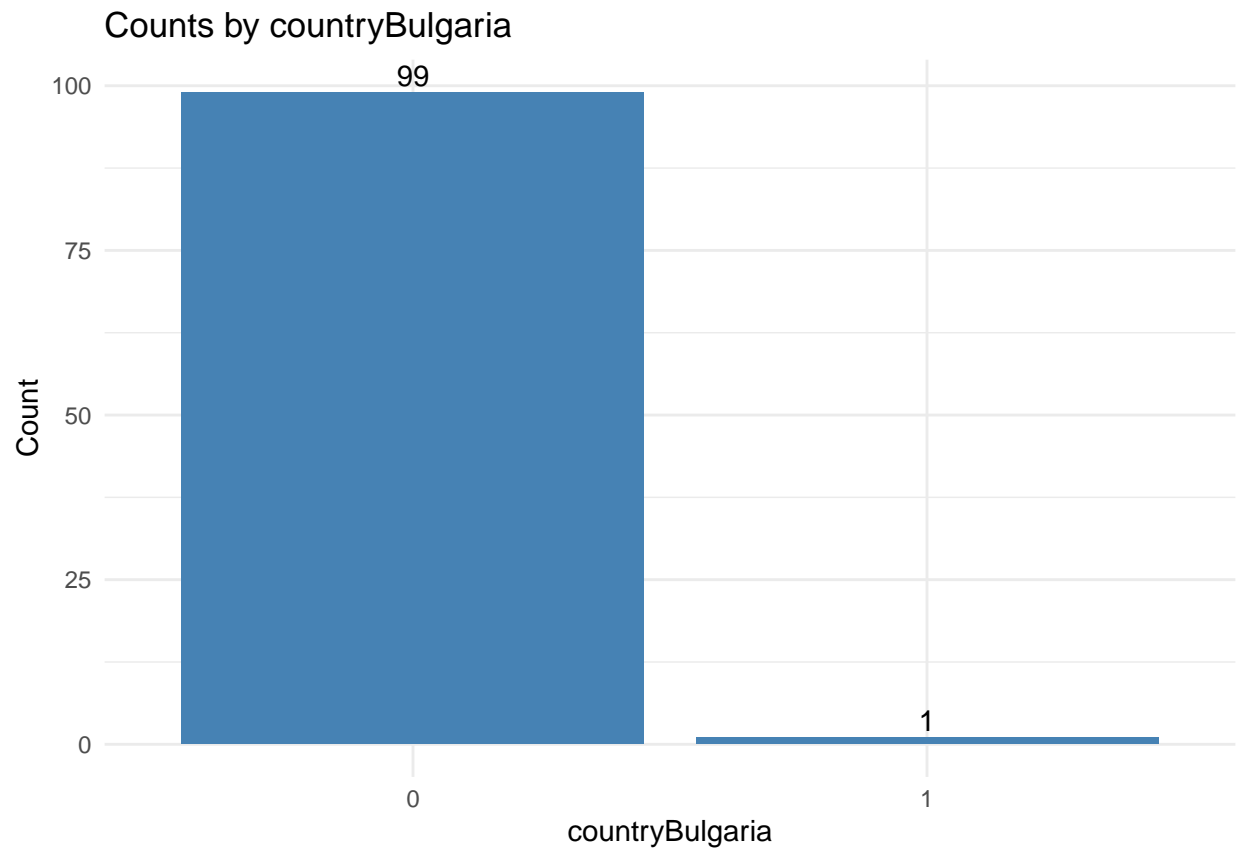


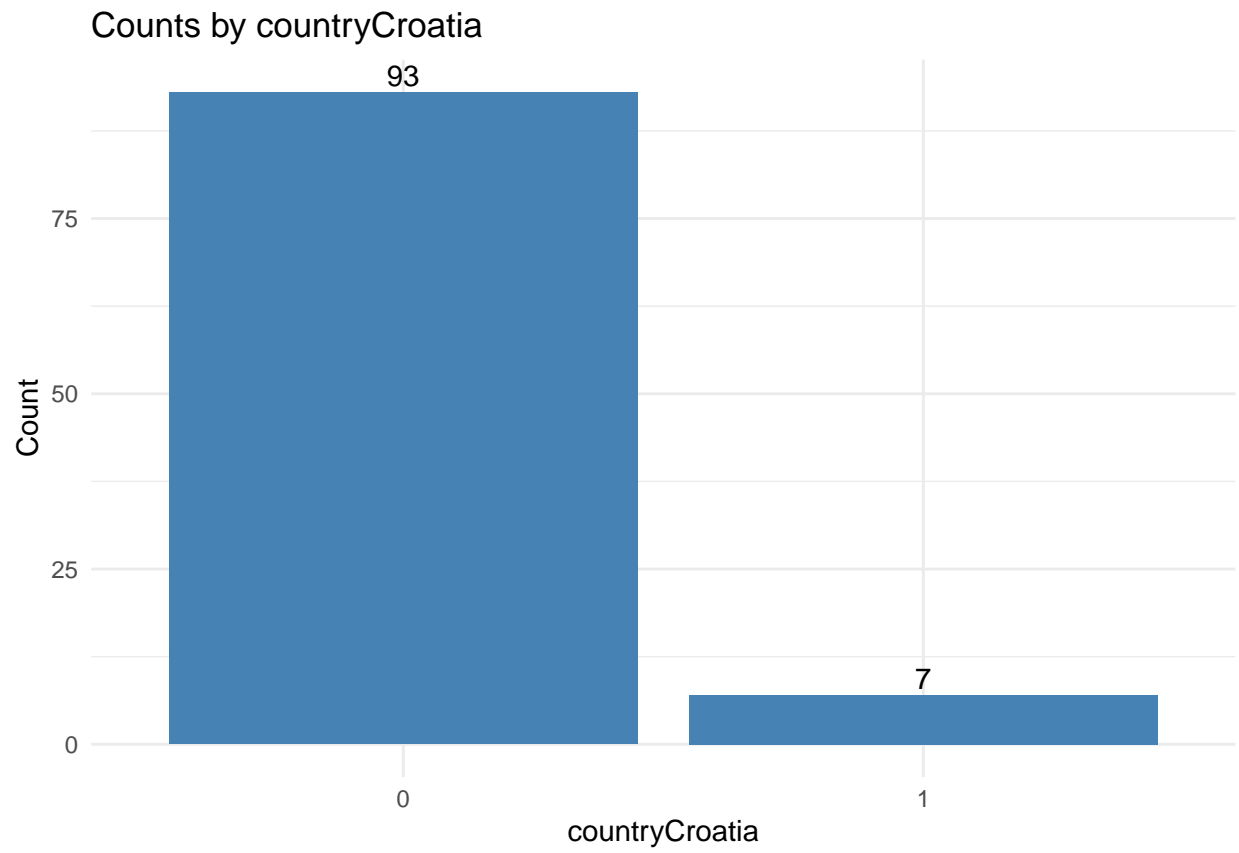


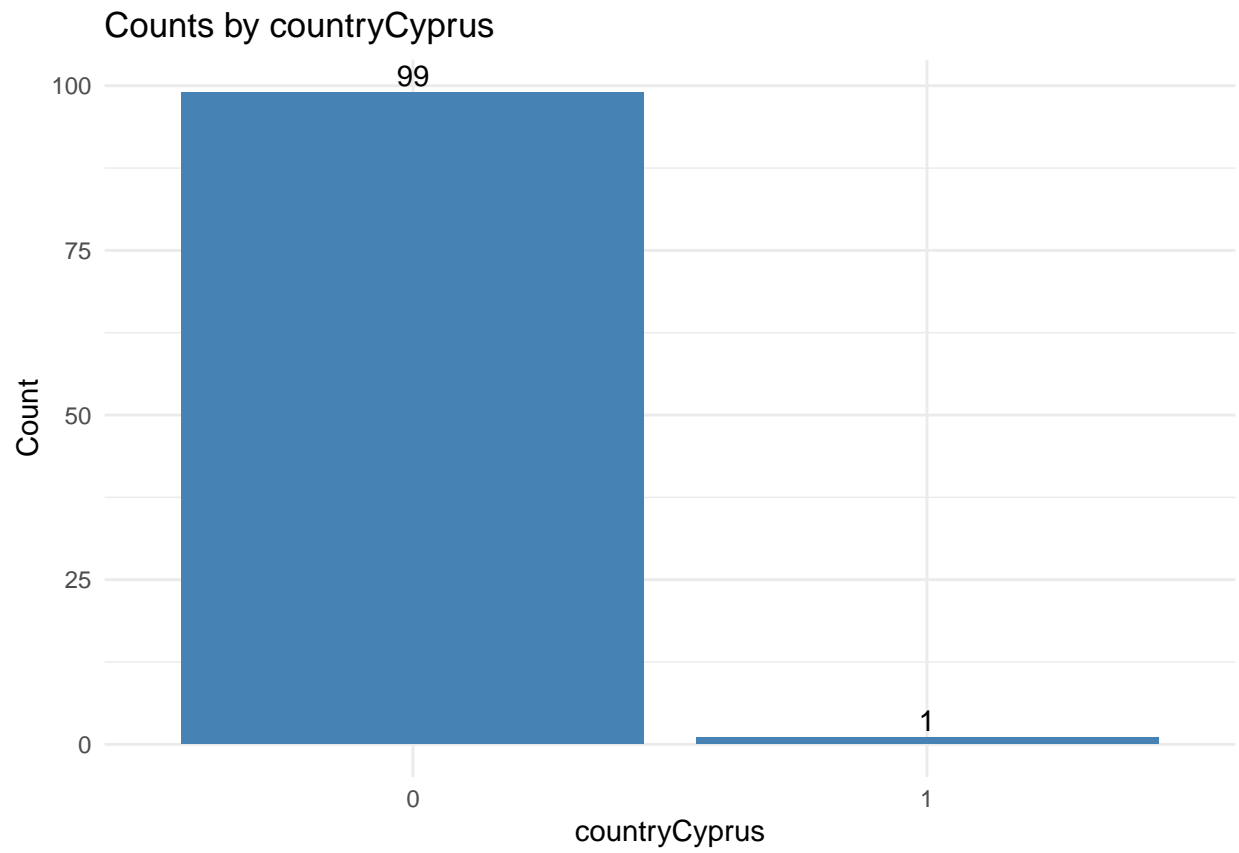


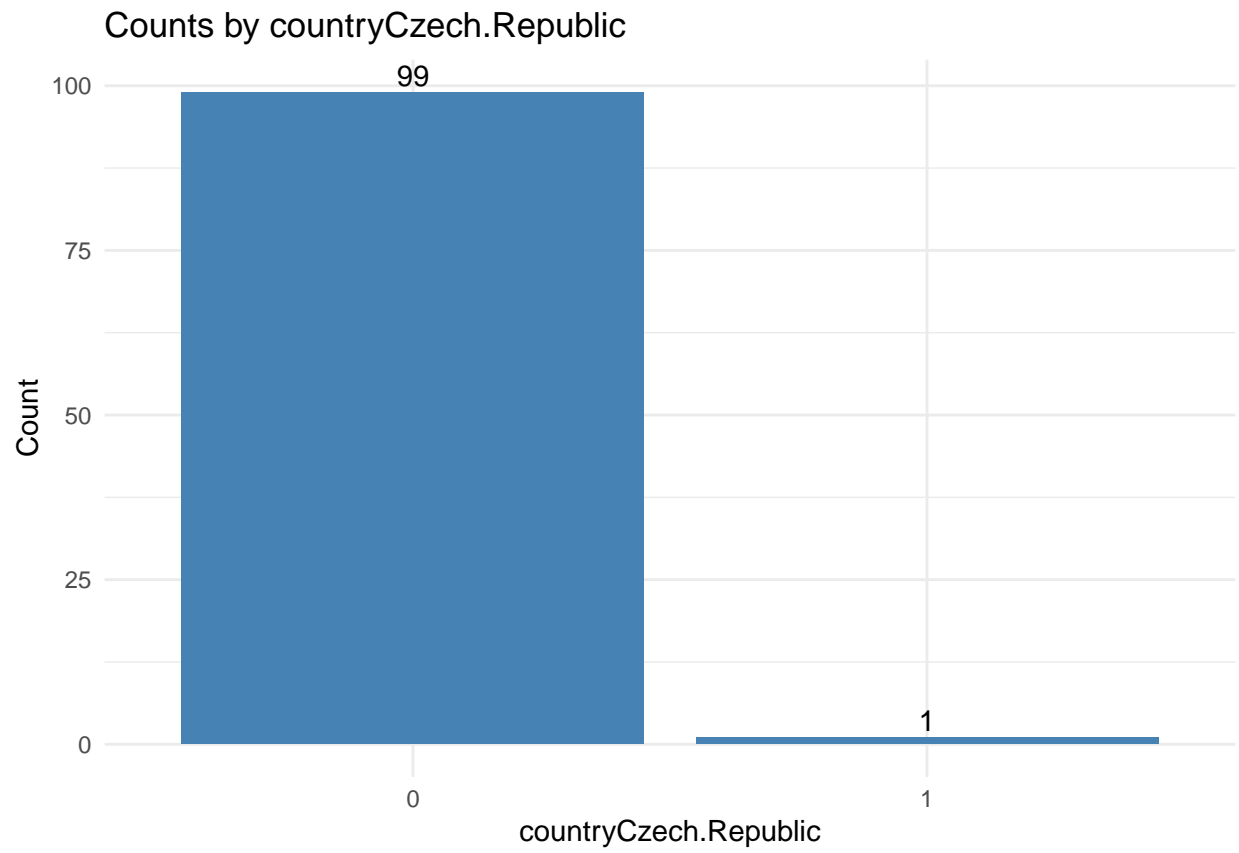


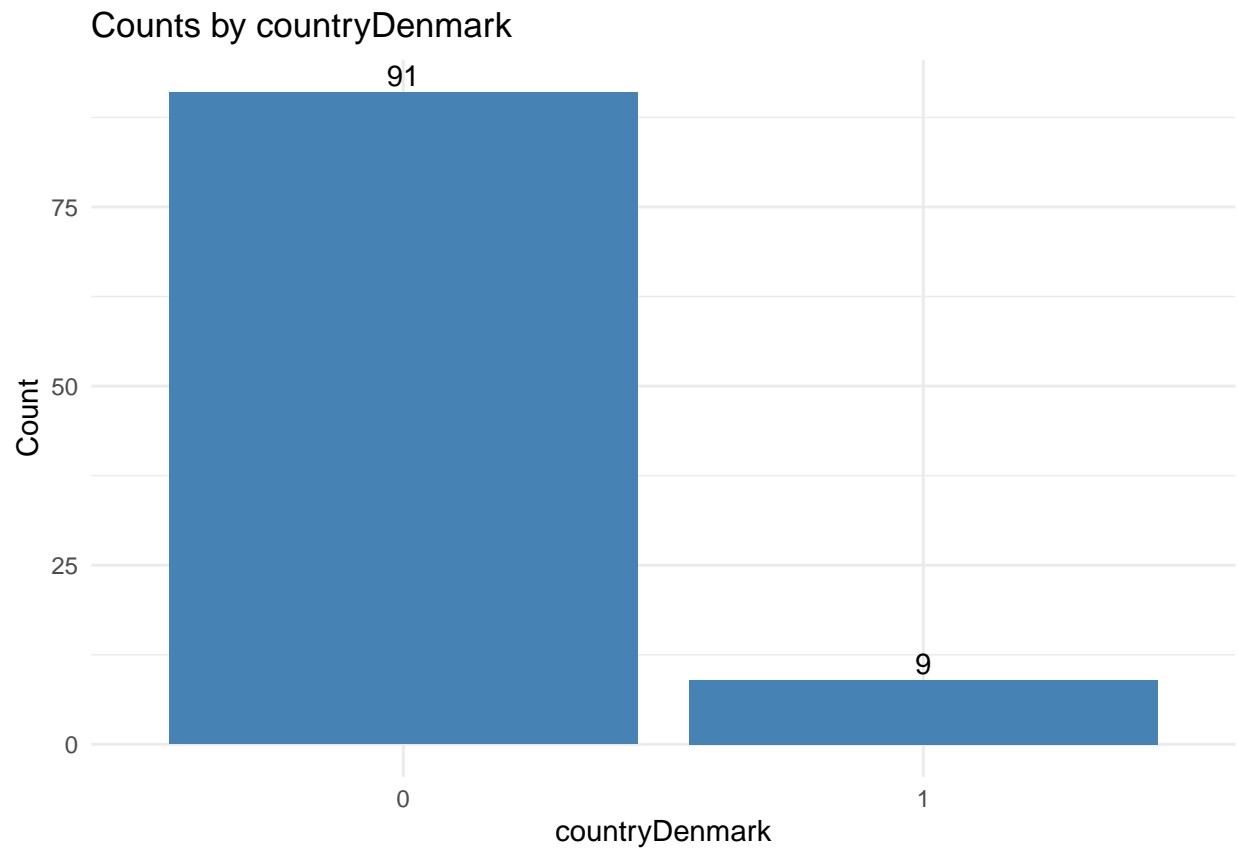


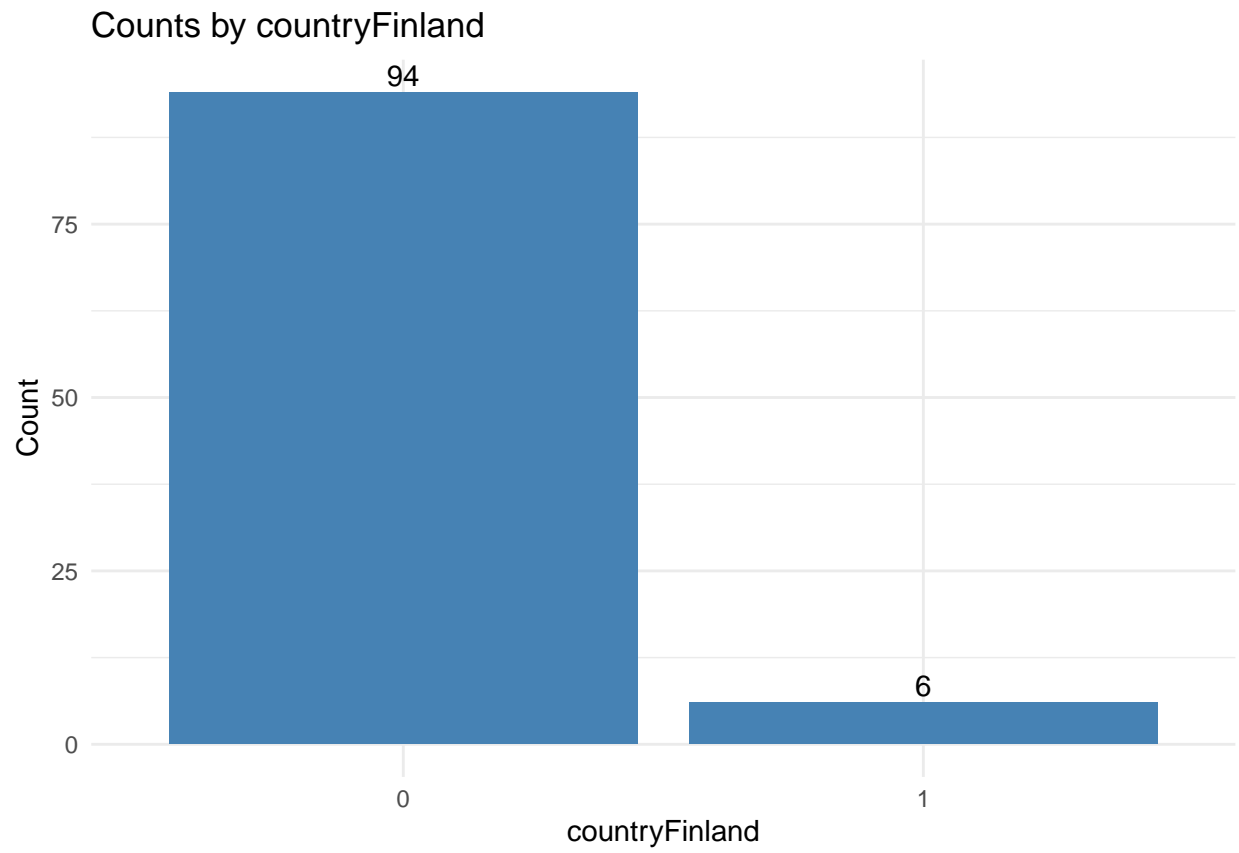


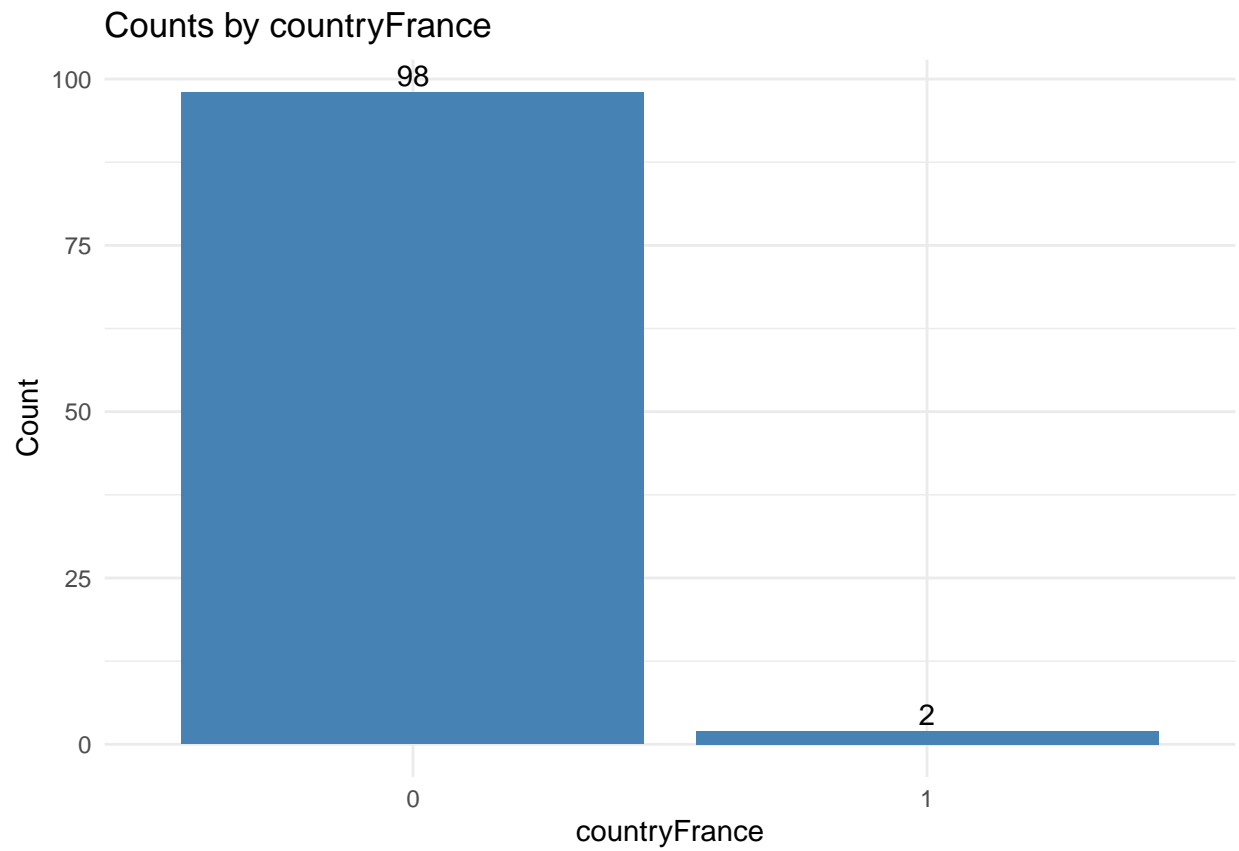


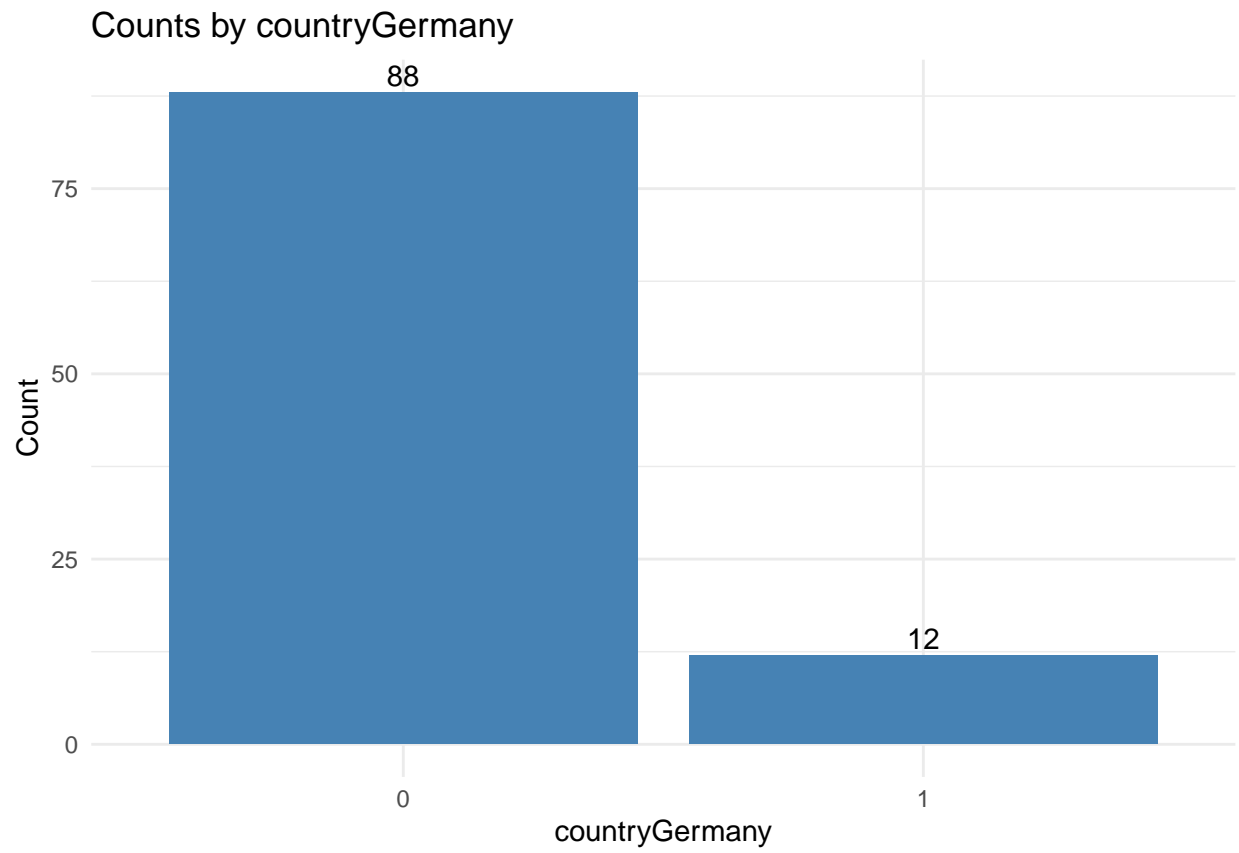


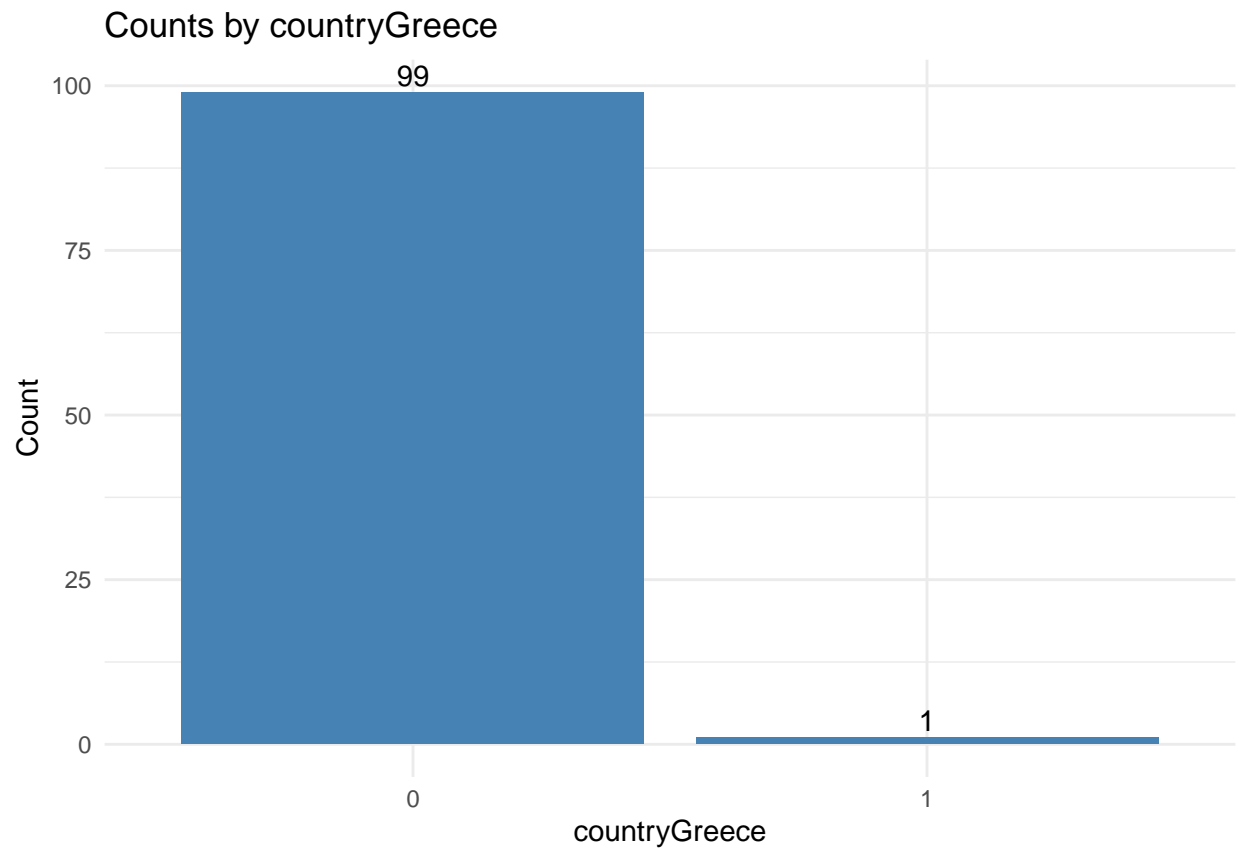


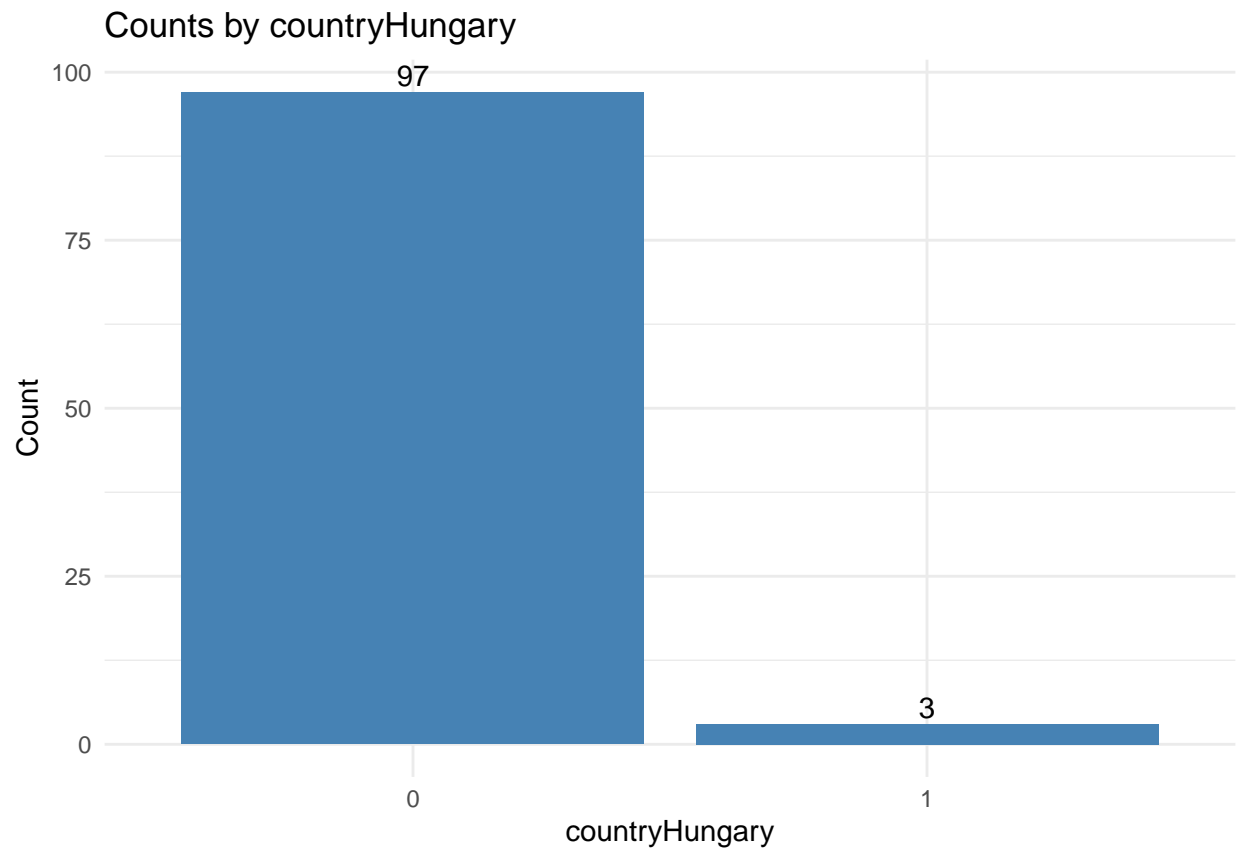


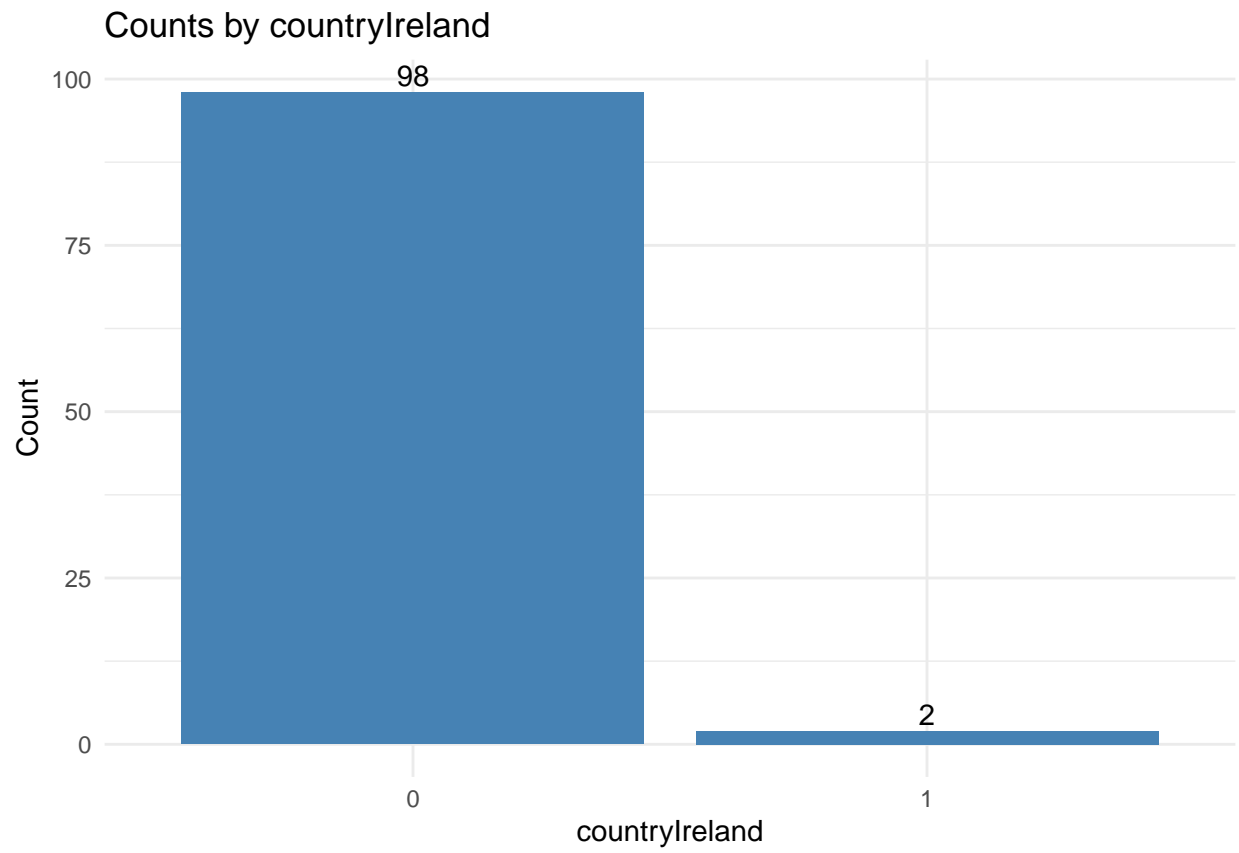


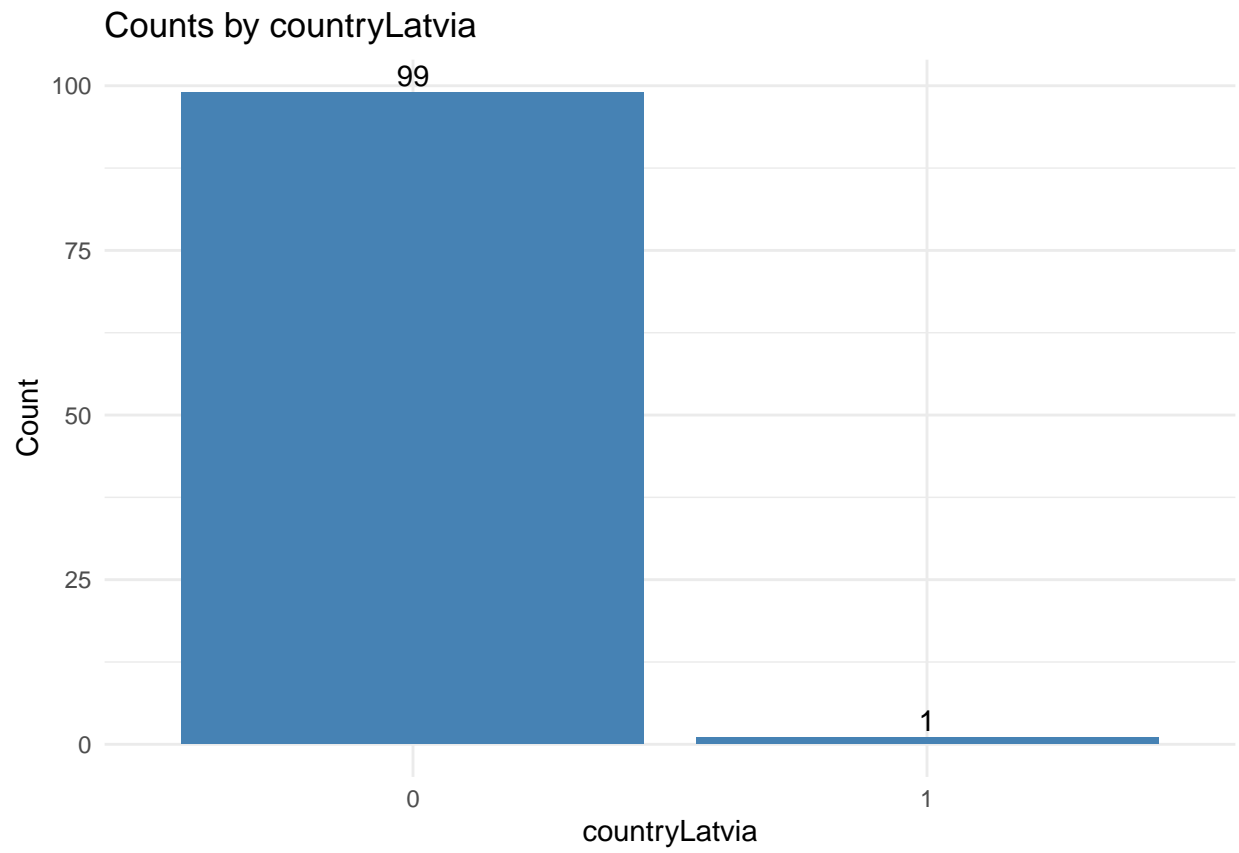


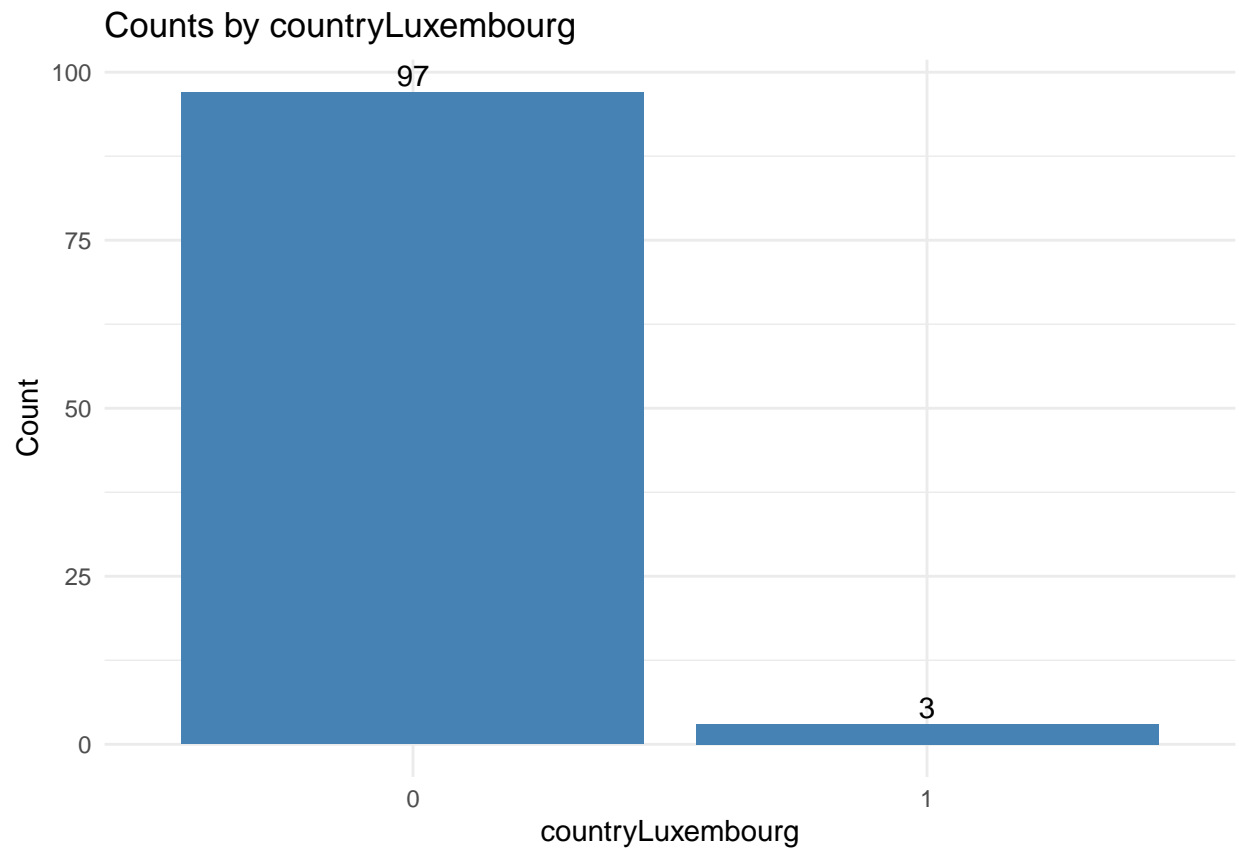


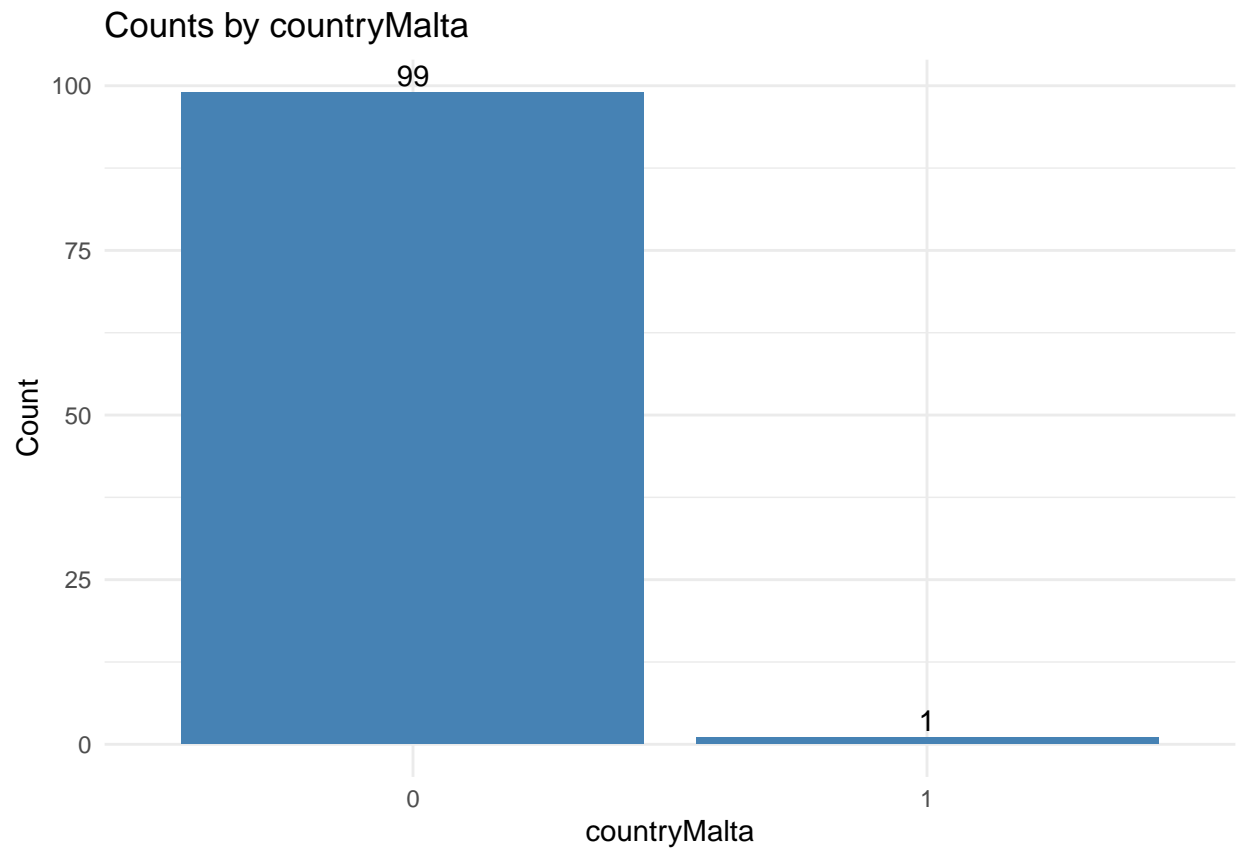


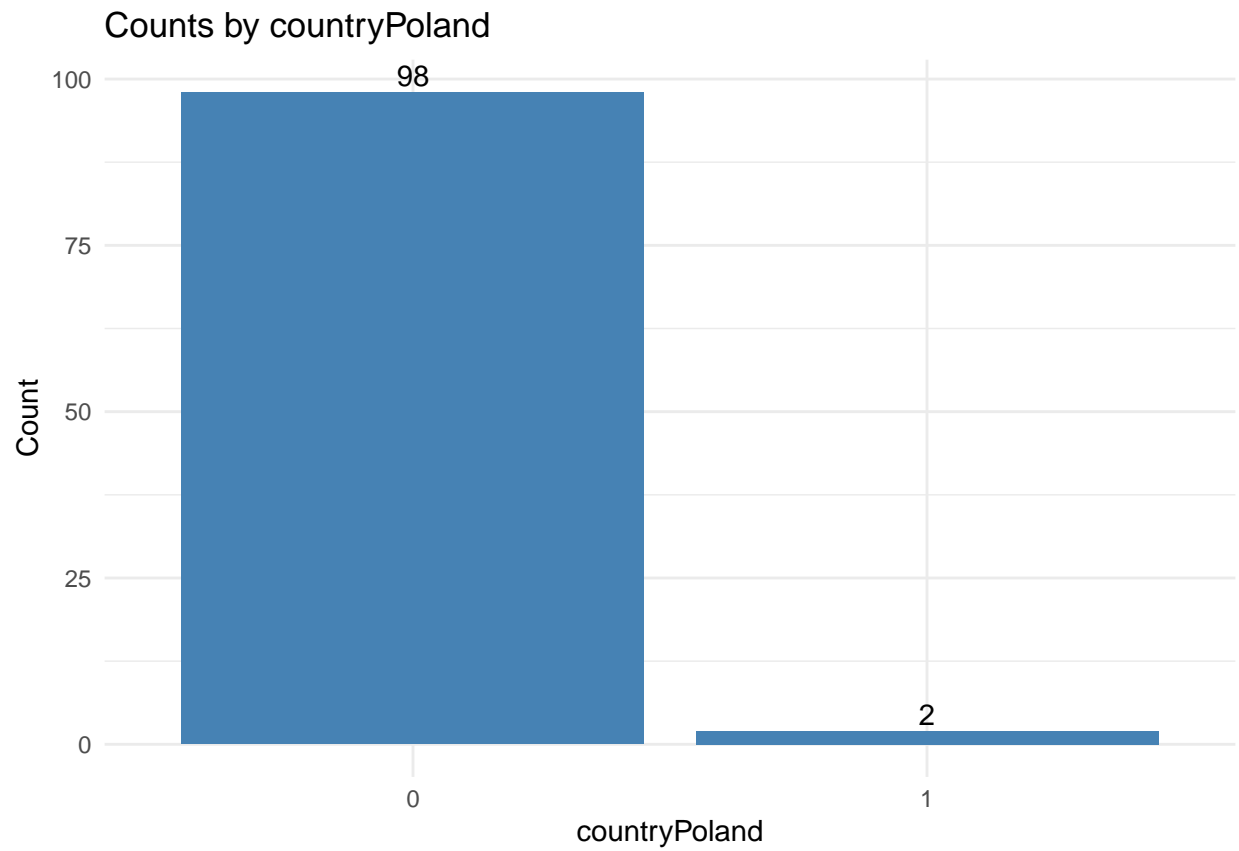


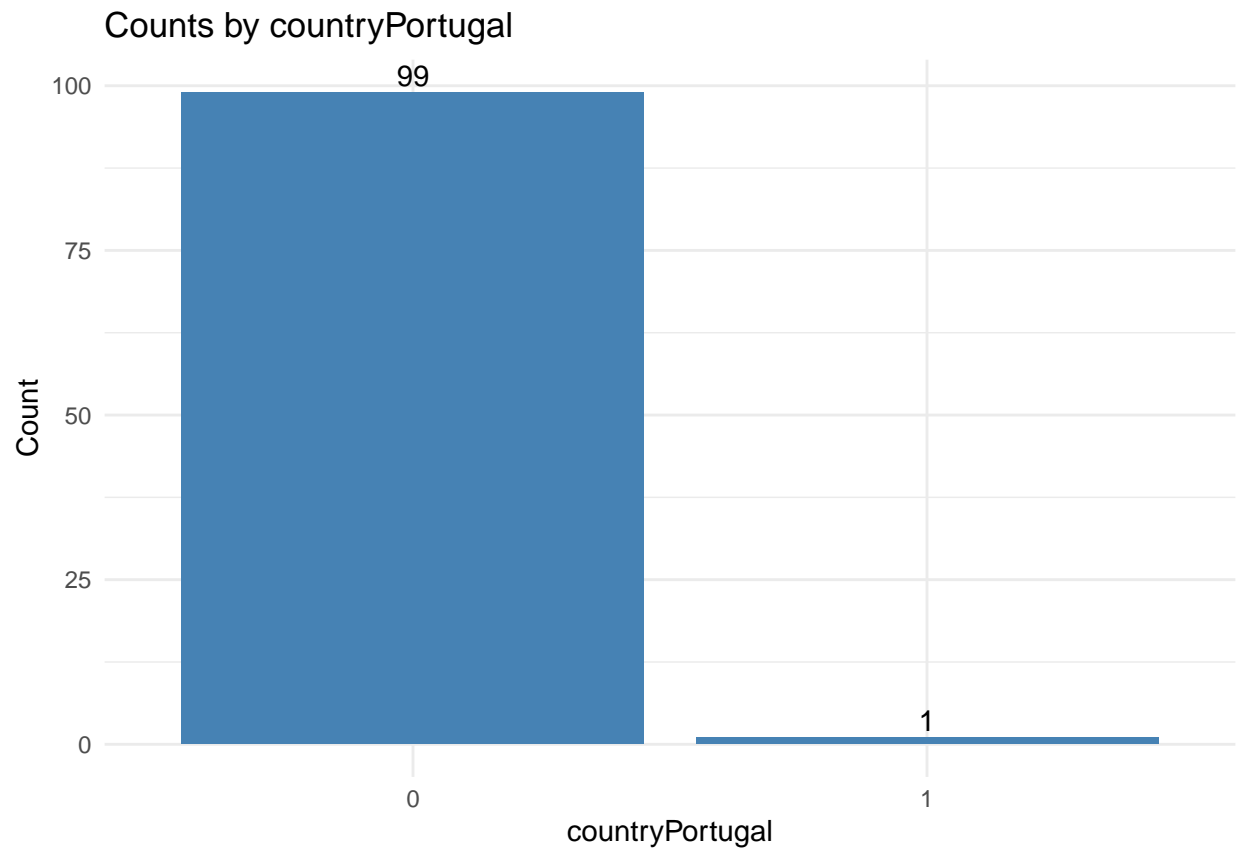


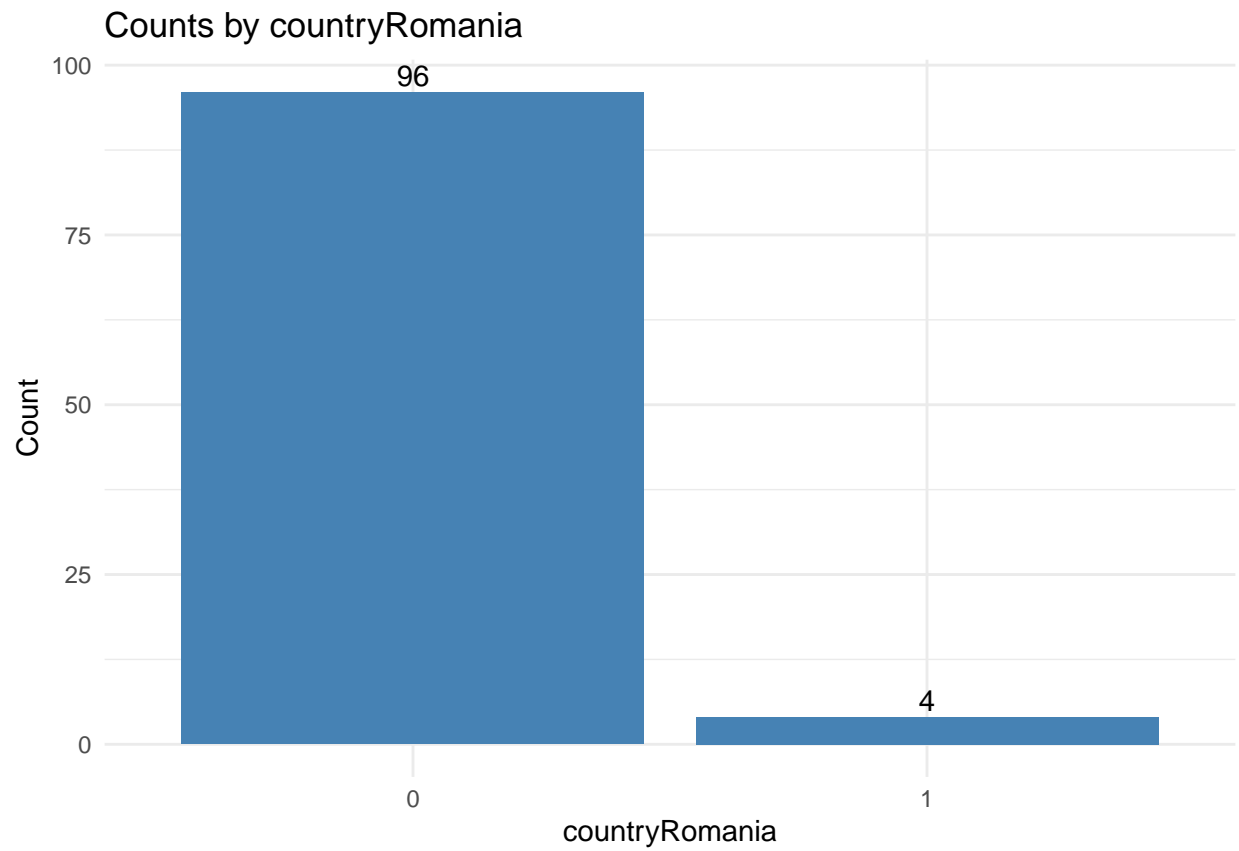


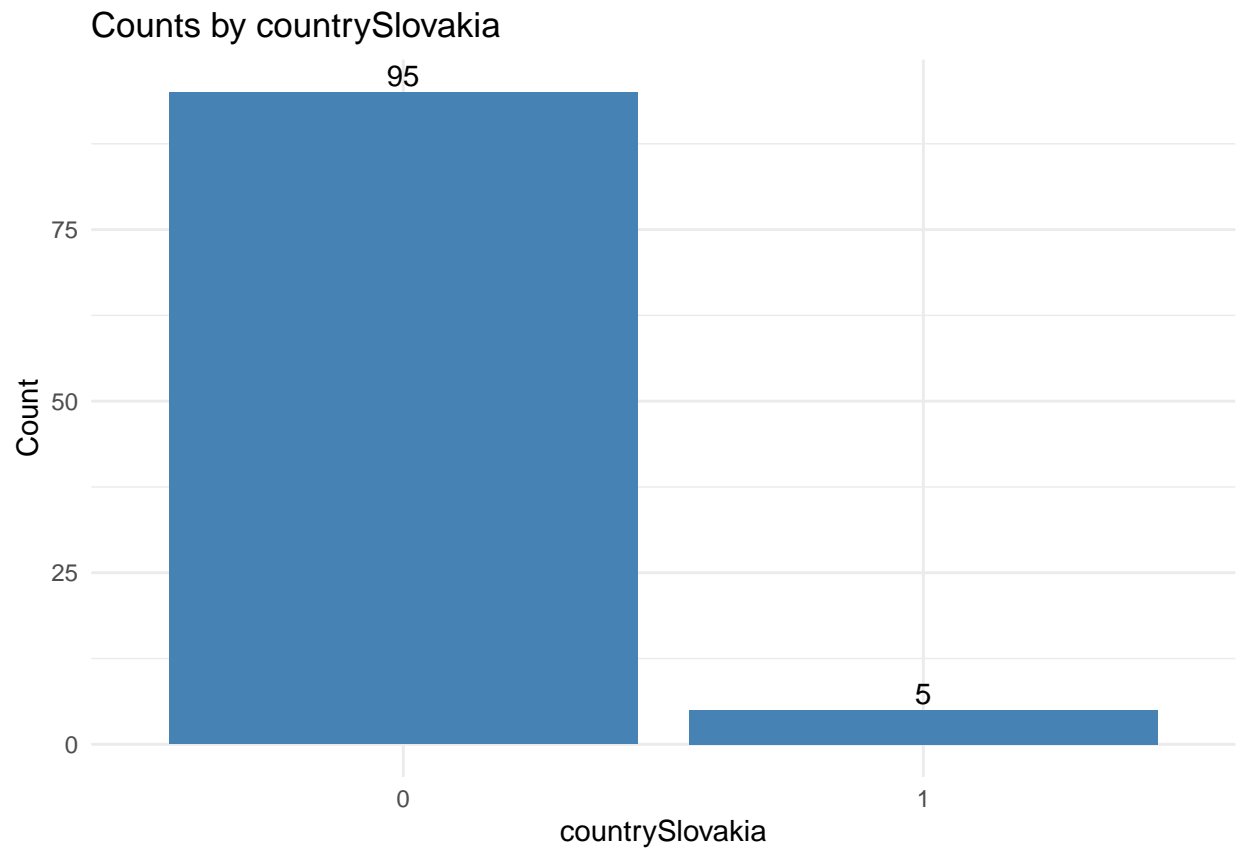


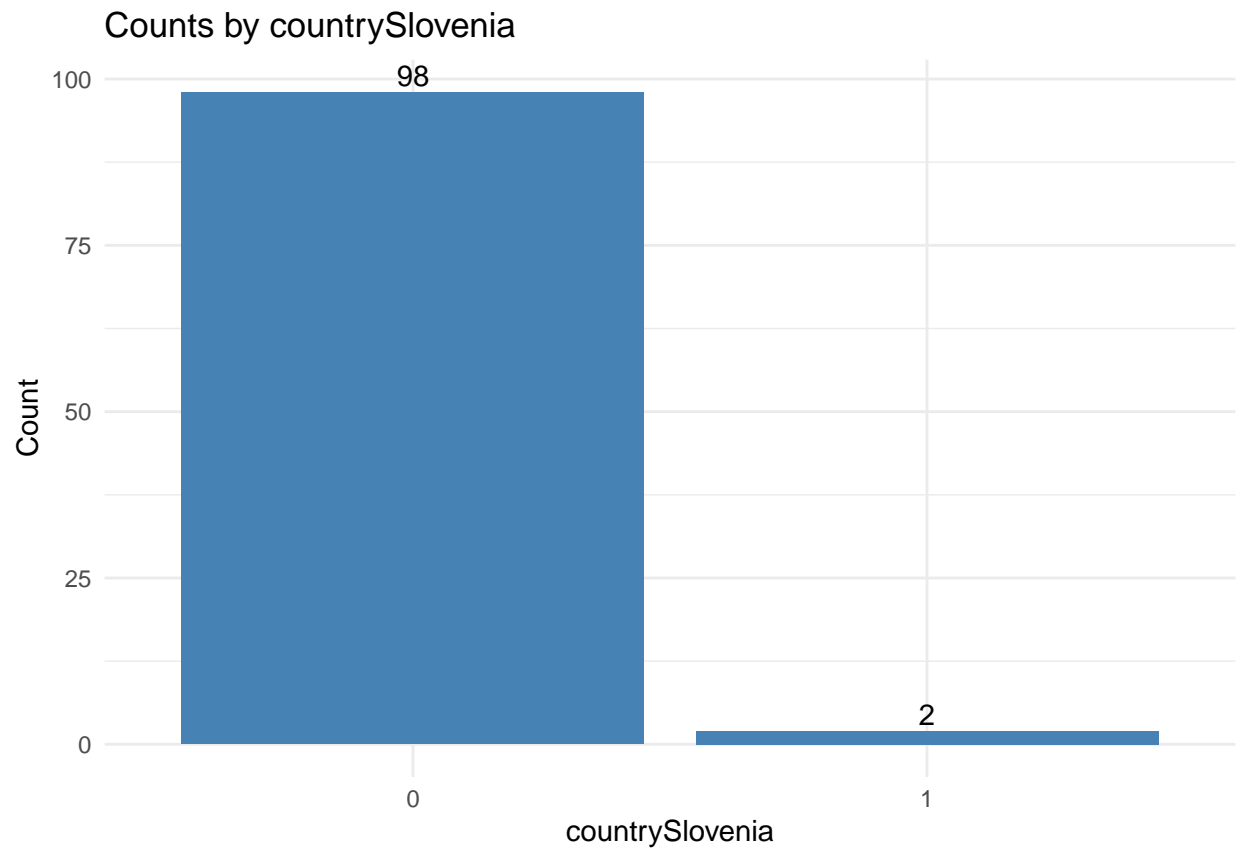


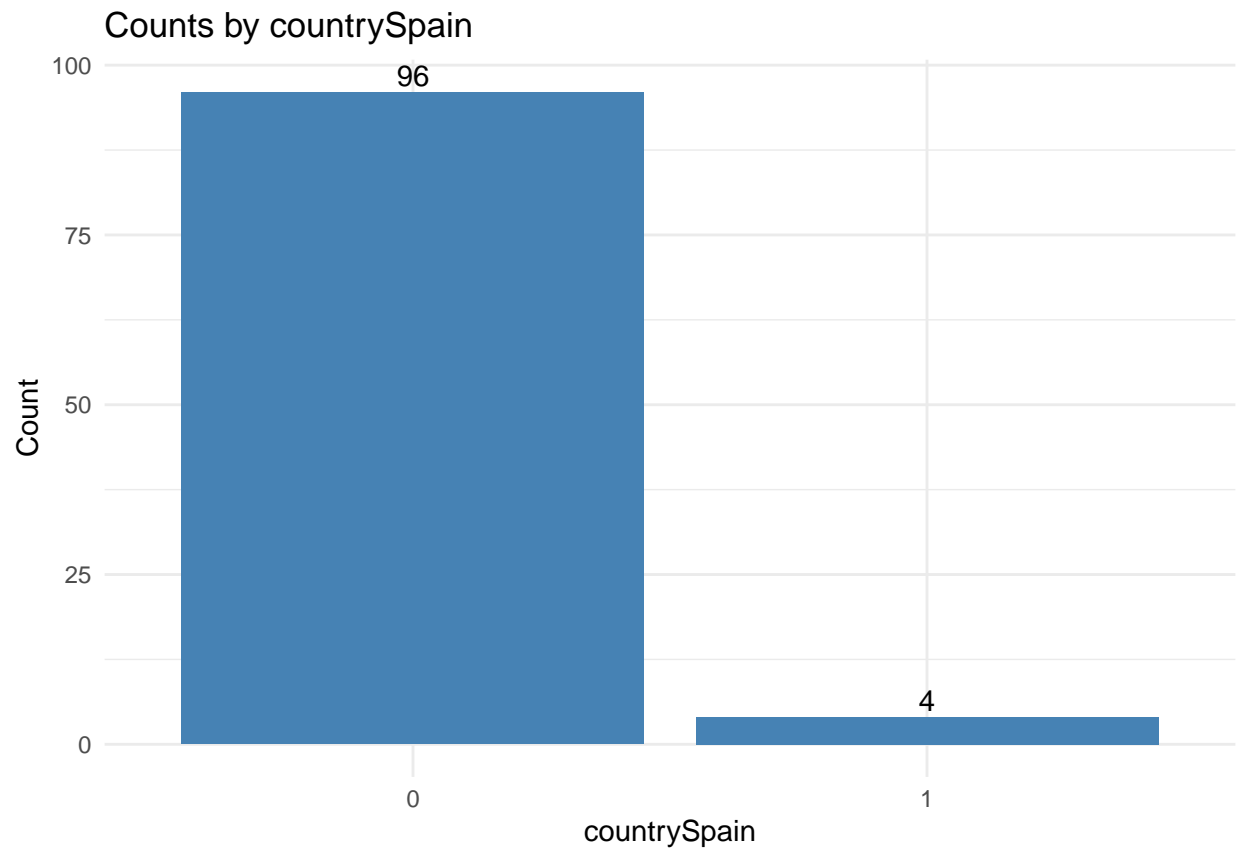


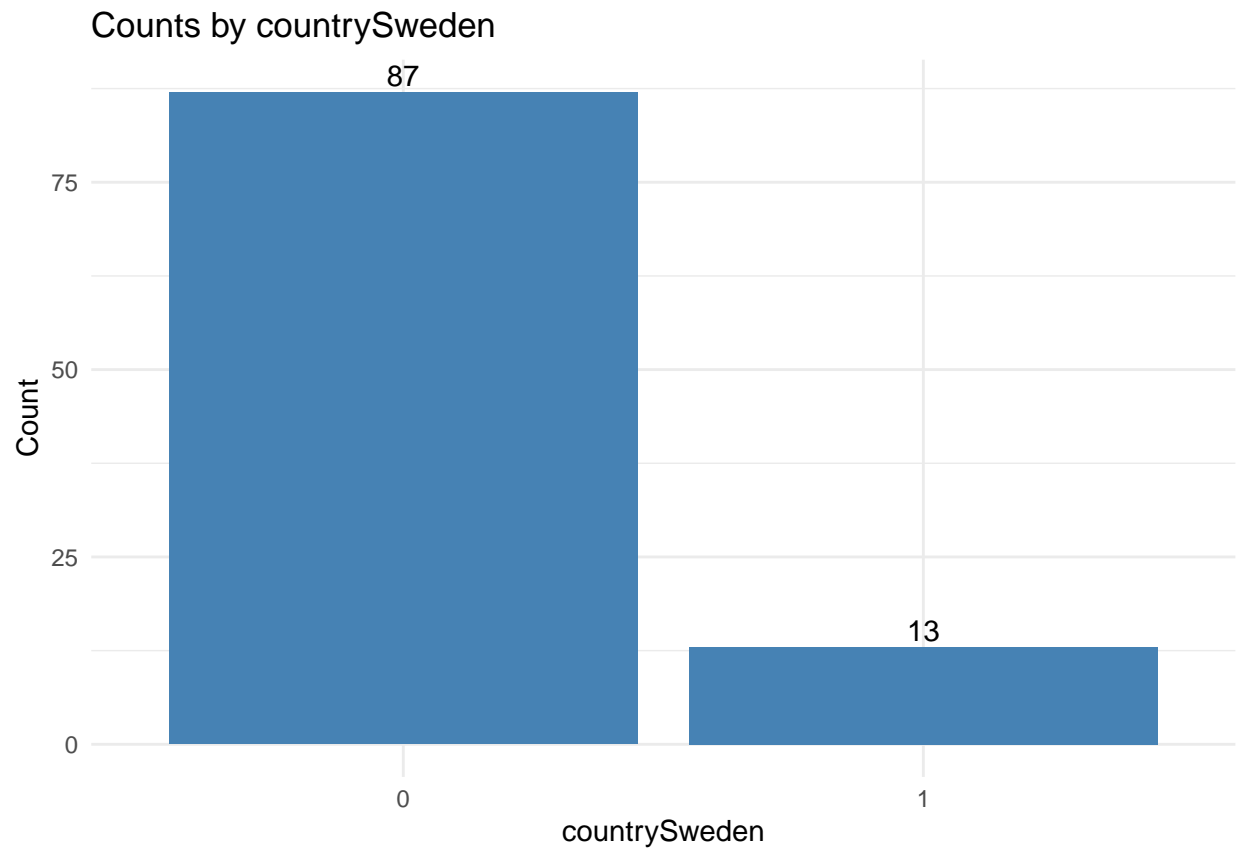


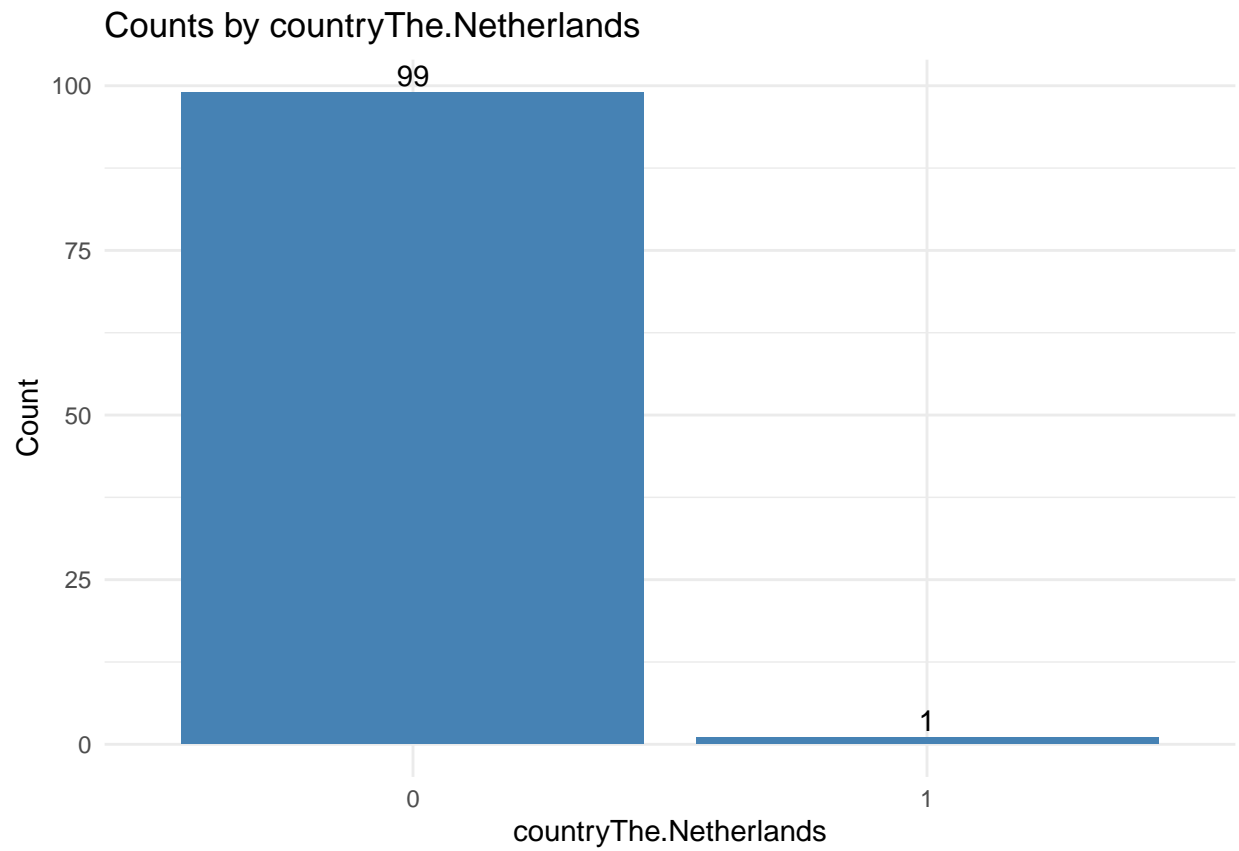


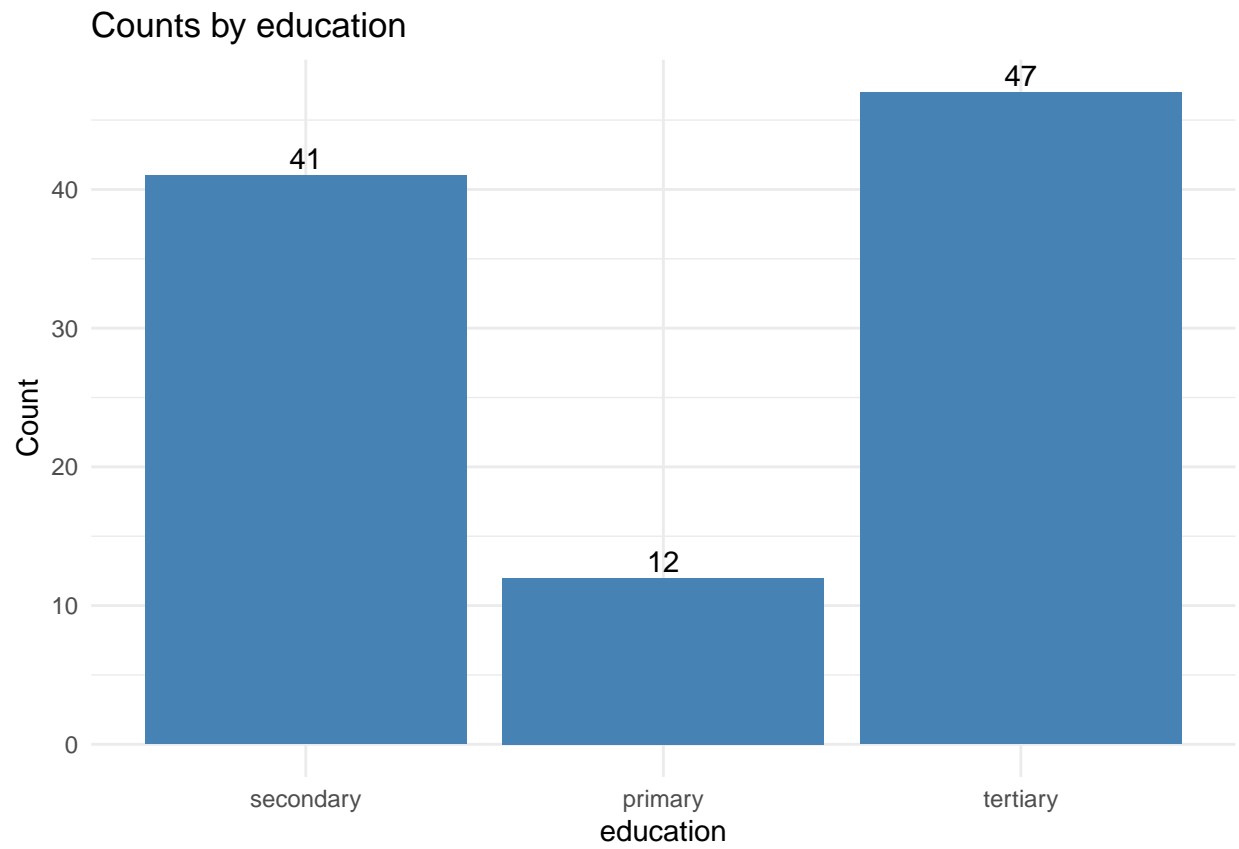


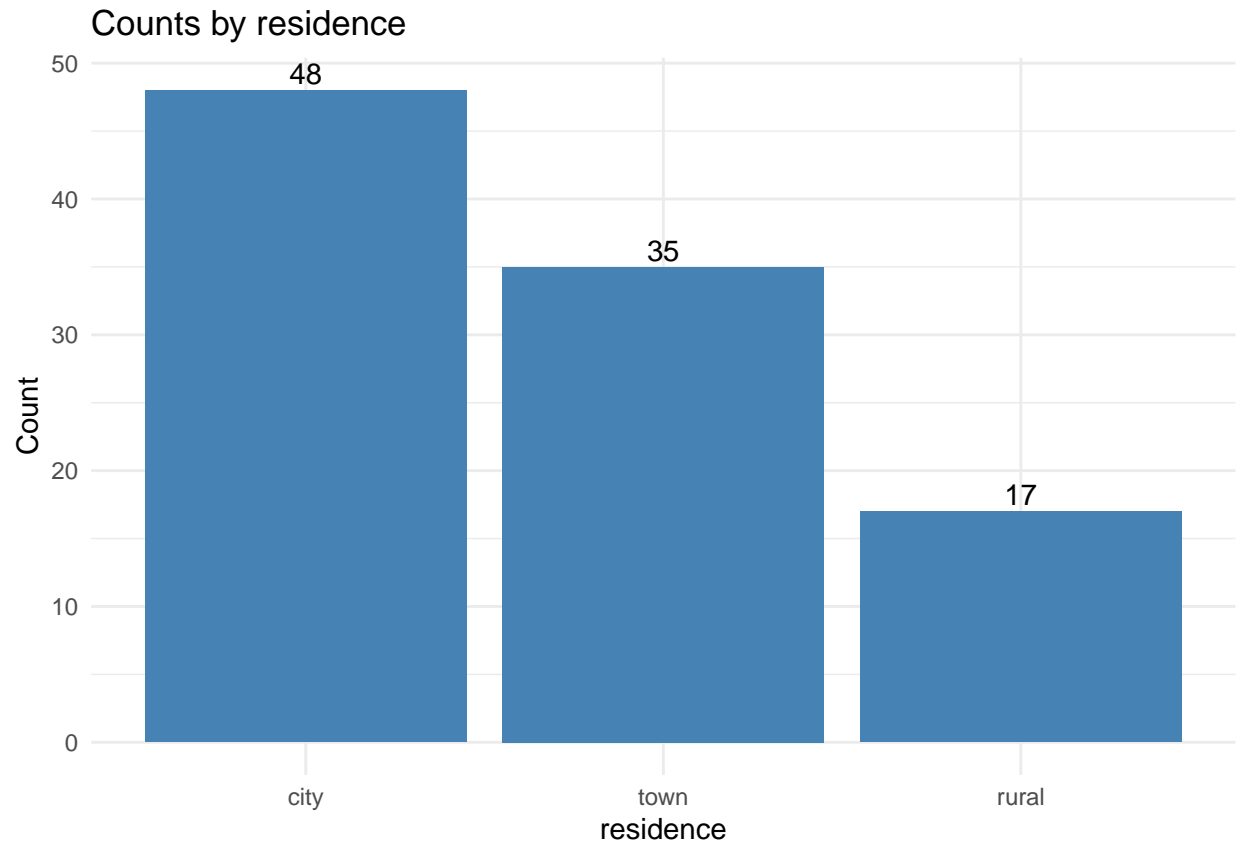












```
library(ggplot2)
library(dplyr)
library(gridExtra)
```

```
##
## Attaching package: 'gridExtra'

## The following object is masked from 'package:dplyr':
##
##   combine

## The following object is masked from 'package:randomForest':
##
##   combine
```

```
compare_model_features_plot <- function(subset_data, full_data, model_formula) {
  model_vars <- all.vars(model_formula)[-1] # exclude response variable

  for (var in model_vars) {
    # Skip if variable is missing in either dataset
    if (!(var %in% names(subset_data)) || !(var %in% names(full_data))) next

    x_subset <- subset_data[[var]]
    x_full <- full_data[[var]]
  }
}
```



```

# Strict binary check (only 0 and 1, numeric)
unique_vals <- unique(na.omit(c(x_subset, x_full)))
is_binary <- is.numeric(x_subset) && length(unique_vals) == 2 && all(unique_vals %in% c(0, 1))
is_categorical <- is.factor(x_subset) || is.character(x_subset) || is_binary

if (is_categorical) {
  # Treat all as factors (including binary)
  all_levels <- union(unique(as.character(x_subset)), unique(as.character(x_full)))
  subset_data[[var]] <- factor(as.character(x_subset), levels = all_levels)
  full_data[[var]] <- factor(as.character(x_full), levels = all_levels)

  p1 <- ggplot(subset_data, aes_string(x = var)) +
    geom_bar(fill = "skyblue") +
    labs(title = paste(var, "(SRM Data)"), x = NULL, y = "Count") +
    theme_minimal() +
    theme(axis.text.x = element_text(angle = 45, hjust = 1))

  p2 <- ggplot(full_data, aes_string(x = var)) +
    geom_bar(fill = "salmon") +
    labs(title = paste(var, "(Full Data)"), x = NULL, y = "Count") +
    theme_minimal() +
    theme(axis.text.x = element_text(angle = 45, hjust = 1))

} else if (is.numeric(x_subset)) {
  # Continuous numeric variable
  df_subset <- data.frame(value = x_subset, group = "SRM Data")
  df_full <- data.frame(value = x_full, group = "Full Data")
  combined <- bind_rows(df_subset, df_full)

  p1 <- ggplot(combined, aes(x = value, fill = group, color = group)) +
    geom_density(alpha = 0.4) +
    labs(title = paste(var, "- Density Comparison"), x = var, y = "Density") +
    theme_minimal()

  p2 <- ggplot(combined, aes(x = group, y = value, fill = group)) +
    geom_boxplot(alpha = 0.6) +
    labs(title = paste(var, "- Boxplot"), x = NULL, y = var) +
    theme_minimal()
} else {
  next
}

# Display side-by-side plots
grid.arrange(p1, p2, ncol = 2)
}
}

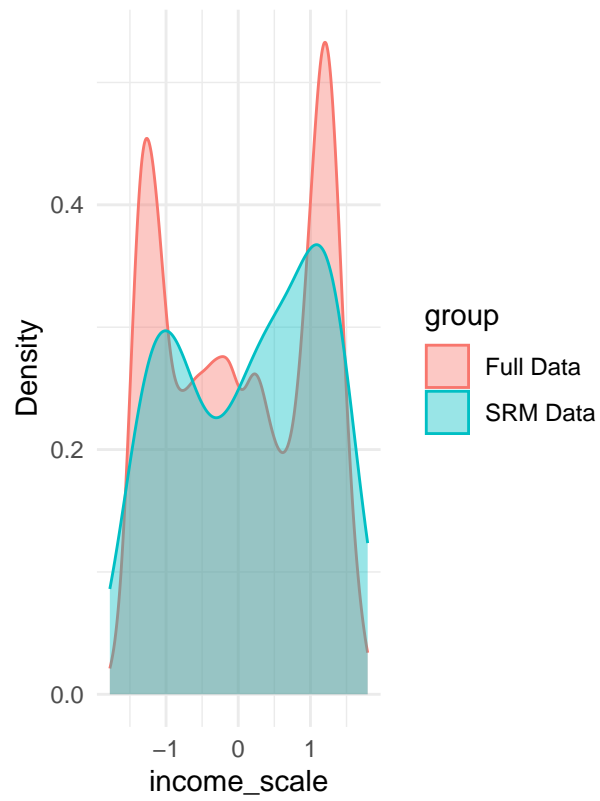
```

```

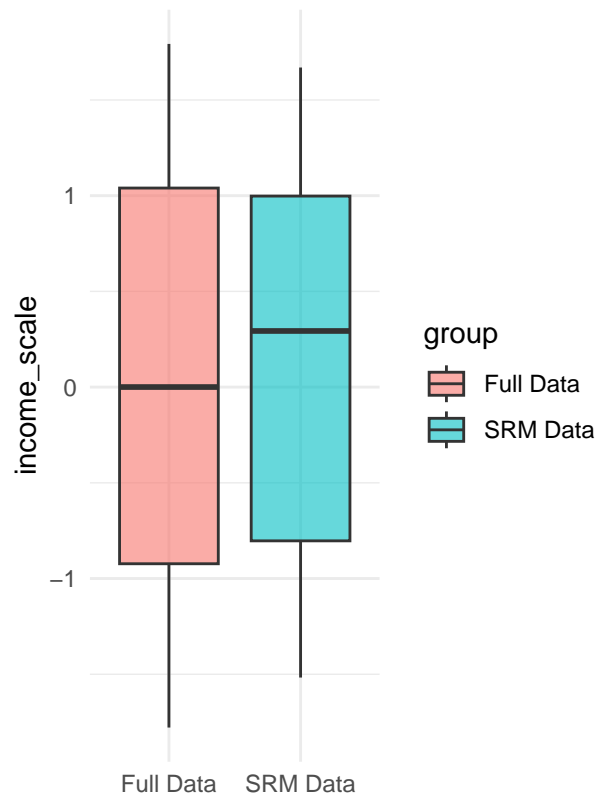
compare_model_features_plot(srm_final, df_clean, model_formula = full_formula)

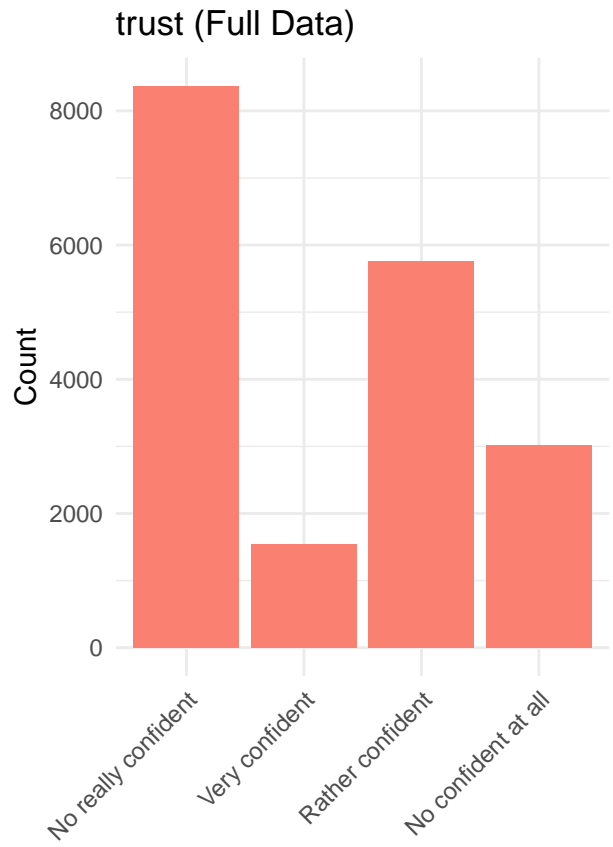
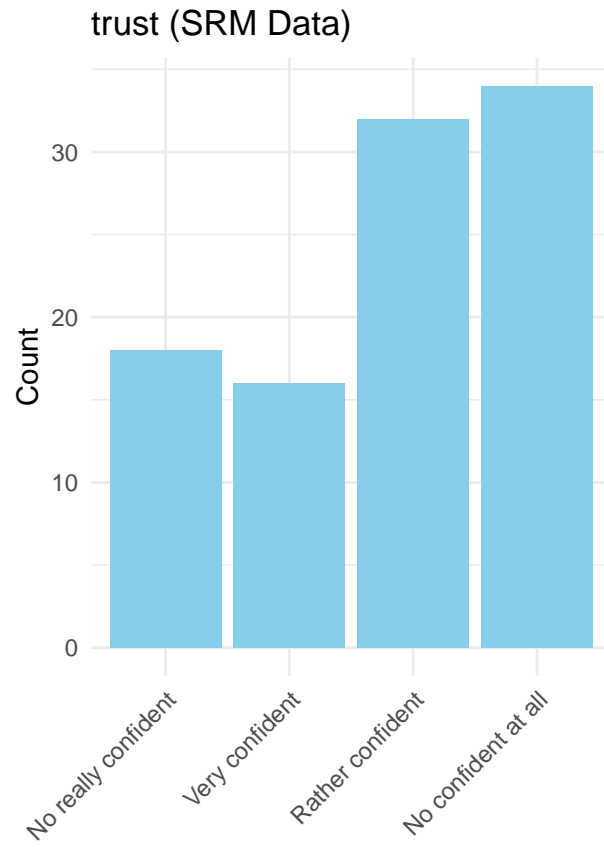
```

income_scale – Density Comparison

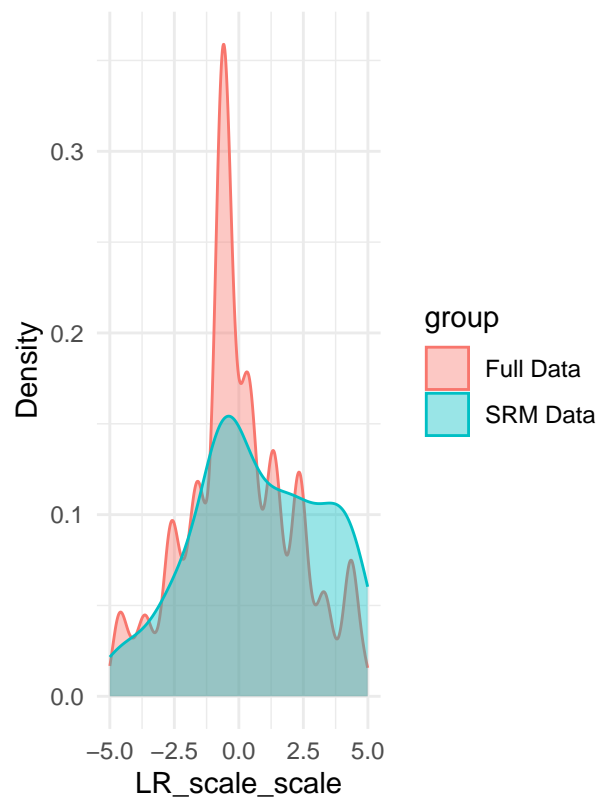


income_scale – Boxplot

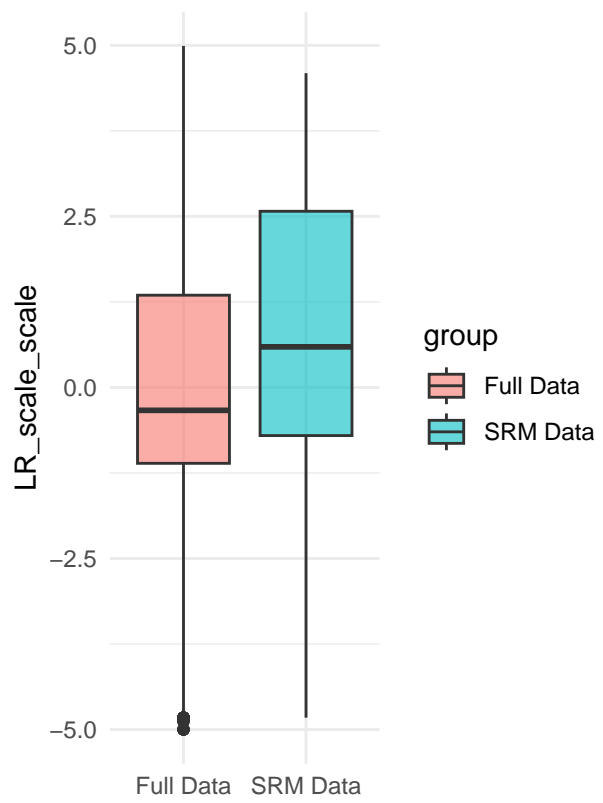


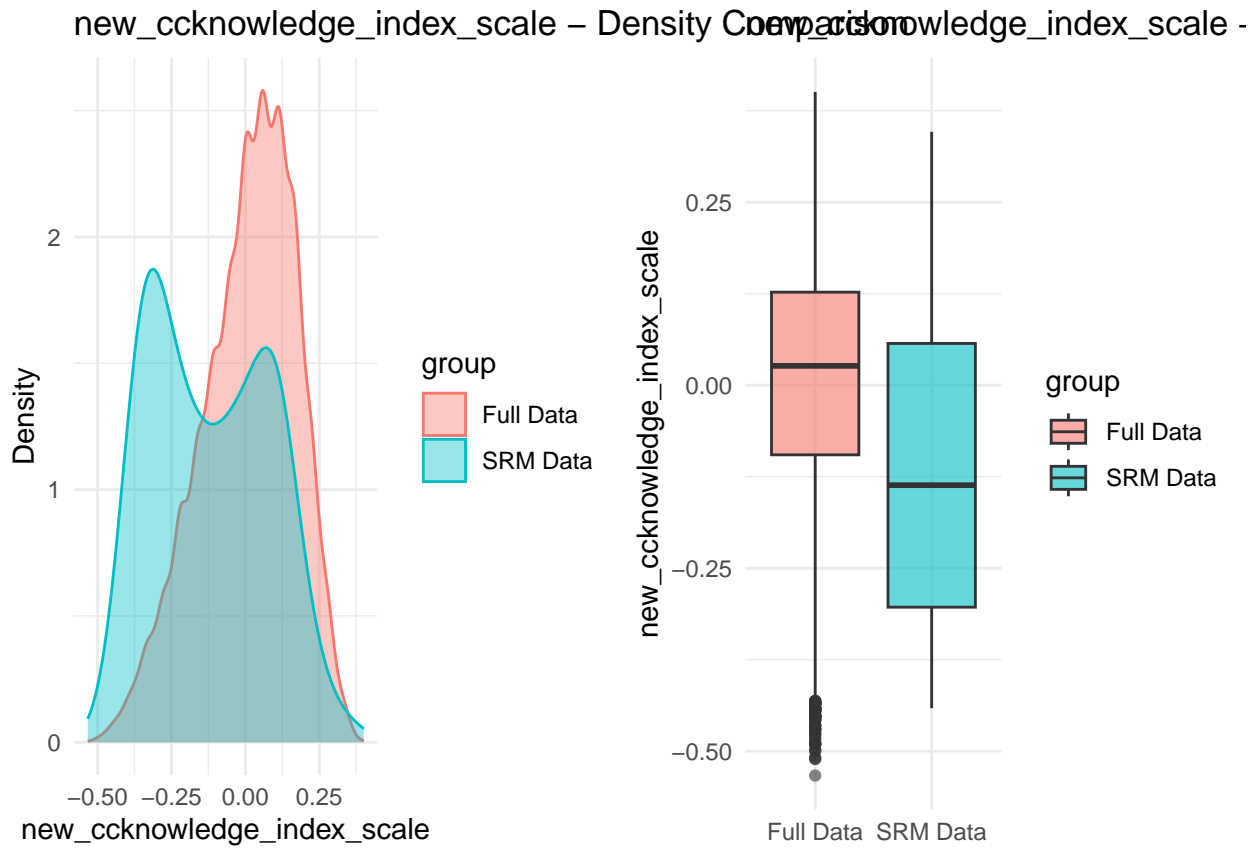


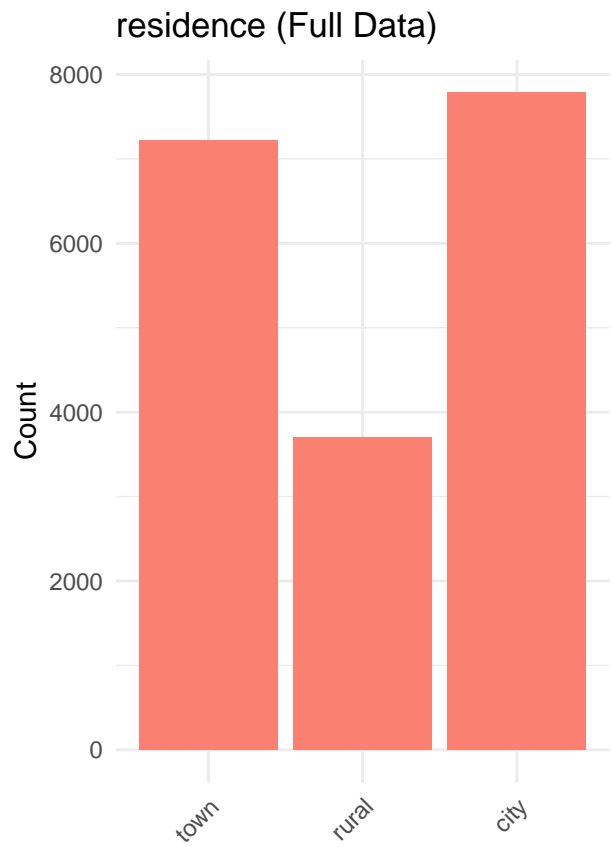
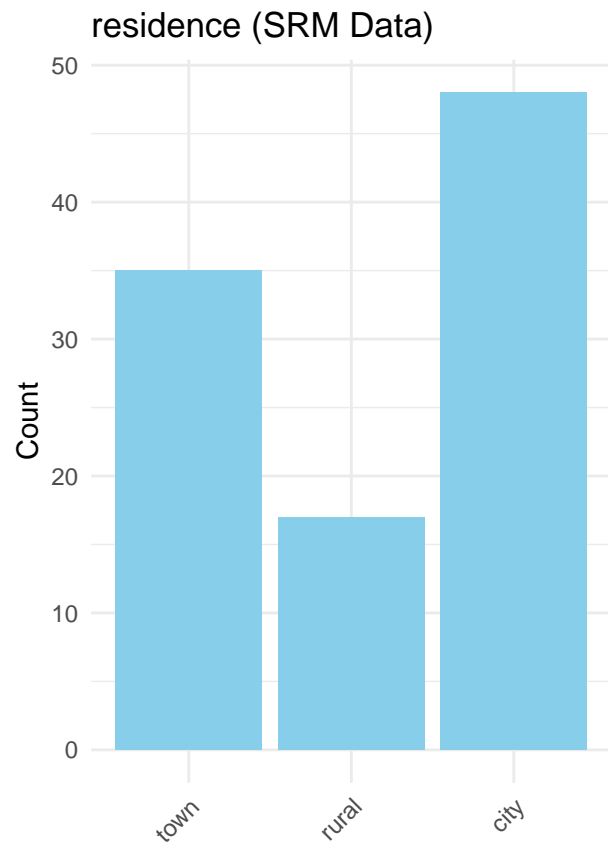
LR_scale_scale – Density Comparison

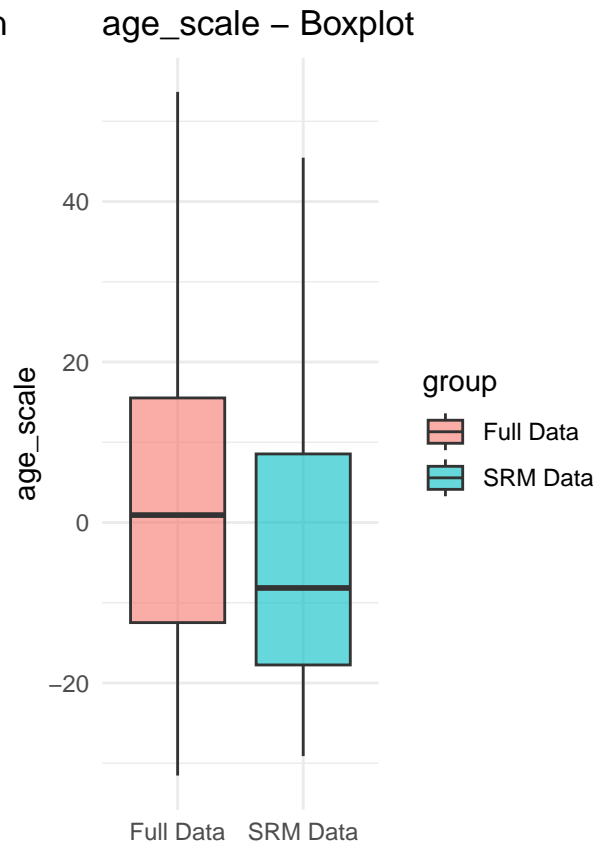
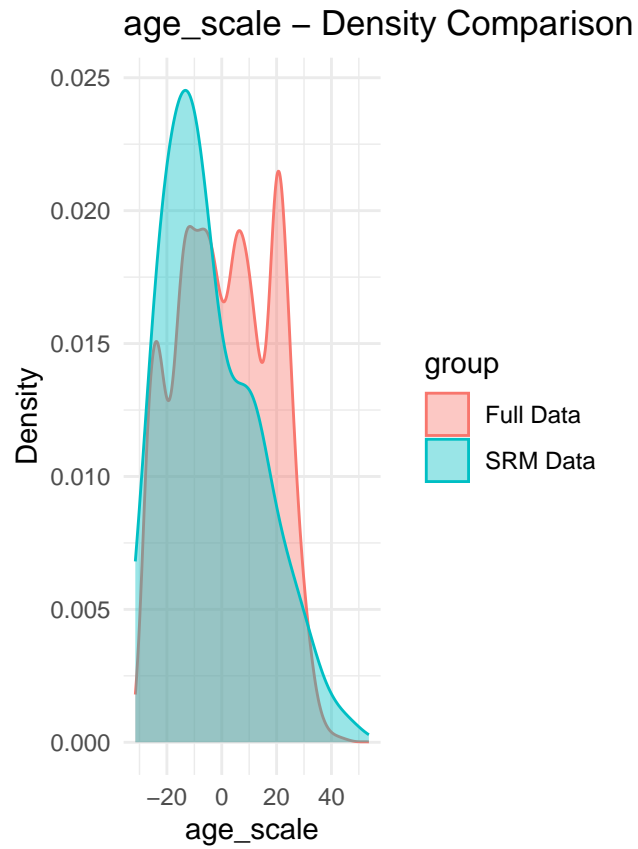


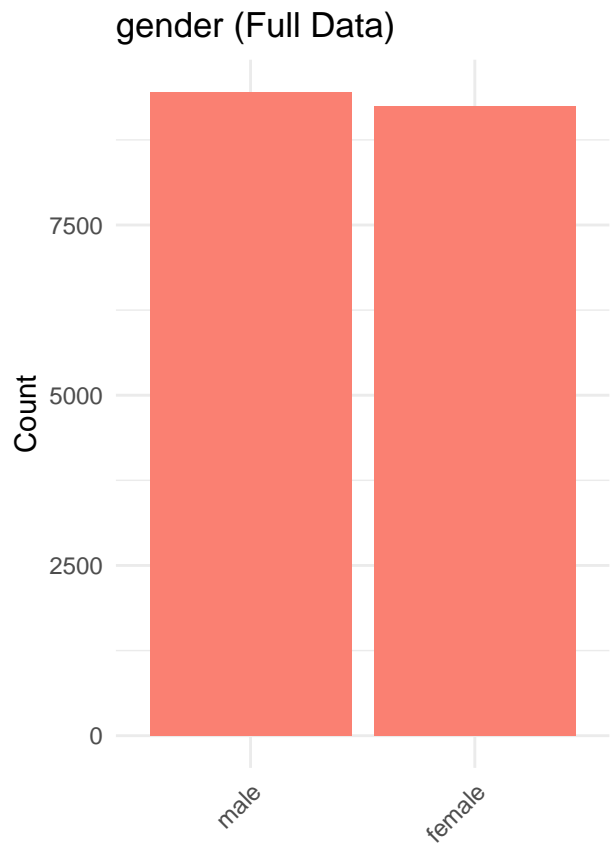
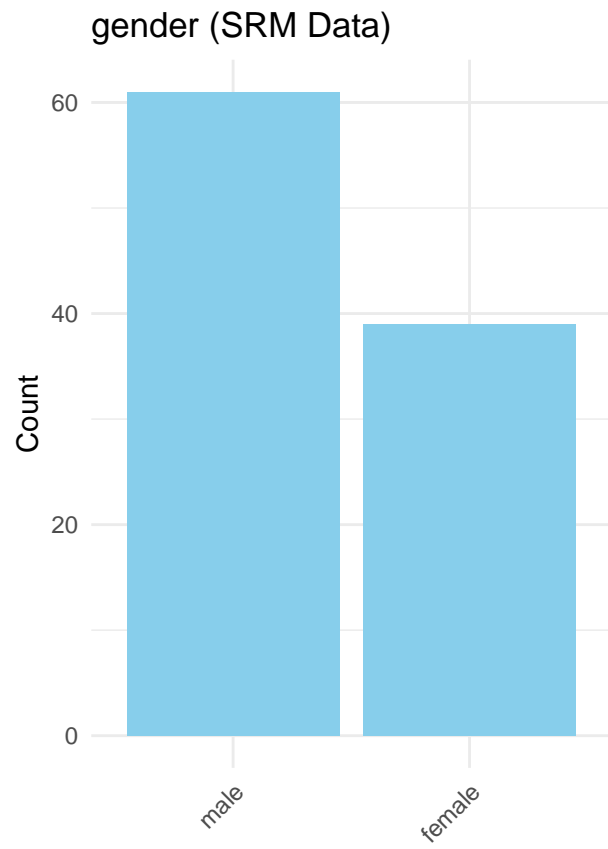
LR_scale_scale – Boxplot

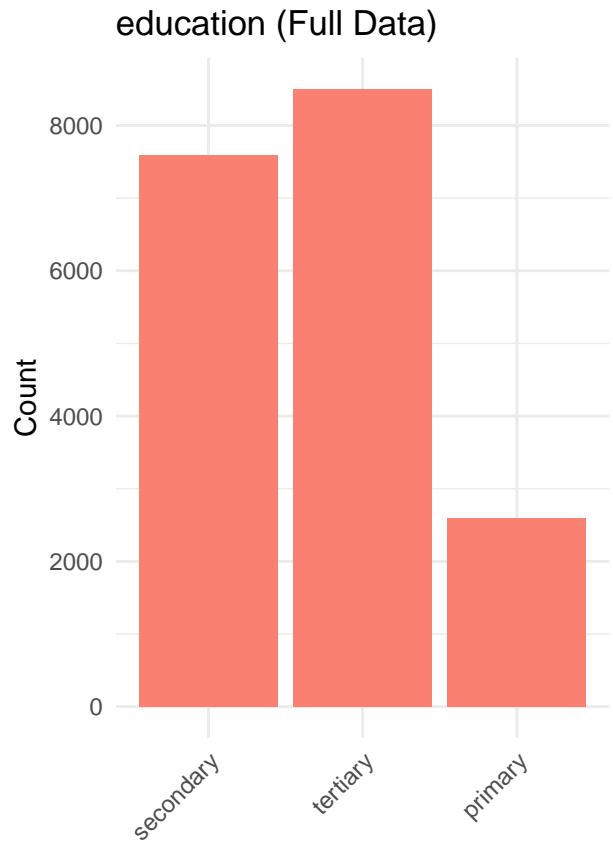
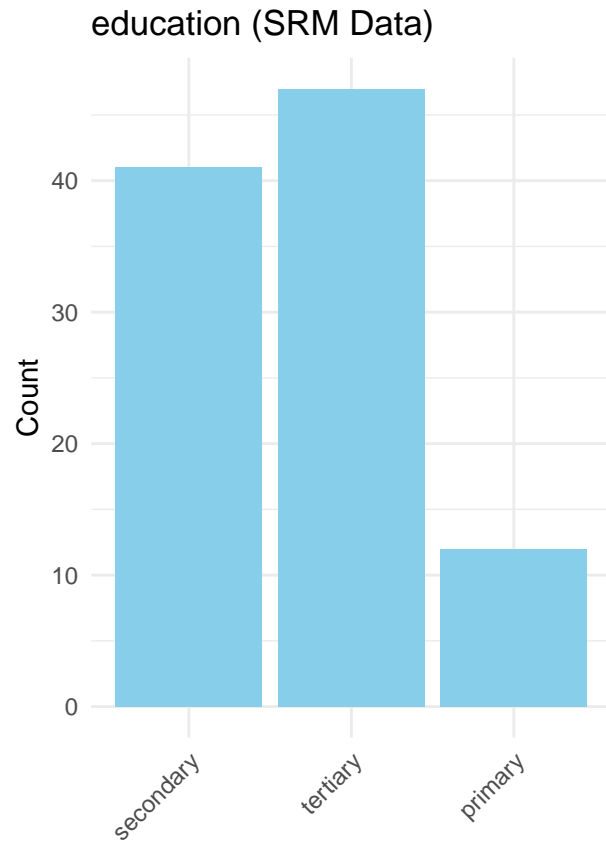


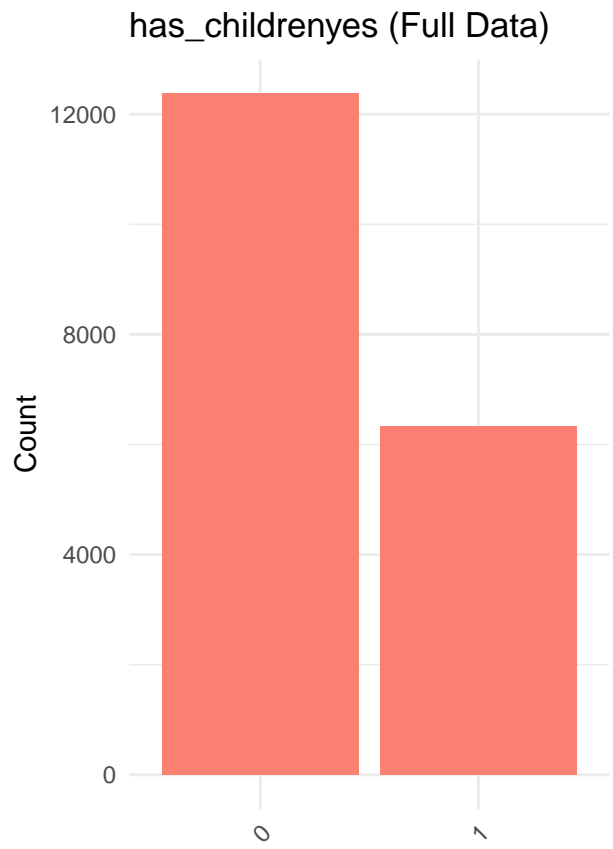
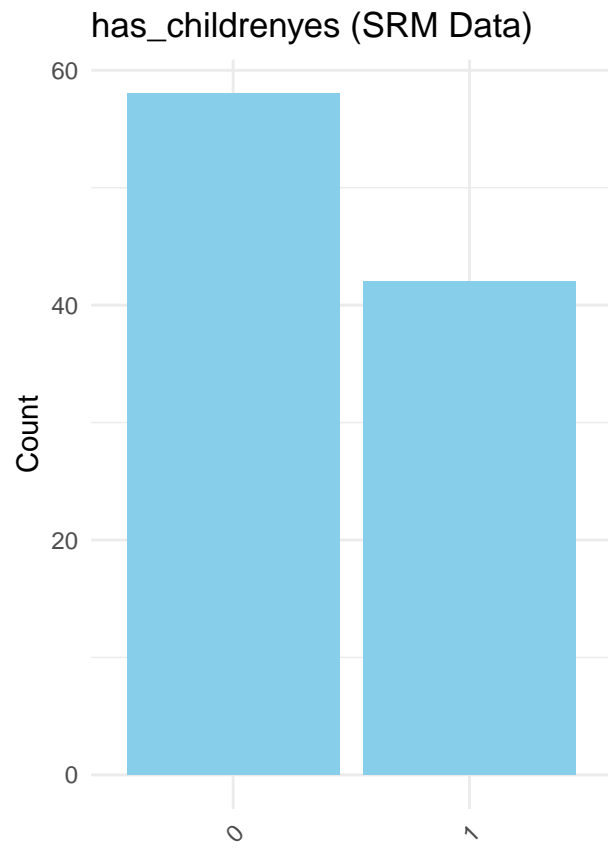


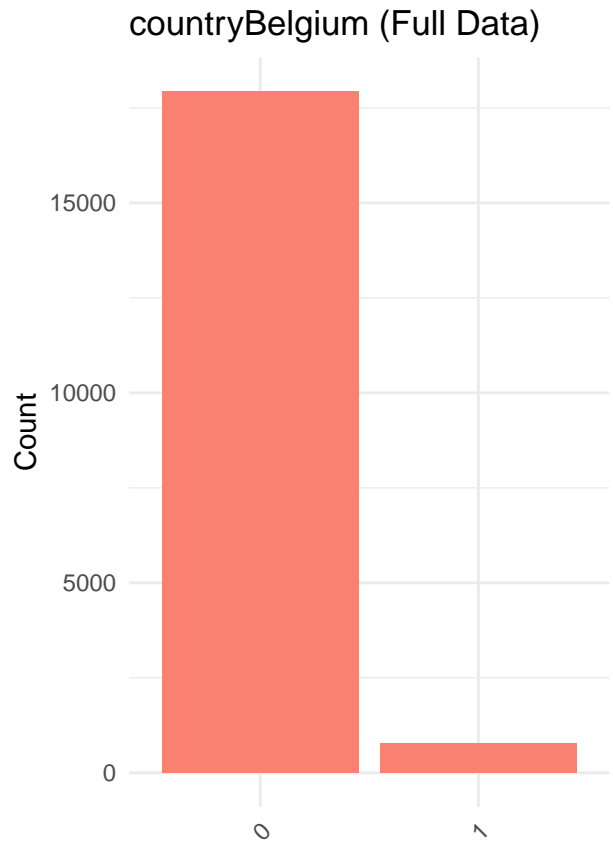
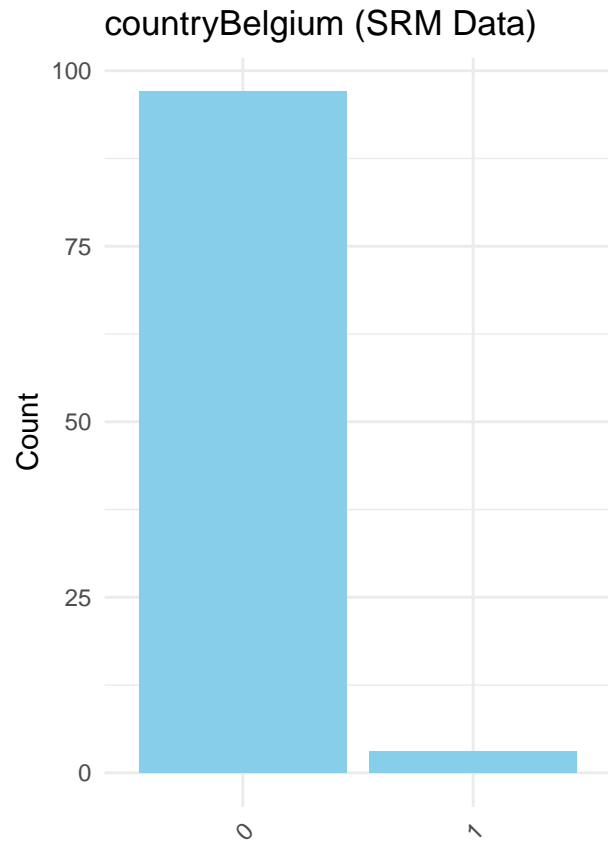


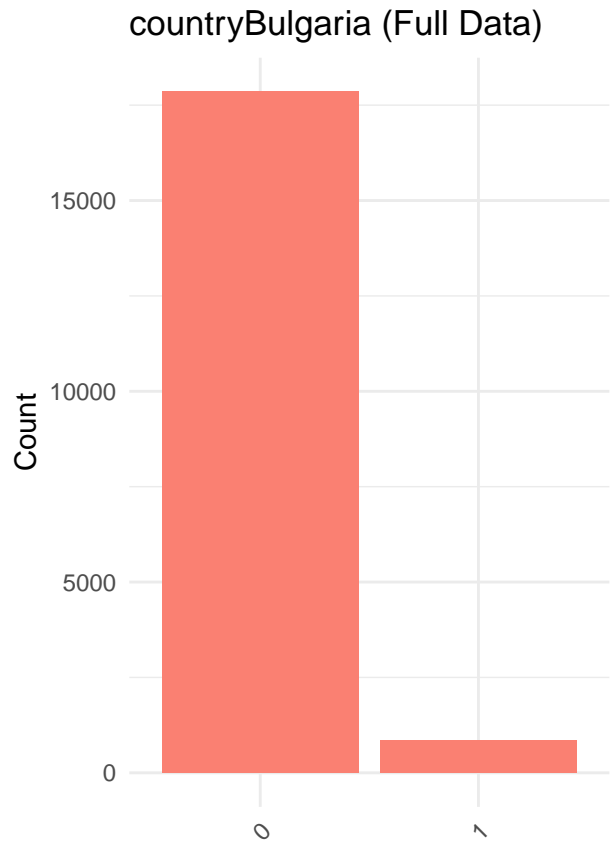
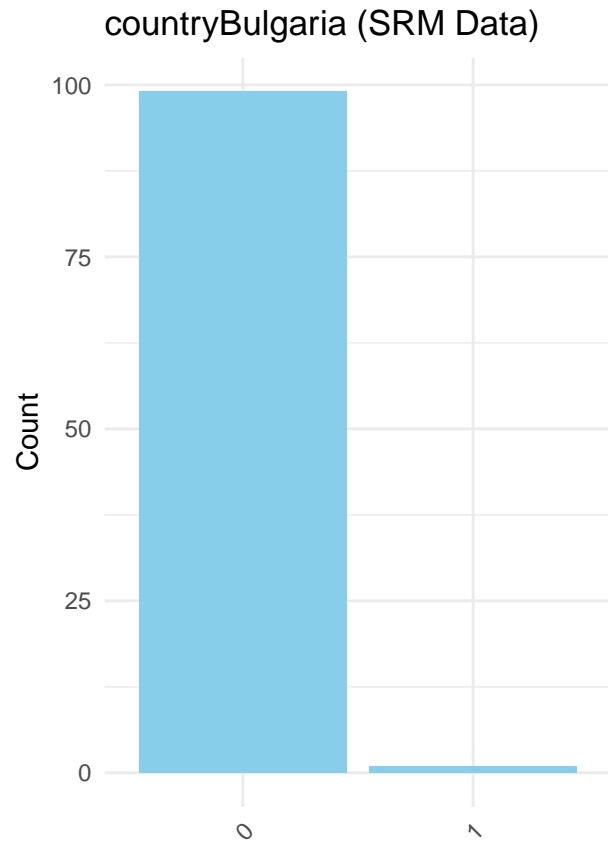


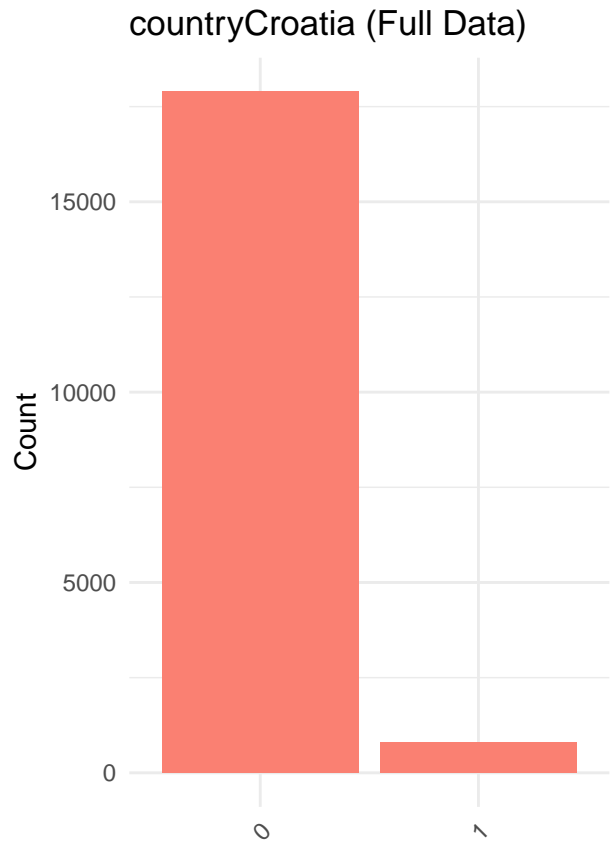
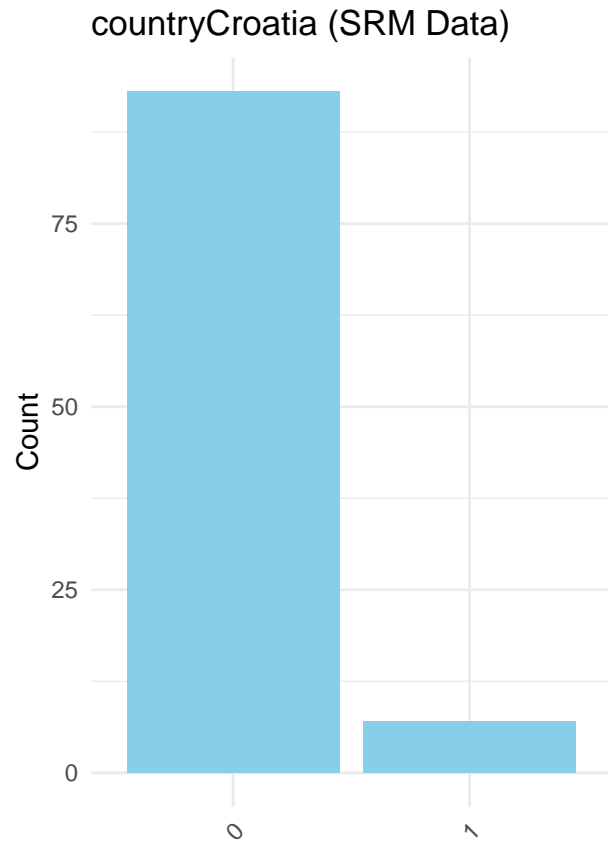


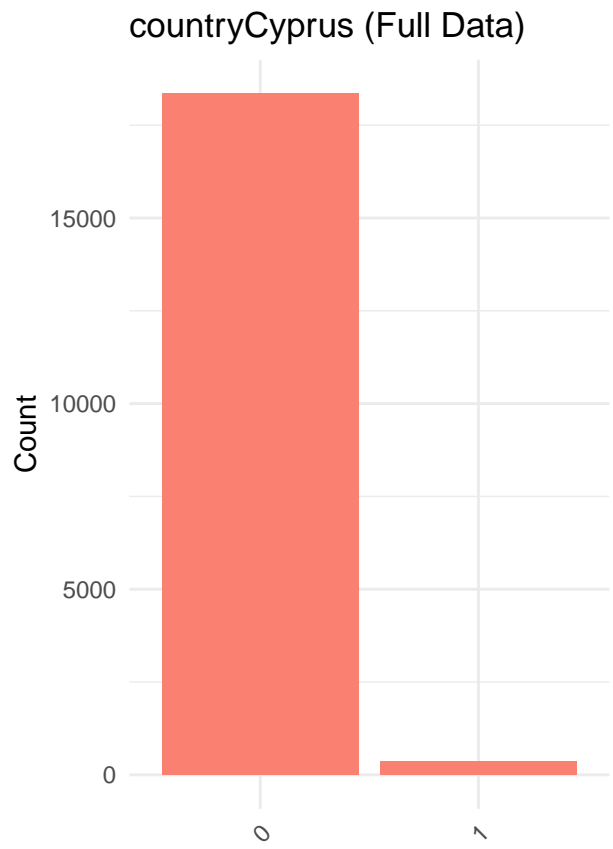
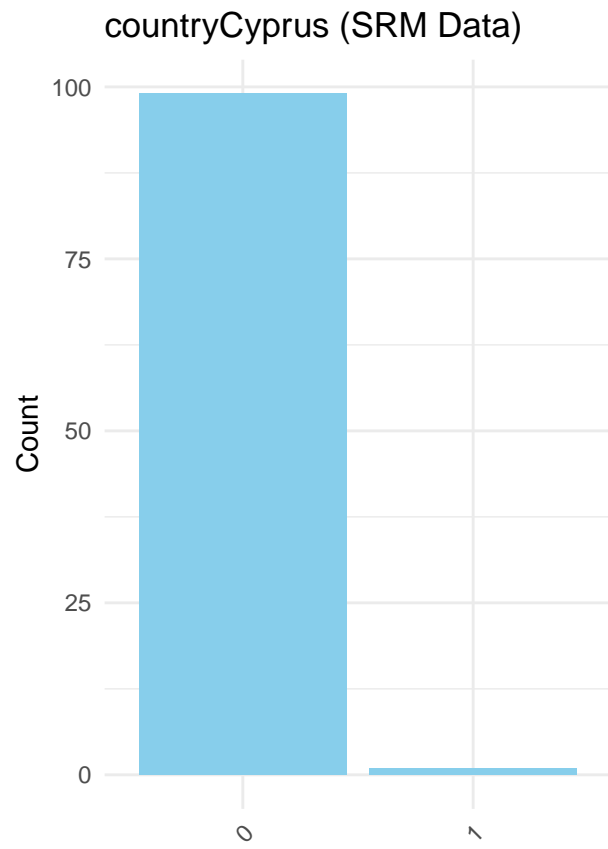


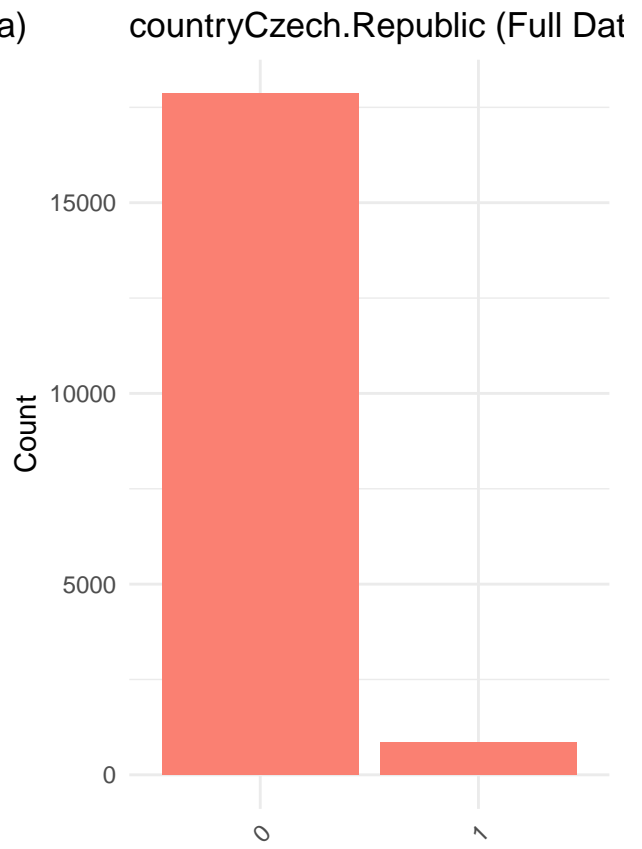
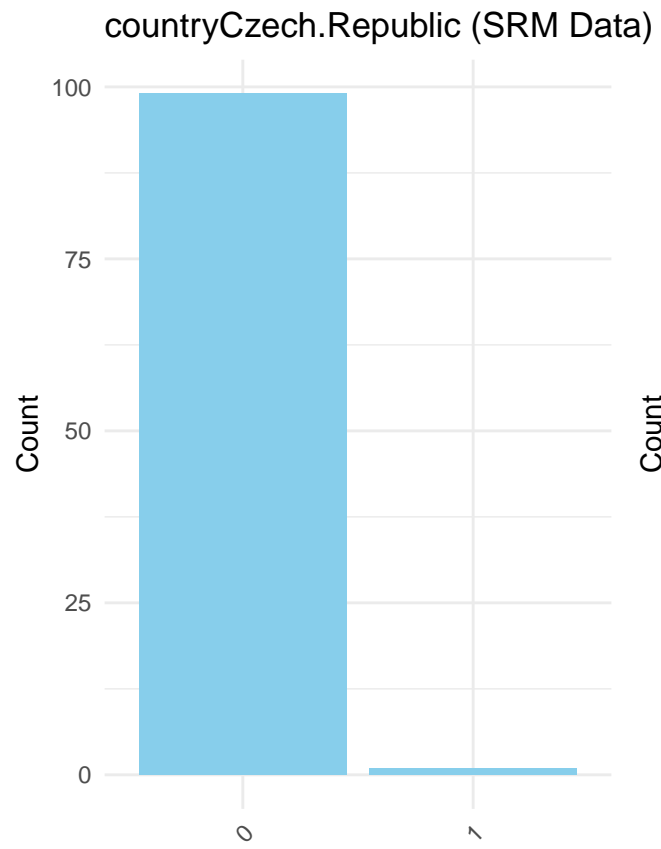


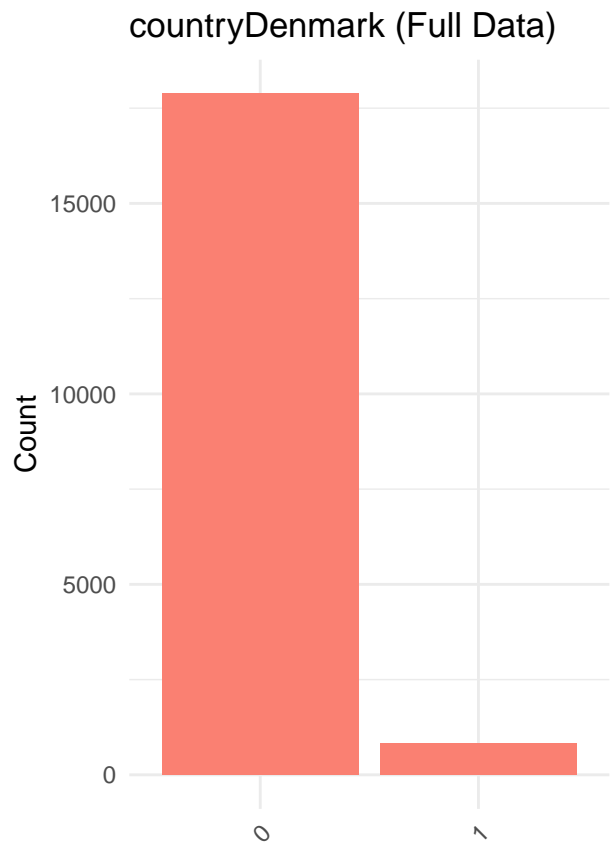
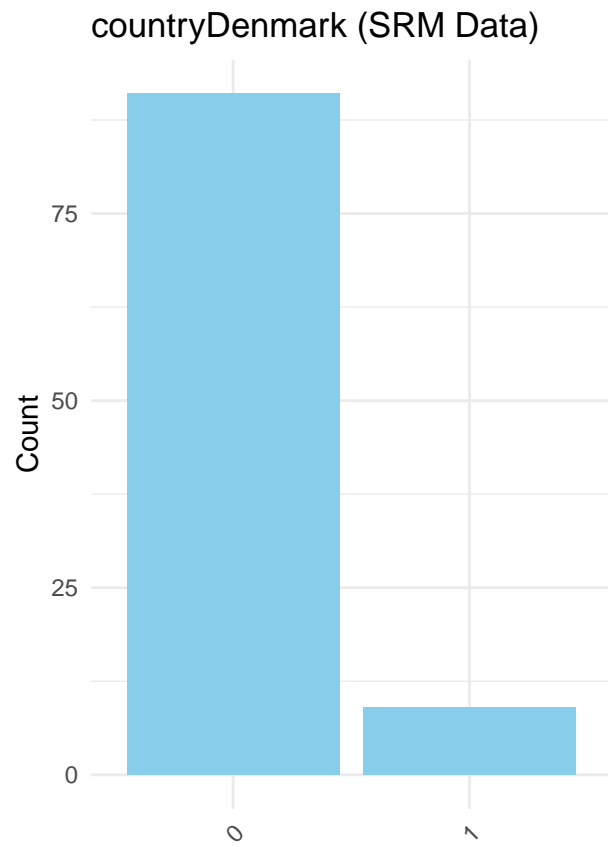


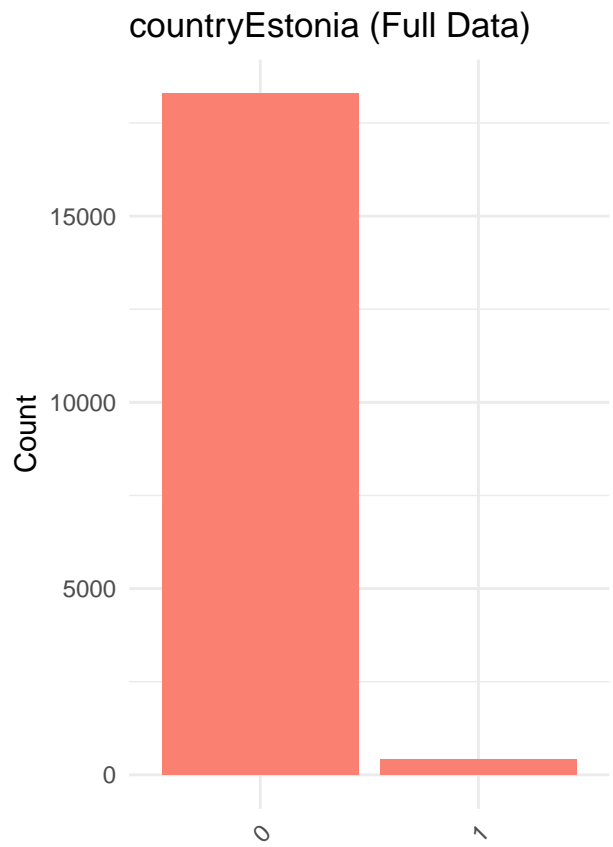
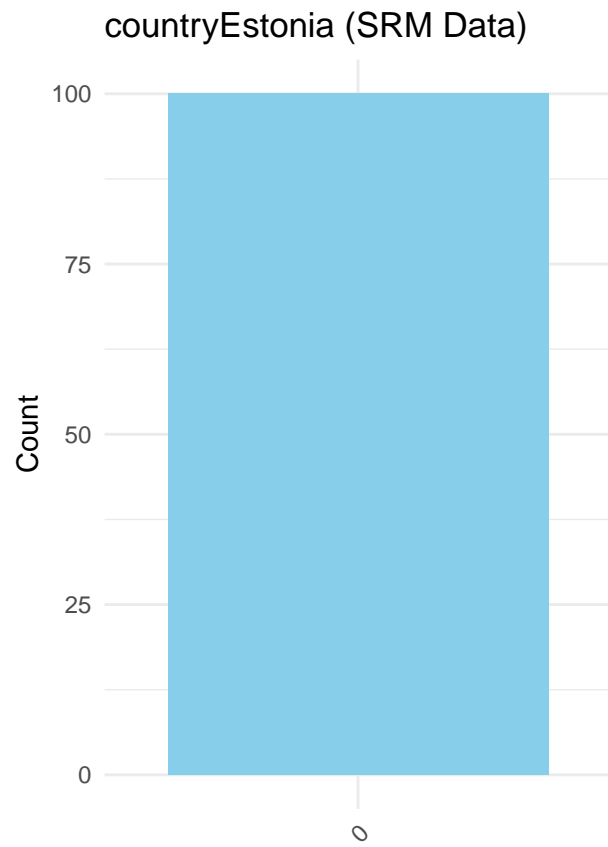


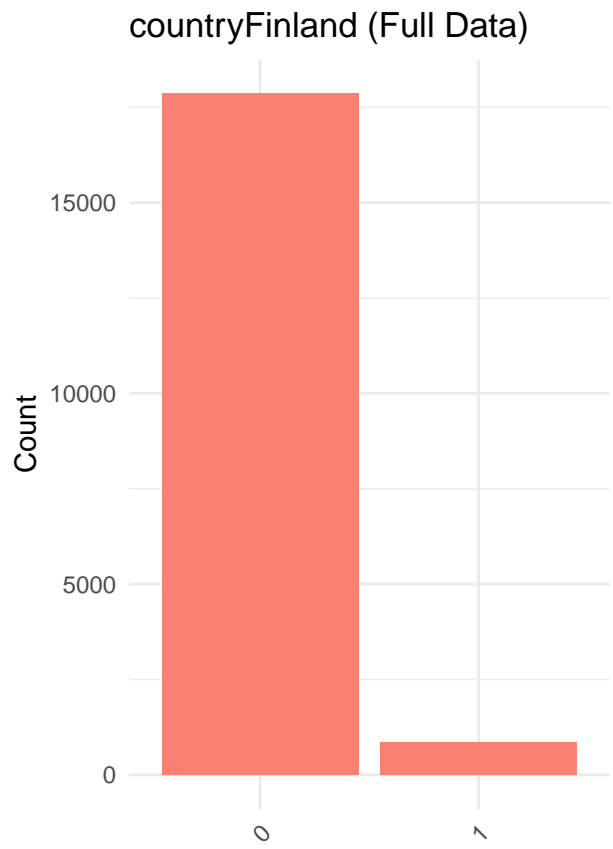
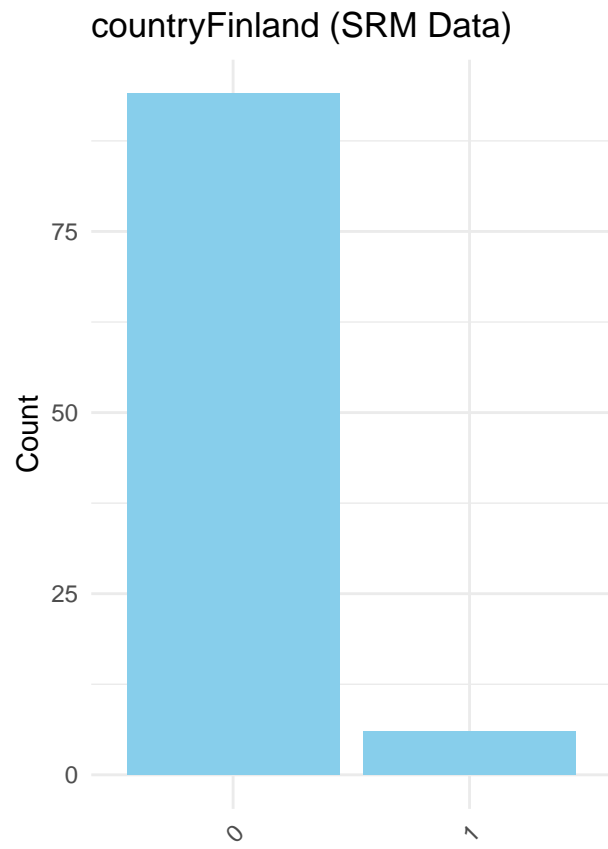


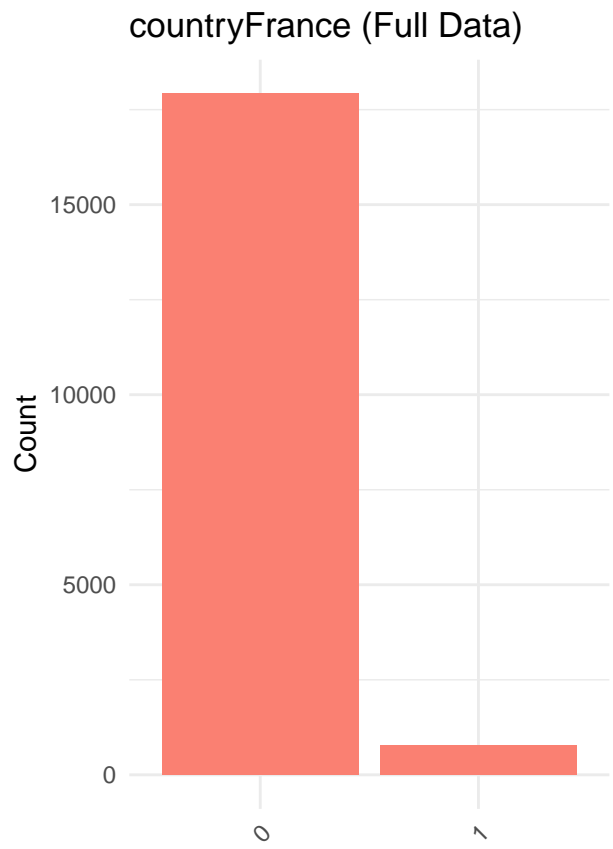
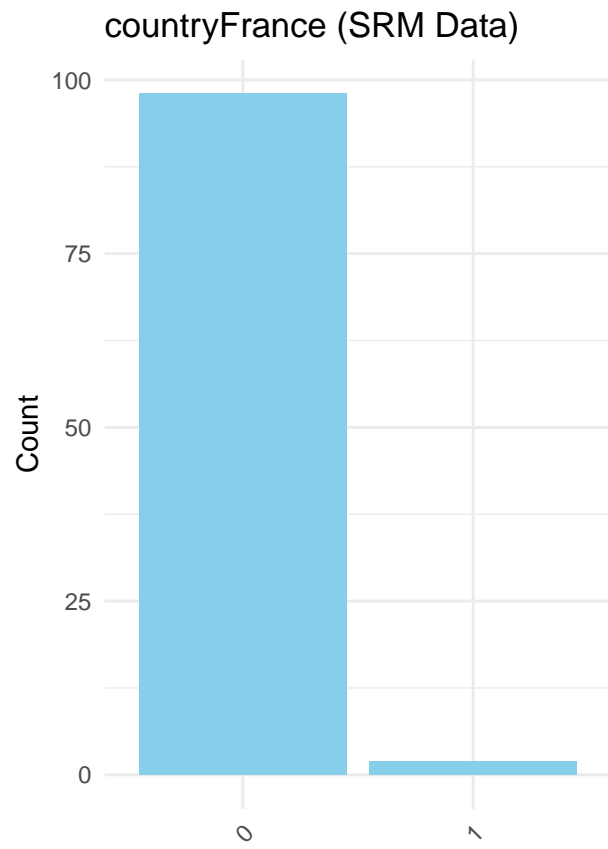


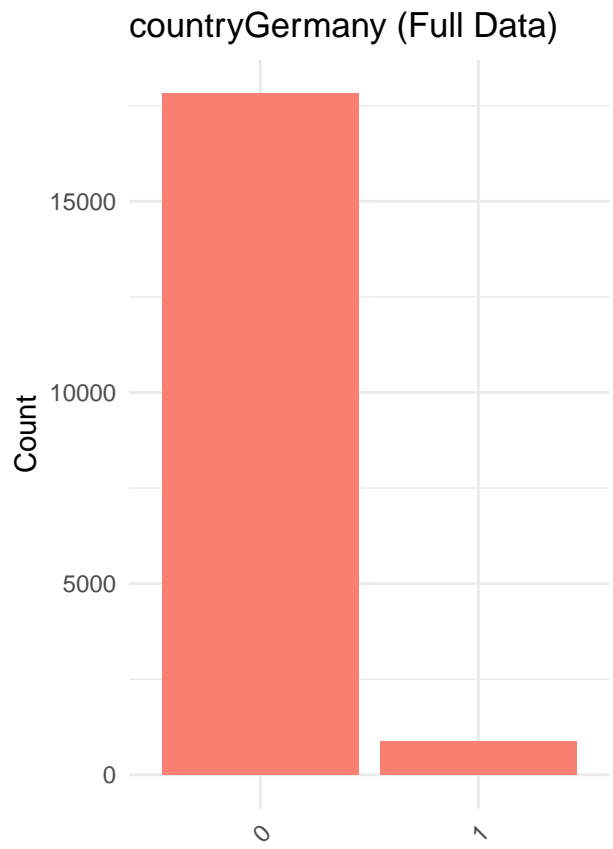
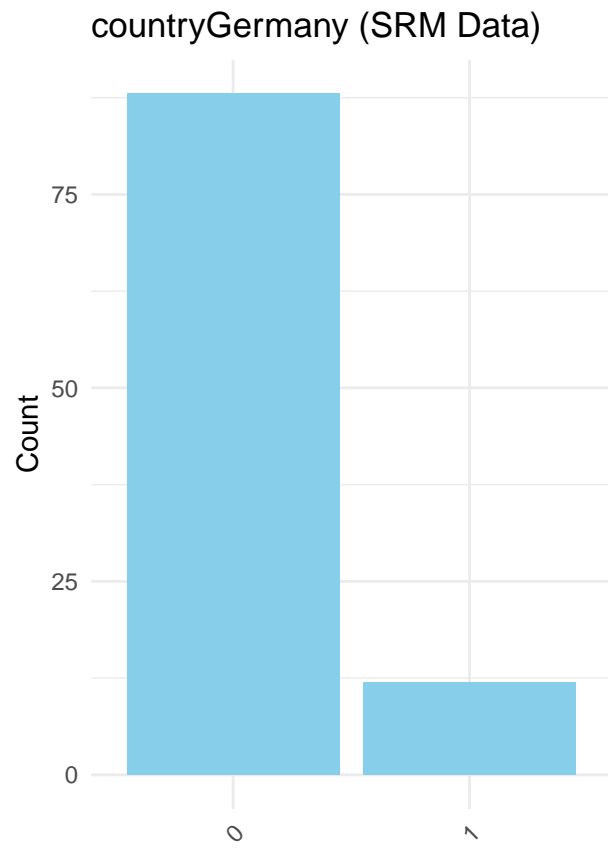


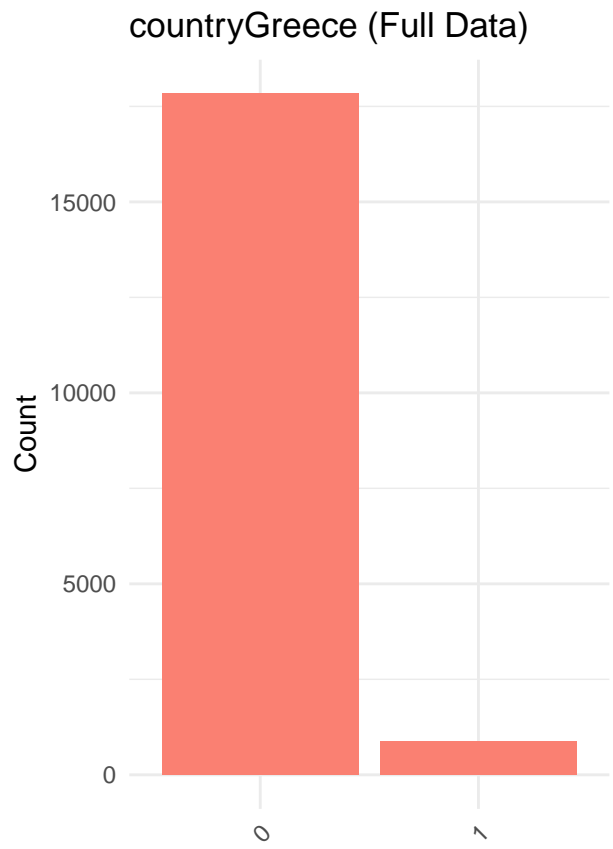
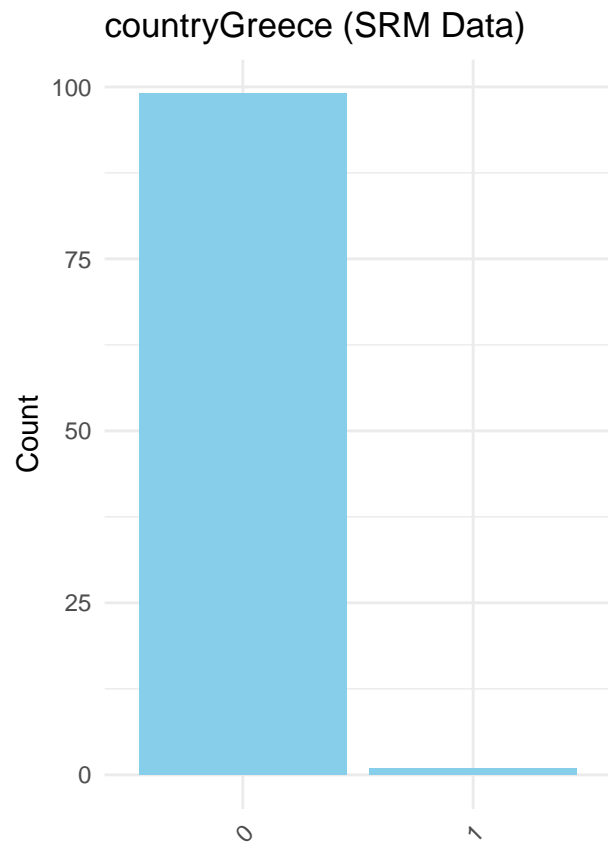


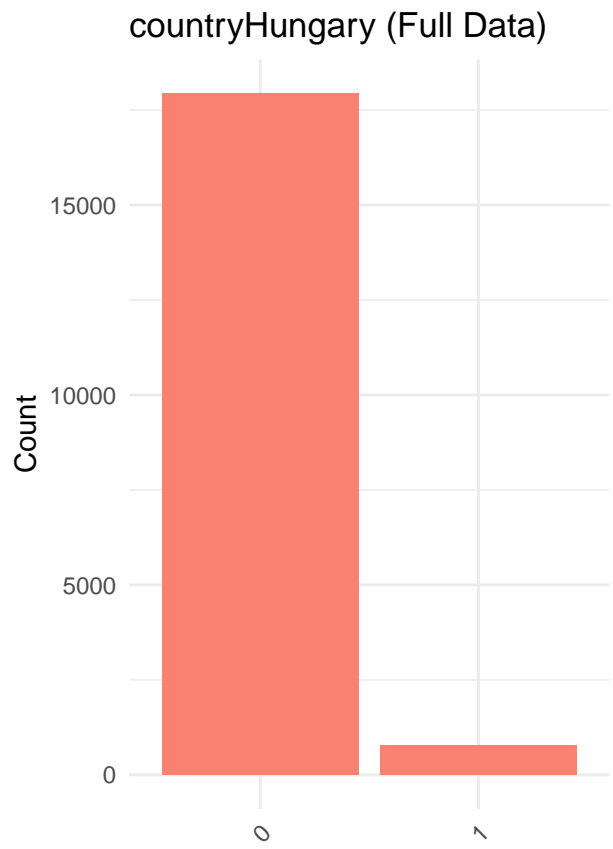
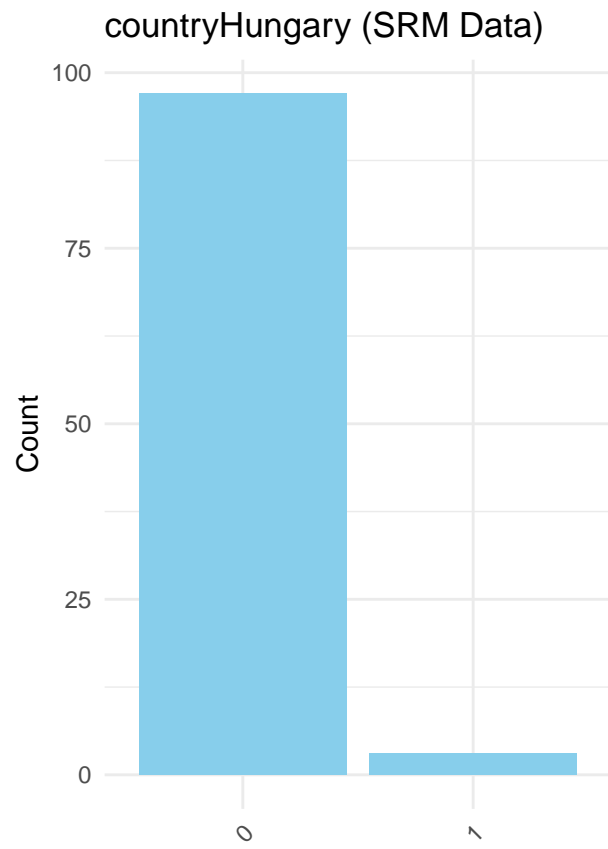


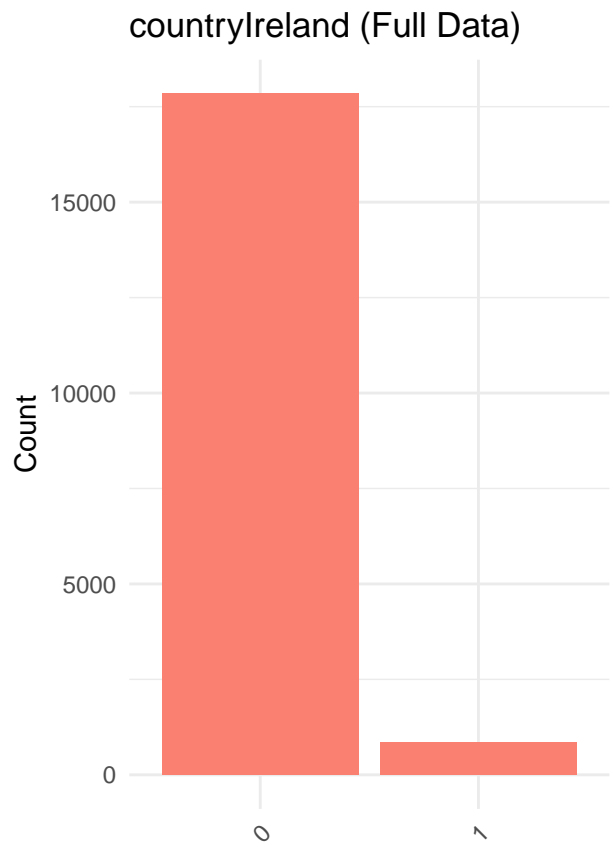
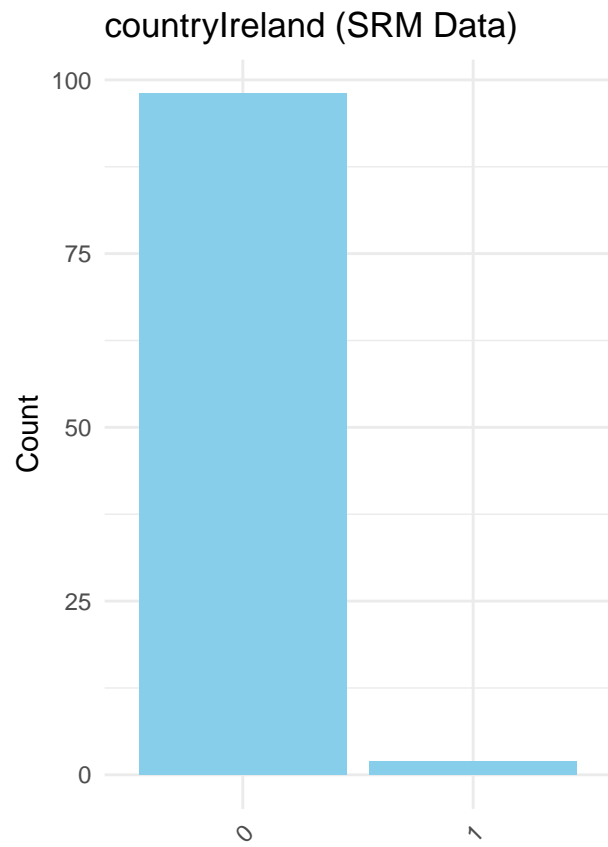


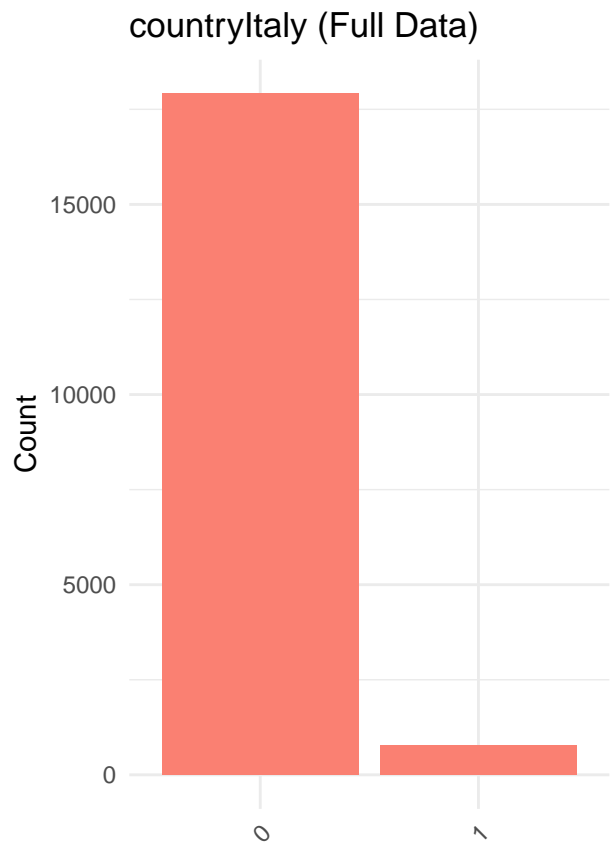
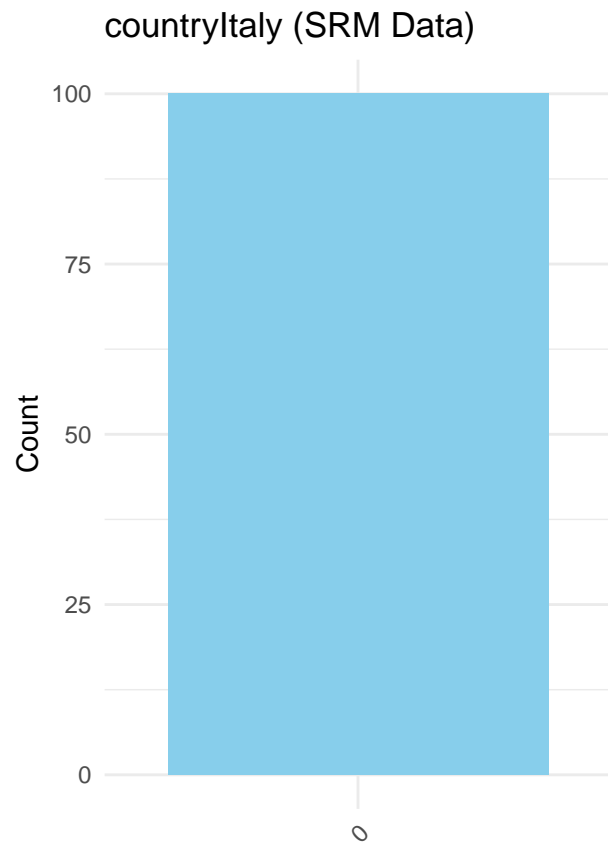


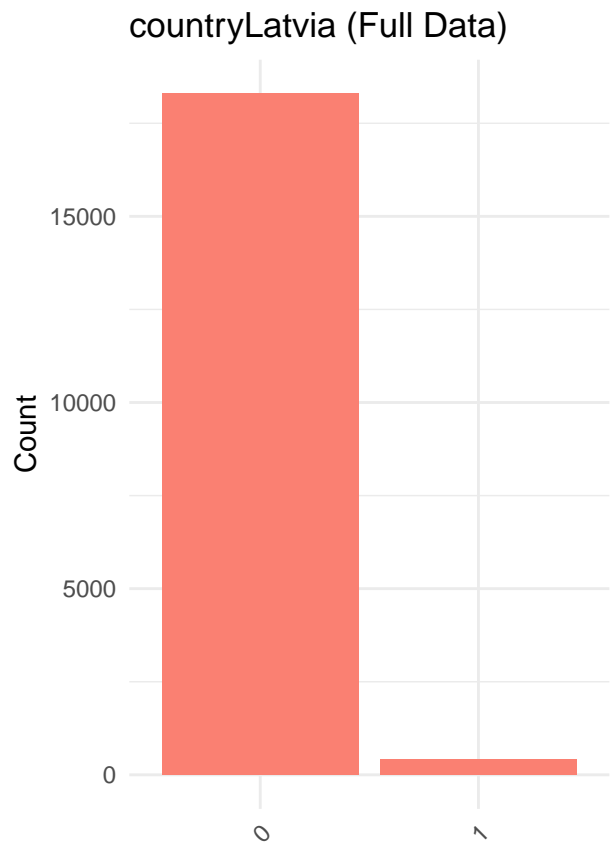
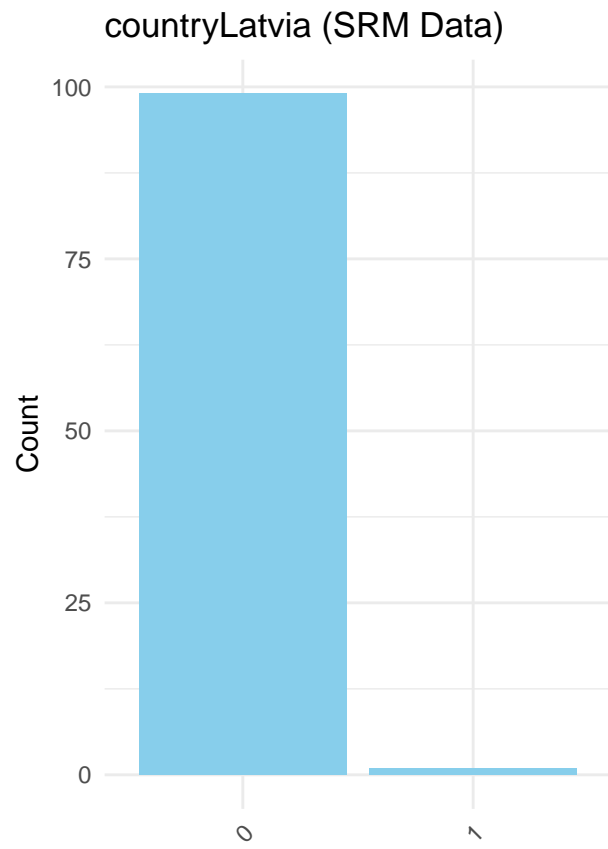


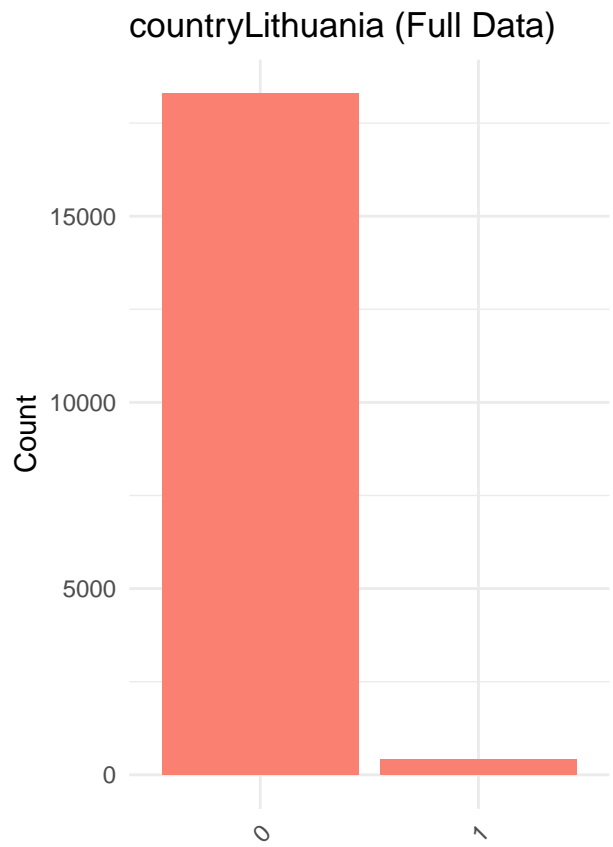
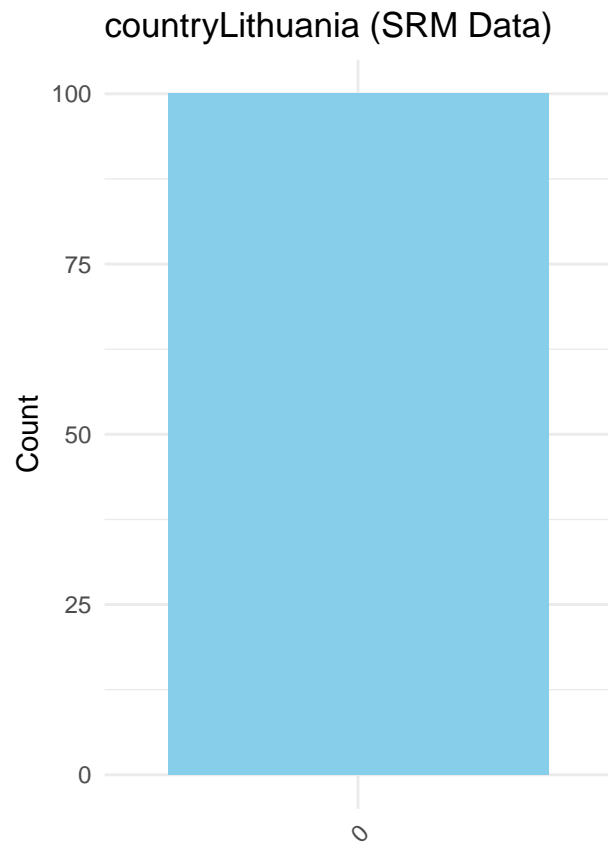


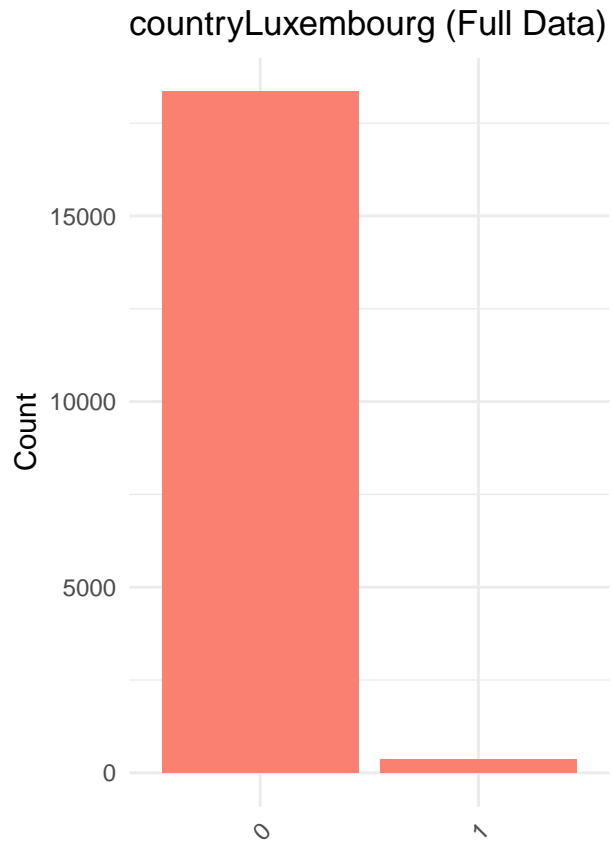
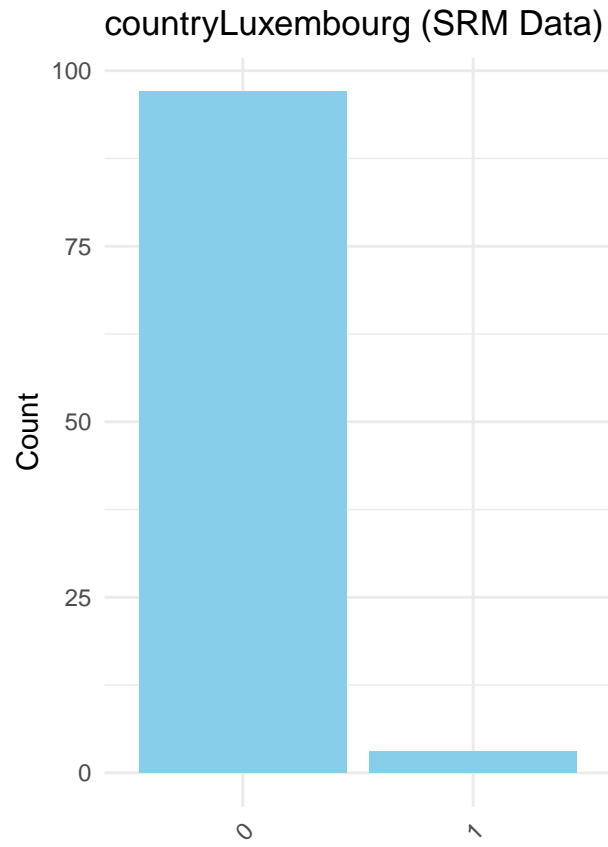


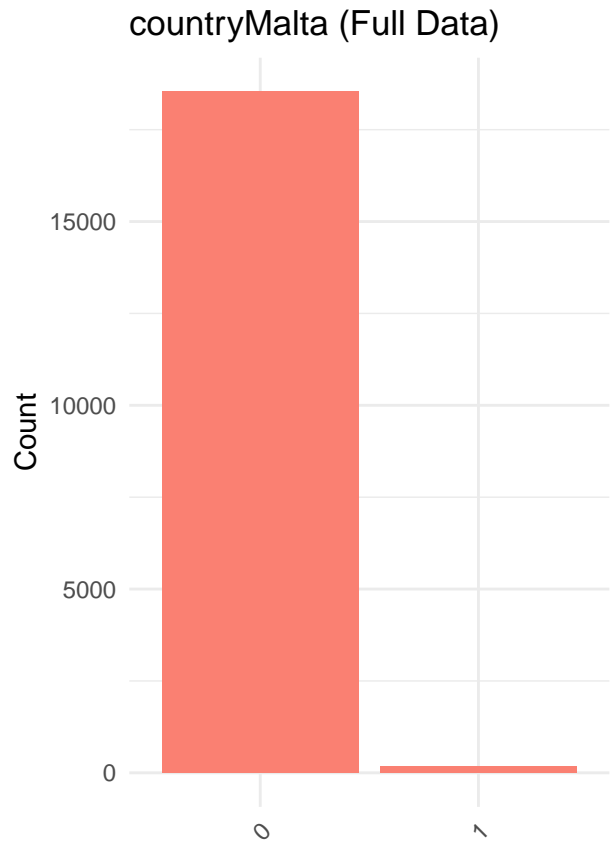
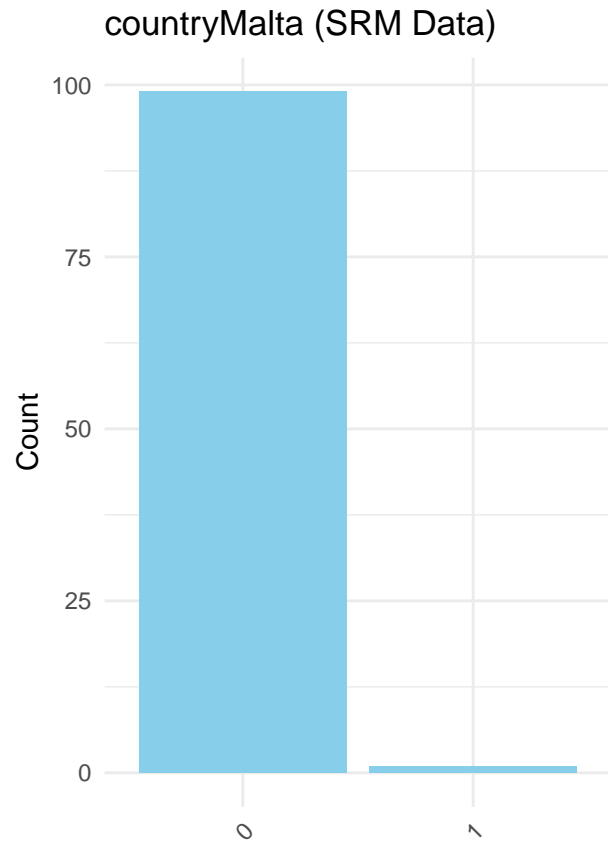


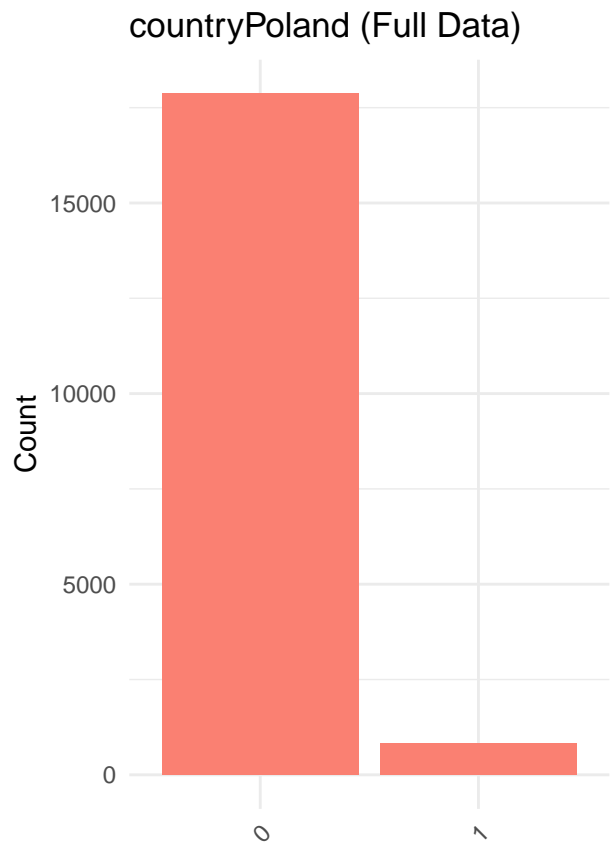
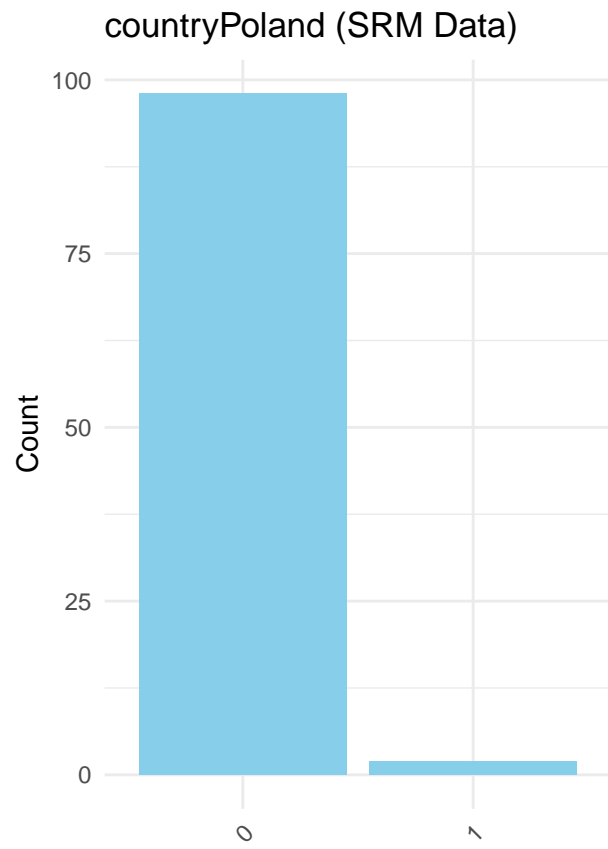


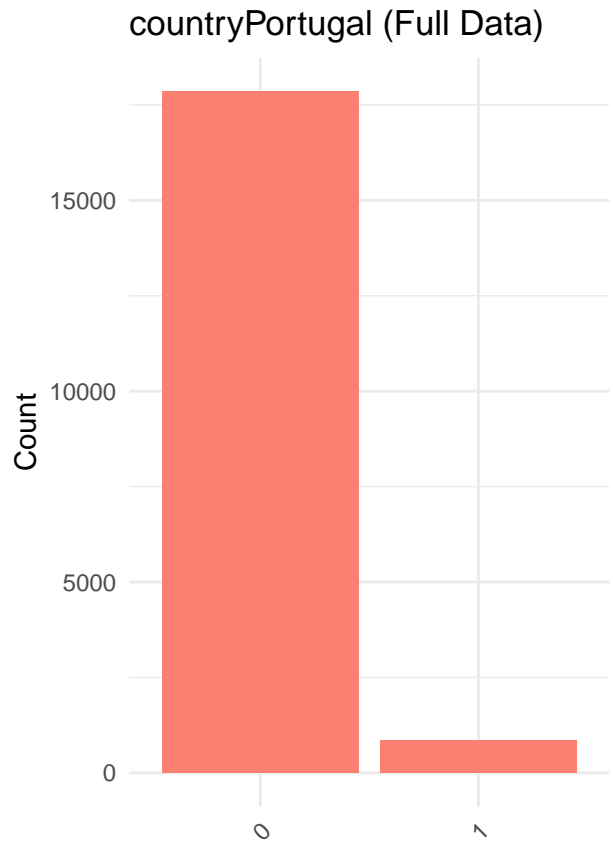
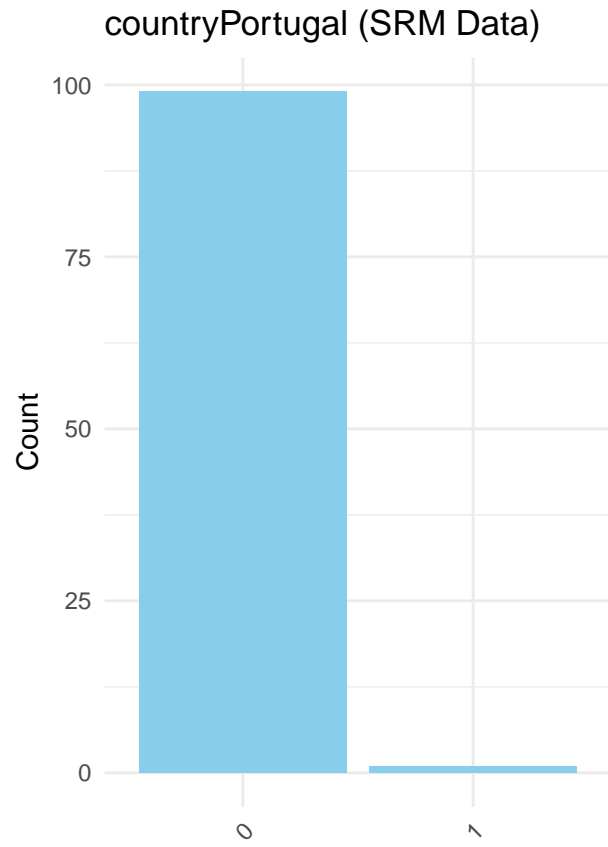


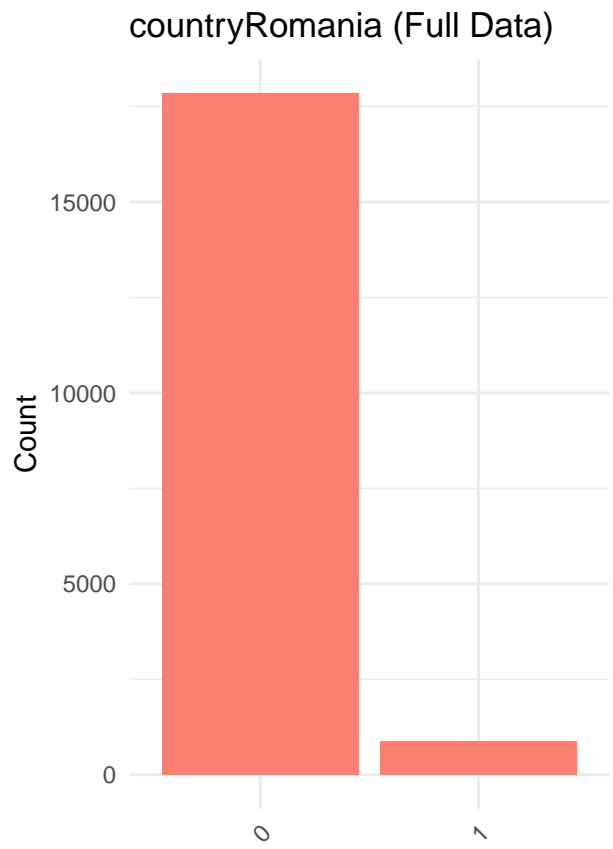
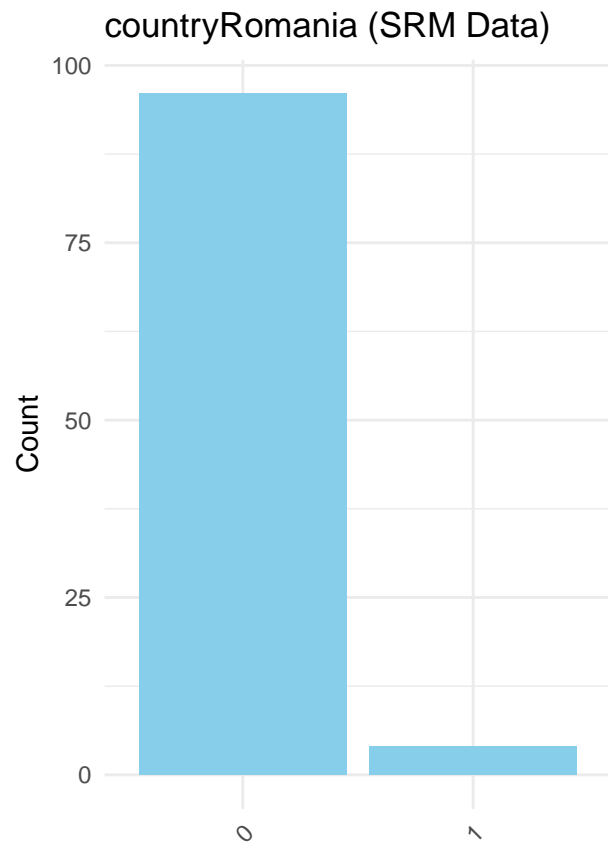


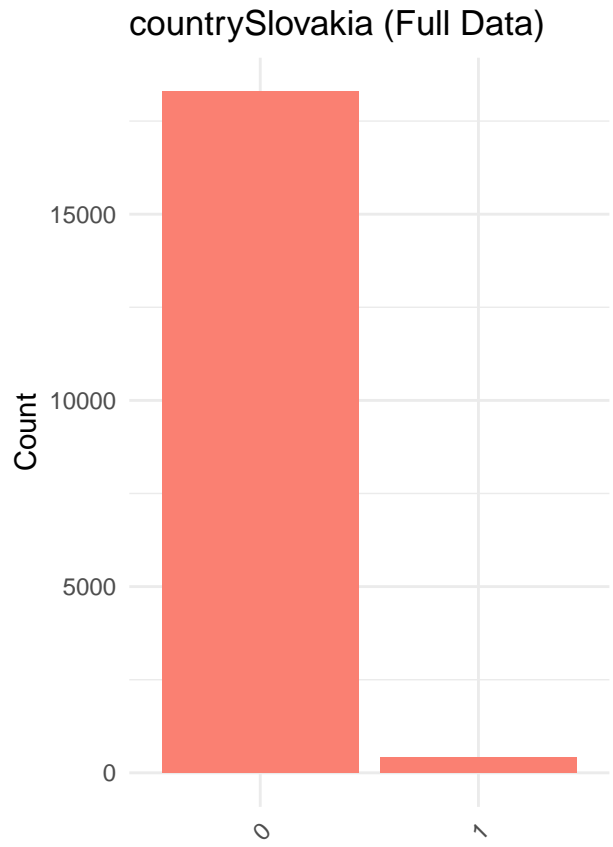
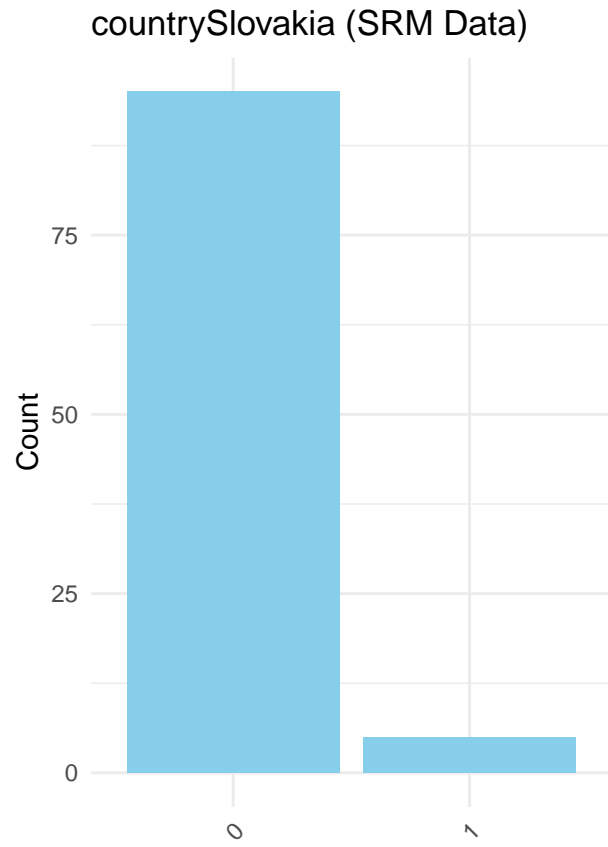


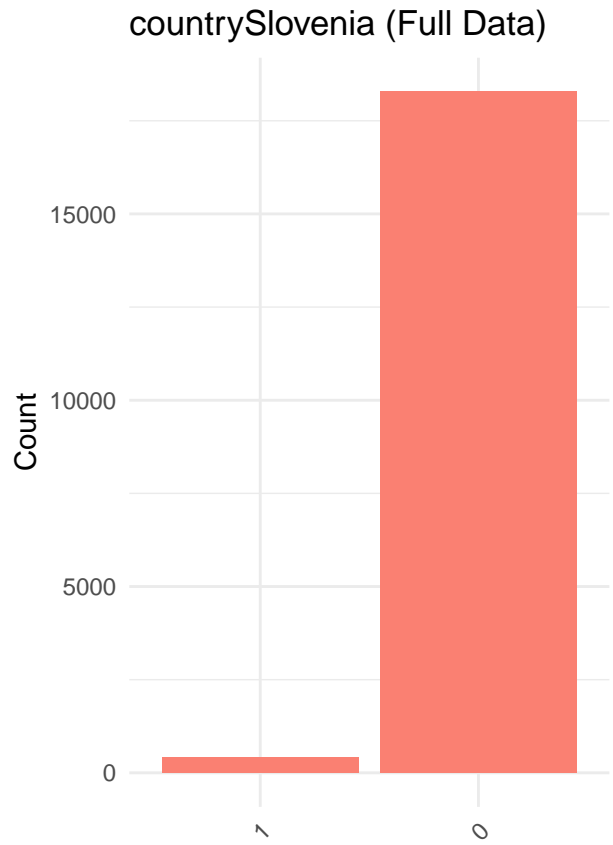
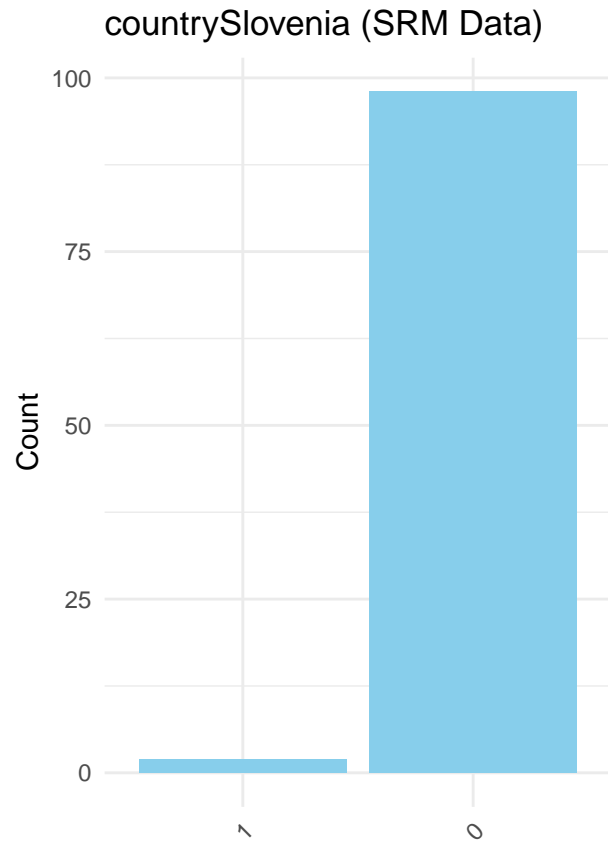


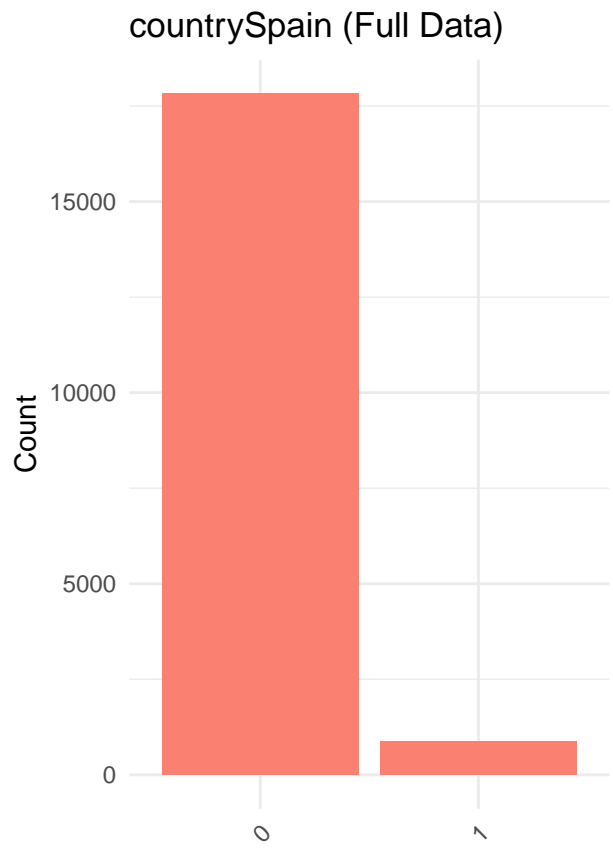
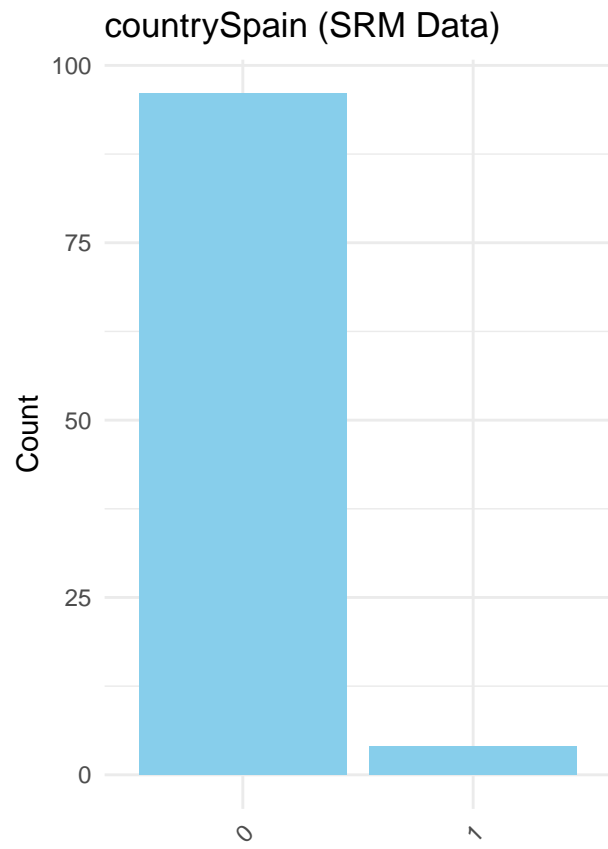


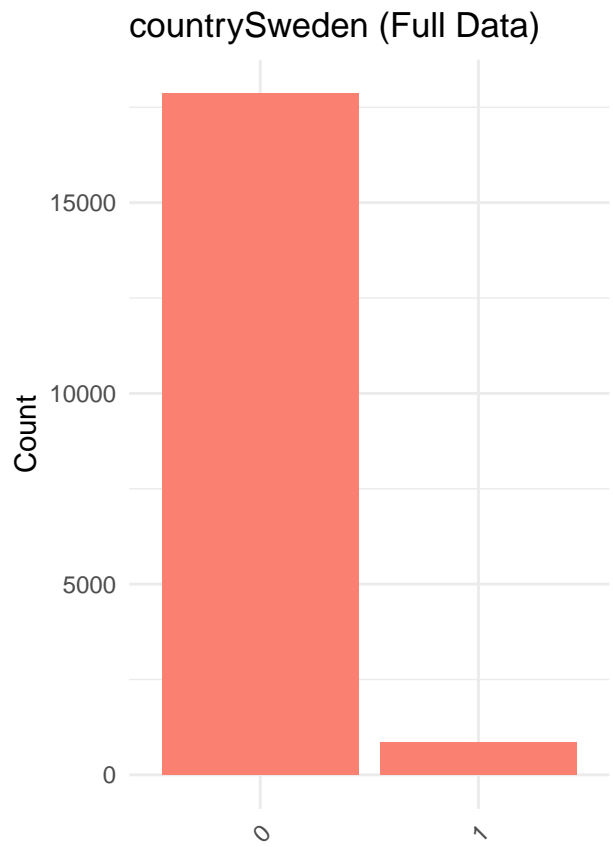
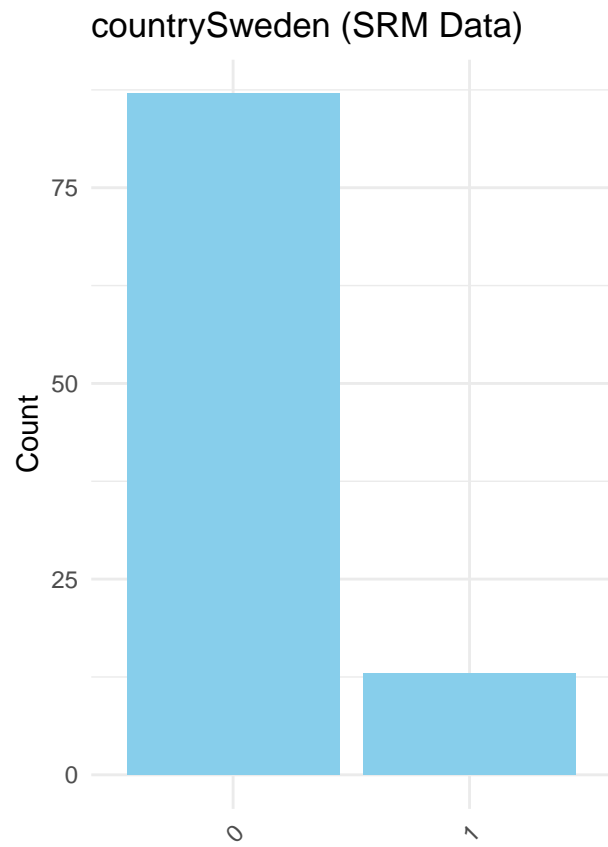


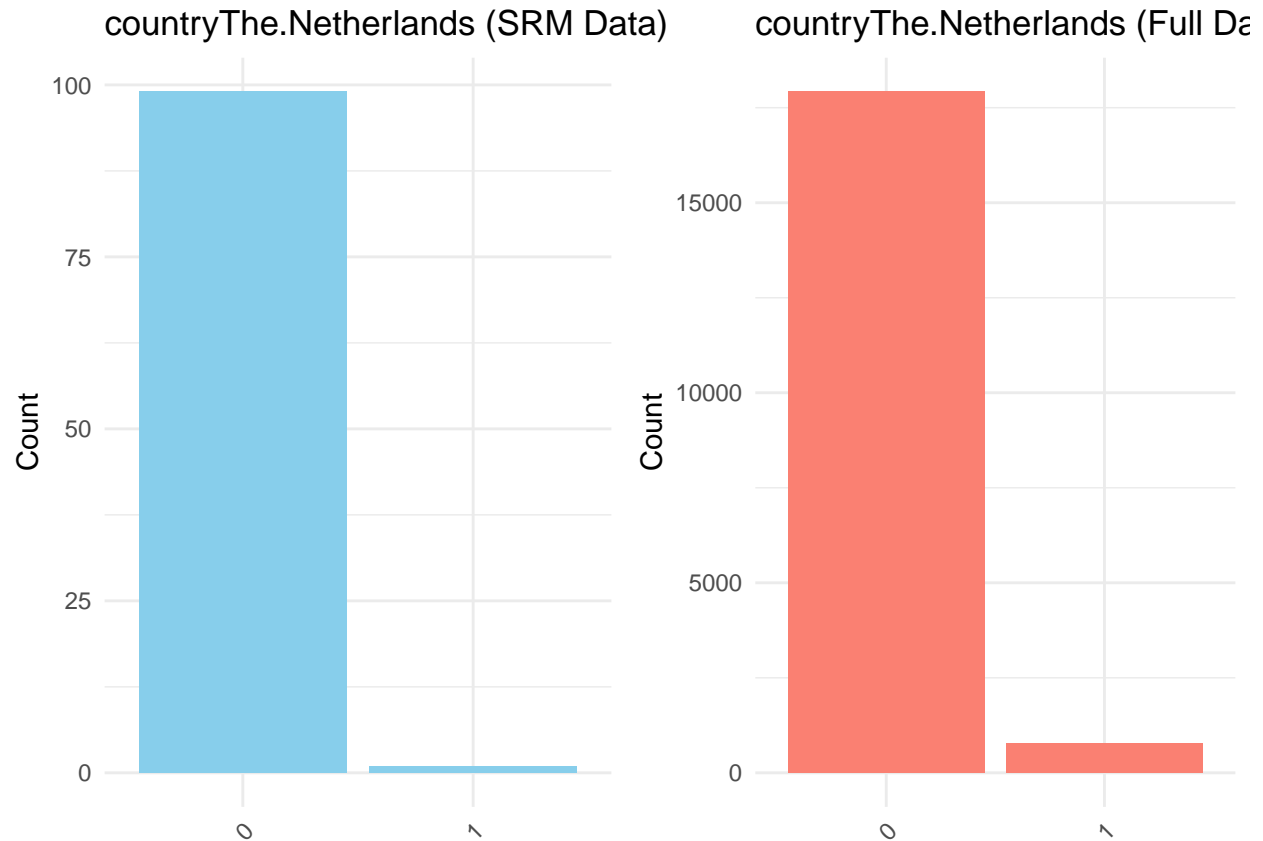












Creating the SRM critique model:

All possible combinations of the interactions respect to SRM critique is looks like :

```
. ~ . + age_scale:trust . ~ . + education:trust . ~ . + LR_scale_scale:trust . ~ . + income_scale:trust . ~ . + gender:trust . ~ . + age_scale:trust + education:trust . ~ . + age_scale:trust + LR_scale_scale:trust . ~ . + age_scale:trust + income_scale:trust . ~ . + age_scale:trust + gender:trust . ~ . + education:trust + LR_scale_scale:trust . ~ . + education:trust + income_scale:trust . ~ . + education:trust + gender:trust . ~ . + LR_scale_scale:trust + income_scale:trust . ~ . + LR_scale_scale:trust + gender:trust . ~ . + income_scale:trust + gender:trust . ~ . + age_scale:trust + education:trust + LR_scale_scale:trust . ~ . + age_scale:trust + education:trust + income_scale:trust . ~ . + age_scale:trust + education:trust + gender:trust . ~ . + age_scale:trust + LR_scale_scale:trust + income_scale:trust . ~ . + age_scale:trust + LR_scale_scale:trust + gender:trust . ~ . + age_scale:trust + income_scale:trust + gender:trust . ~ . + education:trust + LR_scale_scale:trust + income_scale:trust . ~ . + education:trust + LR_scale_scale:trust + gender:trust . ~ . + education:trust + income_scale:trust + gender:trust . ~ . + LR_scale_scale:trust + income_scale:trust + gender:trust . ~ . + age_scale:trust + education:trust + LR_scale_scale:trust + income_scale:trust . ~ . + age_scale:trust + education:trust + LR_scale_scale:trust + gender:trust . ~ . + age_scale:trust + education:trust + income_scale:trust + gender:trust . ~ . + age_scale:trust + LR_scale_scale:trust + income_scale:trust + gender:trust . ~ . + education:trust + LR_scale_scale:trust + income_scale:trust + gender:trust . ~ . + age_scale:trust + education:trust + LR_scale_scale:trust + income_scale:trust + gender:trust
```

```
srm_critique_formula=update(full_formula,. ~ . + age_scale:trust)
```

```
srm_critique_model=glm(formula =srm_critique_formula ,family = binomial(link = "logit"),data=train_data
```

```
## Warning in eval(family$initialize): non-integer #successes in a binomial glm!
```

```
summary(srm_critique_model)
```

```
##
## Call:
## glm(formula = srm_critique_formula, family = binomial(link = "logit"),
##      data = train_data, weights = country_w)
##
## Coefficients:
##              Estimate Std. Error z value Pr(>|z|)
## (Intercept)    0.1324446  0.1092200   1.213  0.22527
## income_scale   -0.0435607  0.0227317  -1.916  0.05533 .
## trustNo really confident  0.5163152  0.0575343   8.974 < 2e-16 ***
## trustRather confident  1.0413448  0.0639744  16.278 < 2e-16 ***
## trustVery confident  1.3710265  0.0993701  13.797 < 2e-16 ***
## LR_scale_scale -0.0792614  0.0099683  -7.951 1.84e-15 ***
## new_ccknowledge_index_scale  3.3359968  0.1392912  23.950 < 2e-16 ***
## residencetown   -0.0519311  0.0483008  -1.075  0.28230
## residencerural  -0.1798561  0.0575040  -3.128  0.00176 **
## age_scale       -0.0004908  0.0030227  -0.162  0.87102
## gendermale      -0.0363312  0.0428705  -0.847  0.39674
## educationprimary  0.0972120  0.0629674   1.544  0.12263
## educationtertiary  0.0917967  0.0471744   1.946  0.05167 .
## has_childrenyes  0.0620776  0.0469028   1.324  0.18566
## countryBelgium   0.3550353  0.1312815   2.704  0.00684 **
## countryBulgaria  0.9359661  0.1409107   6.642 3.09e-11 ***
## countryCroatia   1.0805284  0.1410249   7.662 1.83e-14 ***
## countryCyprus     1.1918366  0.2164242   5.507 3.65e-08 ***
## countryCzech.Republic  0.3511166  0.1279172   2.745  0.00605 **
## countryDenmark   0.5881633  0.1360676   4.323 1.54e-05 ***
## countryEstonia   -0.0295715  0.1502871  -0.197  0.84401
## countryFinland   0.3432525  0.1300498   2.639  0.00831 **
## countryFrance     0.5448647  0.1362616   3.999 6.37e-05 ***
## countryGermany   -0.0056905  0.1217131  -0.047  0.96271
## countryGreece     1.3256636  0.1529074   8.670 < 2e-16 ***
## countryHungary    0.3169152  0.1301902   2.434  0.01492 *
## countryIreland    0.9571435  0.1399371   6.840 7.93e-12 ***
## countryItaly      0.8020915  0.1393301   5.757 8.57e-09 ***
## countryLatvia     0.2045608  0.1553457   1.317  0.18790
## countryLithuania  0.1558939  0.1570157   0.993  0.32078
## countryLuxembourg 0.4704085  0.1753281   2.683  0.00730 **
## countryMalta      2.2427501  0.4248783   5.279 1.30e-07 ***
## countryPoland     0.2068465  0.1281218   1.614  0.10643
## countryPortugal   1.3528302  0.1517026   8.918 < 2e-16 ***
## countryRomania    0.5629941  0.1289535   4.366 1.27e-05 ***
## countrySlovakia   0.2108706  0.1545533   1.364  0.17245
## countrySlovenia   0.9014158  0.1747458   5.158 2.49e-07 ***
## countrySpain      0.6076286  0.1326990   4.579 4.67e-06 ***
## countrySweden     0.0796602  0.1253699   0.635  0.52517
```

```
## countryThe.Netherlands      0.3248200  0.1305375   2.488  0.01283 *
## trustNo really confident:age_scale  0.0080621  0.0035162   2.293  0.02186 *
## trustRather confident:age_scale    0.0049407  0.0038273   1.291  0.19673
## trustVery confident:age_scale     -0.0078091  0.0058785  -1.328  0.18404
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## (Dispersion parameter for binomial family taken to be 1)
##
## Null deviance: 15803  on 14958  degrees of freedom
## Residual deviance: 14302  on 14916  degrees of freedom
## AIC: 14916
##
## Number of Fisher Scoring iterations: 5
```

```
yhat_updated_train=predict(srm_critique_model,newdata=train_data,type="response")
yhat_updated_test=predict(srm_critique_model,newdata=test_data,type="response")

class_pred_updated_train <- ifelse(yhat_updated_train > 0.5, 1, 0)
class_pred_updated_test  <- ifelse(yhat_updated_test > 0.5, 1, 0)
```

Checking the Accuracy Metrics respect to the Train Data for the Updated Model:

```
confusionMatrix(as.factor(train_data$ctax_binary),as.factor(class_pred_updated_train))
```

```
## Confusion Matrix and Statistics
##
##           Reference
## Prediction      0      1
##           0  438  2931
##           1  260 11330
##
##           Accuracy : 0.7867
##           95% CI : (0.78, 0.7932)
##           No Information Rate : 0.9533
##           P-Value [Acc > NIR] : 1
##
##           Kappa : 0.1497
##
## Mcnemar's Test P-Value : <2e-16
##
##           Sensitivity : 0.62751
##           Specificity : 0.79447
##           Pos Pred Value : 0.13001
##           Neg Pred Value : 0.97757
##           Prevalence : 0.04666
##           Detection Rate : 0.02928
##           Detection Prevalence : 0.22522
##           Balanced Accuracy : 0.71099
```

```

##
##      'Positive' Class : 0
##

confusionMatrix(as.factor(test_data$ctax_binary),as.factor(class_pred_updated_test))

## Confusion Matrix and Statistics
##
##           Reference
## Prediction    0    1
##           0  117  733
##           1   51 2838
##
##           Accuracy : 0.7903
##           95% CI : (0.7769, 0.8033)
##      No Information Rate : 0.9551
##      P-Value [Acc > NIR] : 1
##
##           Kappa : 0.1674
##
##  Mcnemar's Test P-Value : <2e-16
##
##           Sensitivity : 0.69643
##           Specificity : 0.79474
##           Pos Pred Value : 0.13765
##           Neg Pred Value : 0.98235
##           Prevalence : 0.04493
##           Detection Rate : 0.03129
##      Detection Prevalence : 0.22733
##           Balanced Accuracy : 0.74558
##
##      'Positive' Class : 0
##

accuracy_df <- data.frame(
  model = c("Baseline", "SRM Critique", "ML Model"),
  accuracy = c(
    mean(class_pred_test == test_data$ctax_binary, na.rm = TRUE),
    mean(class_pred_updated_test == test_data$ctax_binary, na.rm = TRUE),
    mean(yhat_rf_test == test_data$ctax_binary, na.rm = TRUE)
  )
)

accuracy_df$label <- substr(as.character(accuracy_df$accuracy), 1, 6)

library(ggplot2)

ggplot(accuracy_df, aes(x = model, y = accuracy, fill = model)) +
  geom_bar(stat = "identity") +
  geom_text(aes(label = label), vjust = -0.5) +
  labs(
    title = "Model Accuracy Comparison",

```

```

x = "Model",
y = "Accuracy"
) +
scale_fill_manual(values = c(
  "Baseline" = "lightgreen",
  "Baseline + SRM" = "steelblue",
  "ML Model" = "darkgreen"
)) +
ylim(0, 1) +
theme_minimal()

```

